

15-441/641: Computer Networks

Intradomain Routing, cont'd

15-441 Spring 2019

Profs Peter Steenkiste & **Justine Sherry**



Carnegie
Mellon
University

I'm sorry

- See Piazza for a discussion of what a diameter of a graph means.
- Here's some candy because I said the wrong thing.
- I do not recommend that you take graph theory from me.



Welcome Back

A's Route Table

	via B	via C	A's DV
	via B	via C	mindist
to B	1	2	?
to C	4	3	?
to D	2	6	?

A

- What values does A announce in its Distance Vector?

C

C's Route Table

	via B	via D
to A	1	5
to B	1	7
to D	2	3

	mindist
to A	4
to B	2
to D	0

C receives the above DV from its neighbor D. How does it change its routing table?
(Assume the link weight from C to D is 3)



Distance Vector Routing: Summary

- Each router knows the links to its neighbors
- Each router has provisional “shortest path” to **every** other router -- its **distance vector (DV)**
- Routers exchange this DV with their neighbors
- Routers look over the set of options offered by their neighbors and select the best one
- Iterative process converges to set of shortest paths



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Distance Vector
e.g RIP

Need to recompute
spanning tree if failure

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

Need to run spanning
tree protocol before
routing

Still sends new
connections everywhere.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Distance Vector
e.g RIP

Need to recompute
spanning tree if failure

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

Need to run spanning
tree protocol before
routing

Need to run DV before
routing — takes length of
longest best path time.

Still sends new
connections everywhere.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Distance Vector
e.g RIP

Need to recompute
spanning tree if failure

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

$O(\# \text{ switches} * \text{max node degree}) + O(\#nodes)$

Need to run spanning
tree protocol before
routing

Need to run DV before
routing — takes length of
longest best path time.

Still sends new
connections everywhere.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Distance Vector
e.g RIP

Need to recompute
spanning tree if failure

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

Yes

$O(\# \text{ switches} * \text{max node degree}) + O(\#nodes)$

Need to run spanning
tree protocol before
routing

Need to run DV before
routing — takes length of
longest best path time.

Still sends new
connections everywhere.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Need to recompute
spanning tree if failure

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

Need to run spanning
tree protocol before
routing

Distance Vector e.g RIP

I have some bad news.

Yes

$O(\# \text{ switches} * \text{max node degree}) + O(\#nodes)$

Need to run DV before
routing — takes length of
longest best path time.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Distance Vector Algorithms suffer long convergence times when link weights increase or when *links go down*.



Watch out:

These slides are brand new, and they are detailed.

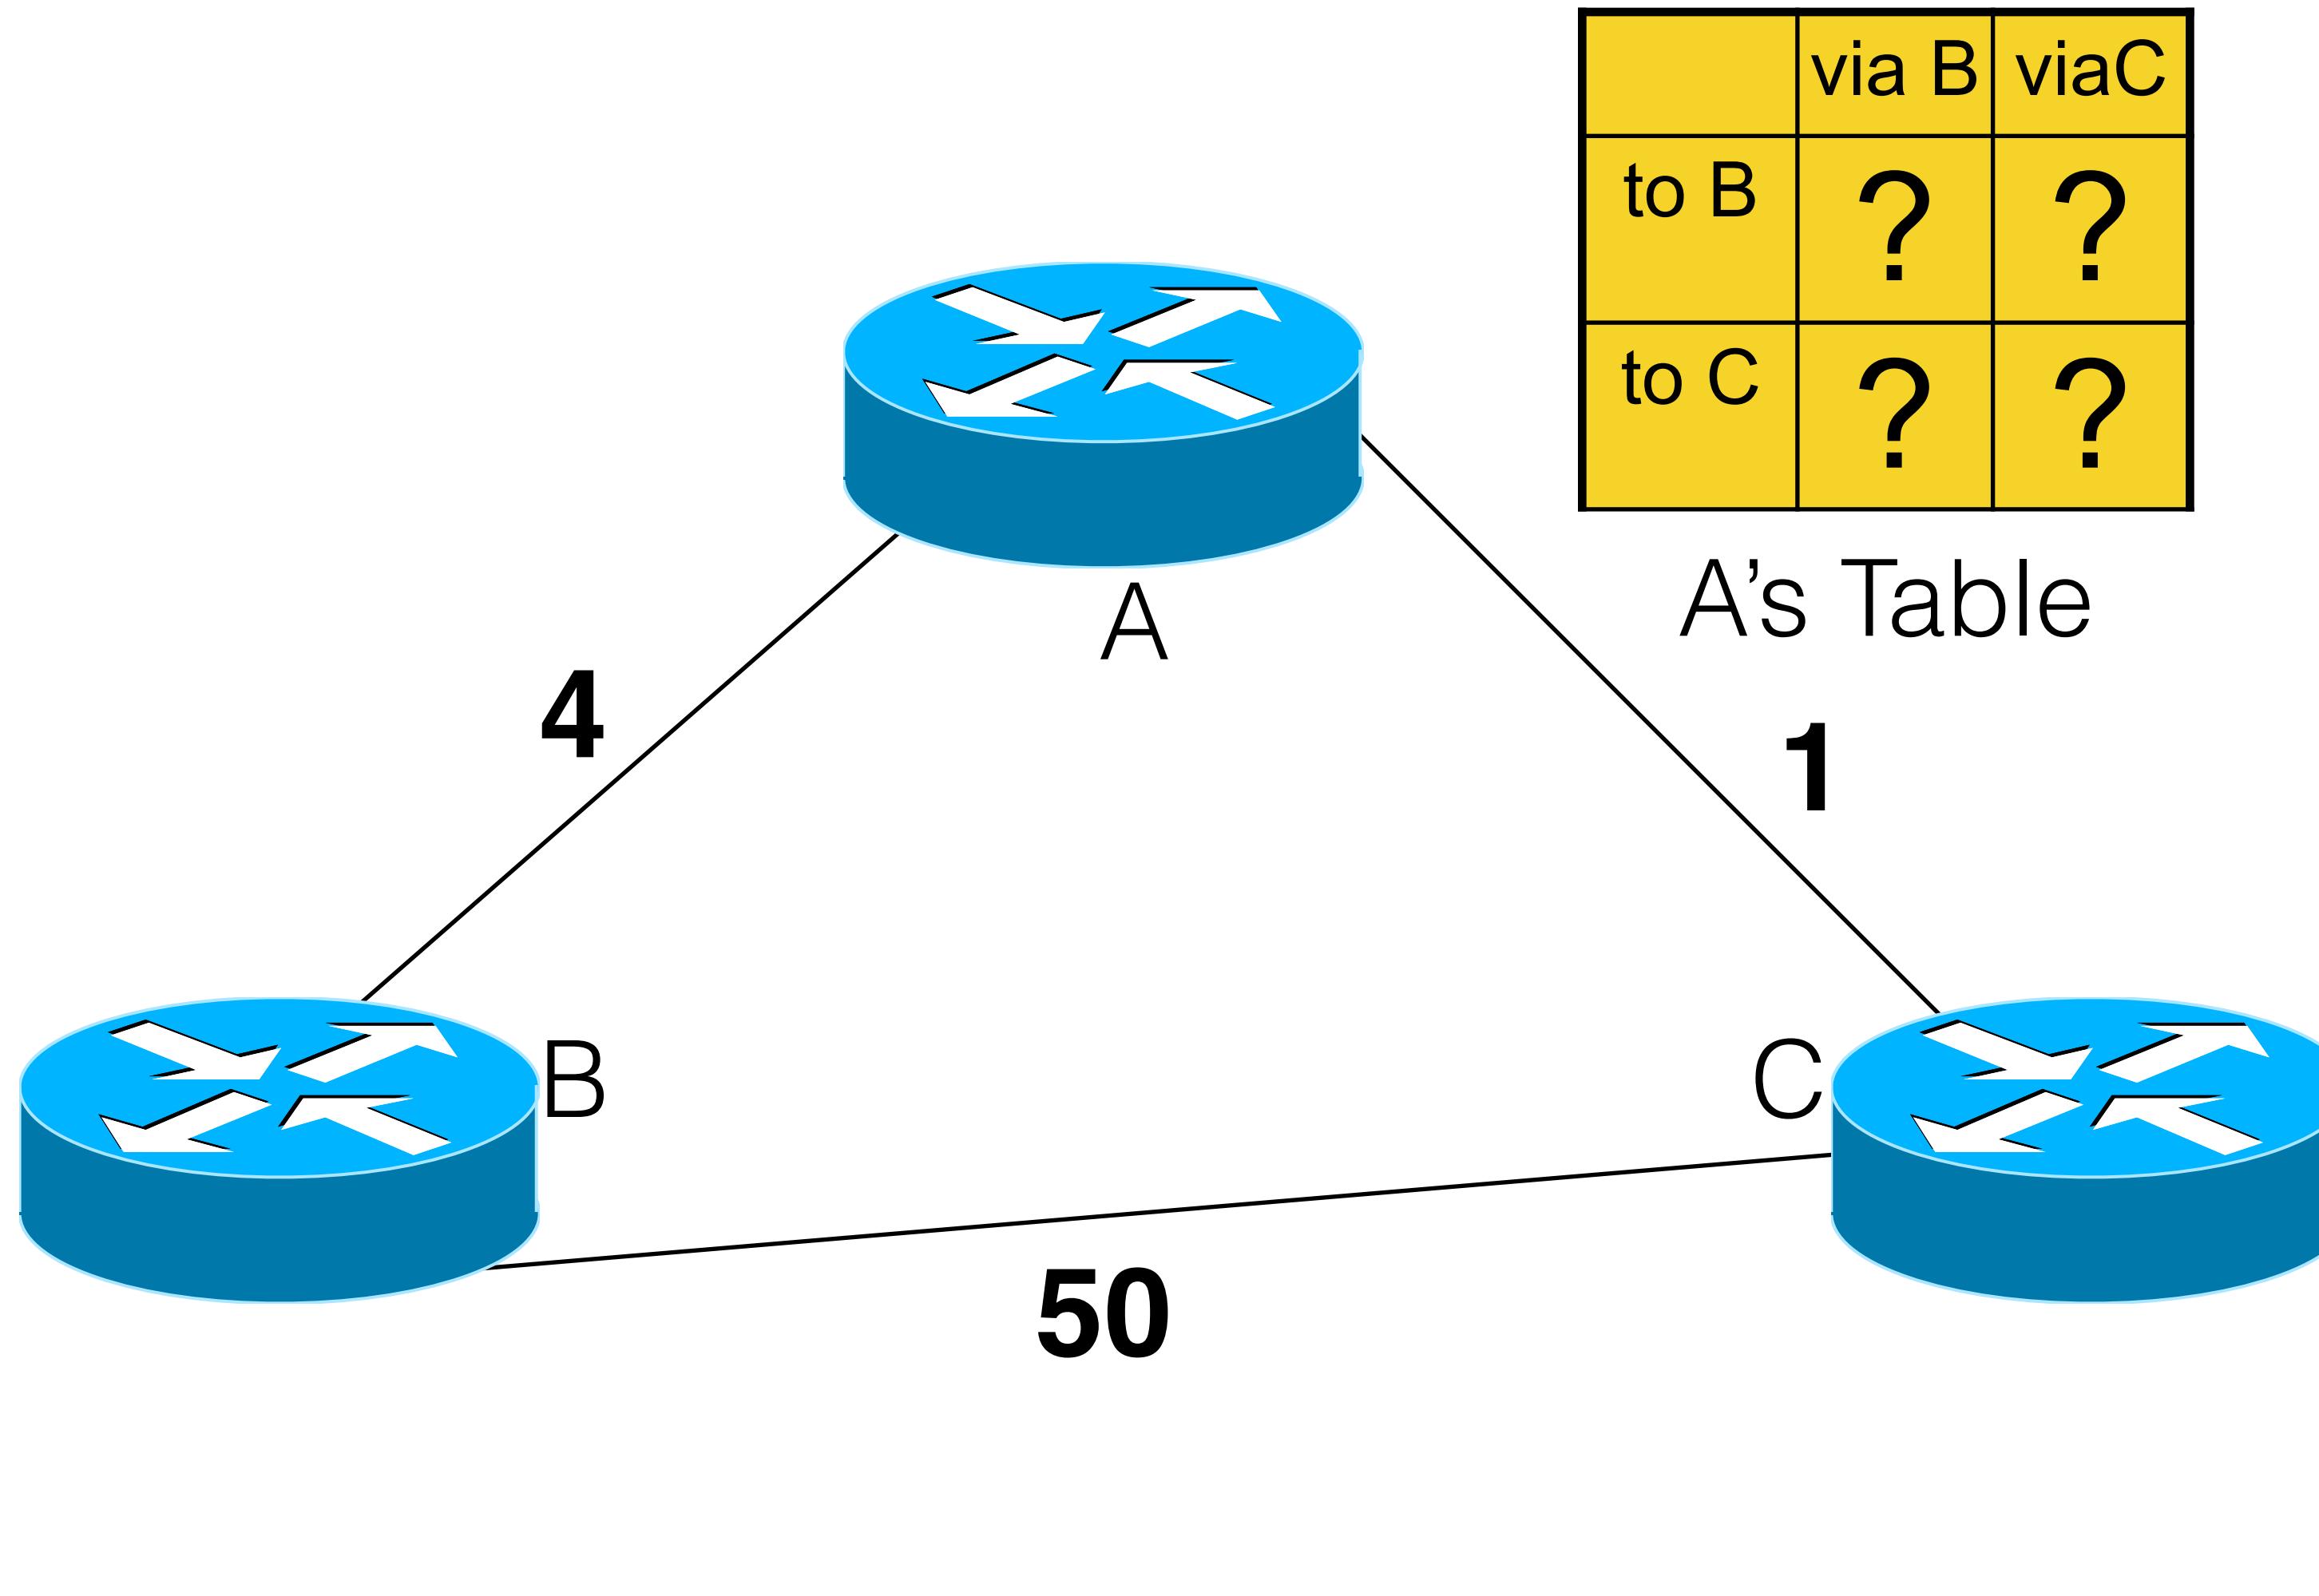
If you find a bug in them I will give you a tee shirt.



Running into trouble....

	via A	via C
to A	?	?
to C	?	?

B's Table



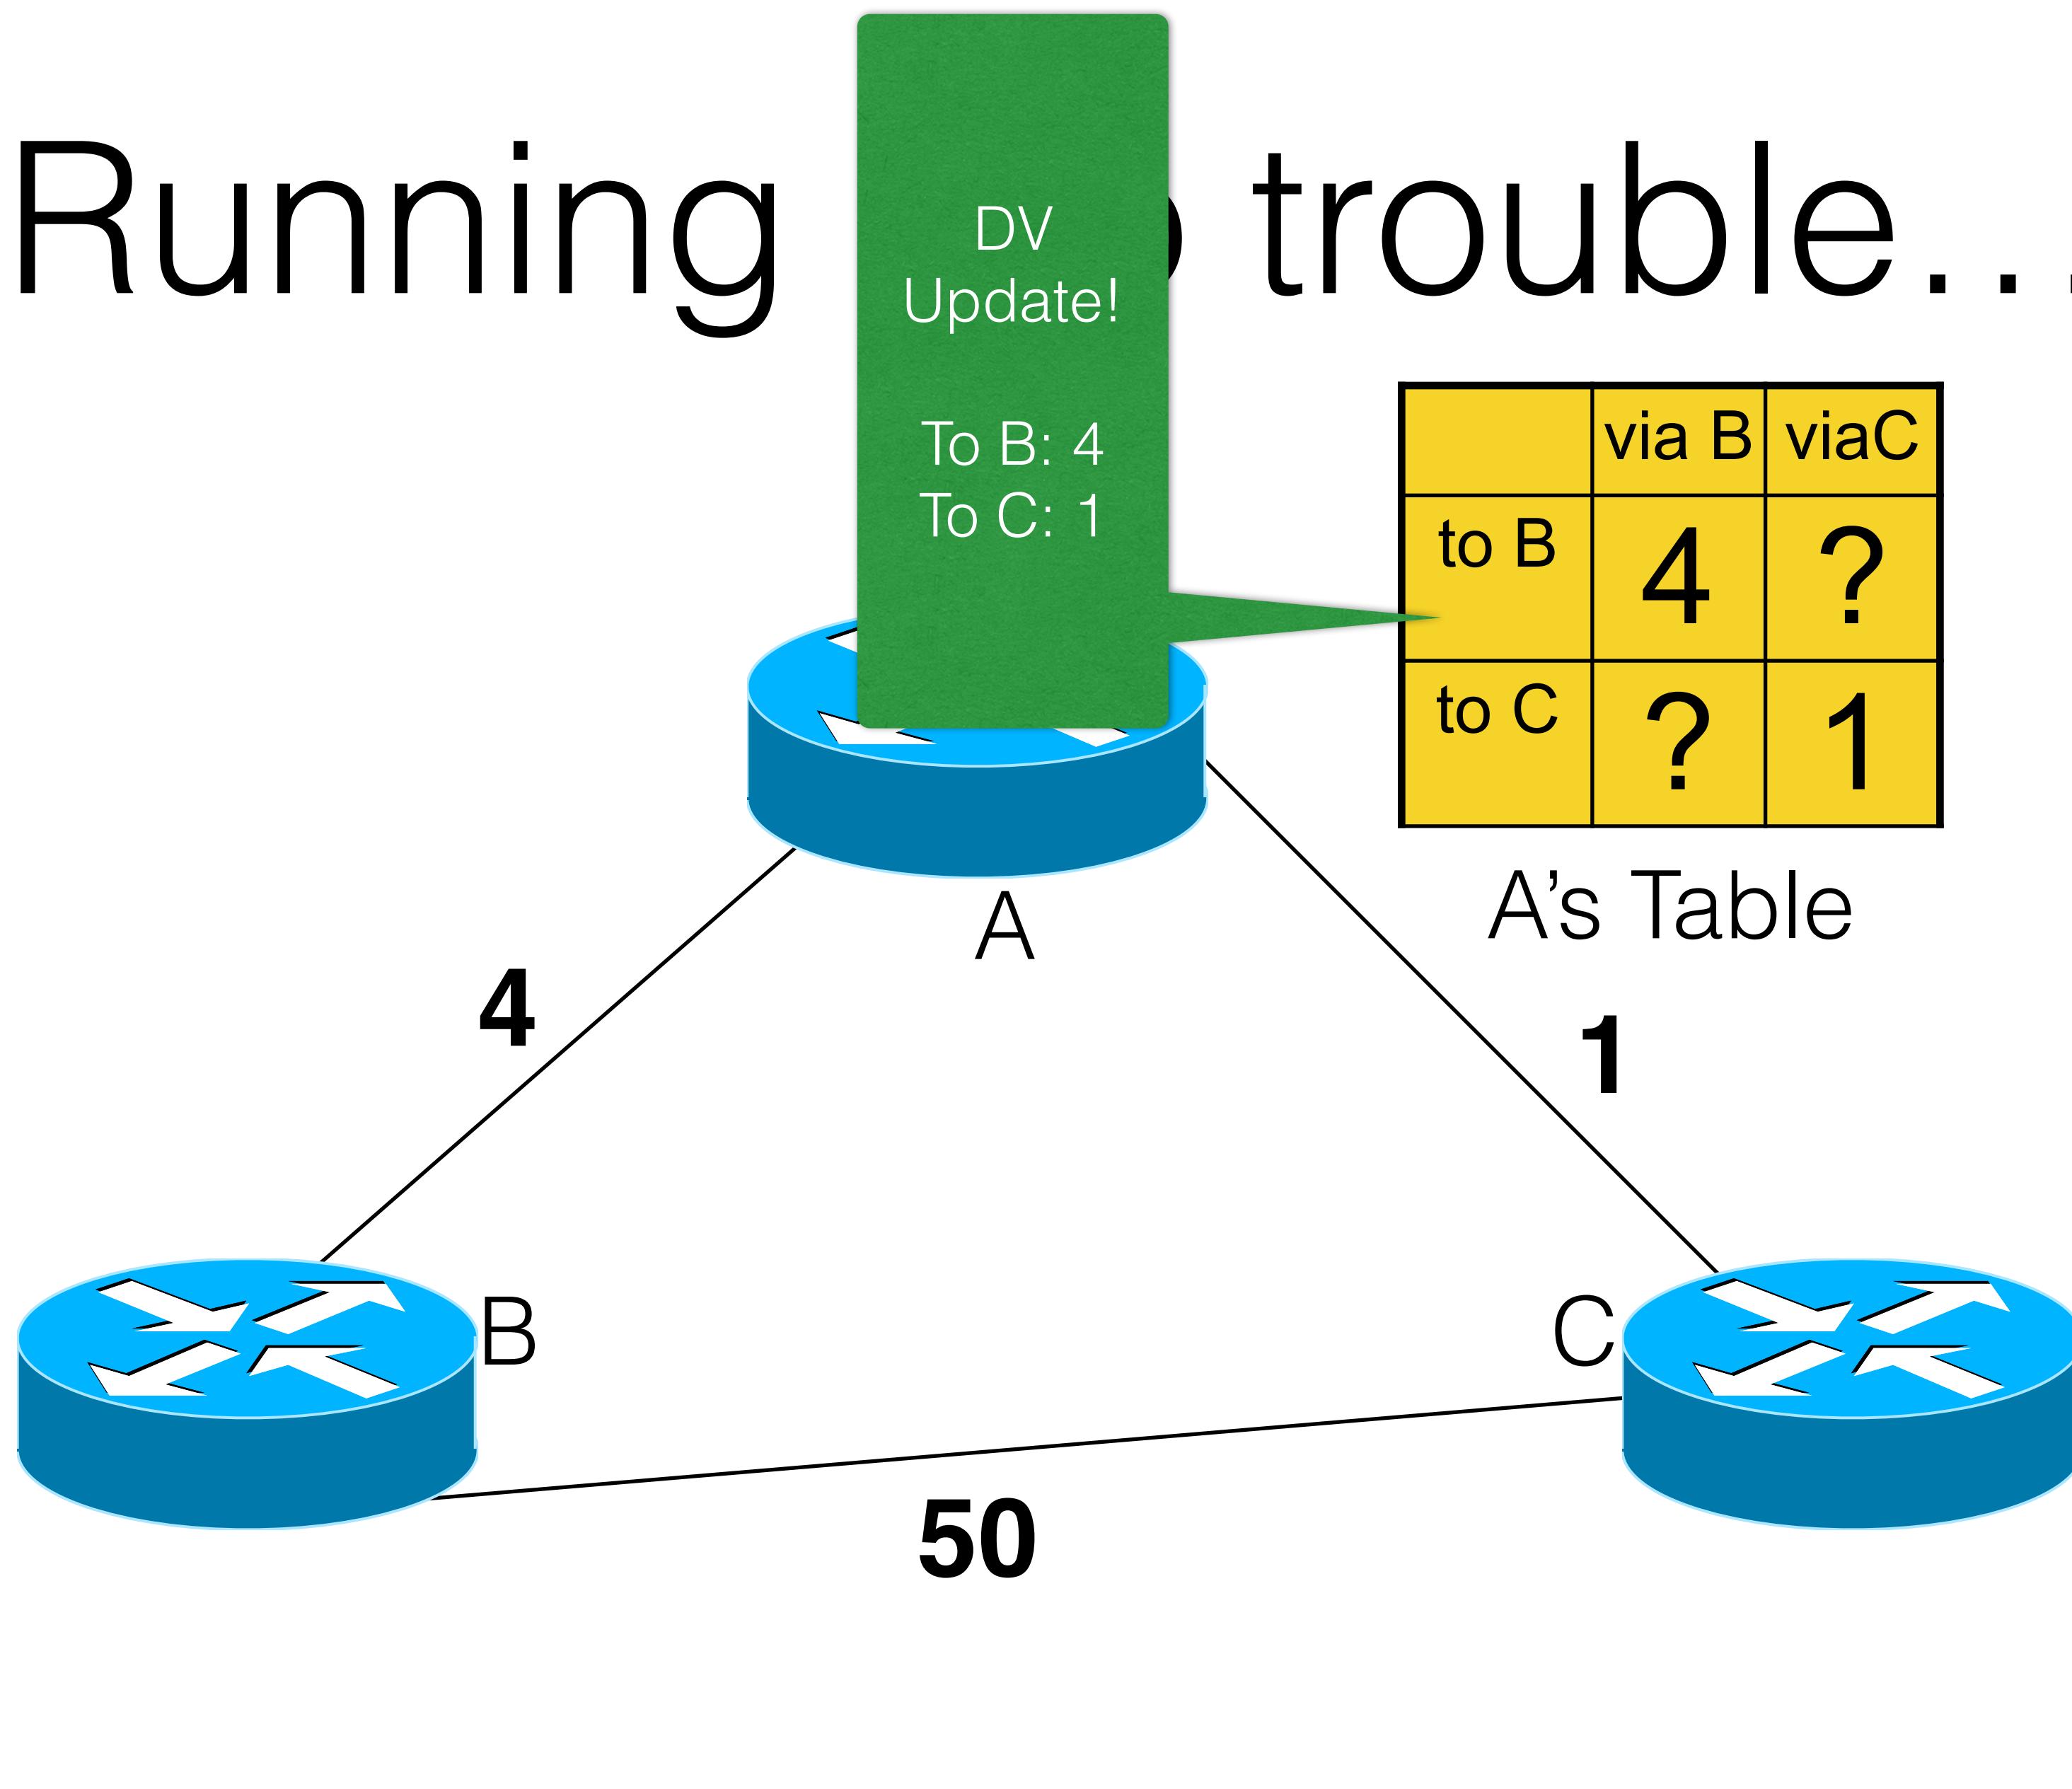
Running trouble....

DV Update!

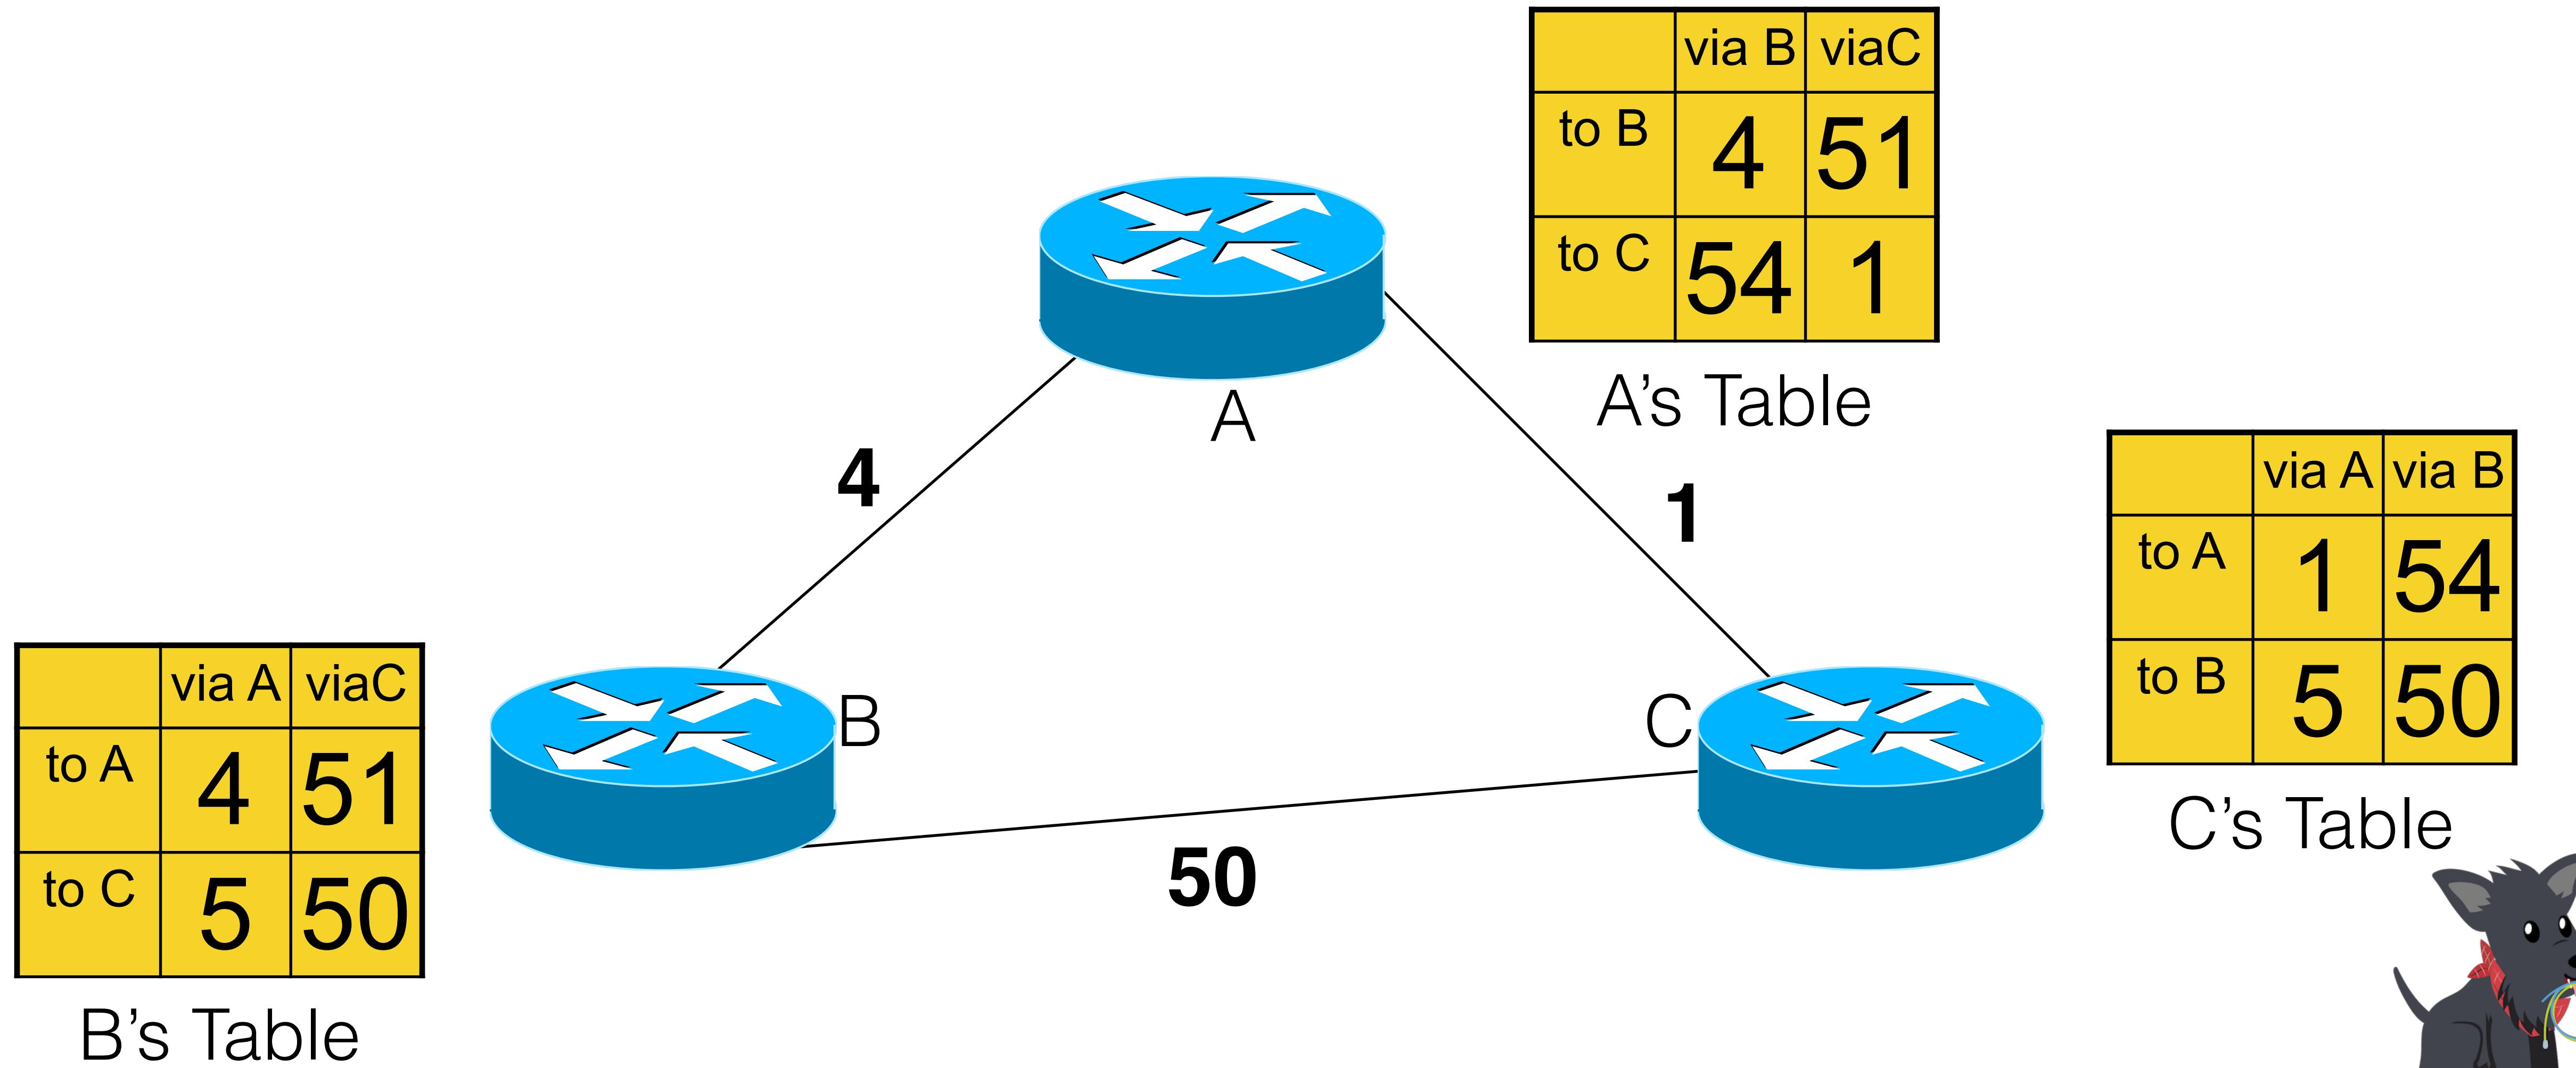
To A: 4
To C: 50

	via A	viaC
to A	4	?
to C	?	50

B's Table



Running into trouble....



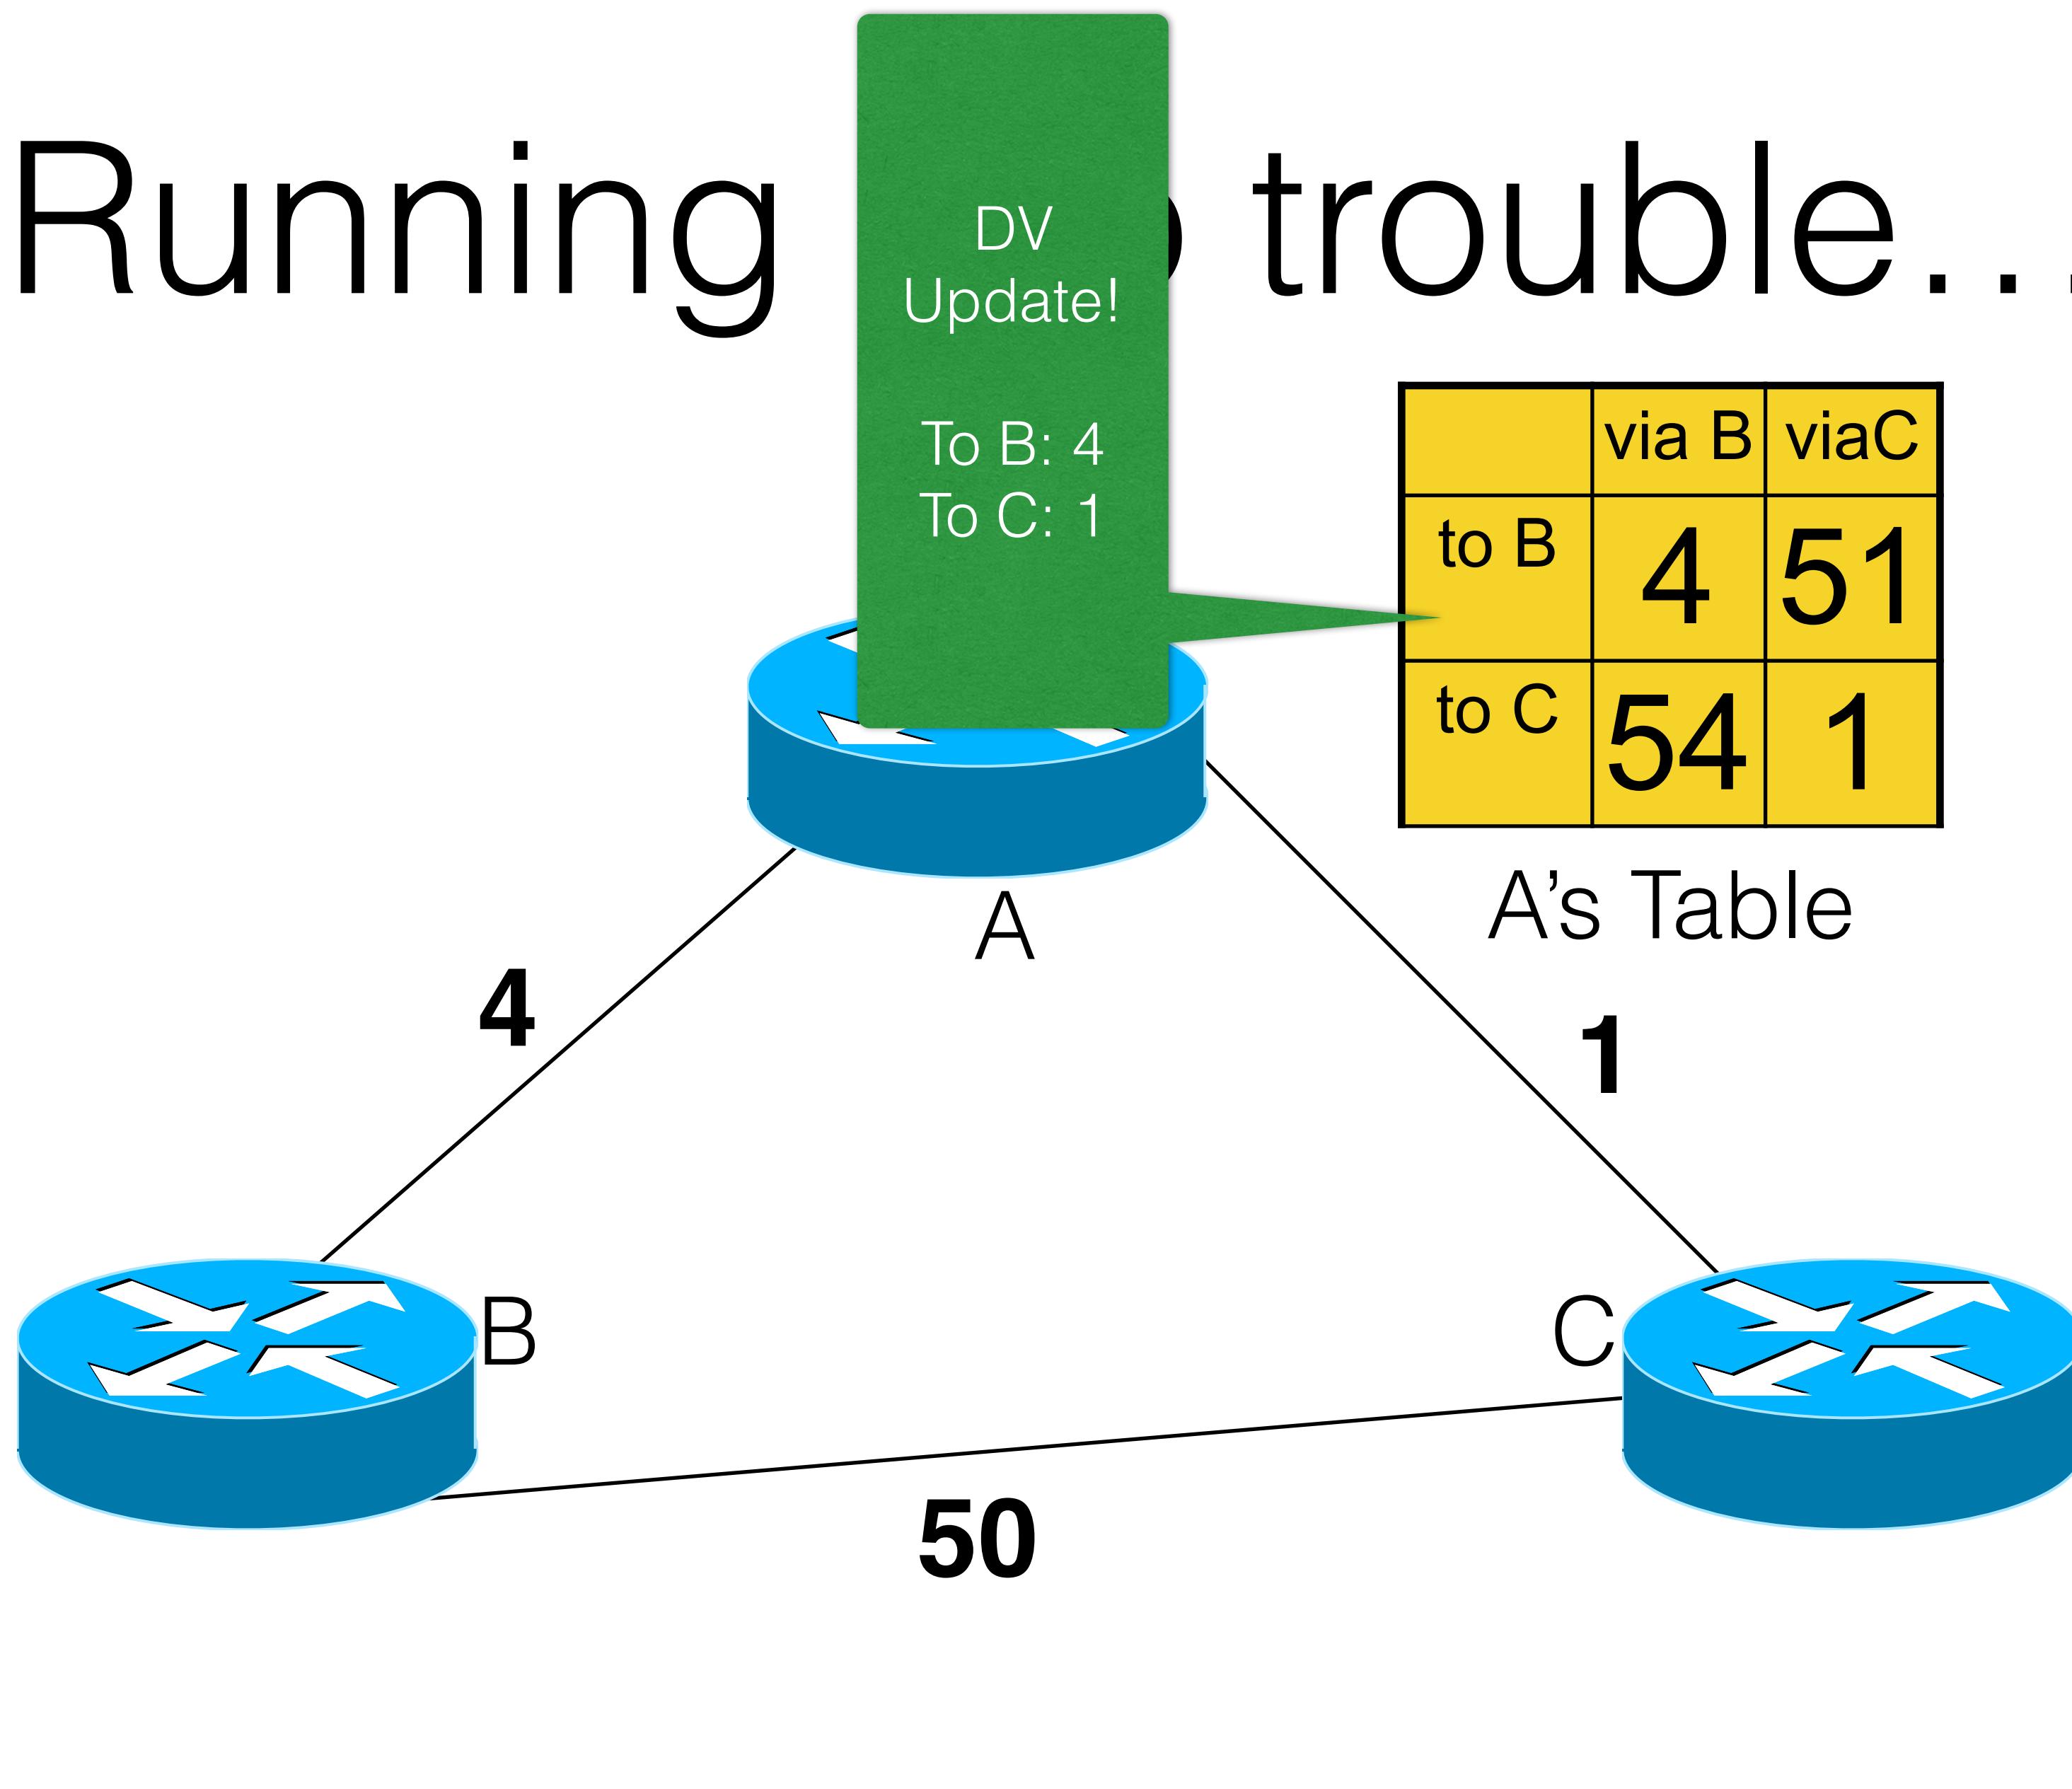
Running trouble....

DV Update!

To A: 4
To C: 5

	via A	viaC
to A	4	51
to C	5	50

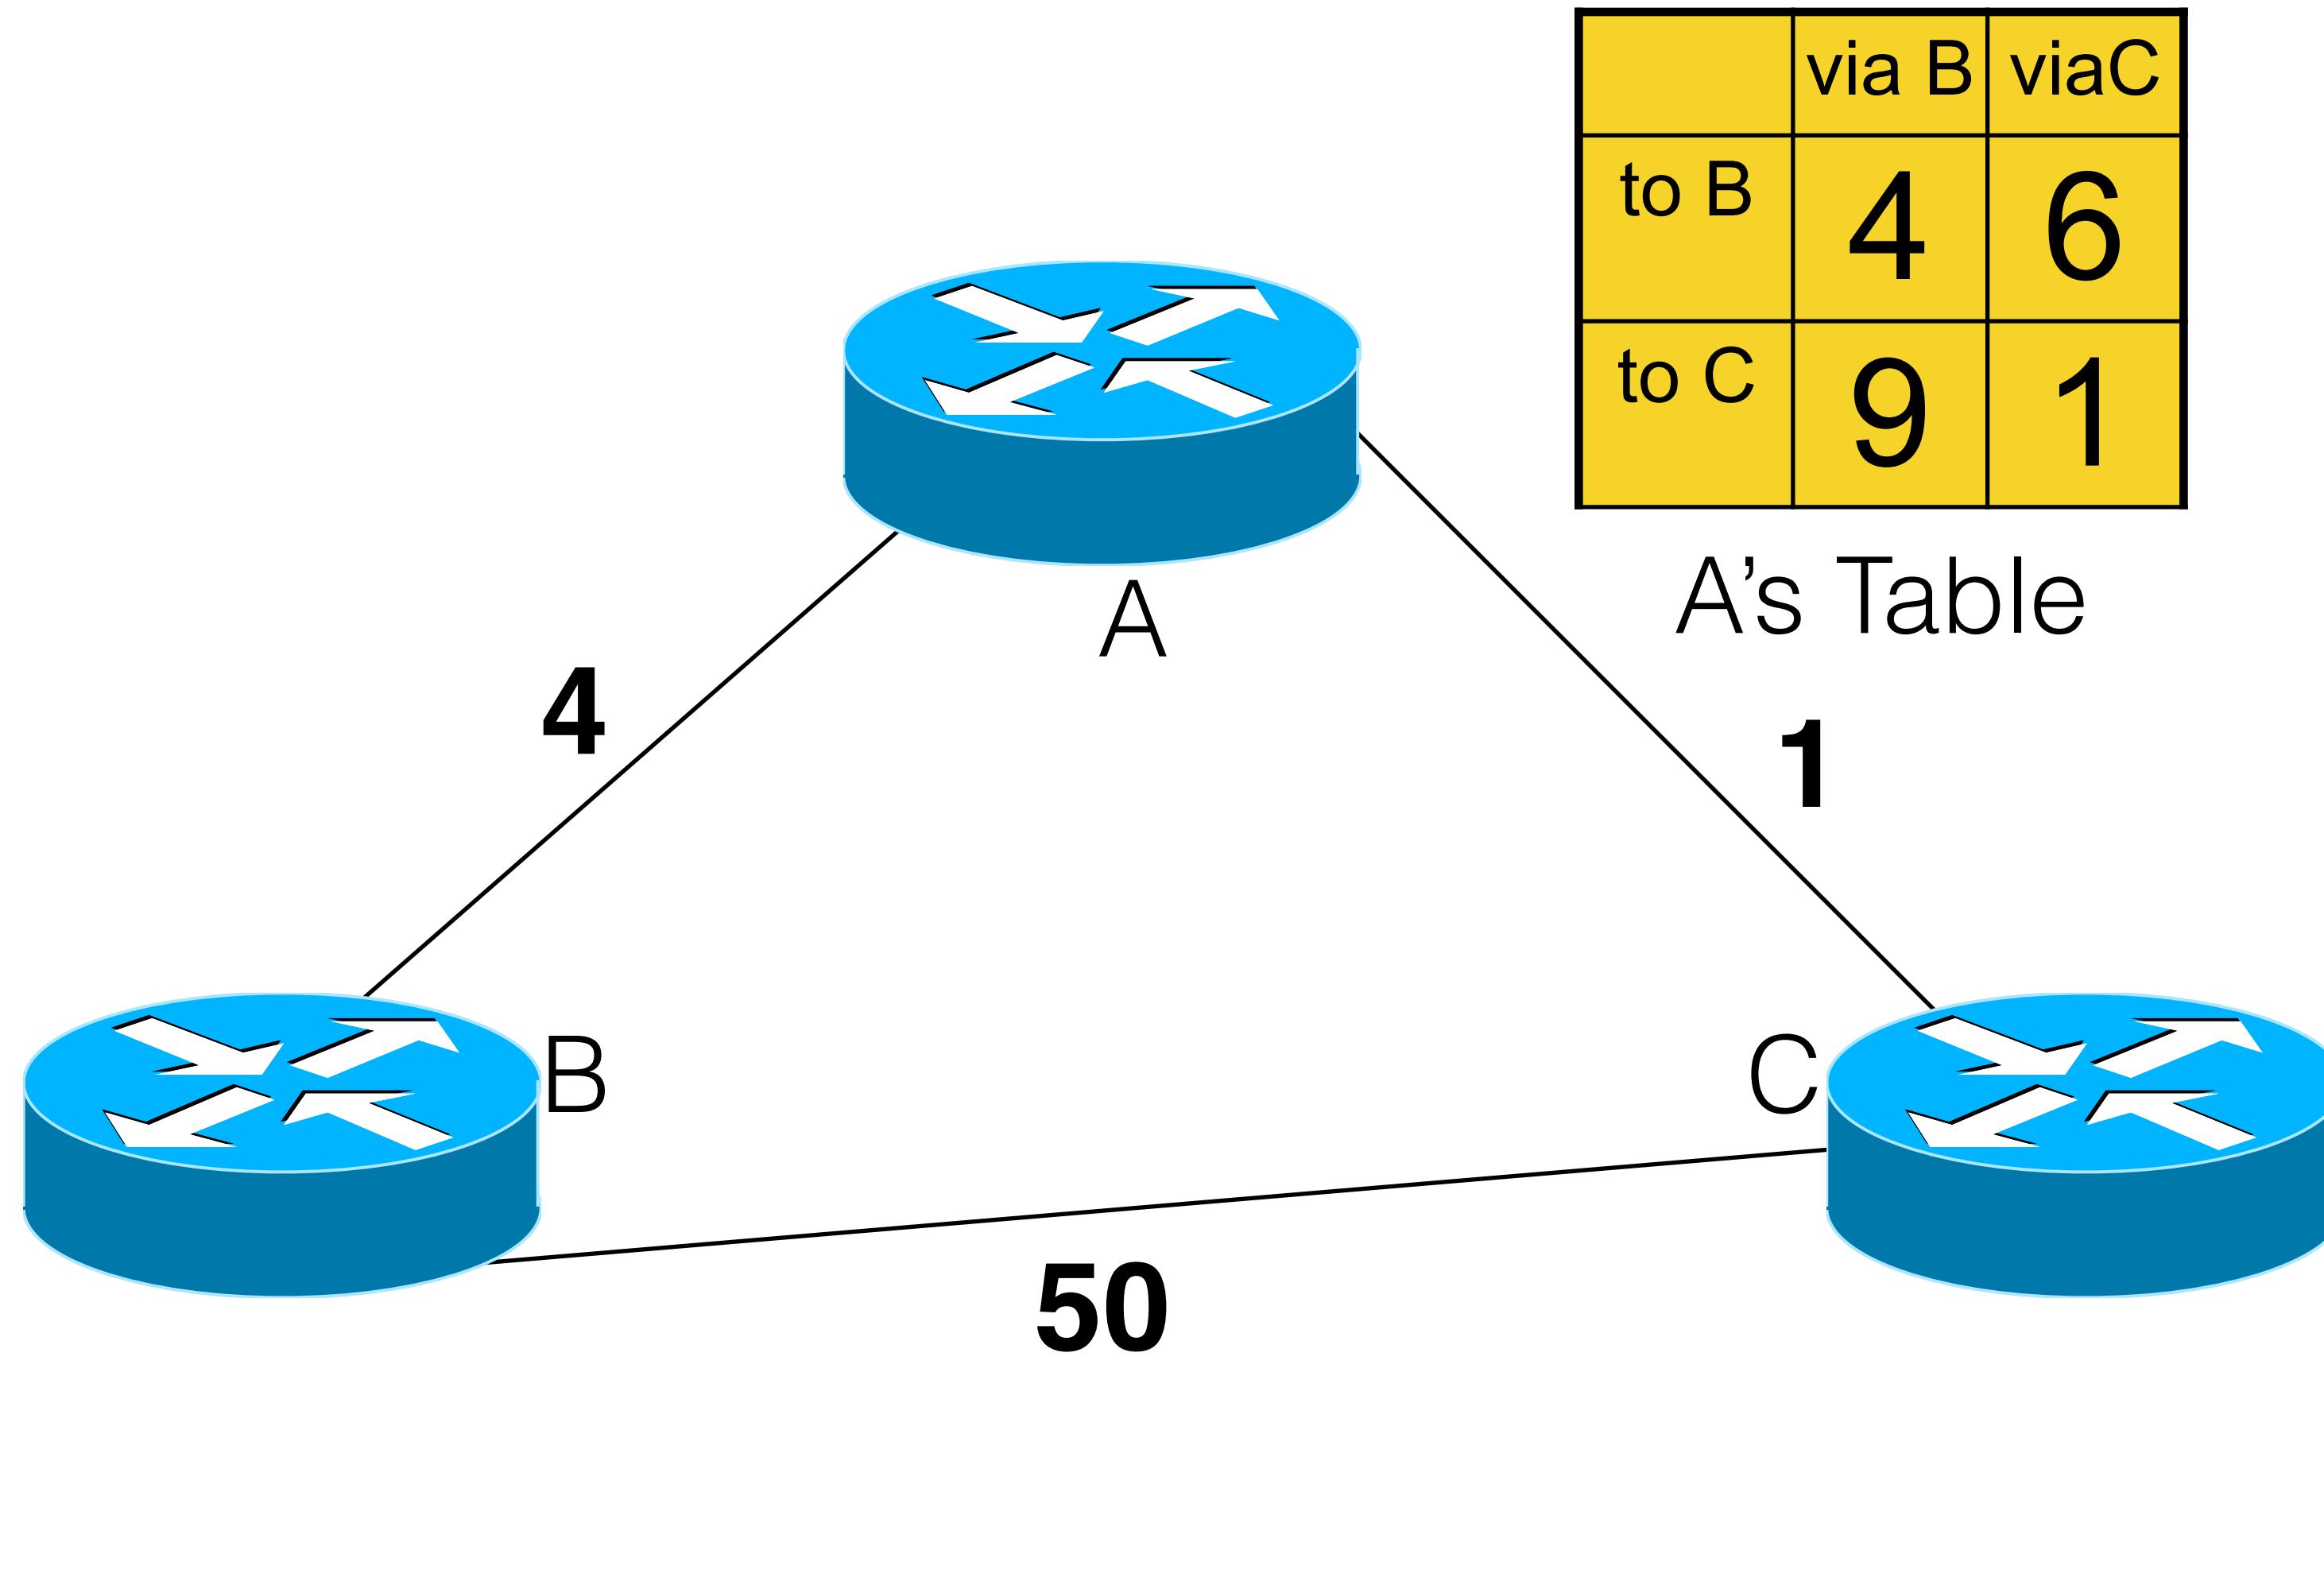
B's Table



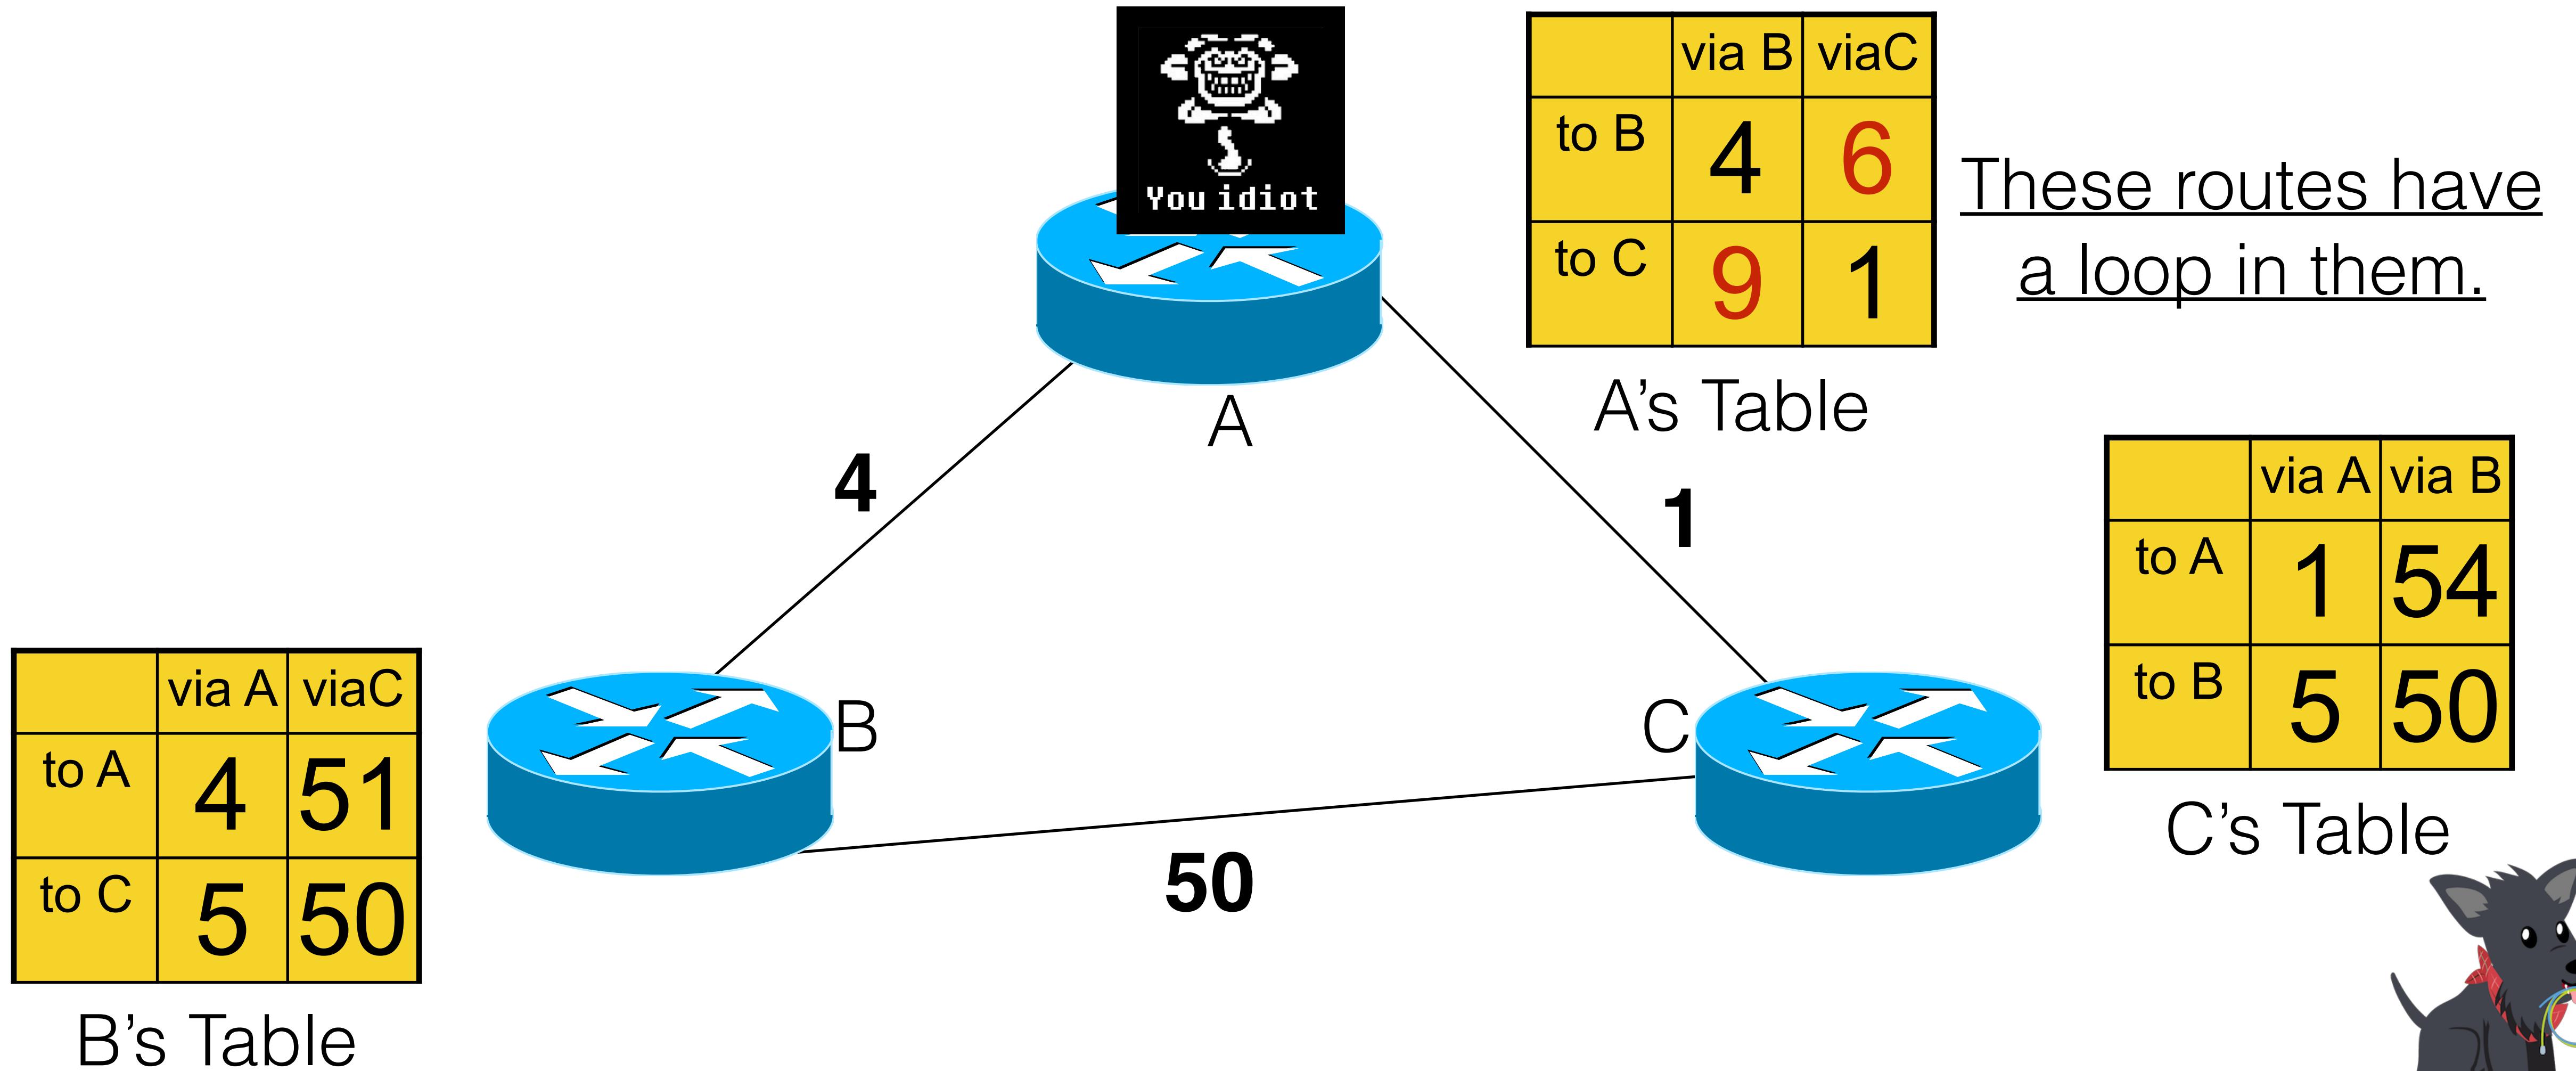
Running into trouble....

	via A	via C
to A	4	51
to C	5	50

B's Table



Running into trouble....



Loopy Routes

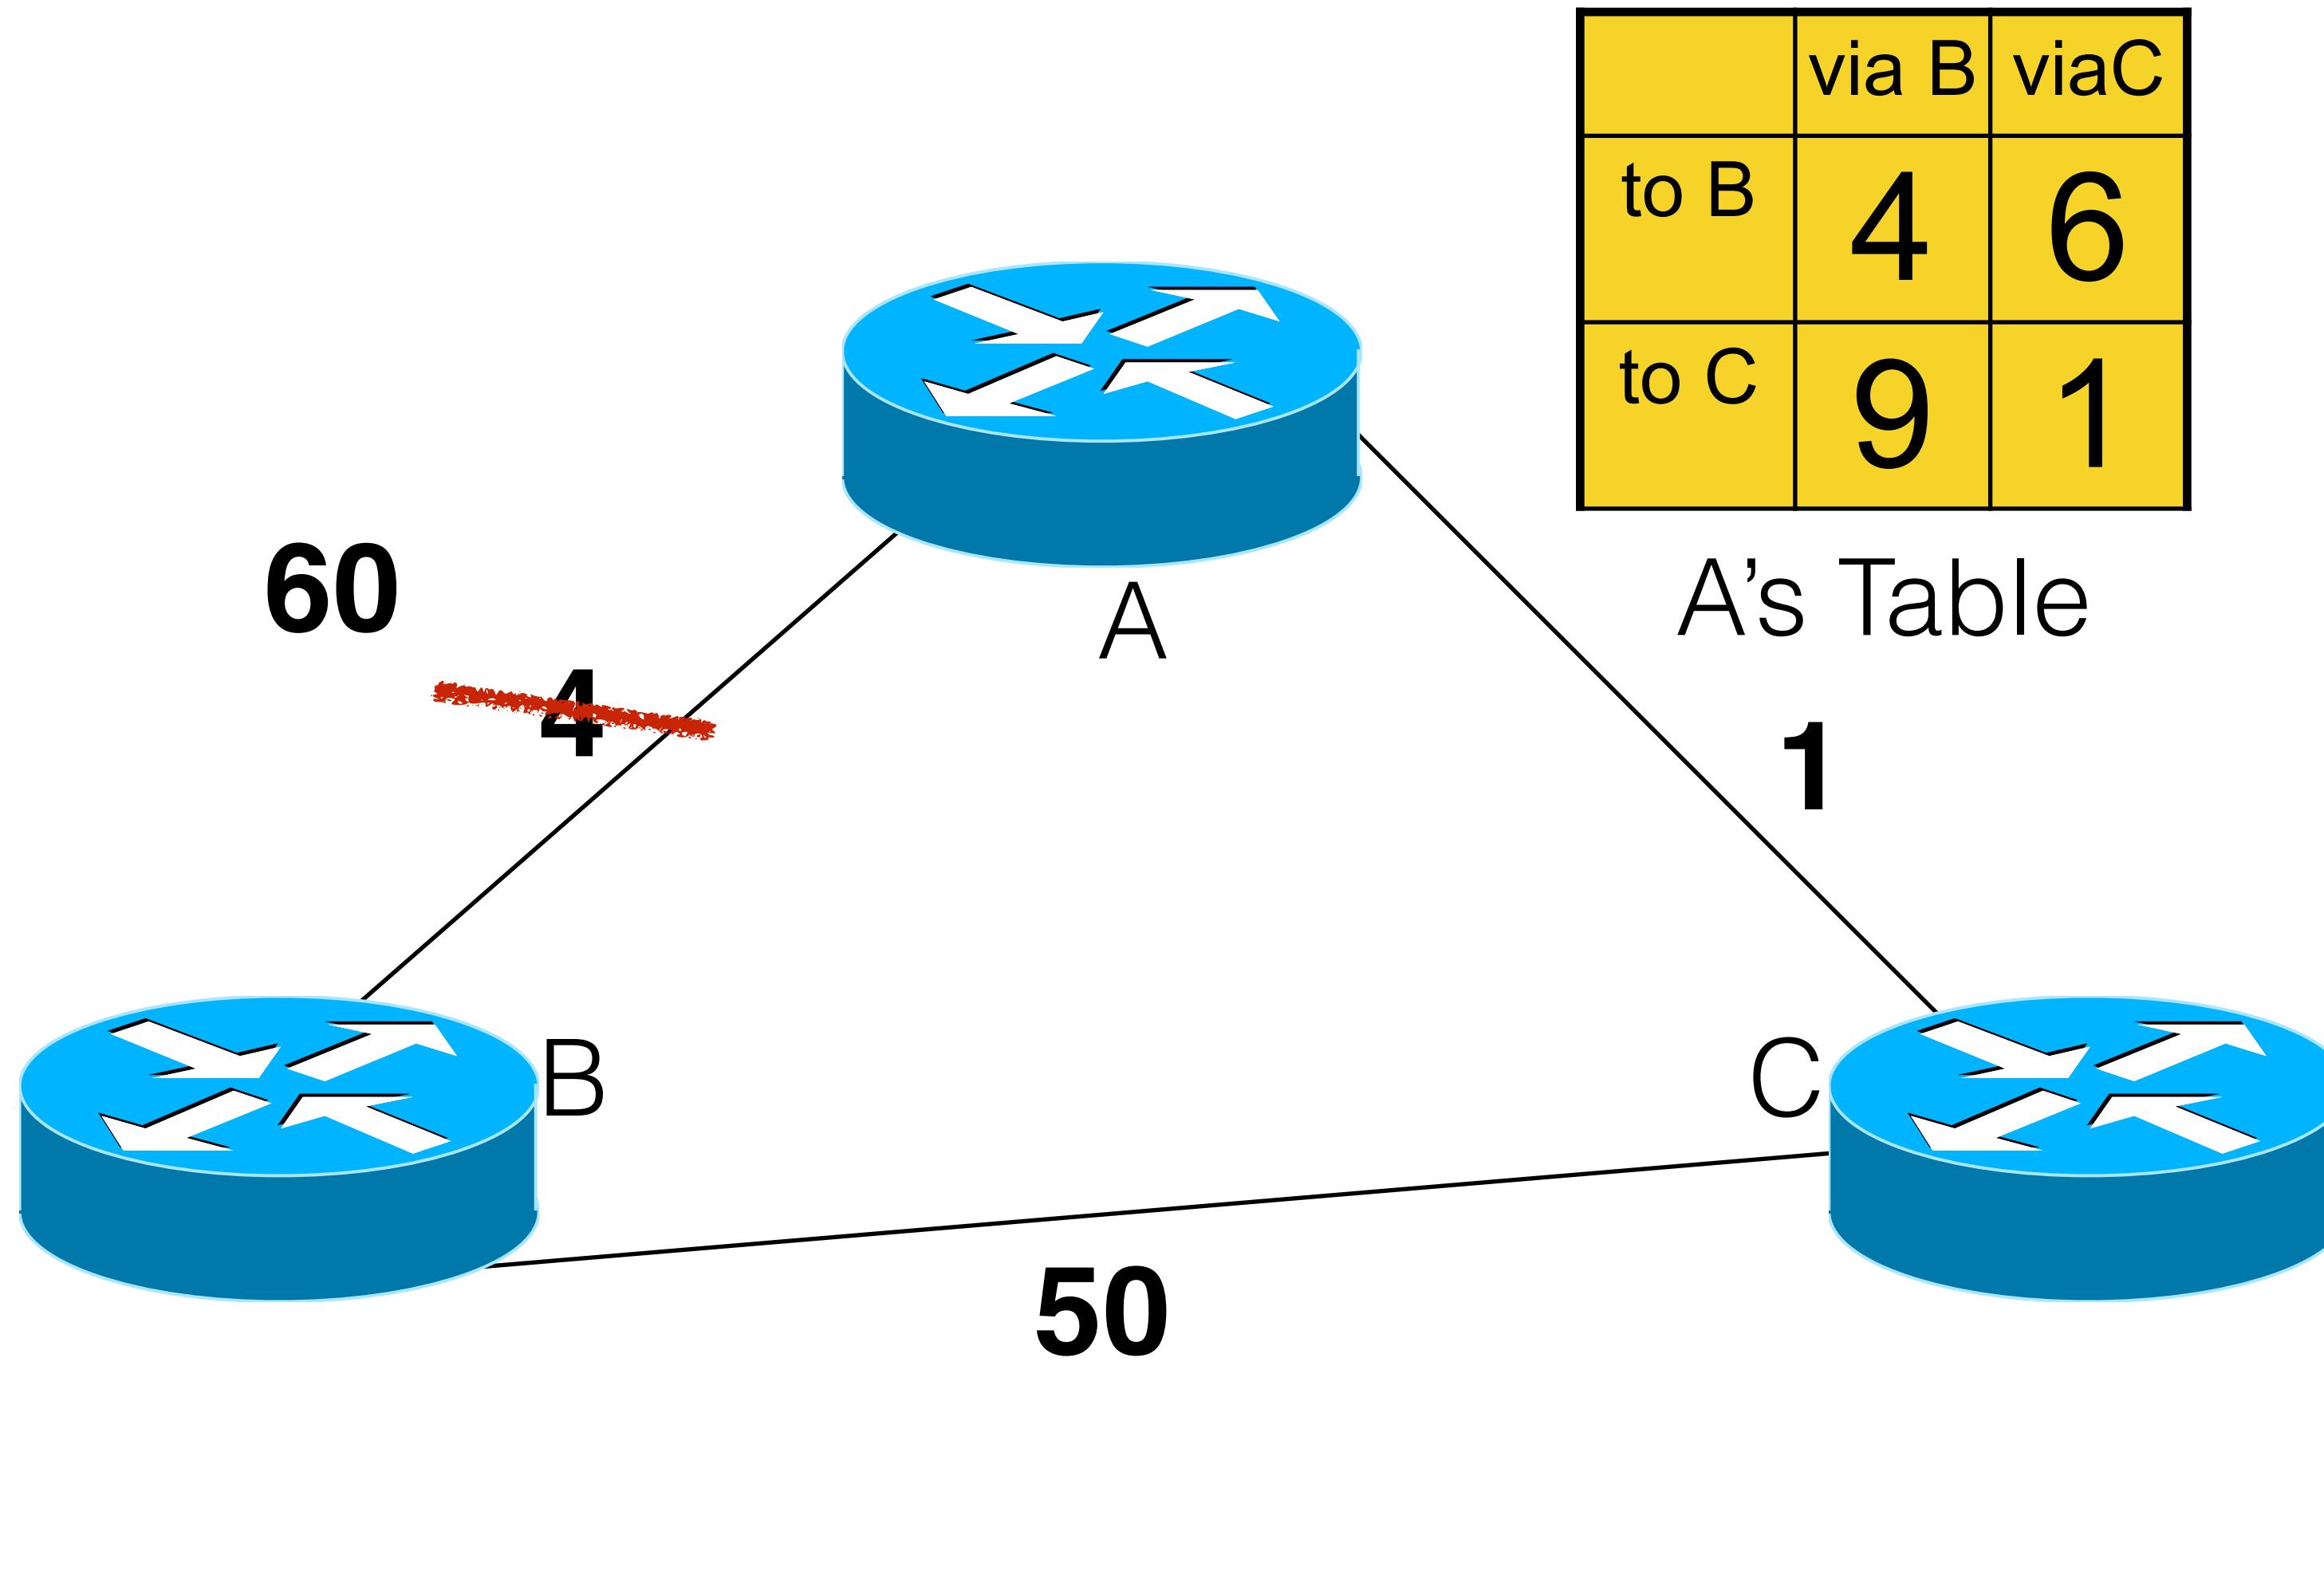
- These routes are *fine* under normal operations — because they don't get used.
 - Why take the loopy route when you can take the direct path?
- But when link *updates* happen, bad things can happen.
 - If a link becomes more expensive.
 - If a link fails.



Running into trouble....

	via A	via C
to A	4	51
to C	5	50

B's Table



DO IT YOURSELF

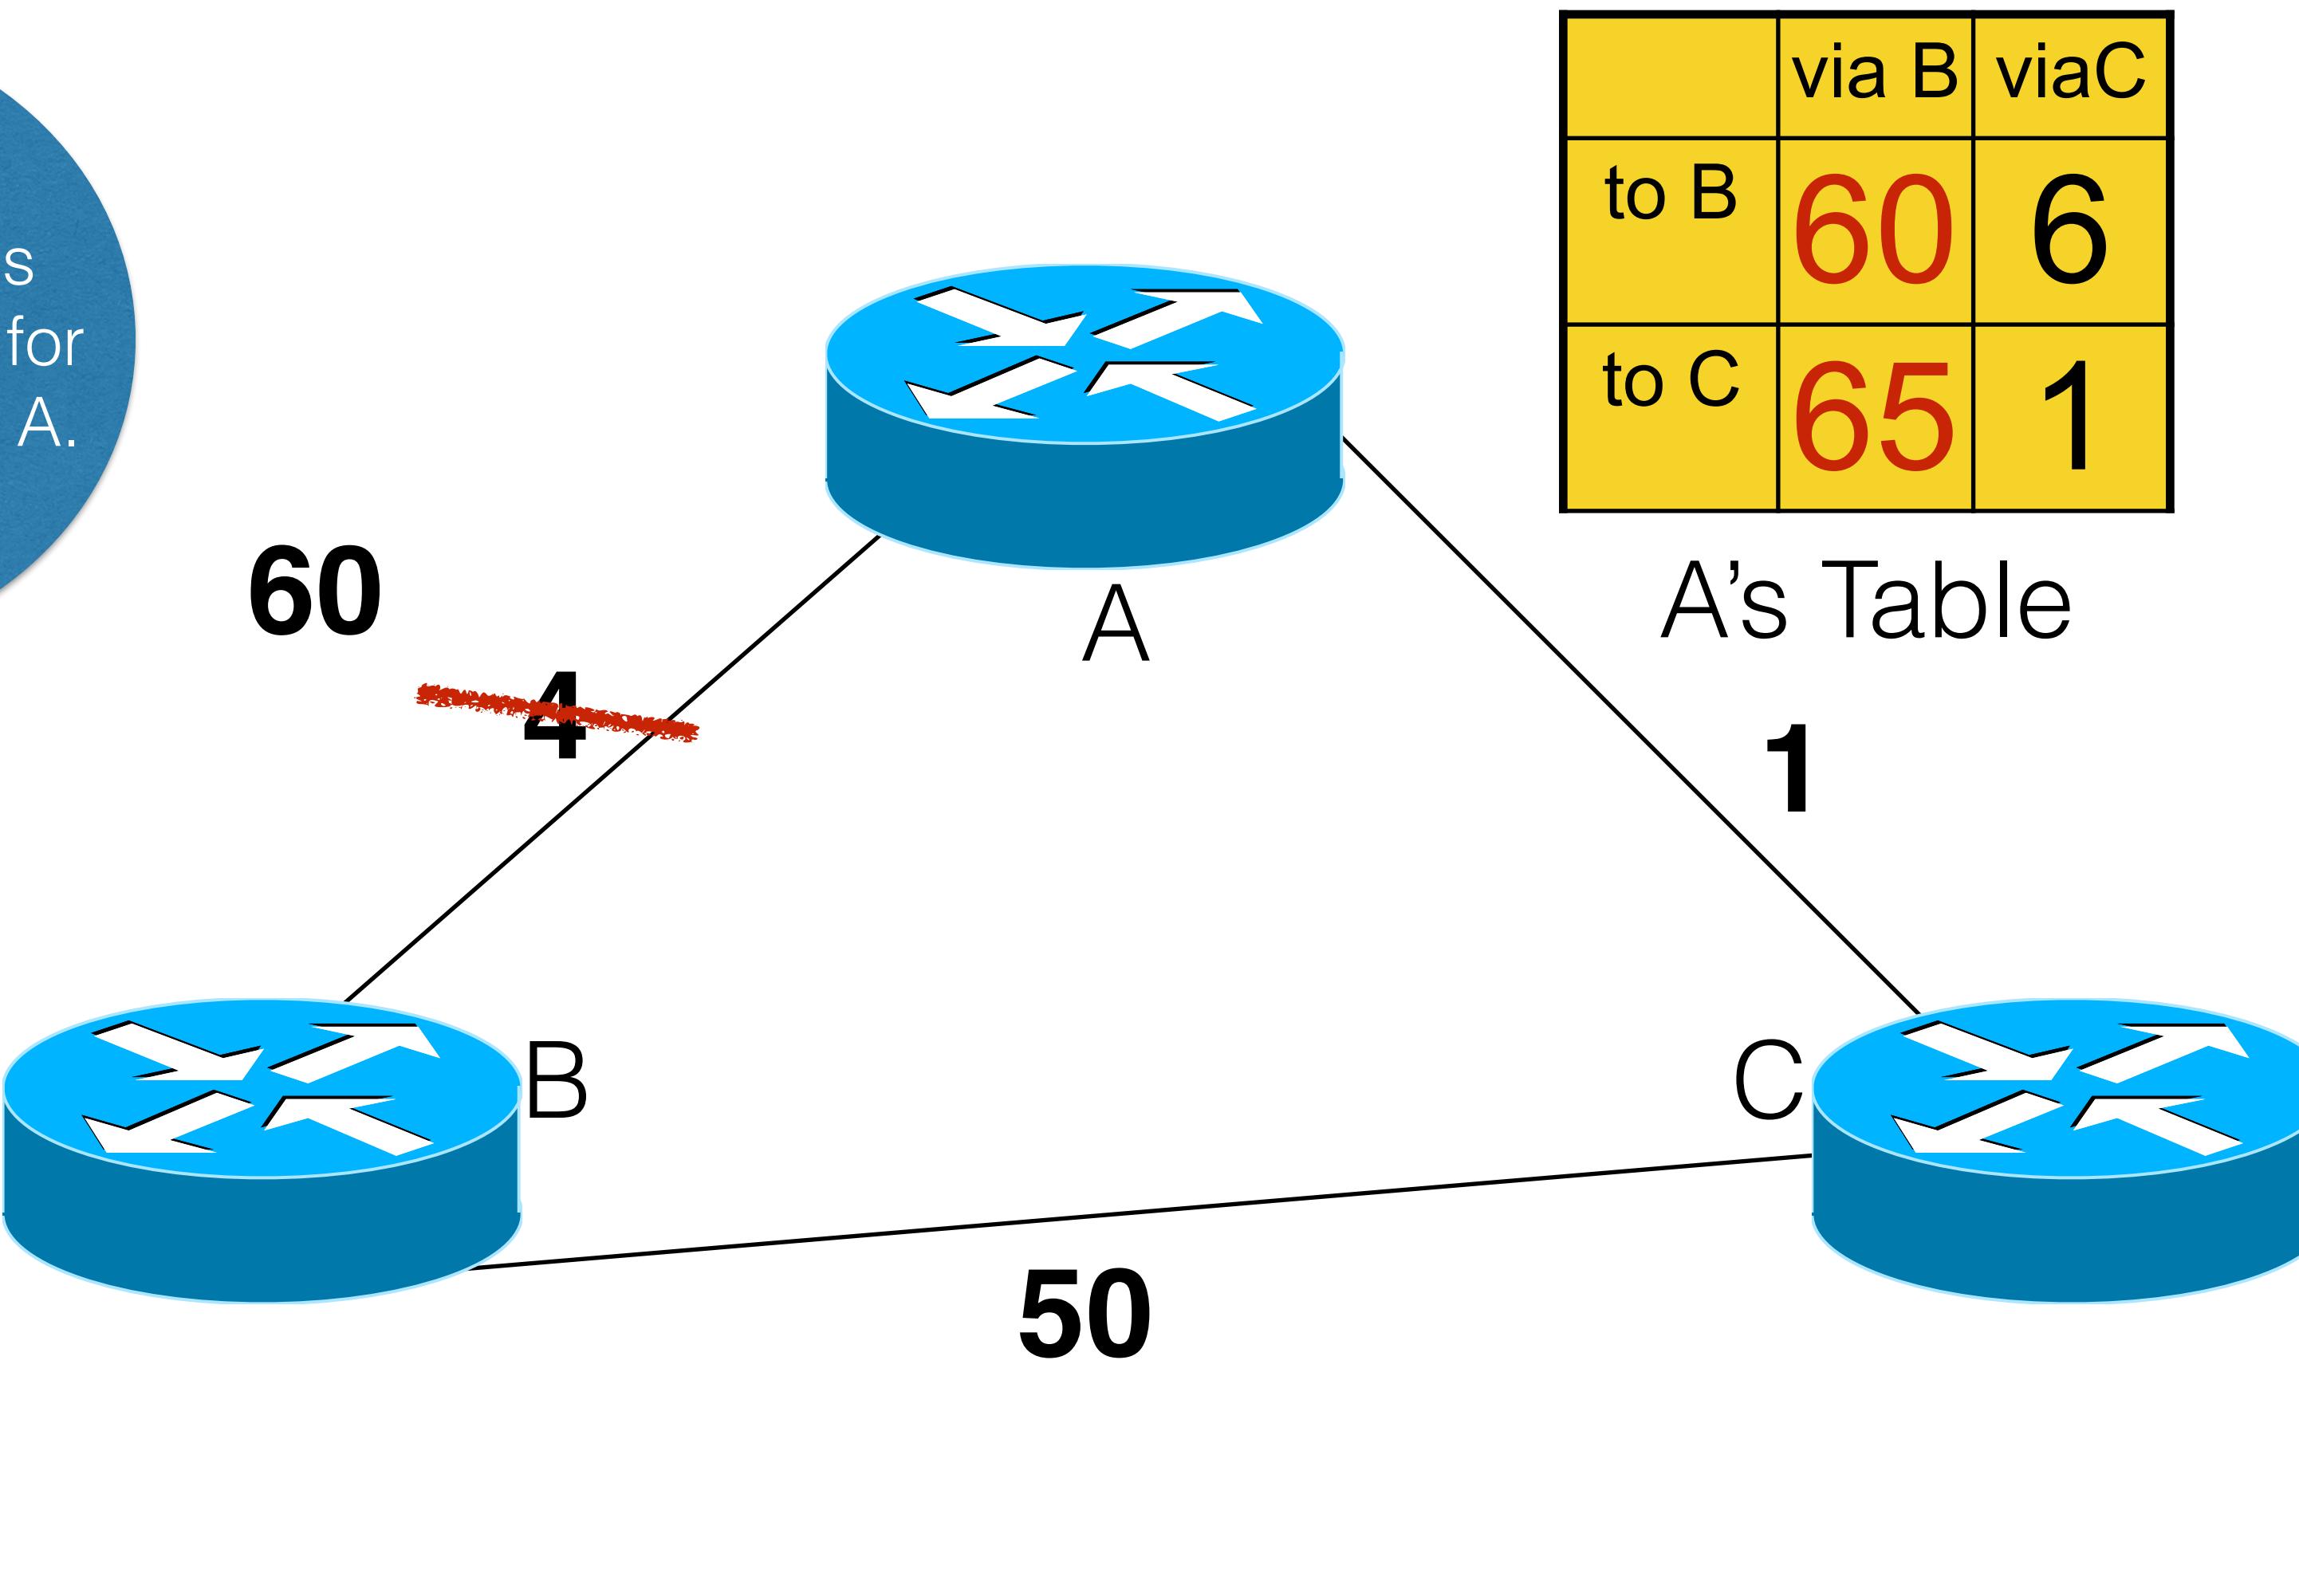


Running into trouble....

Add 56 to routes from A via B and for routes from B via A.

	via A	via C
to A	60	51
to C	61	50

B's Table

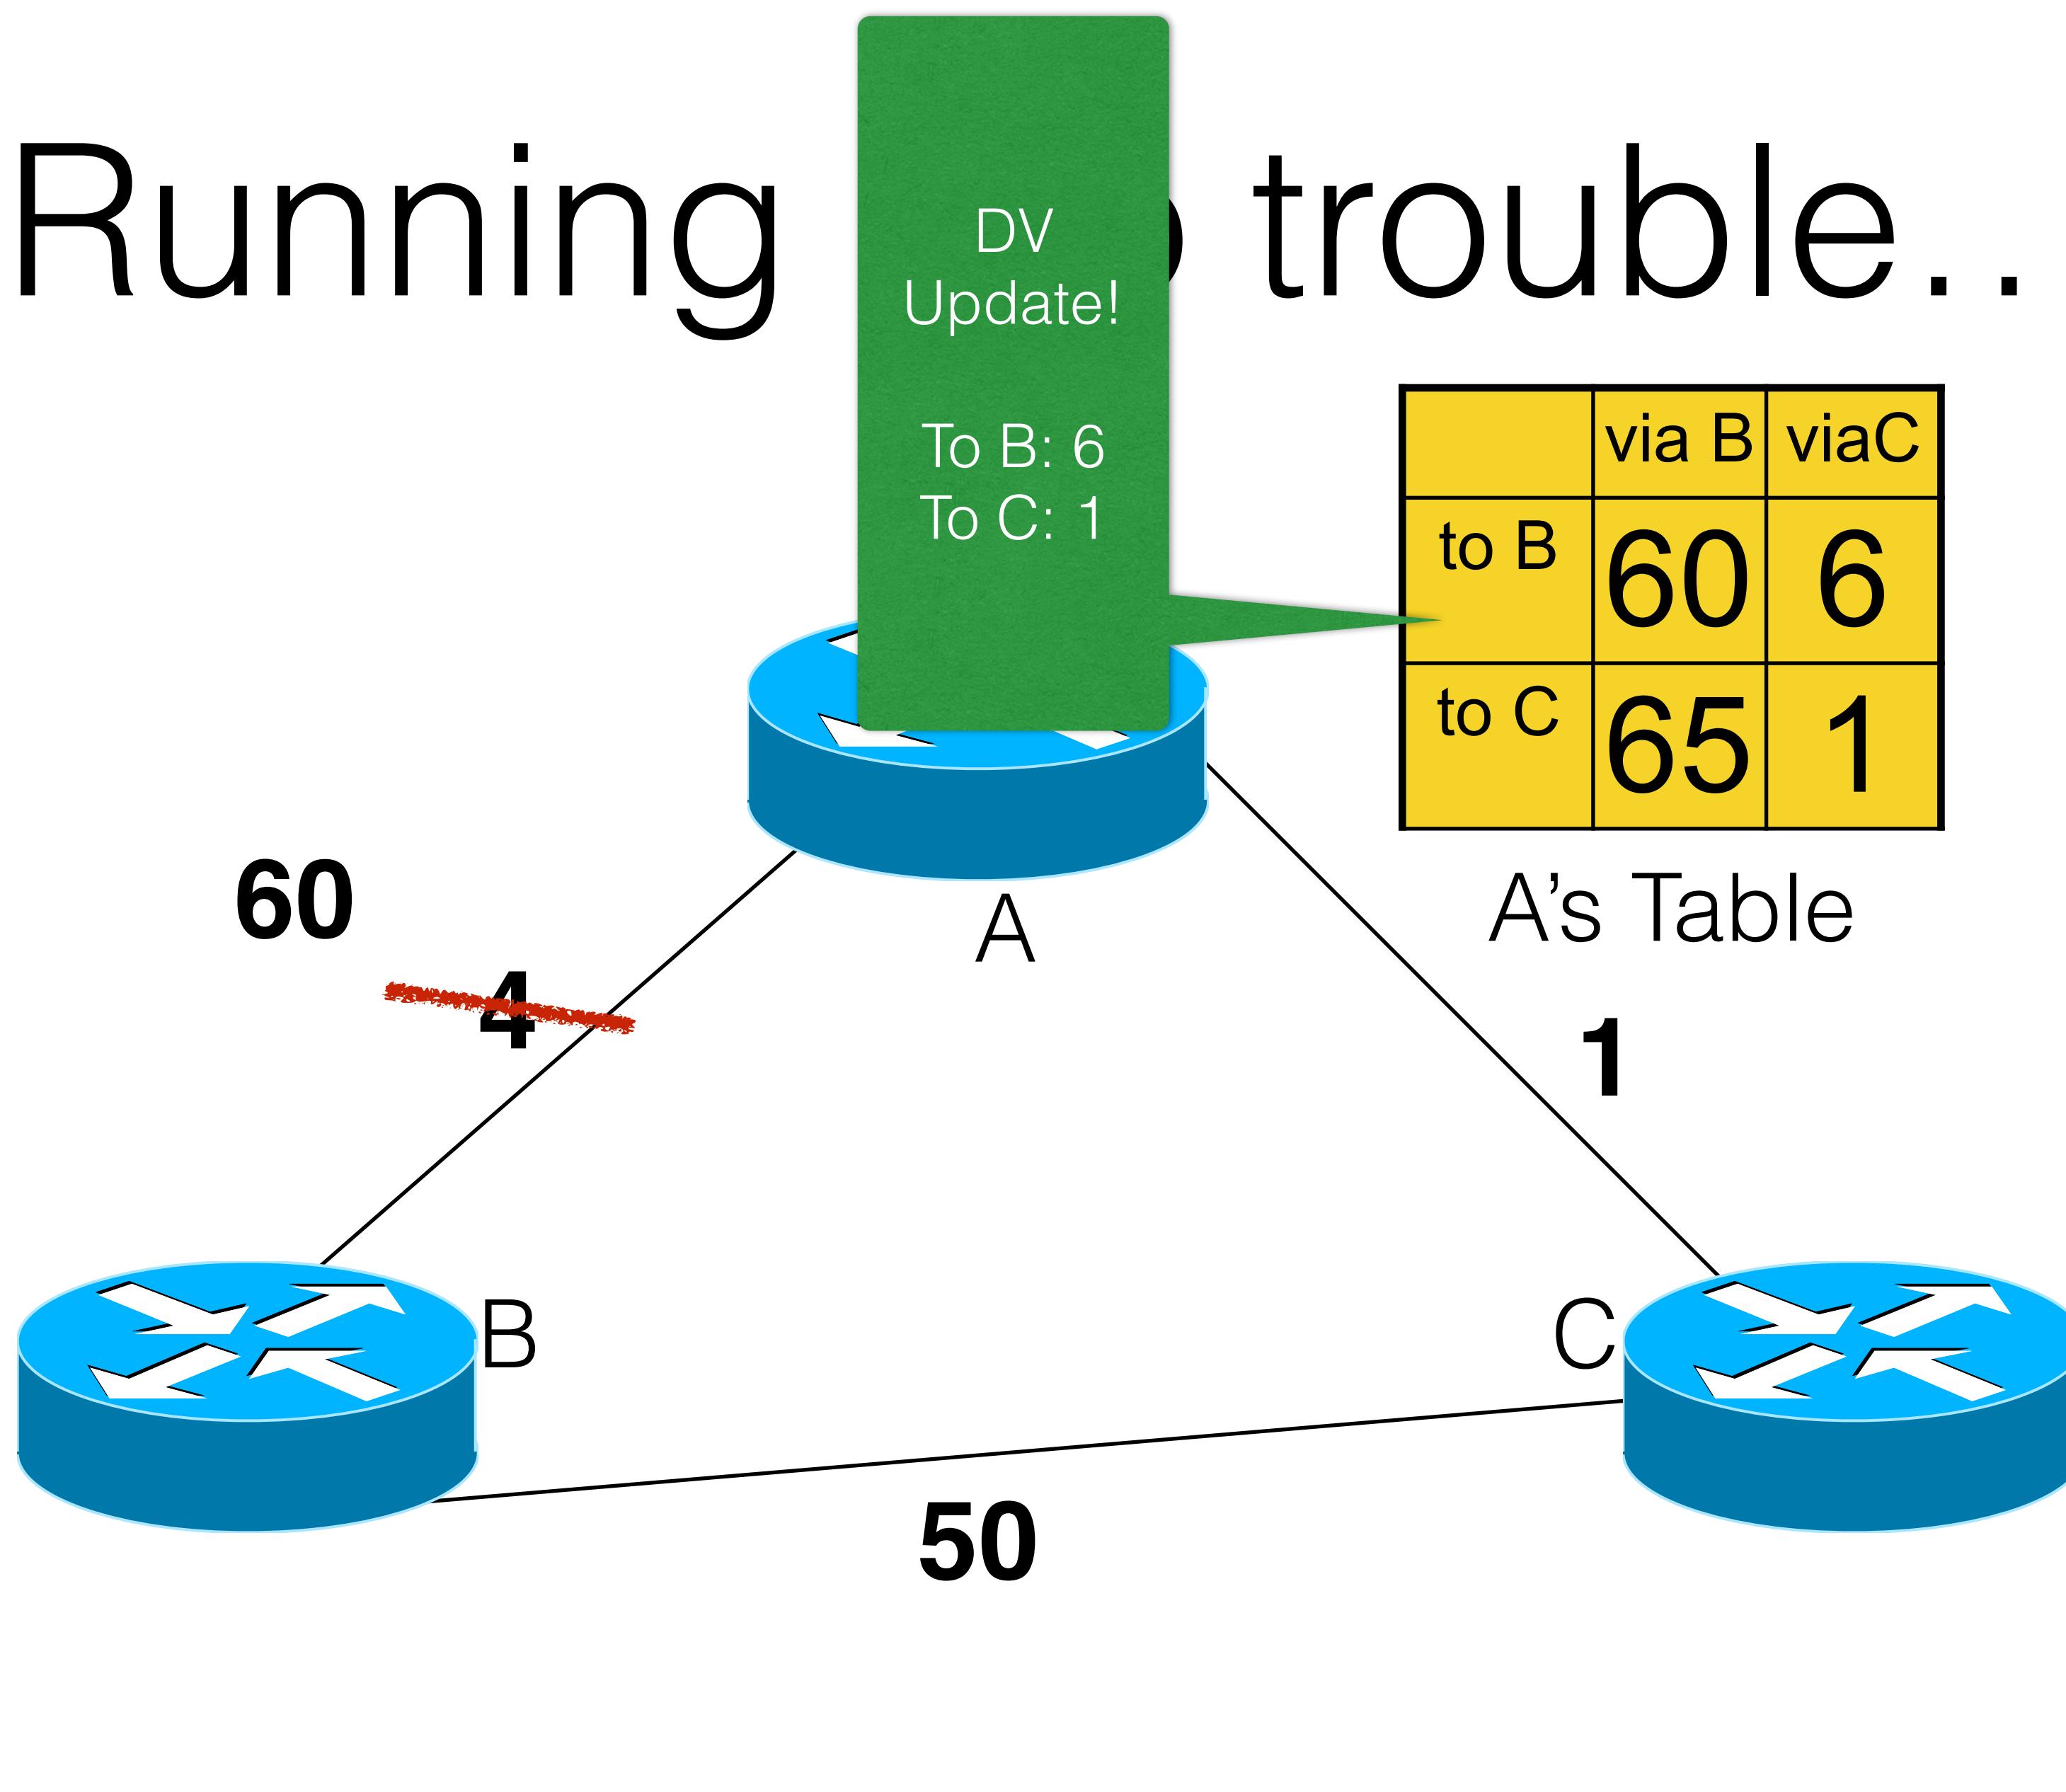


Running trouble....

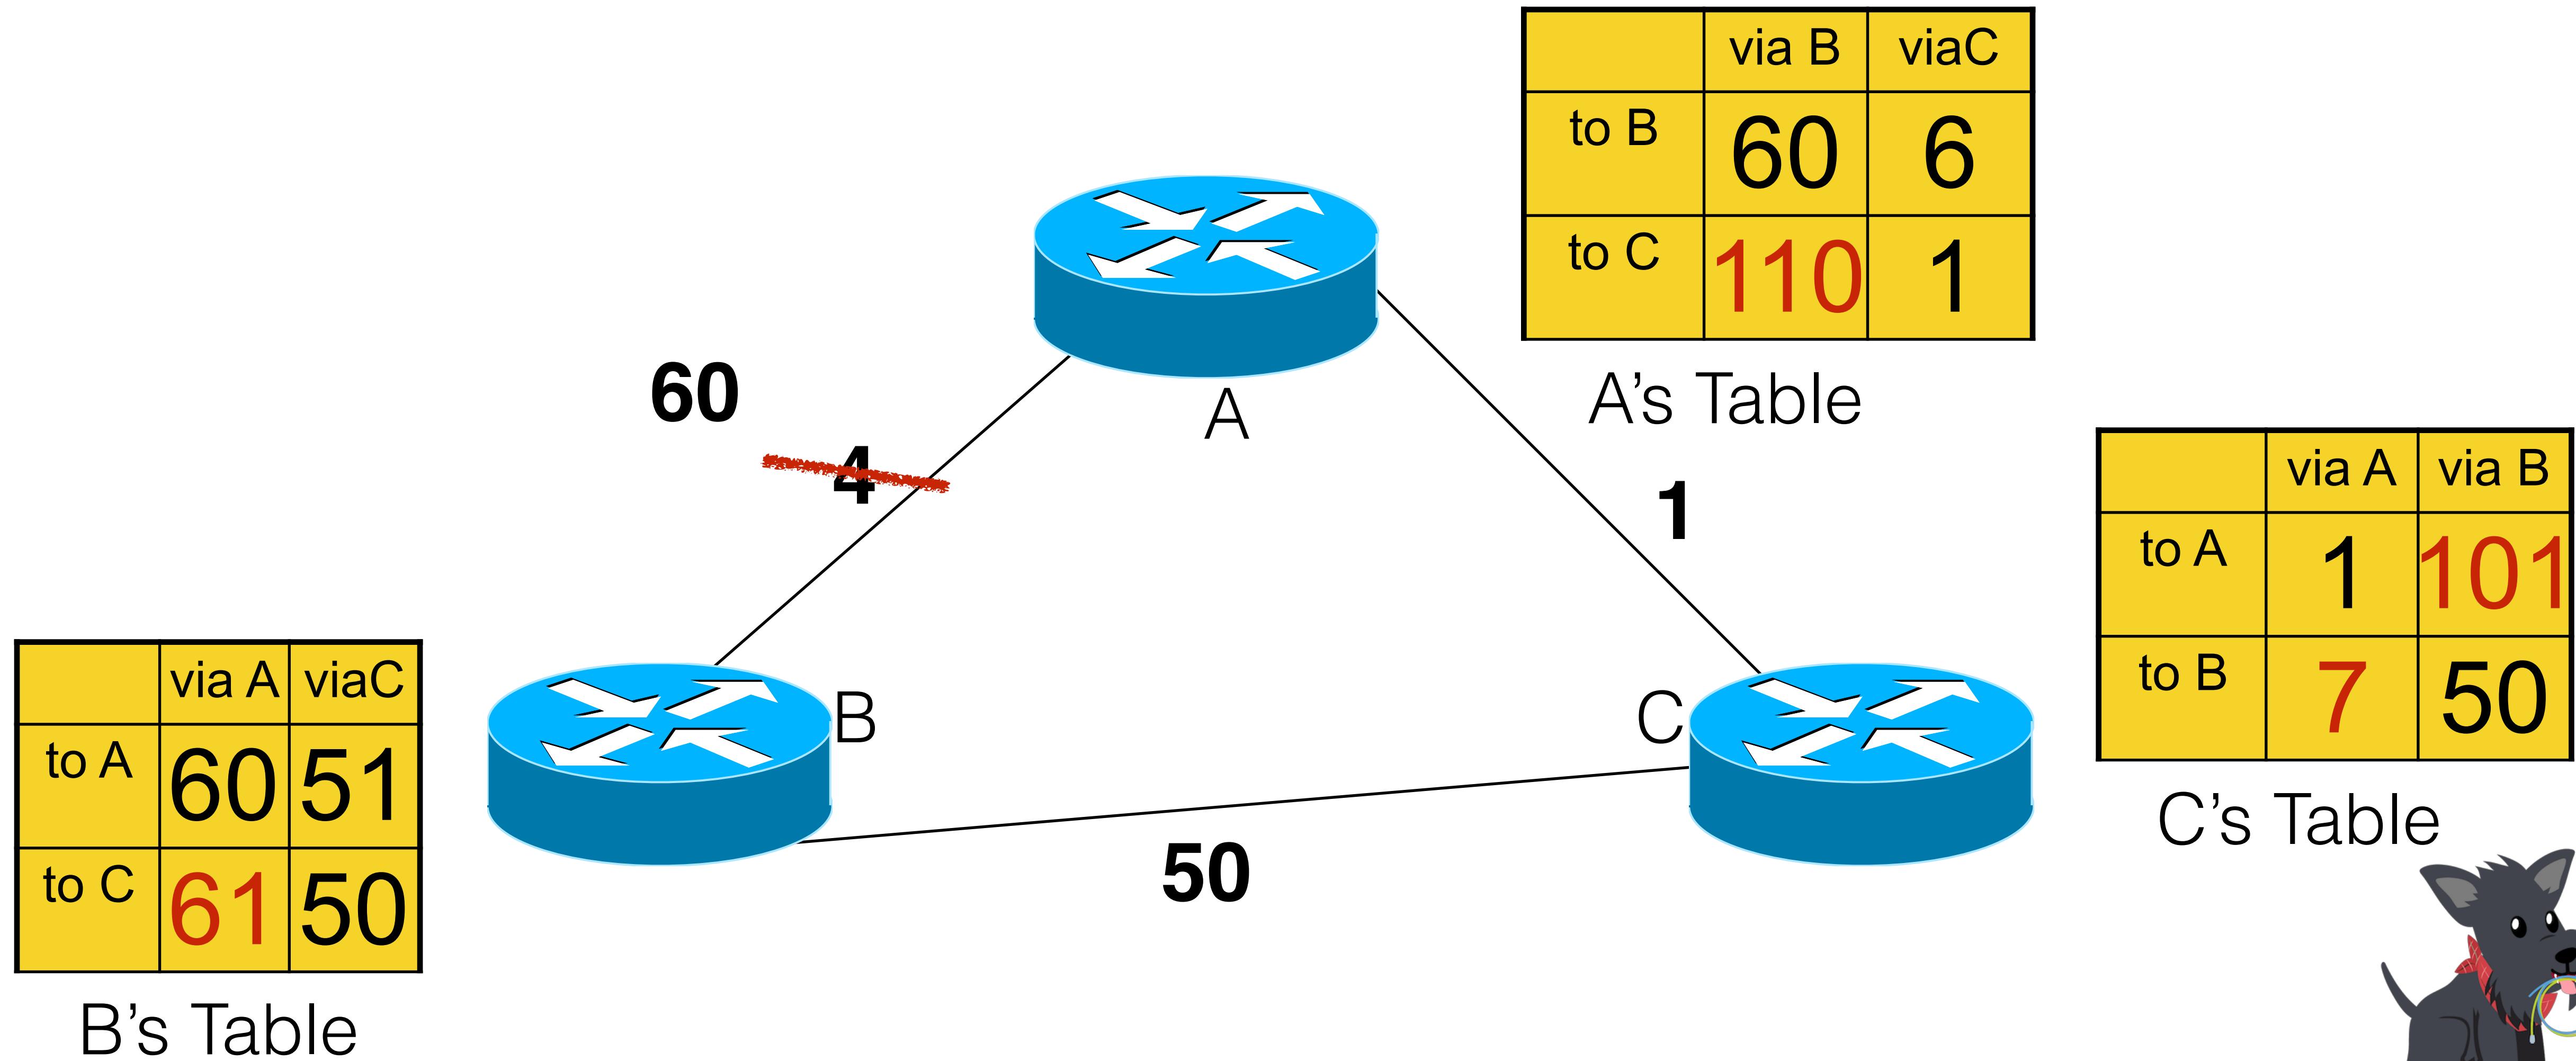
DV Update!
To A: 51
To C: 50

	via A	viaC
to A	60	51
to C	61	50

B's Table



Running into trouble....



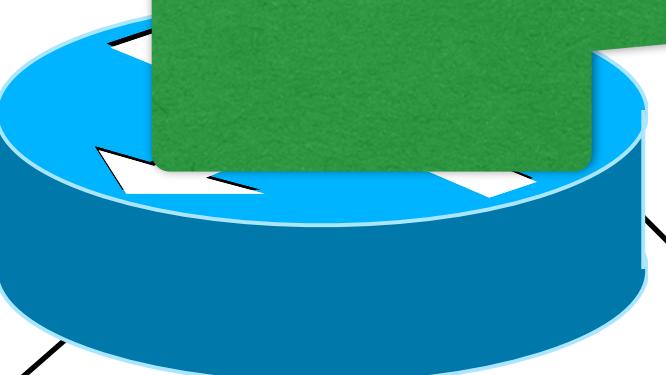
Running trouble...

DV Update!
To A: 51
To C: 50

	via A	viaC
to A	60	51
to C	61	50

B's Table

DV Update!
To B: 6
To C: 1



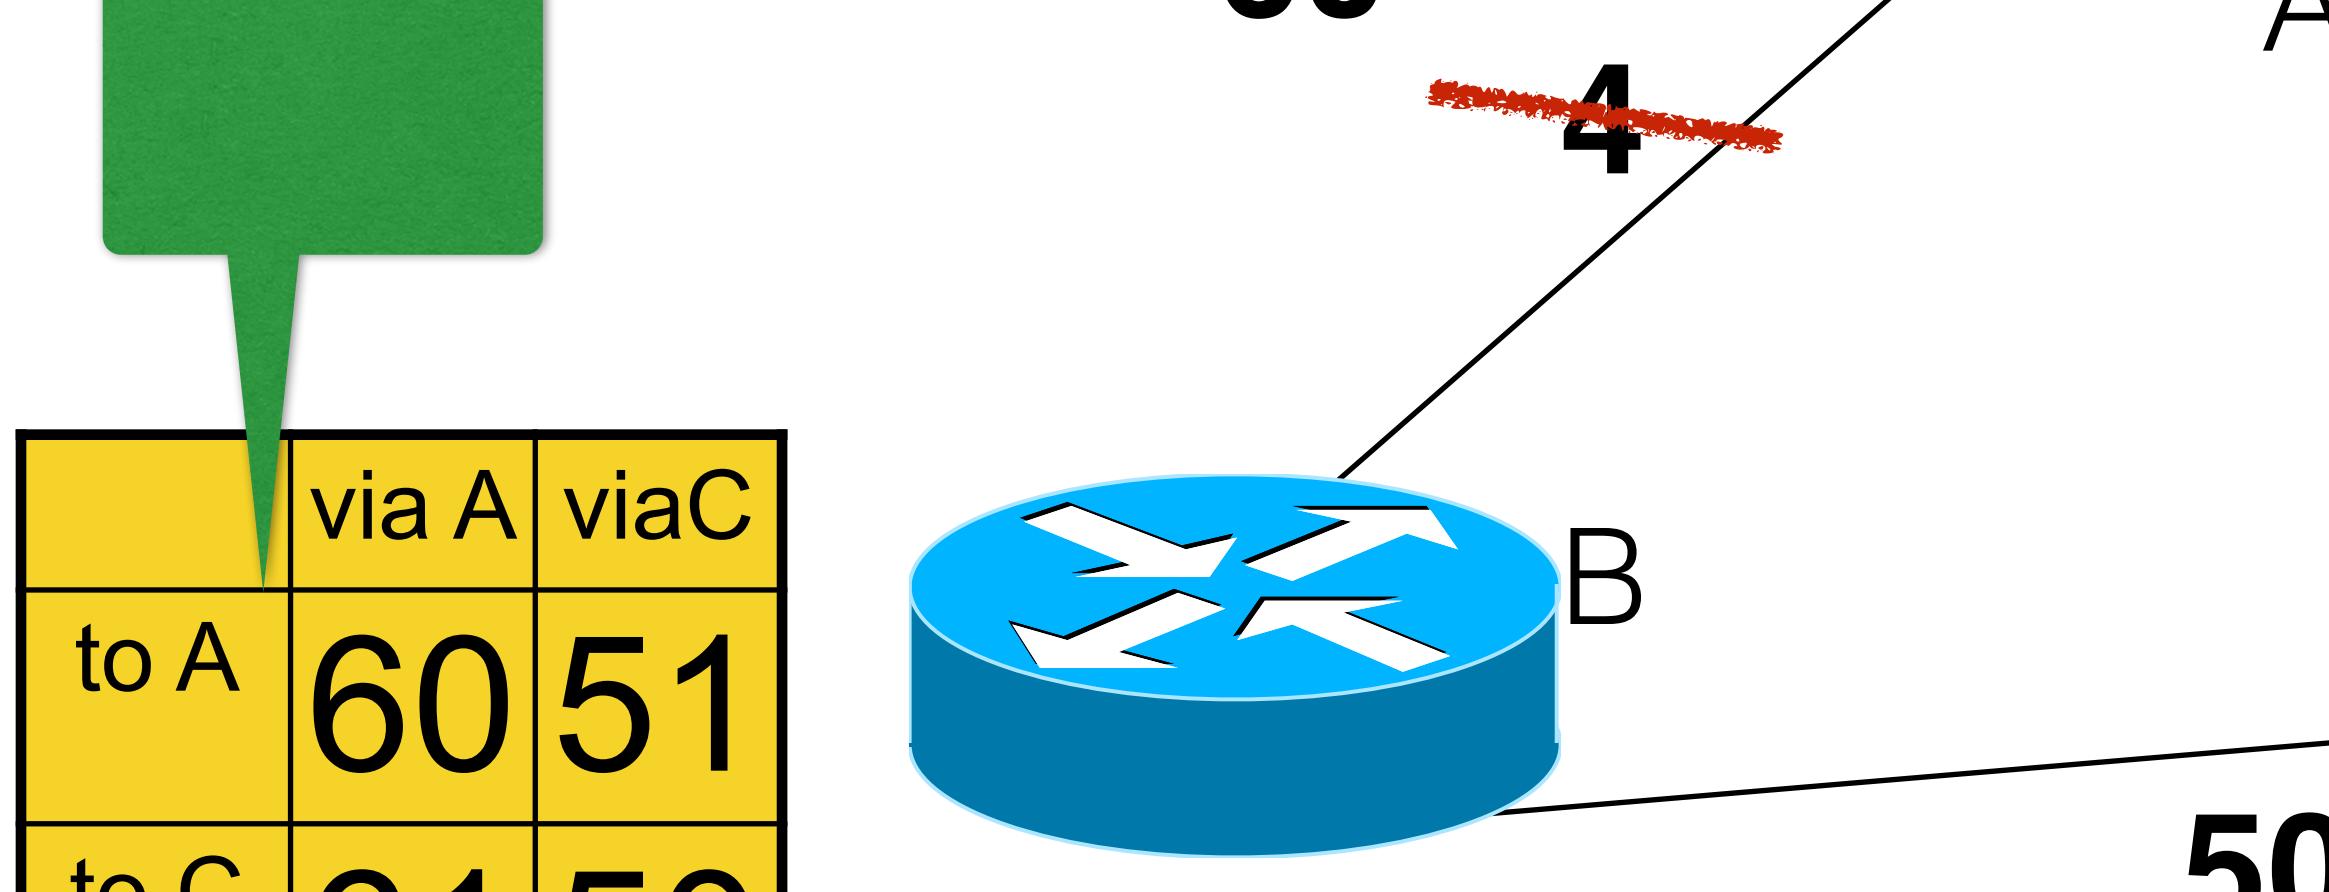
	via B	viaC
to B	60	6
to C	110	1

A's Table

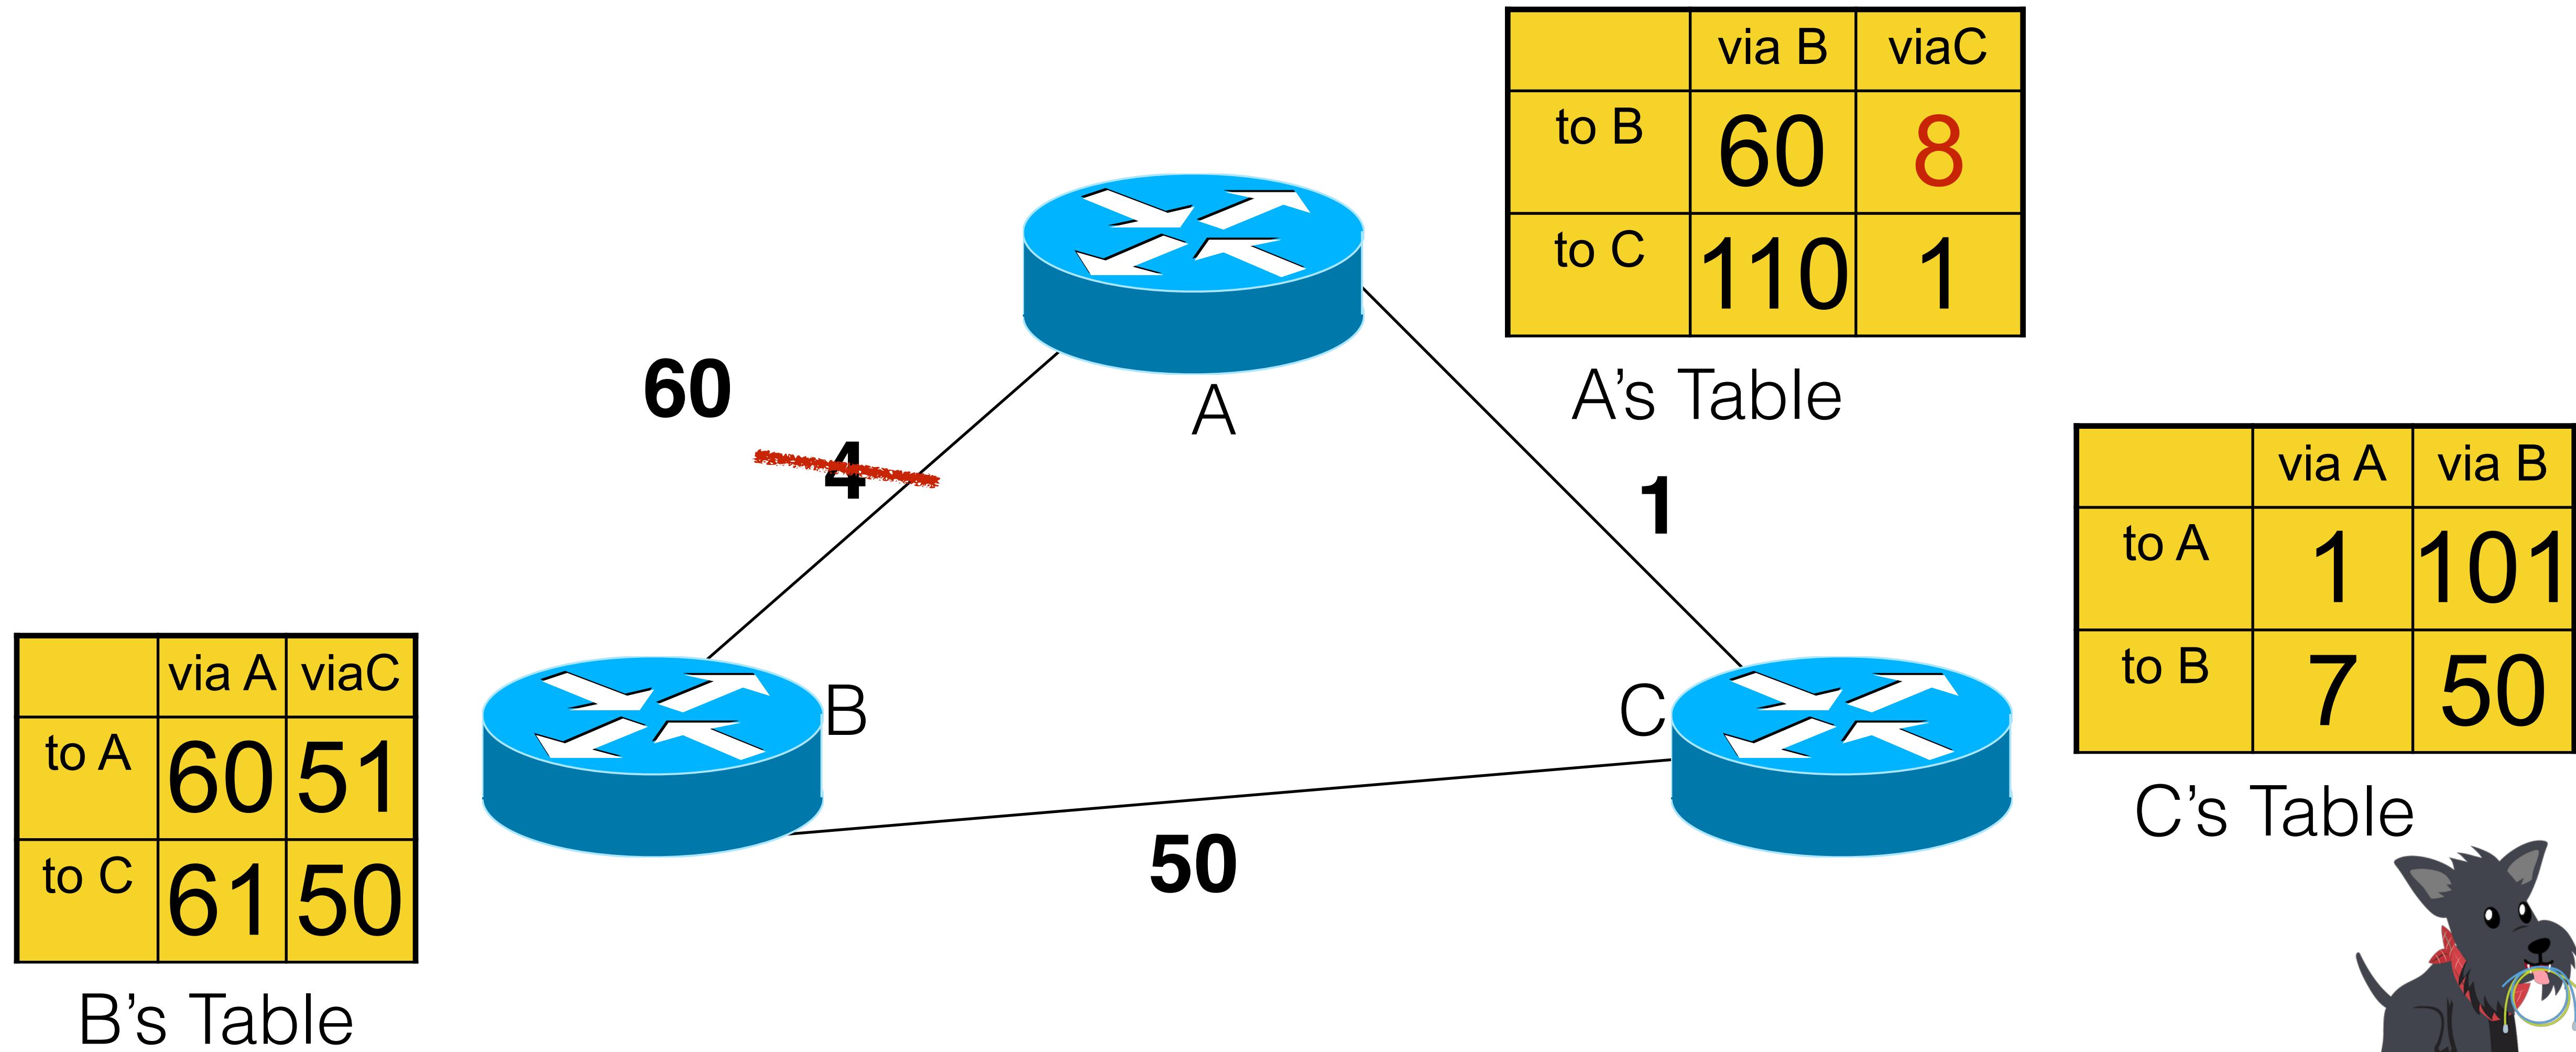
DV Update!
To A: 1
To B: 7

	via A	via B
to A	1	101
to B	7	50

C's Table



Running into trouble....

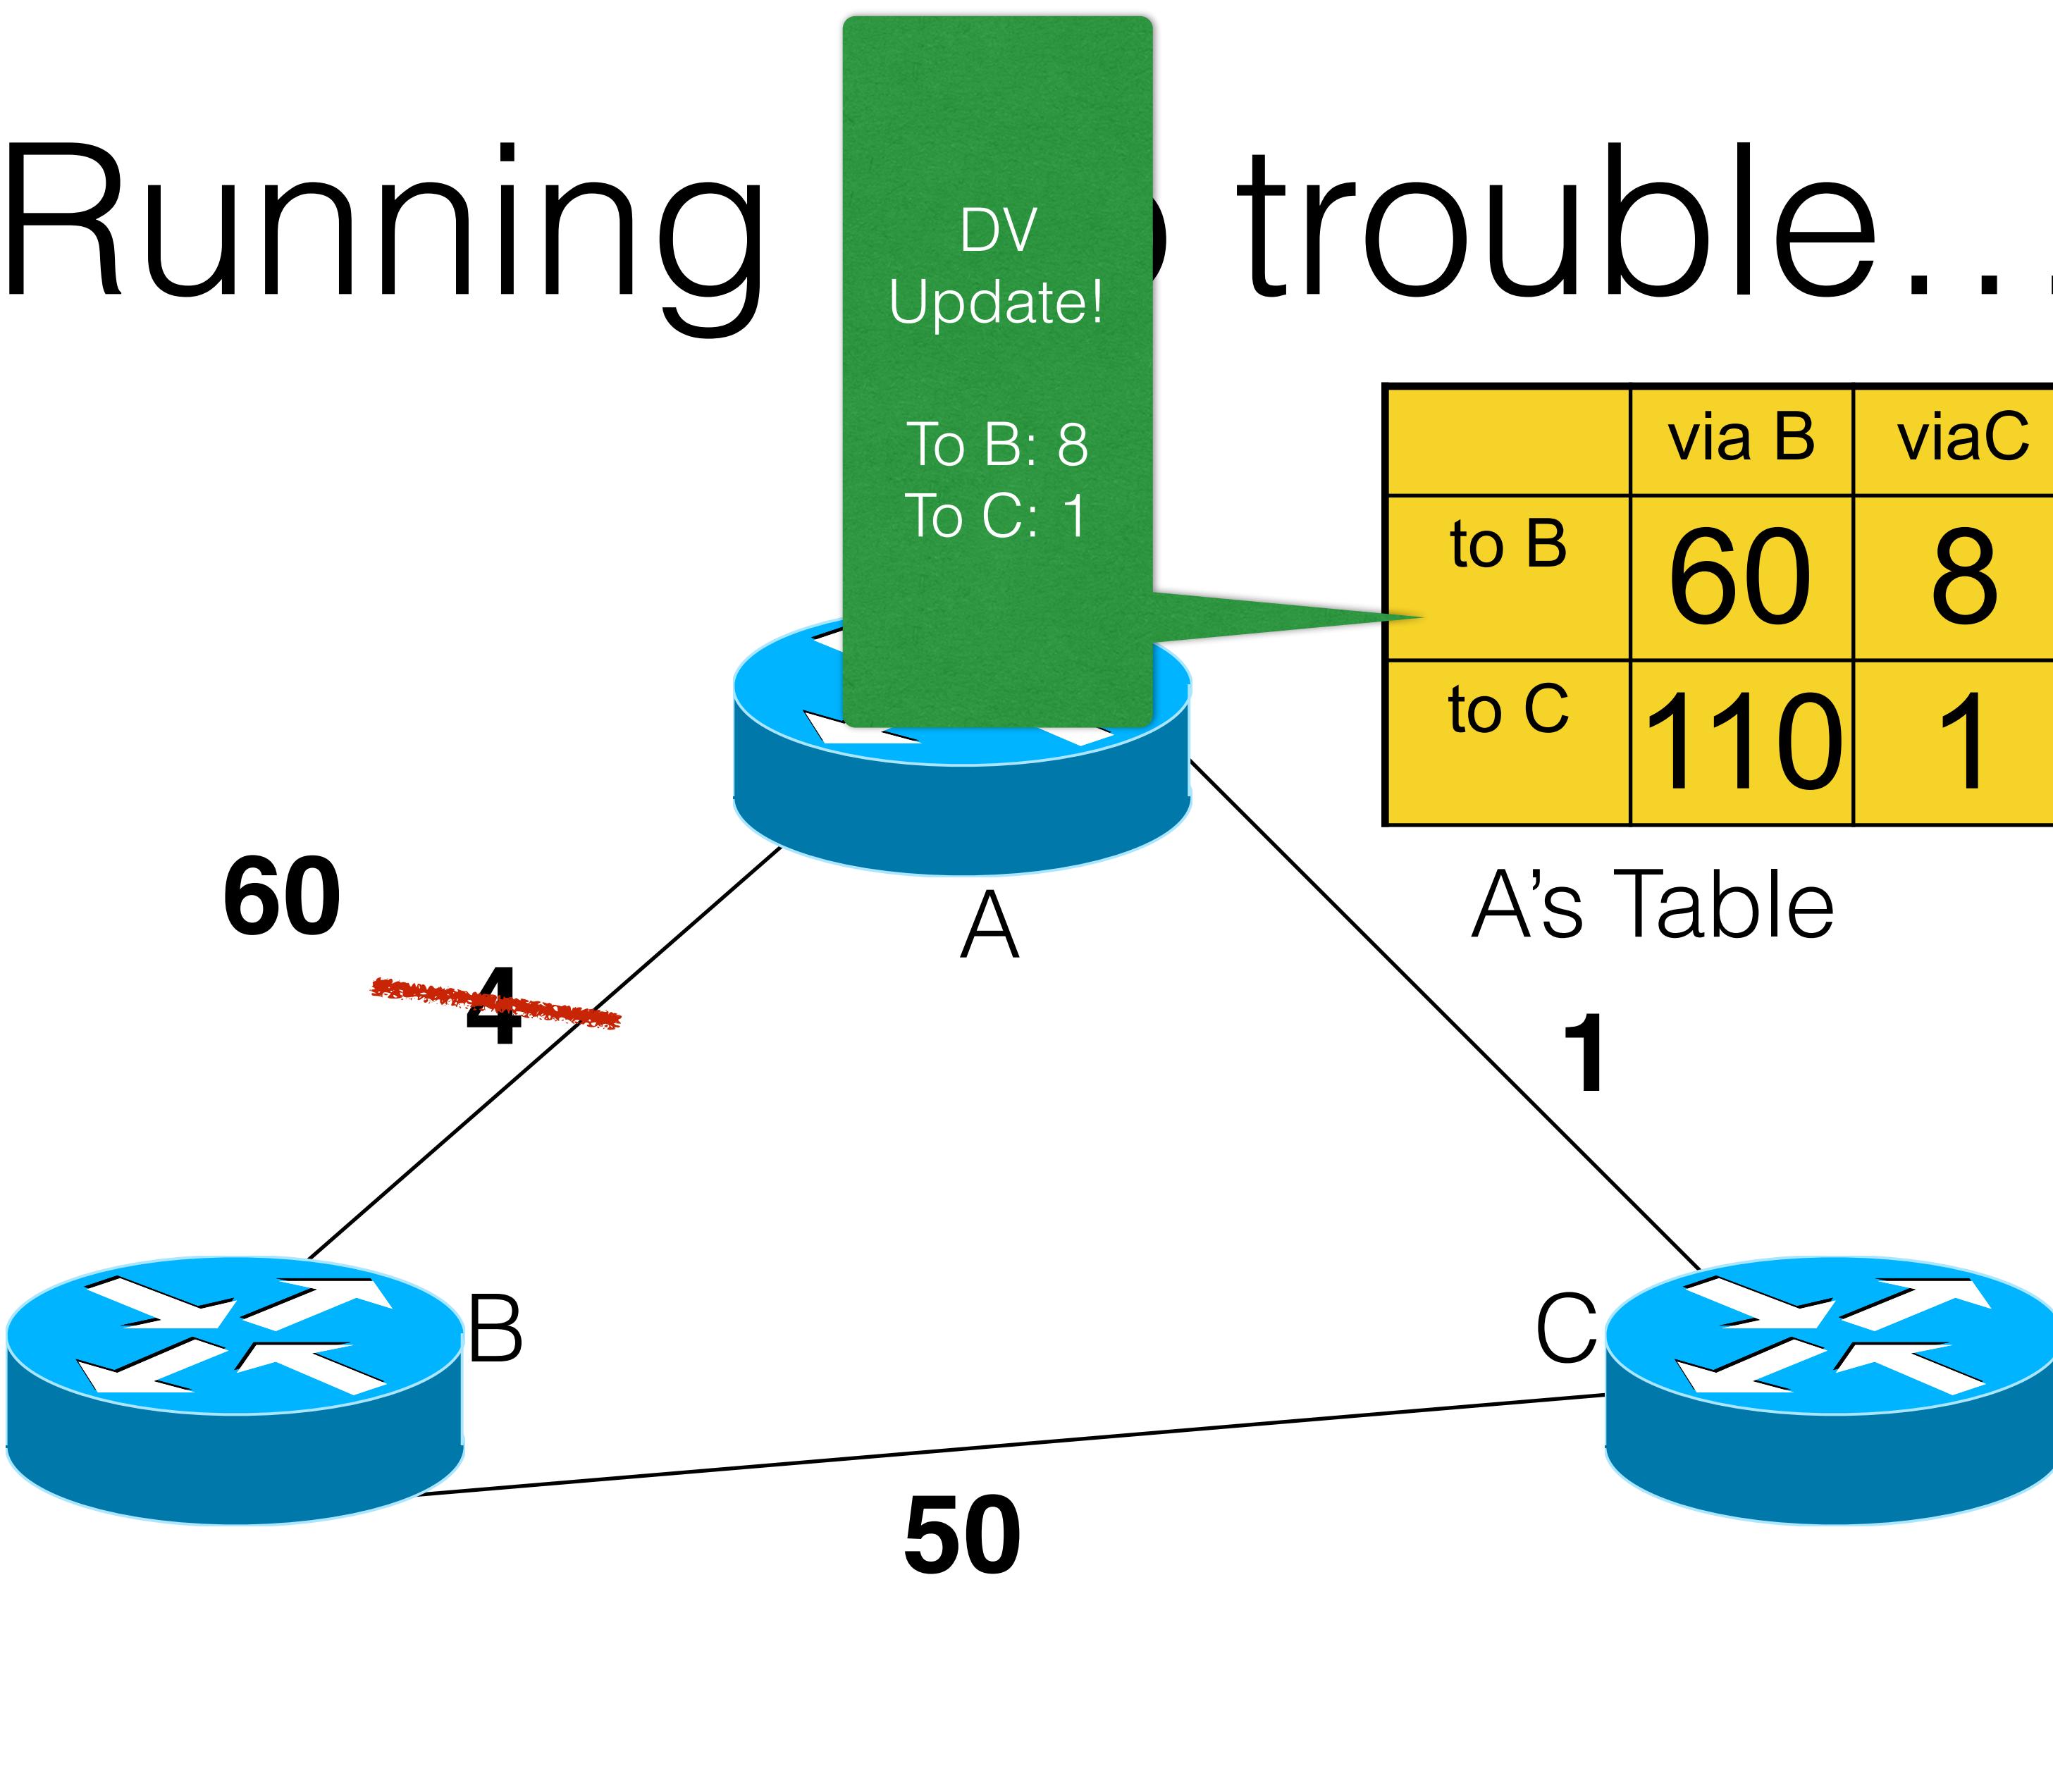


Running trouble....

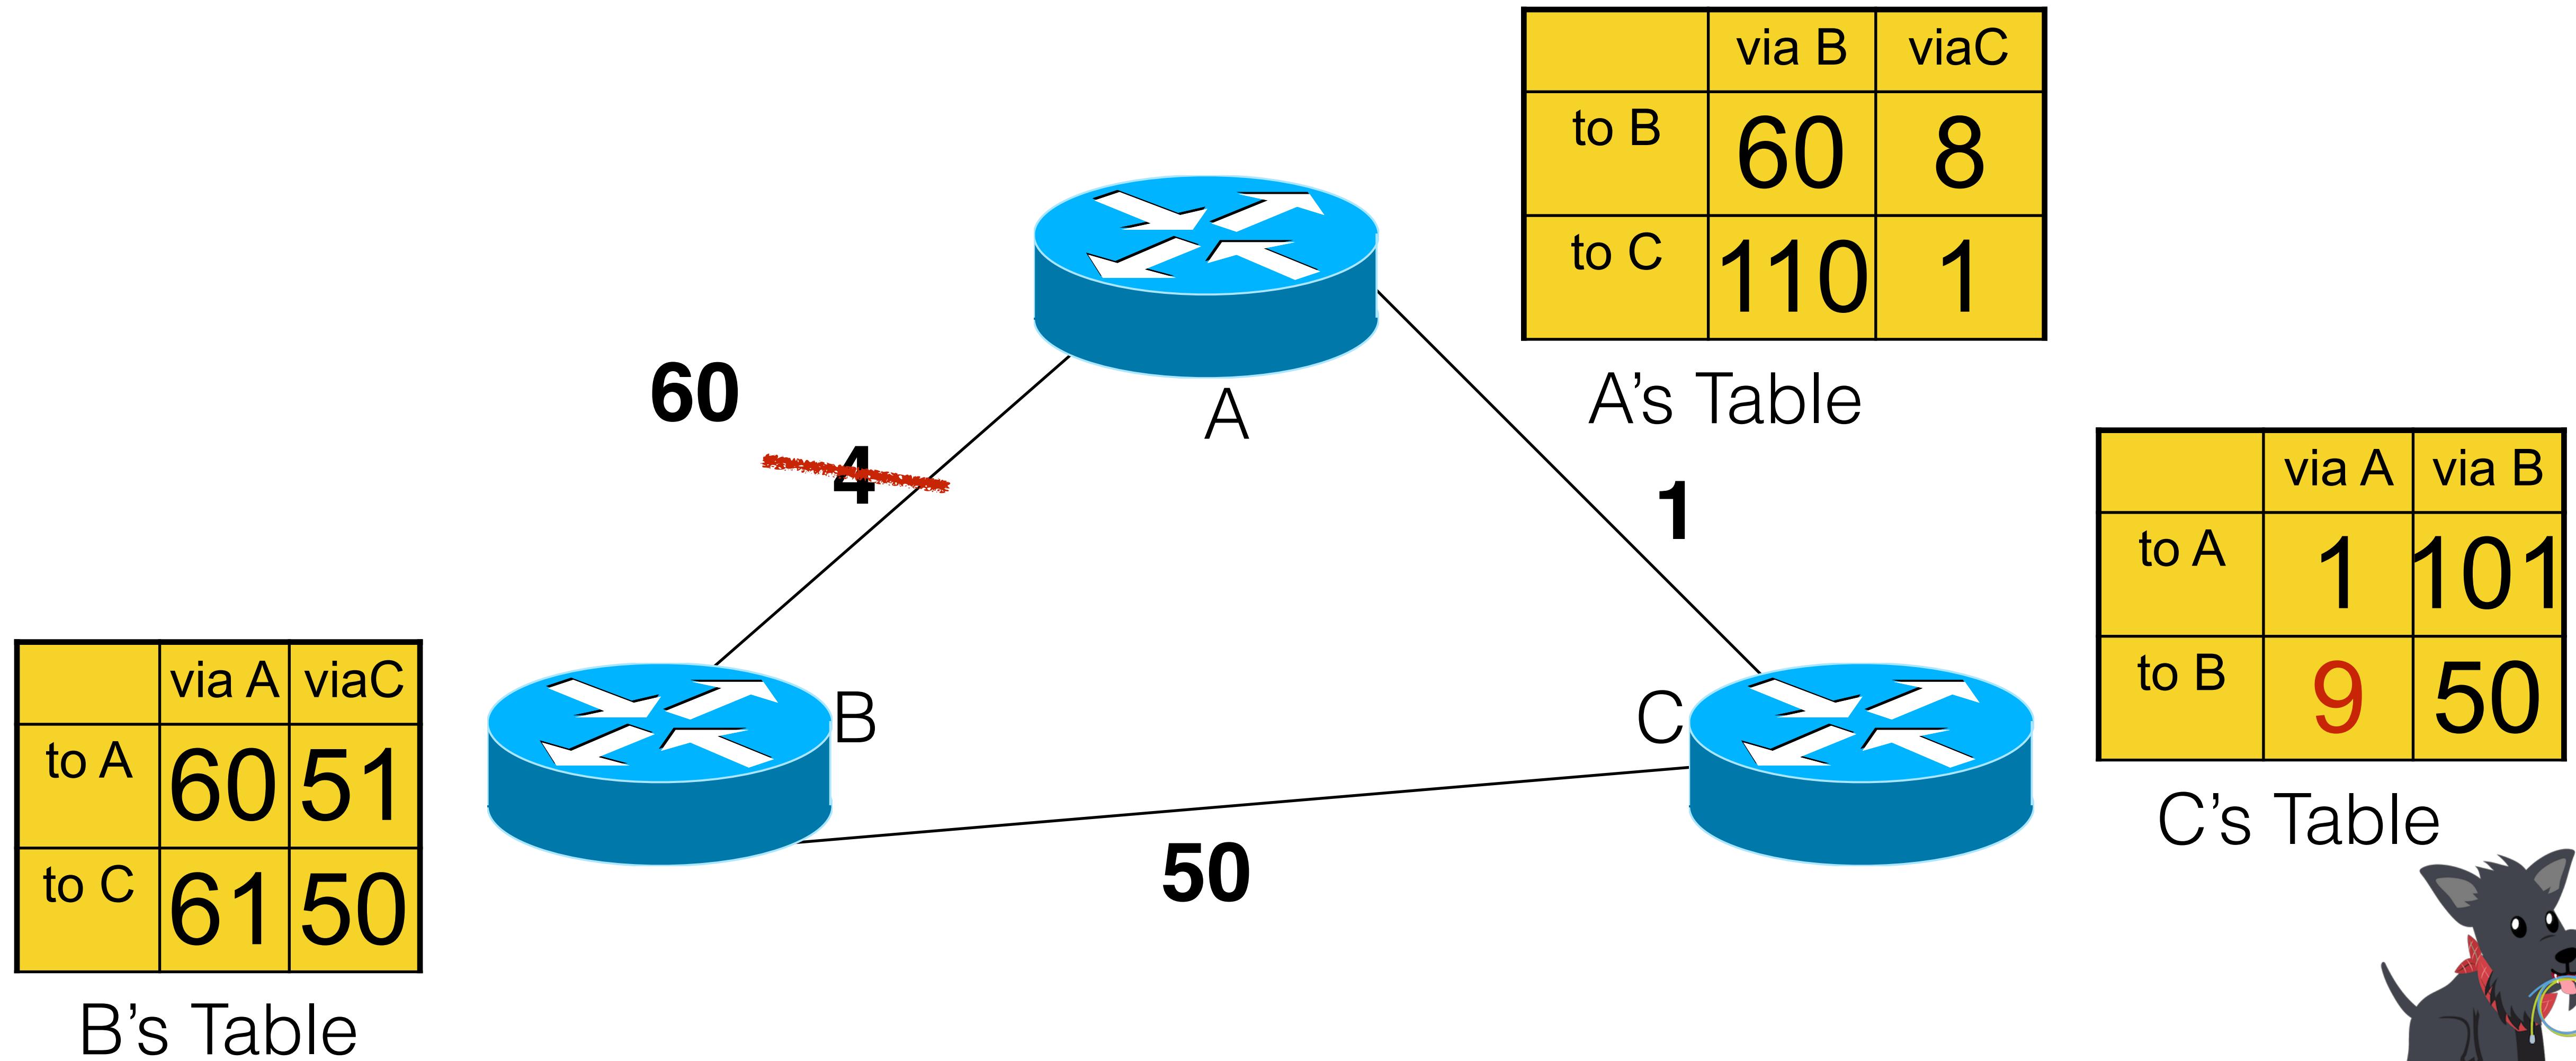
DV Update!
To A: 51
To C: 50

	via A	viaC
to A	60	51
to C	61	50

B's Table



Running into trouble....

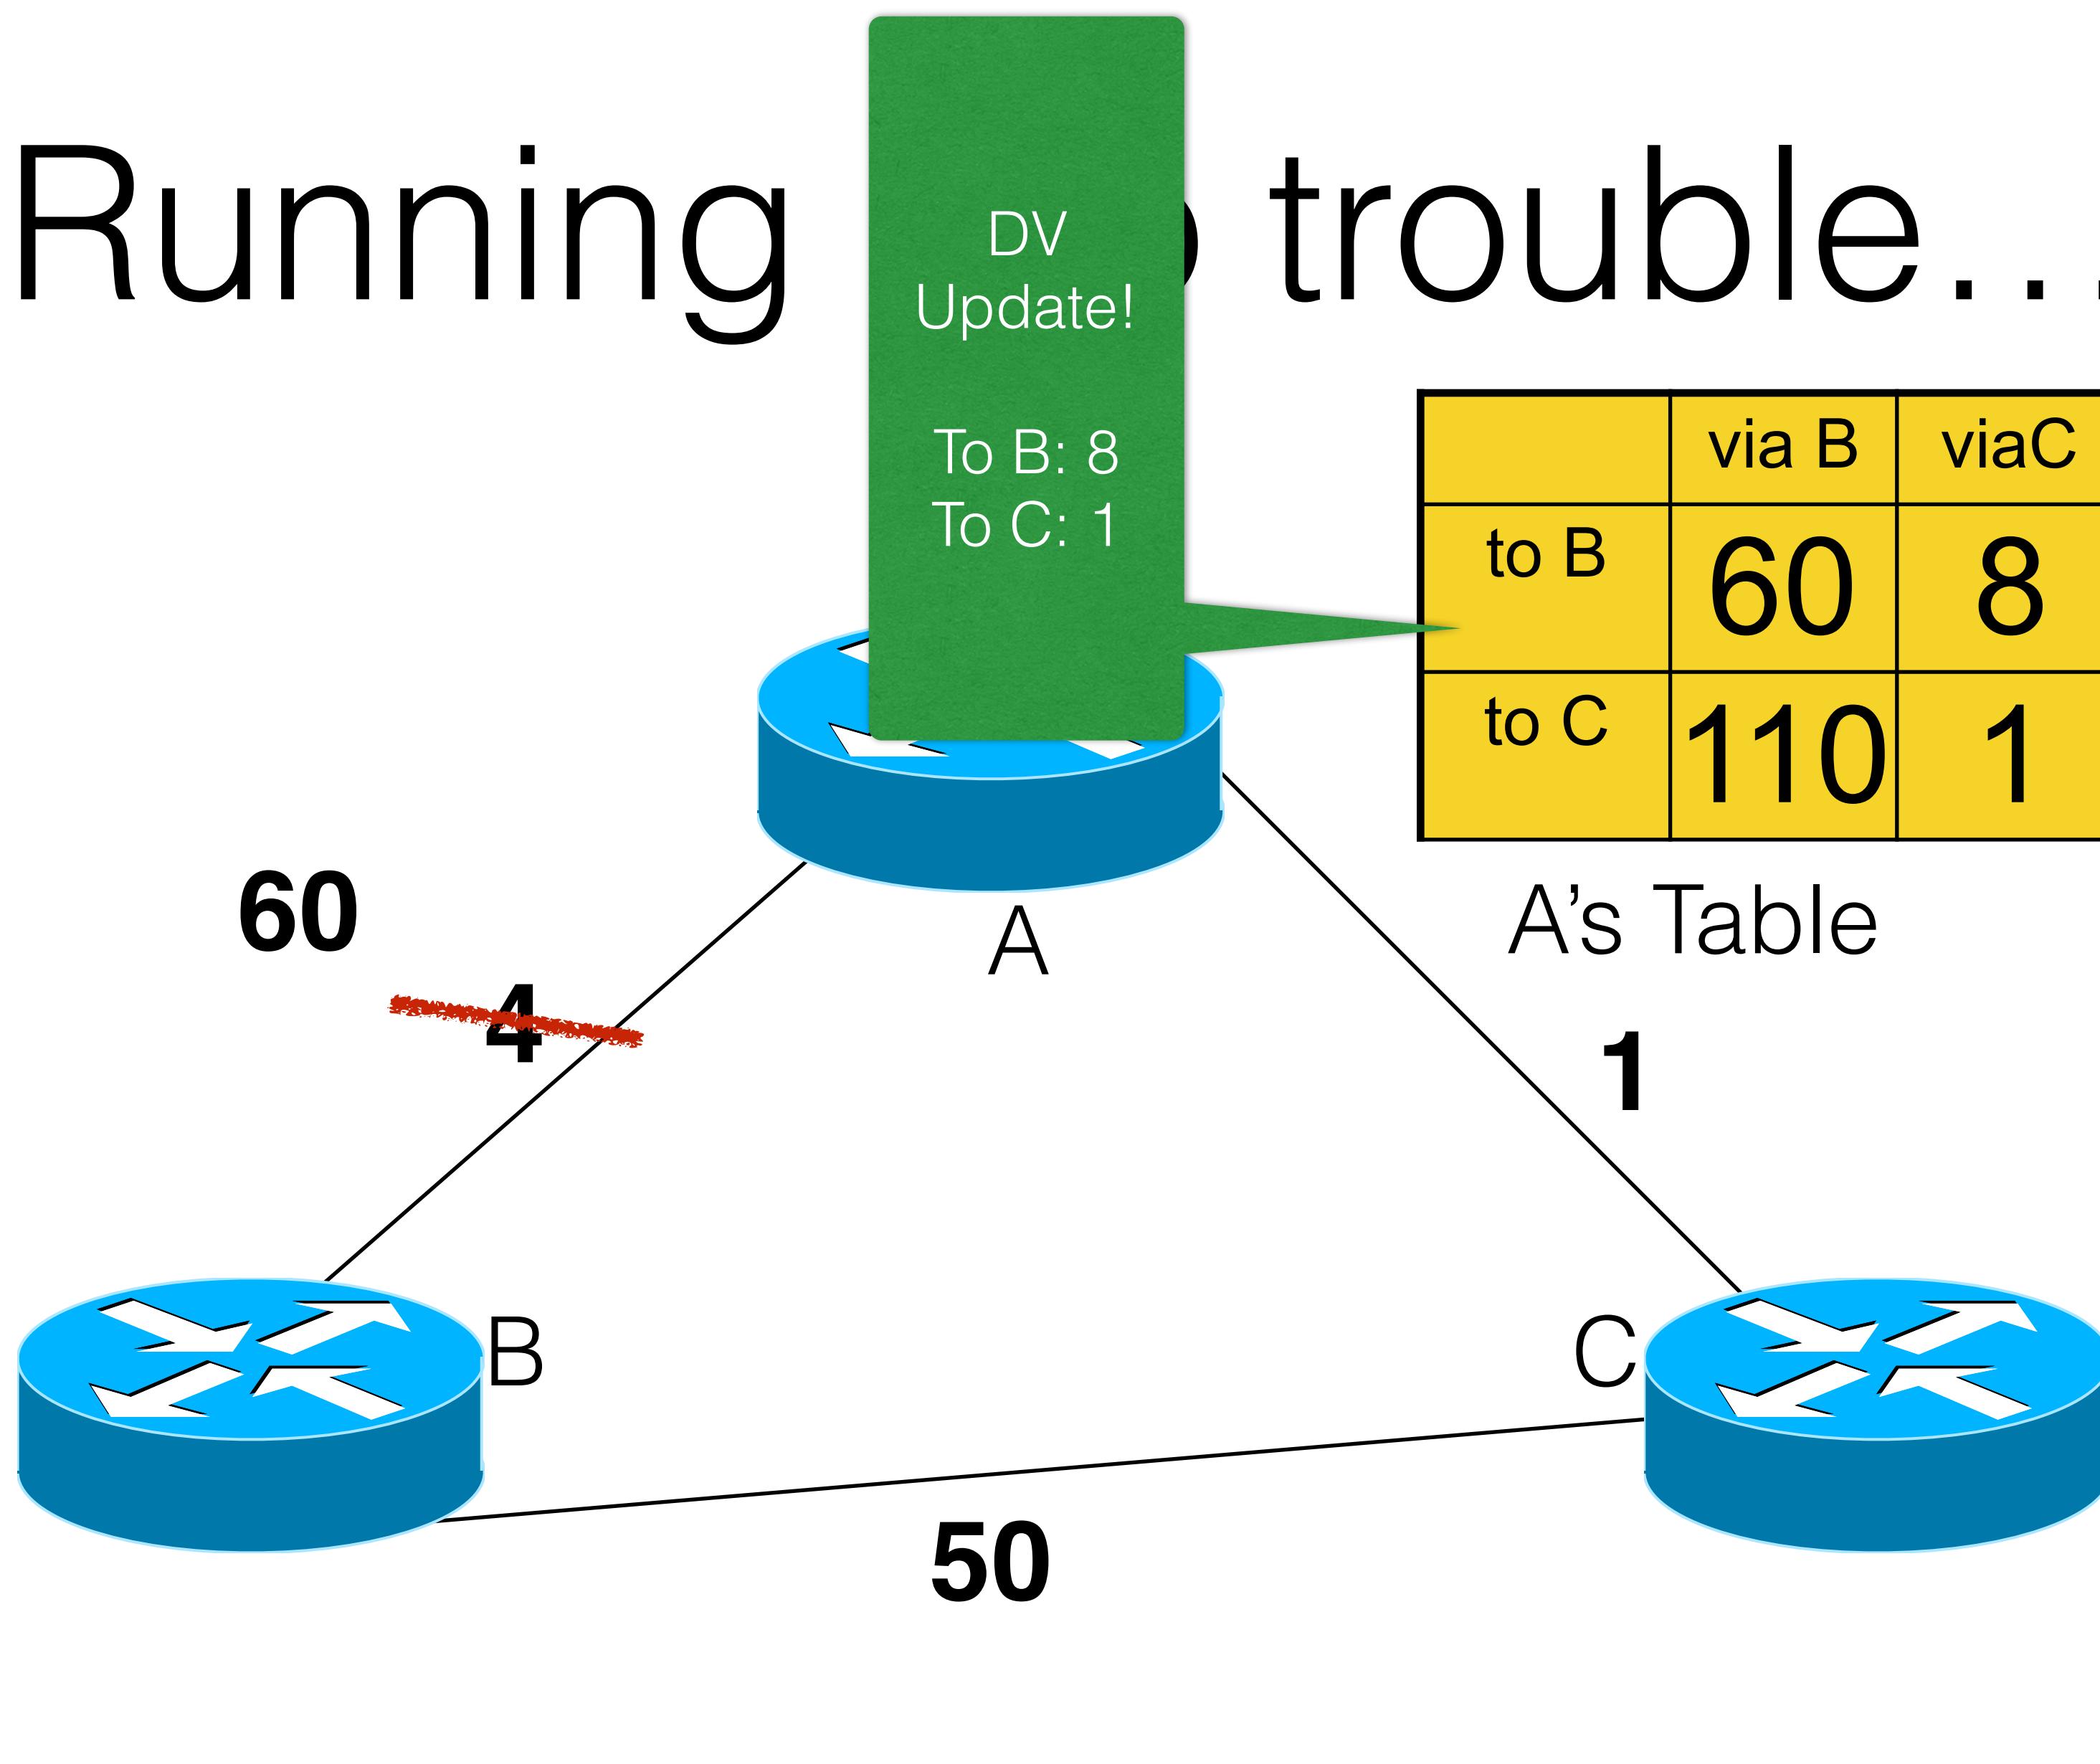


Running trouble....

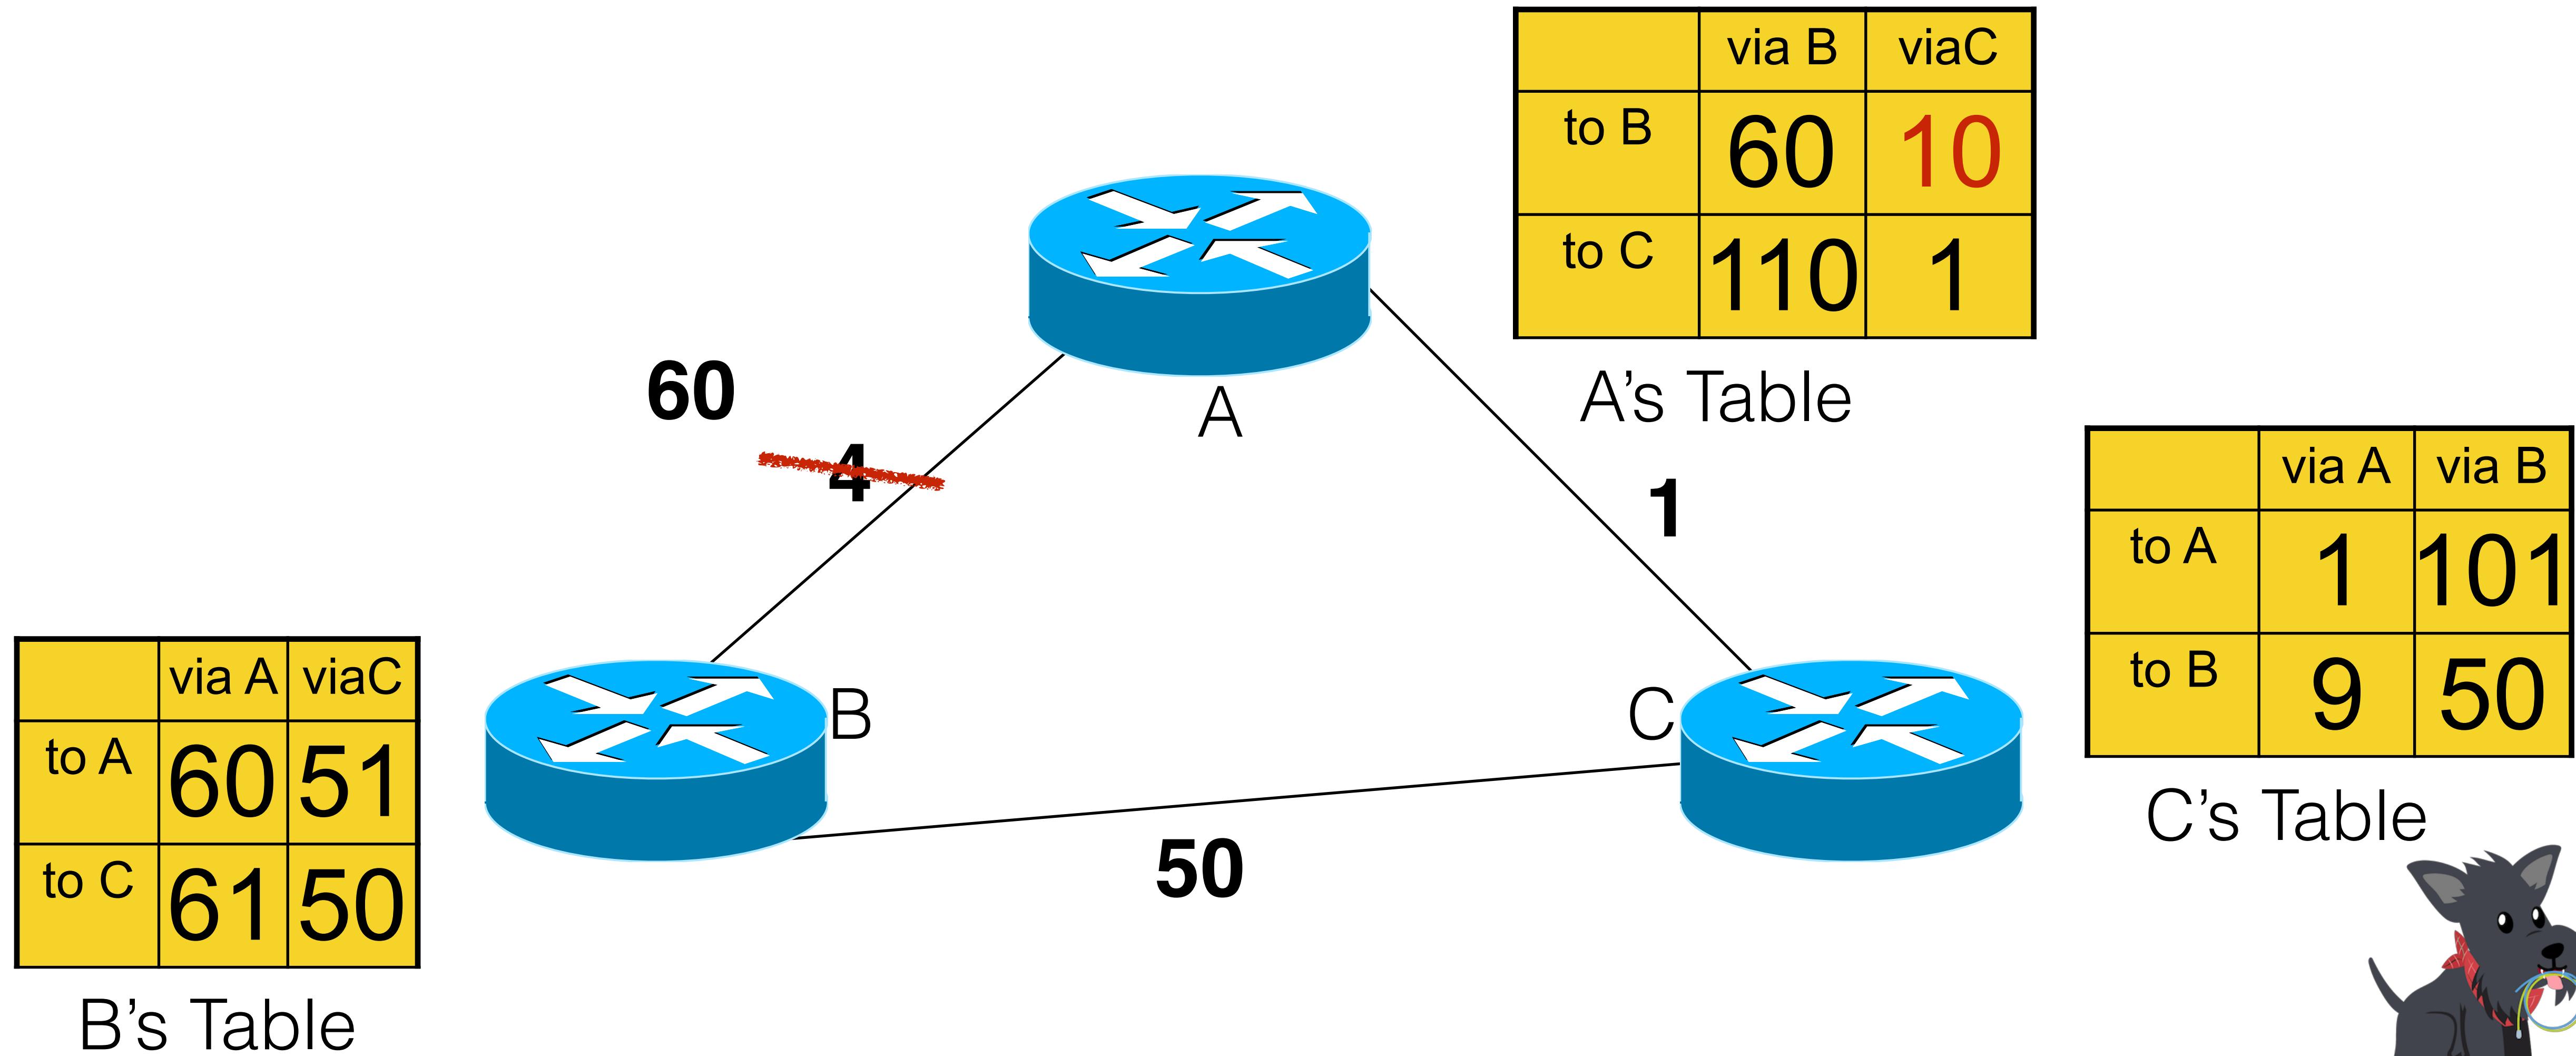
DV Update!
To A: 51
To C: 50

	via A	viaC
to A	60	51
to C	61	50

B's Table



Running into trouble....



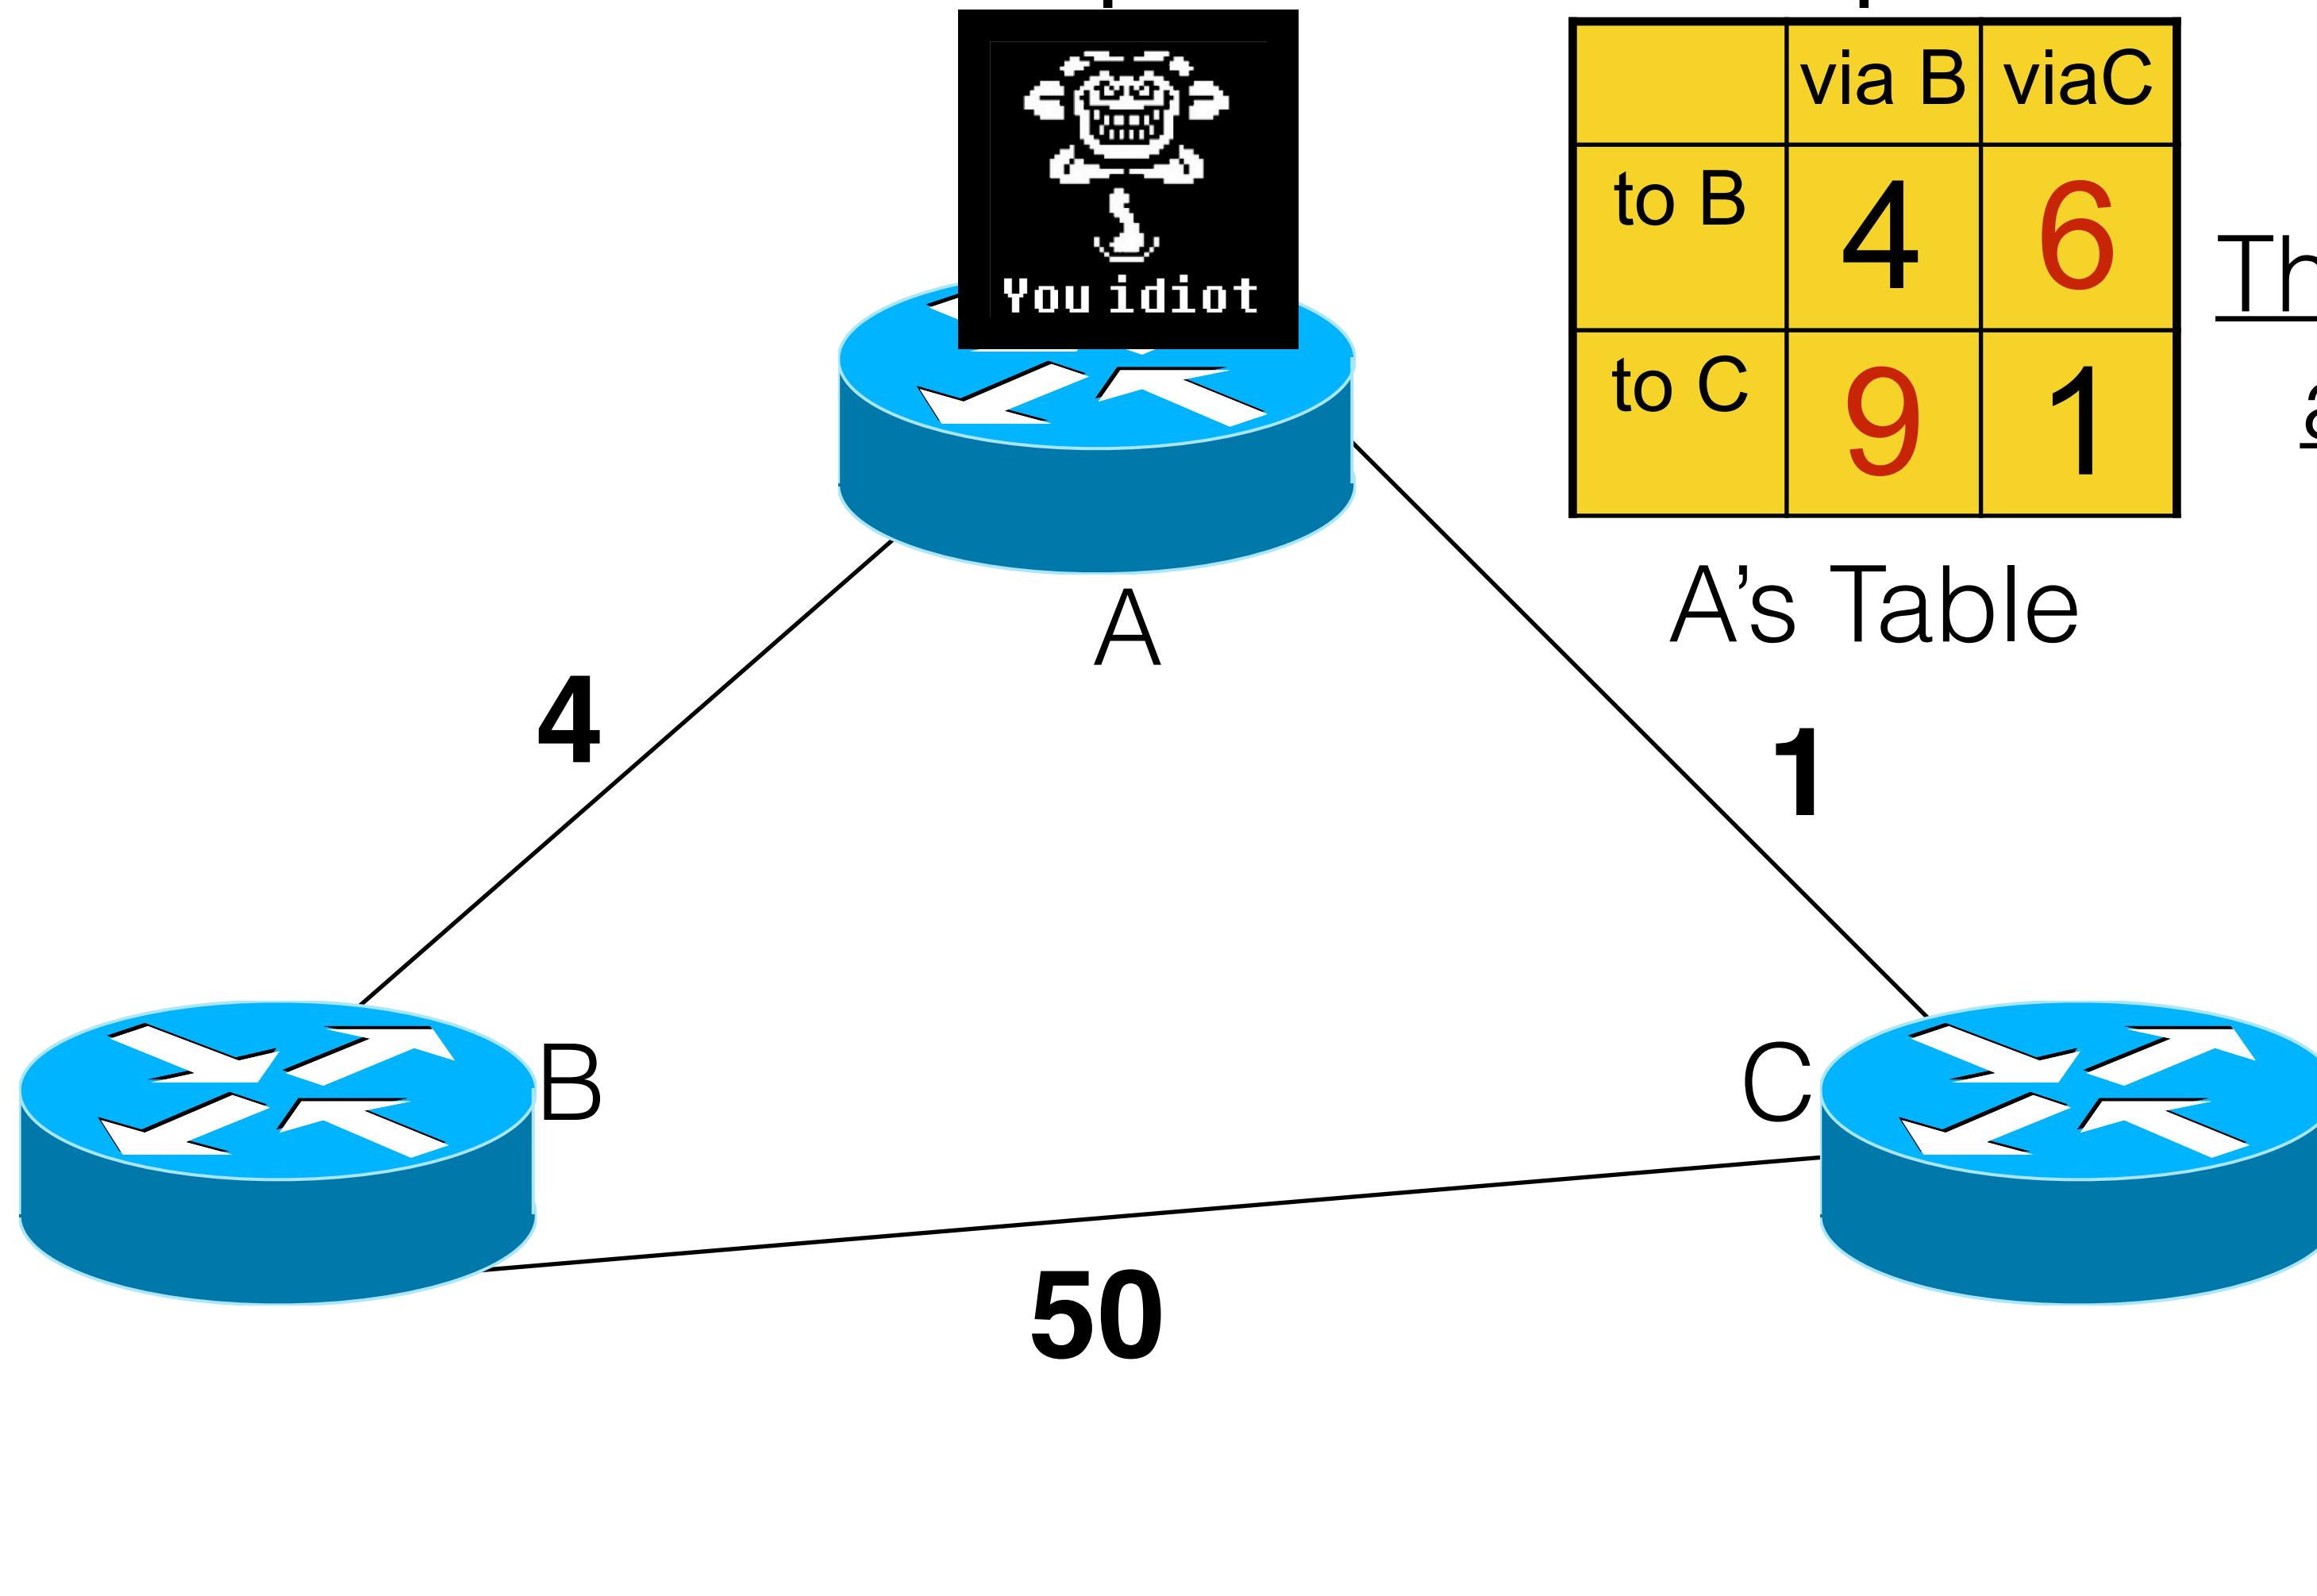
This is called the “count to infinity” problem.



Root of the Problem: DV algorithm has no way to detect and prevent loops.

	via A	via C
to A	4	51
to C	5	50

B's Table



These routes have a loop in them.



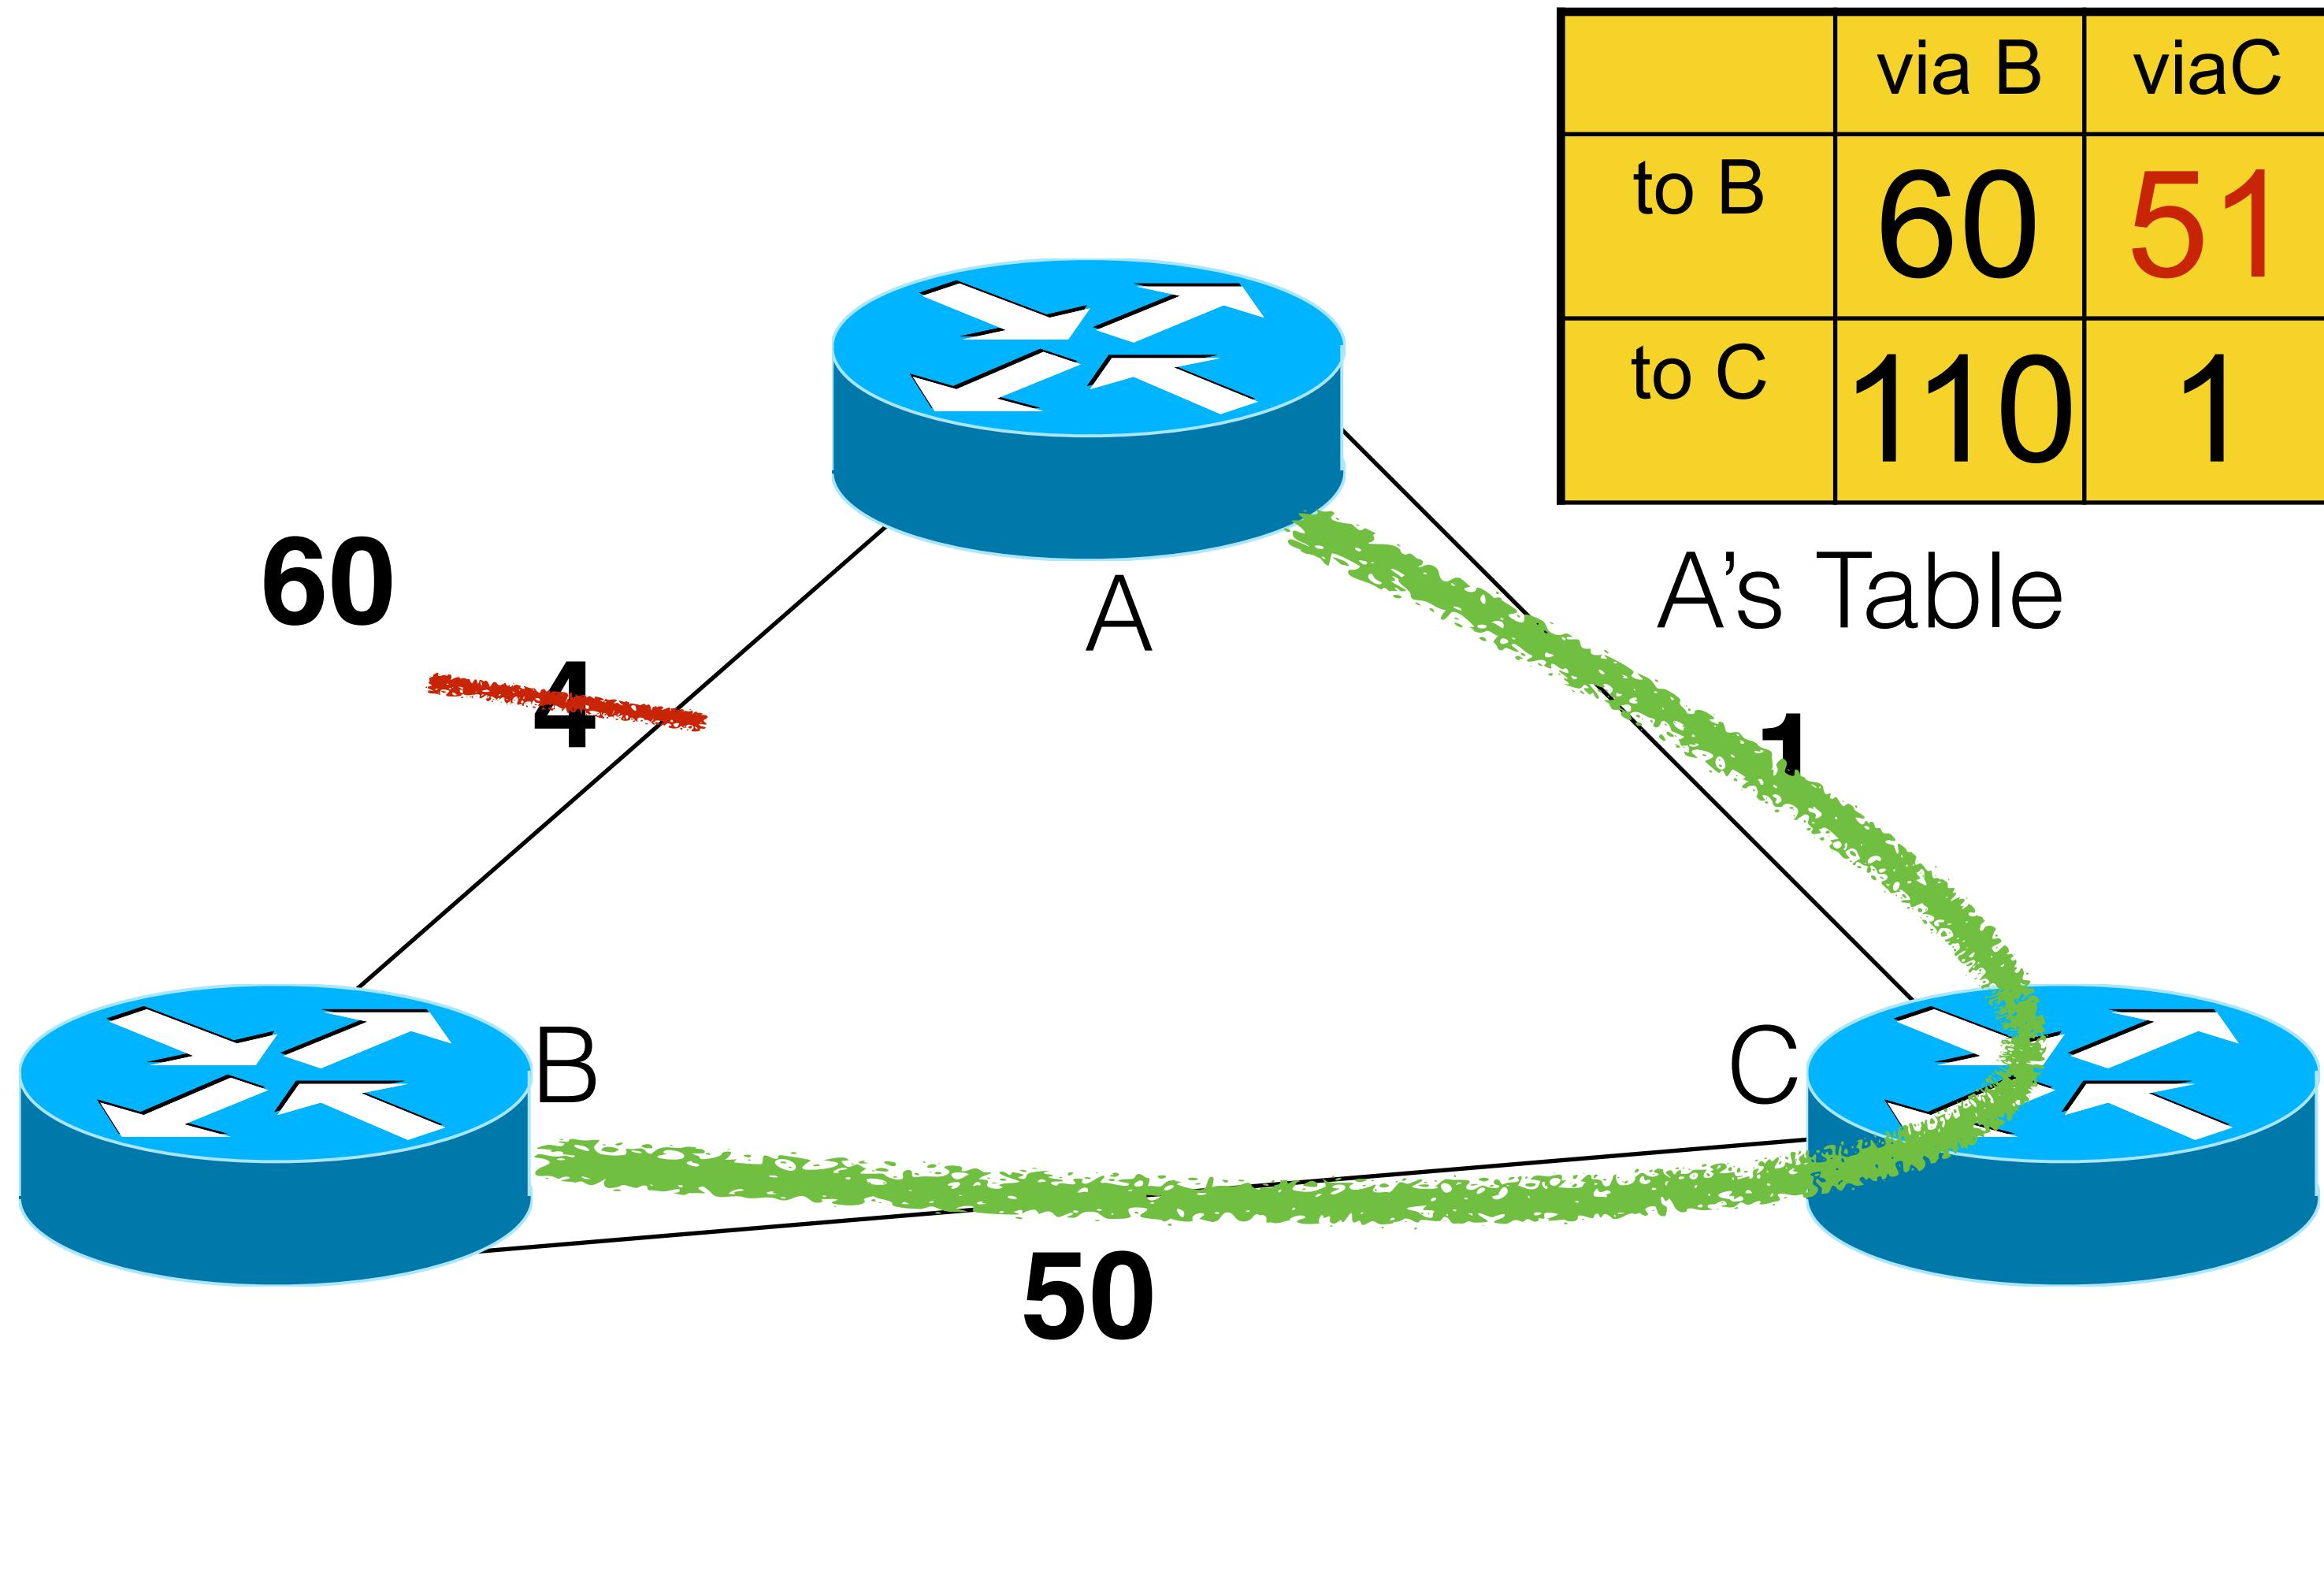
When will this slow-counting
finally end?



Count until we equal the “real” shortest path.

	via A	via C
to A	60	51
to C	61	50

B's Table



	via B	via C
to B	60	51
to C	110	1

A's Table

	via A	via B
to A	1	101
to B	52	50

C's Table



Three Techniques for Mitigating Count to Infinity

- Split Horizon/Poison Reverse
- Maximum Path Lengths
- Pushdown Timers

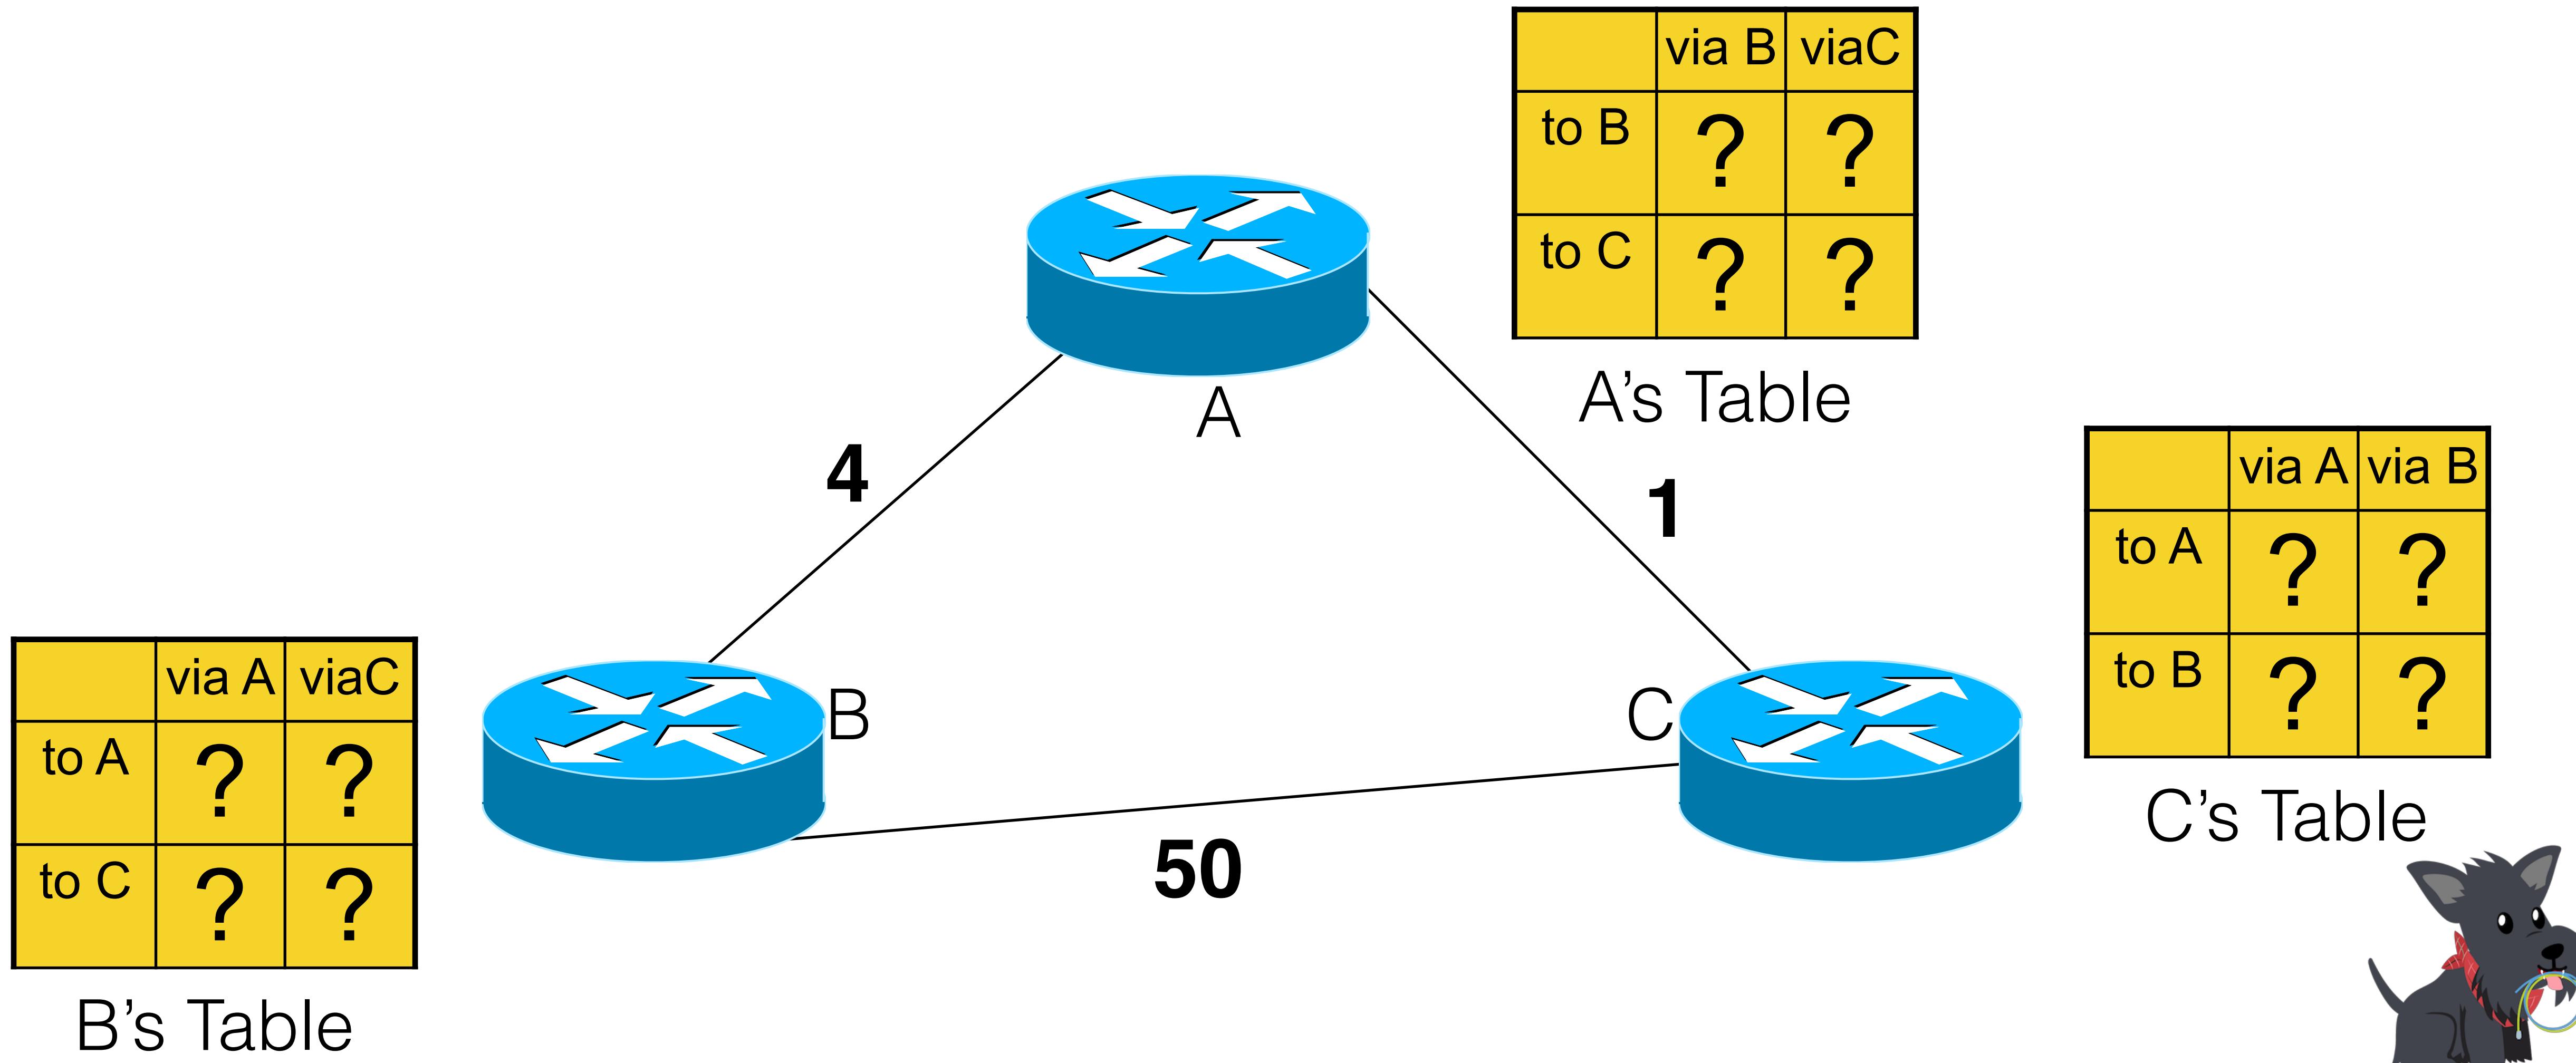


Three Techniques for Mitigating Count to Infinity

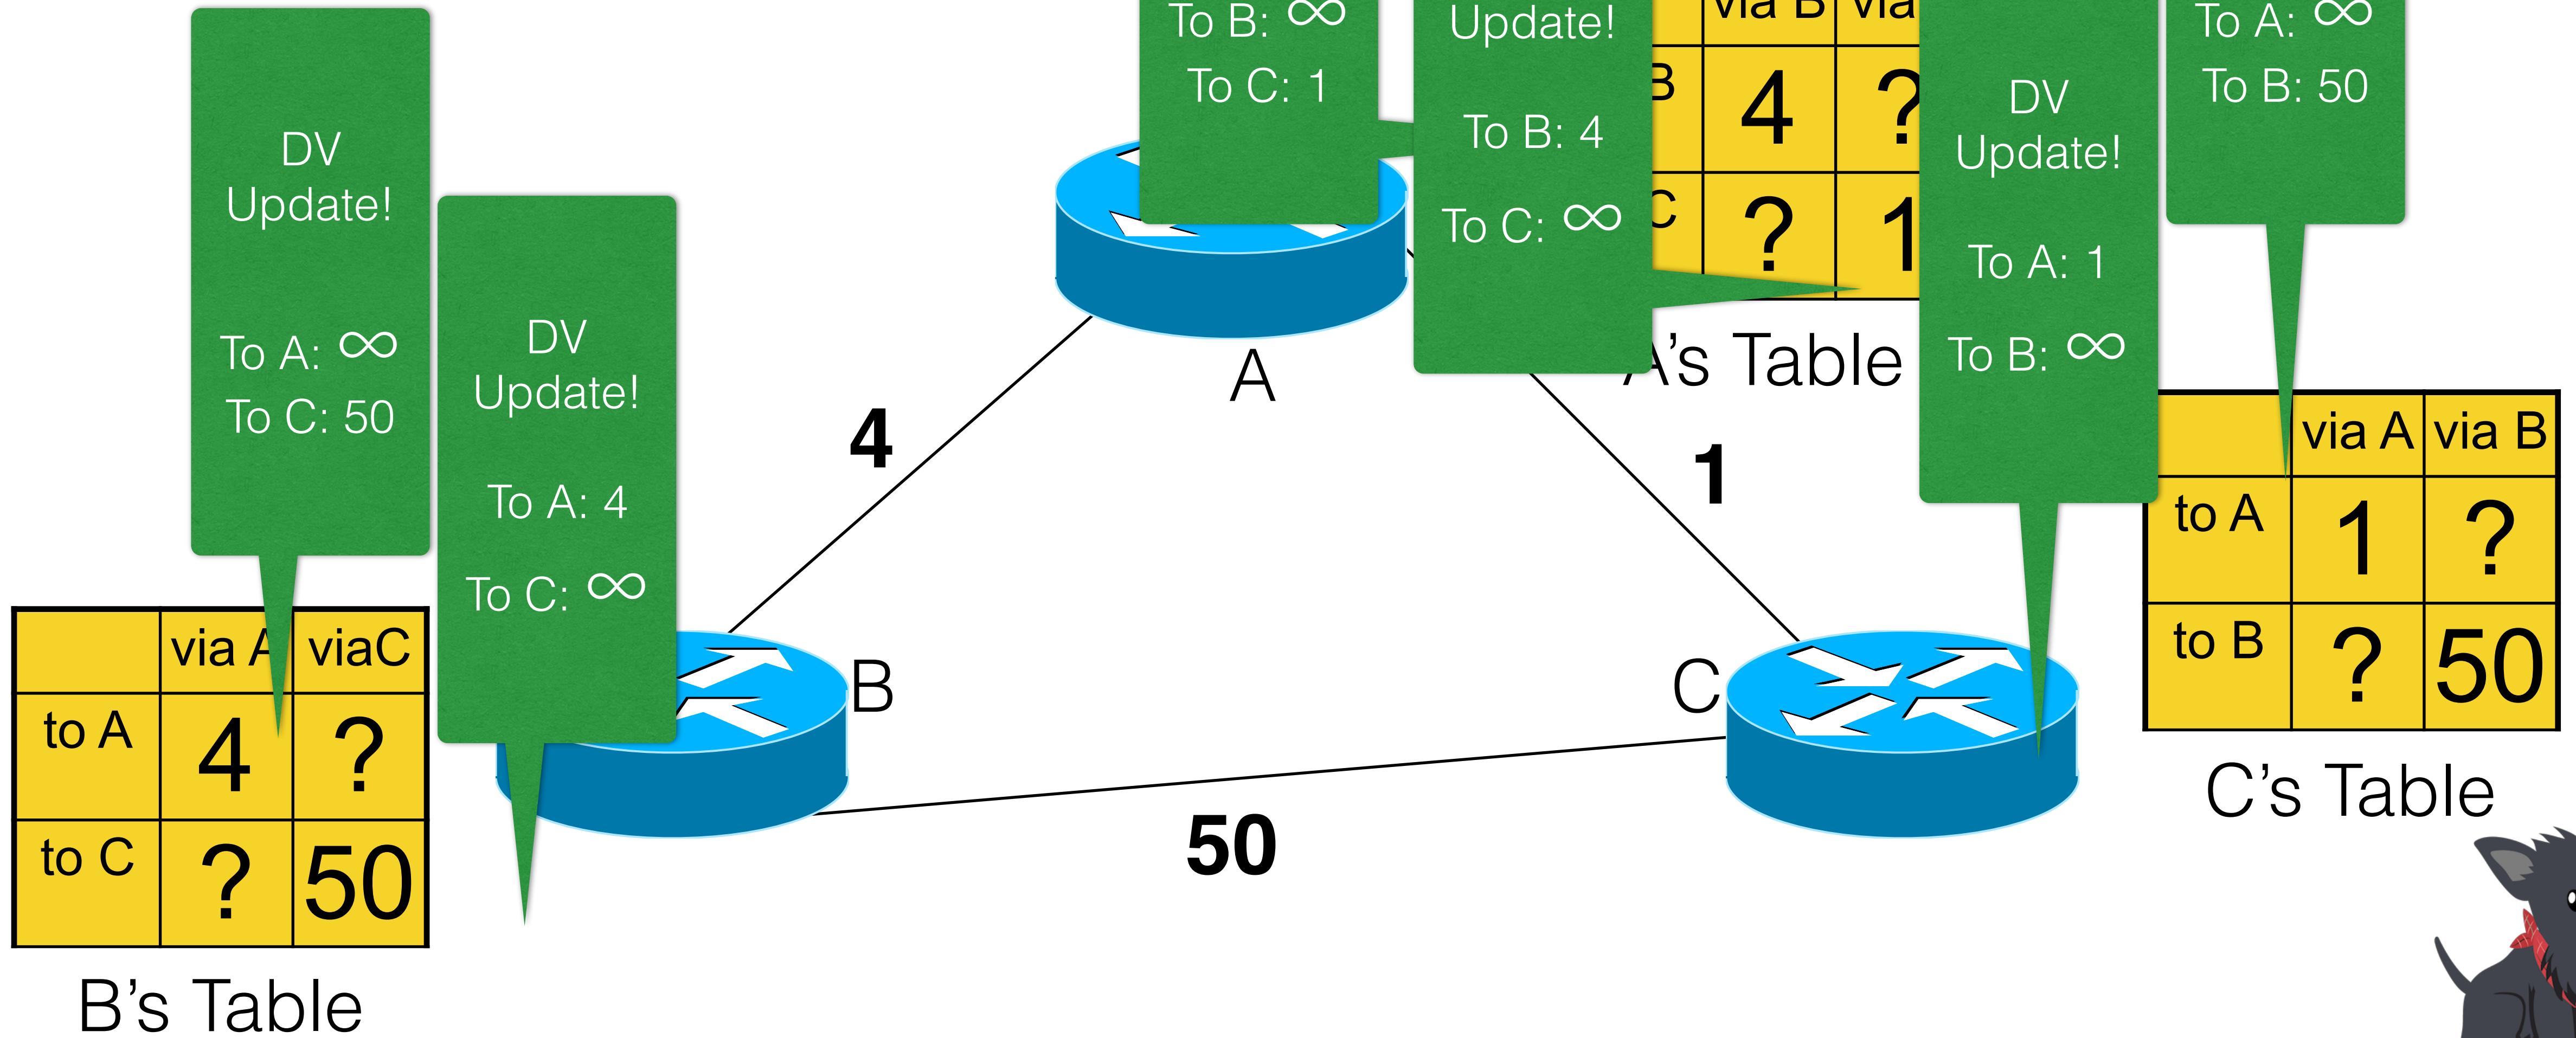
- Split Horizon/Poison Reverse
 - If I select a route I received from you in my distance vector, instead I will report a path length of INFINITY back to you.
- Maximum Path Lengths
- Pushdown Timers



Flashback: What Does That Look Like?



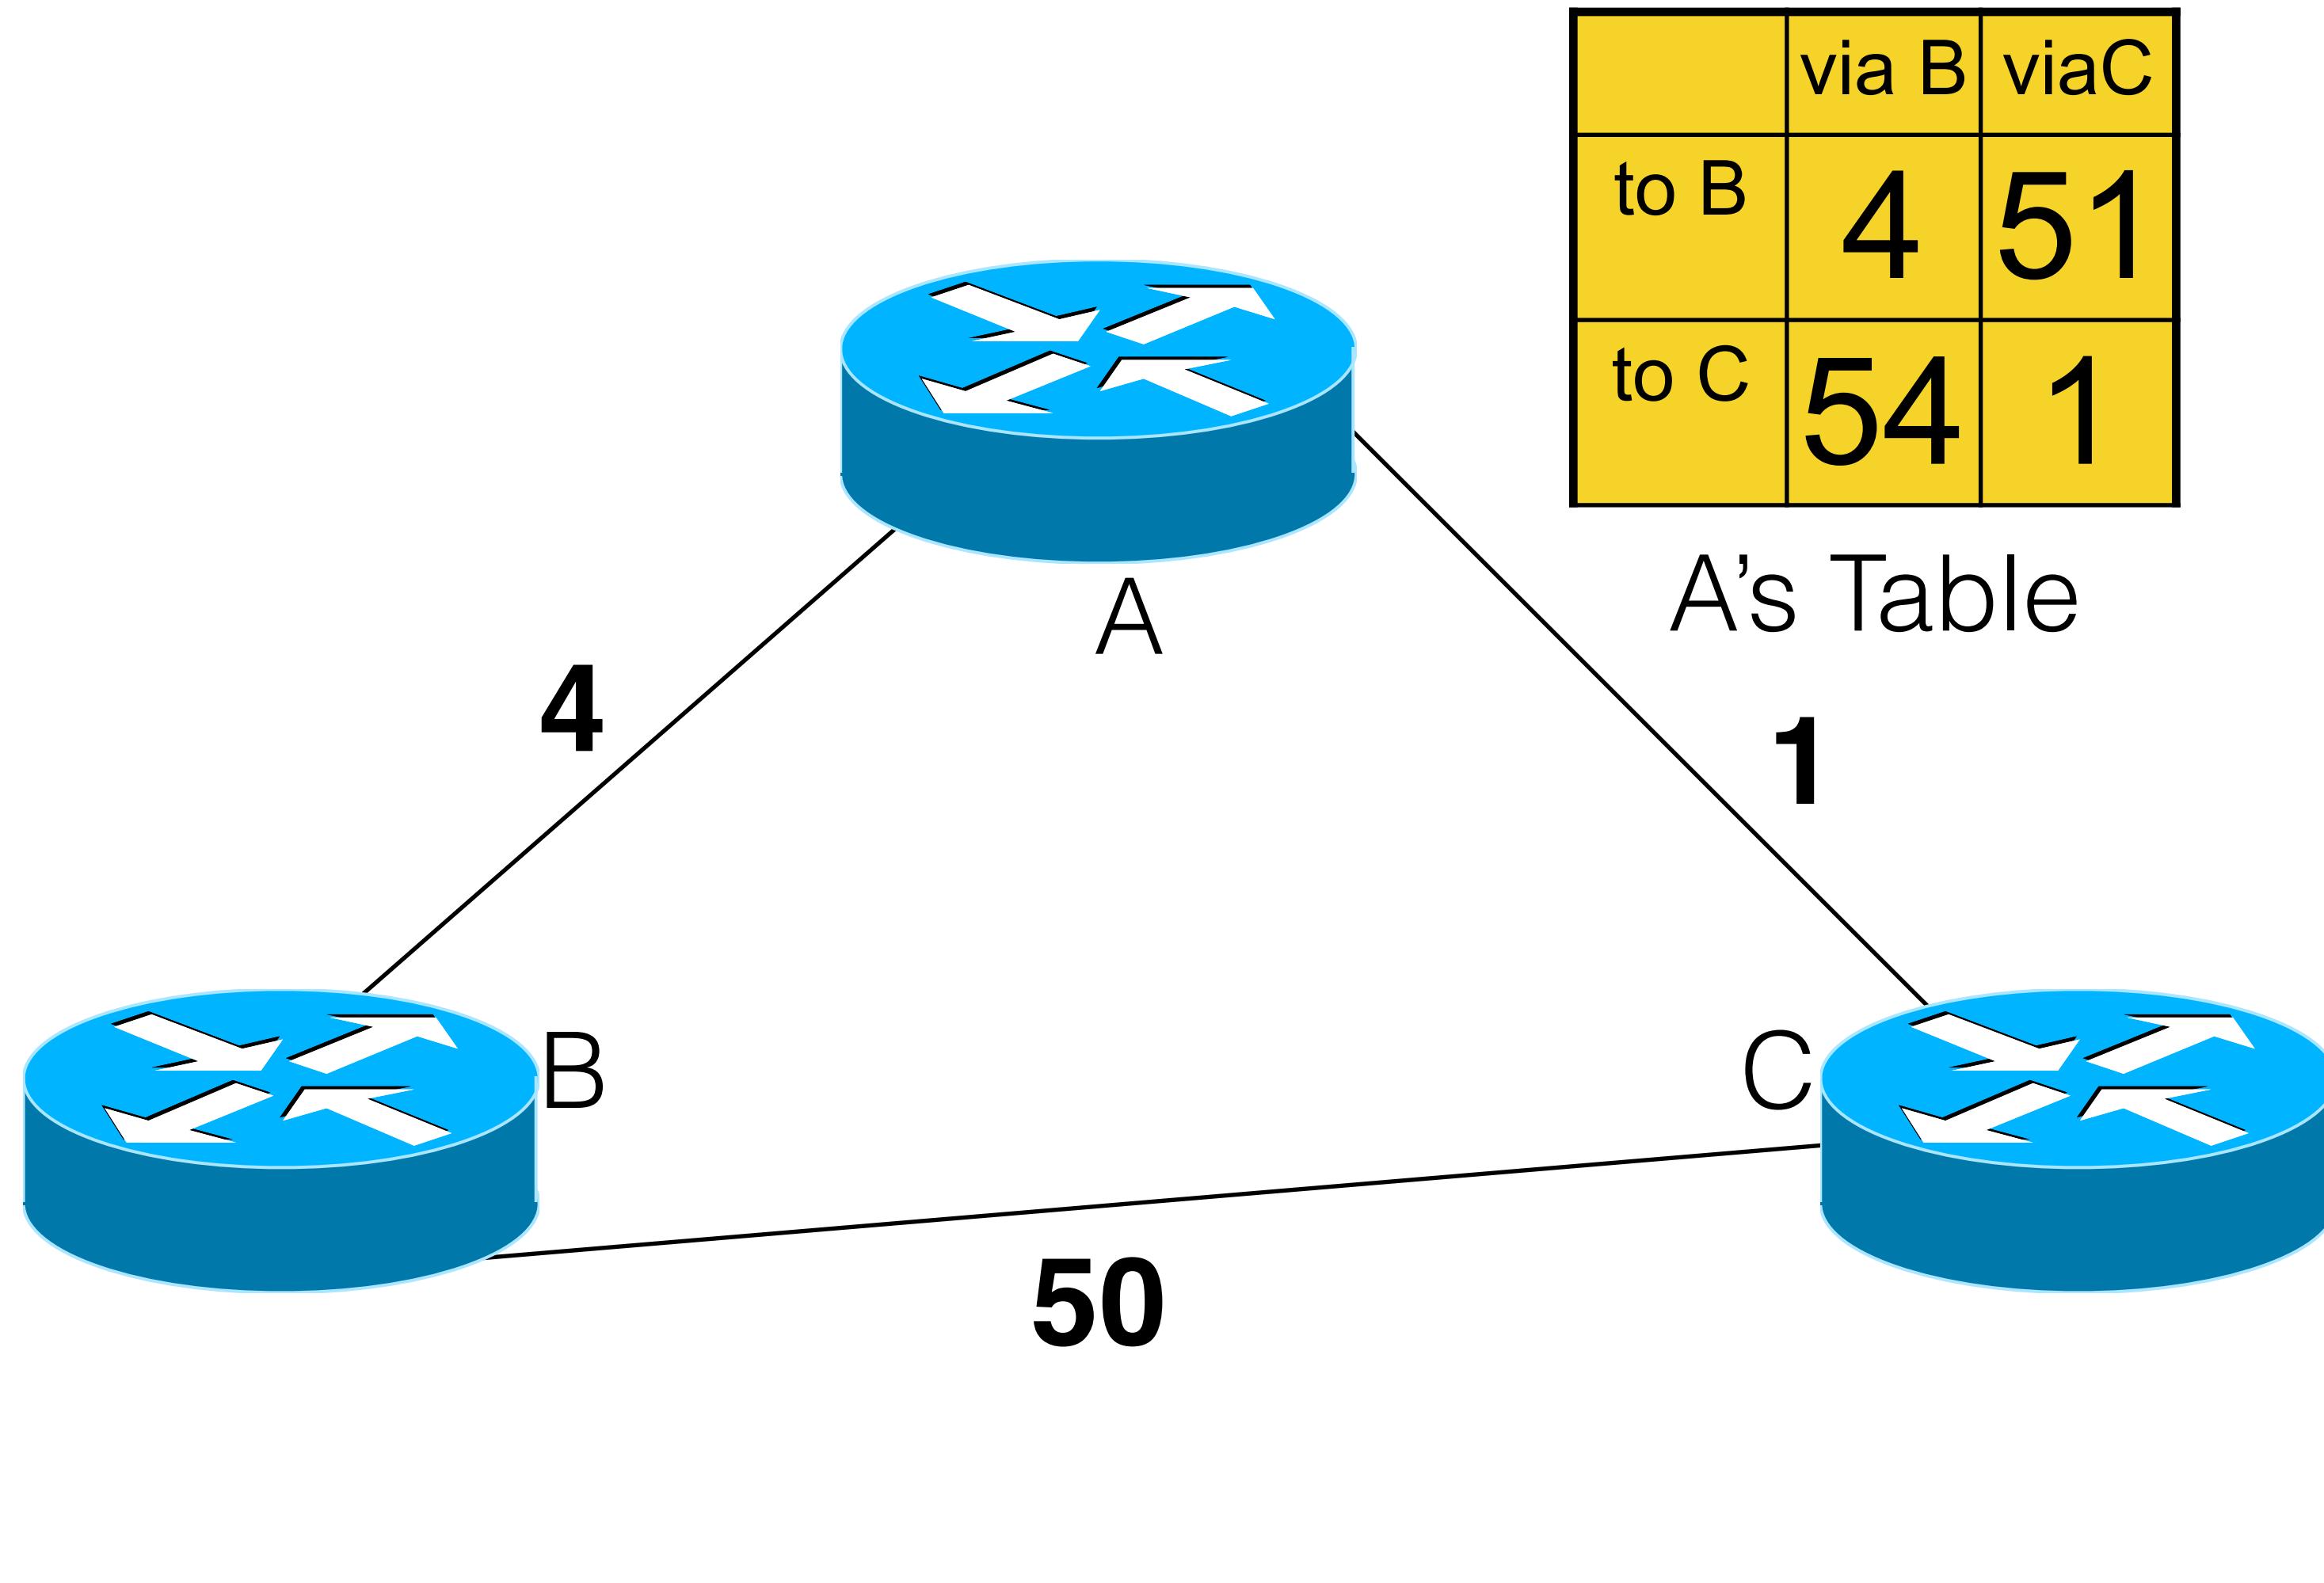
Running trouble...



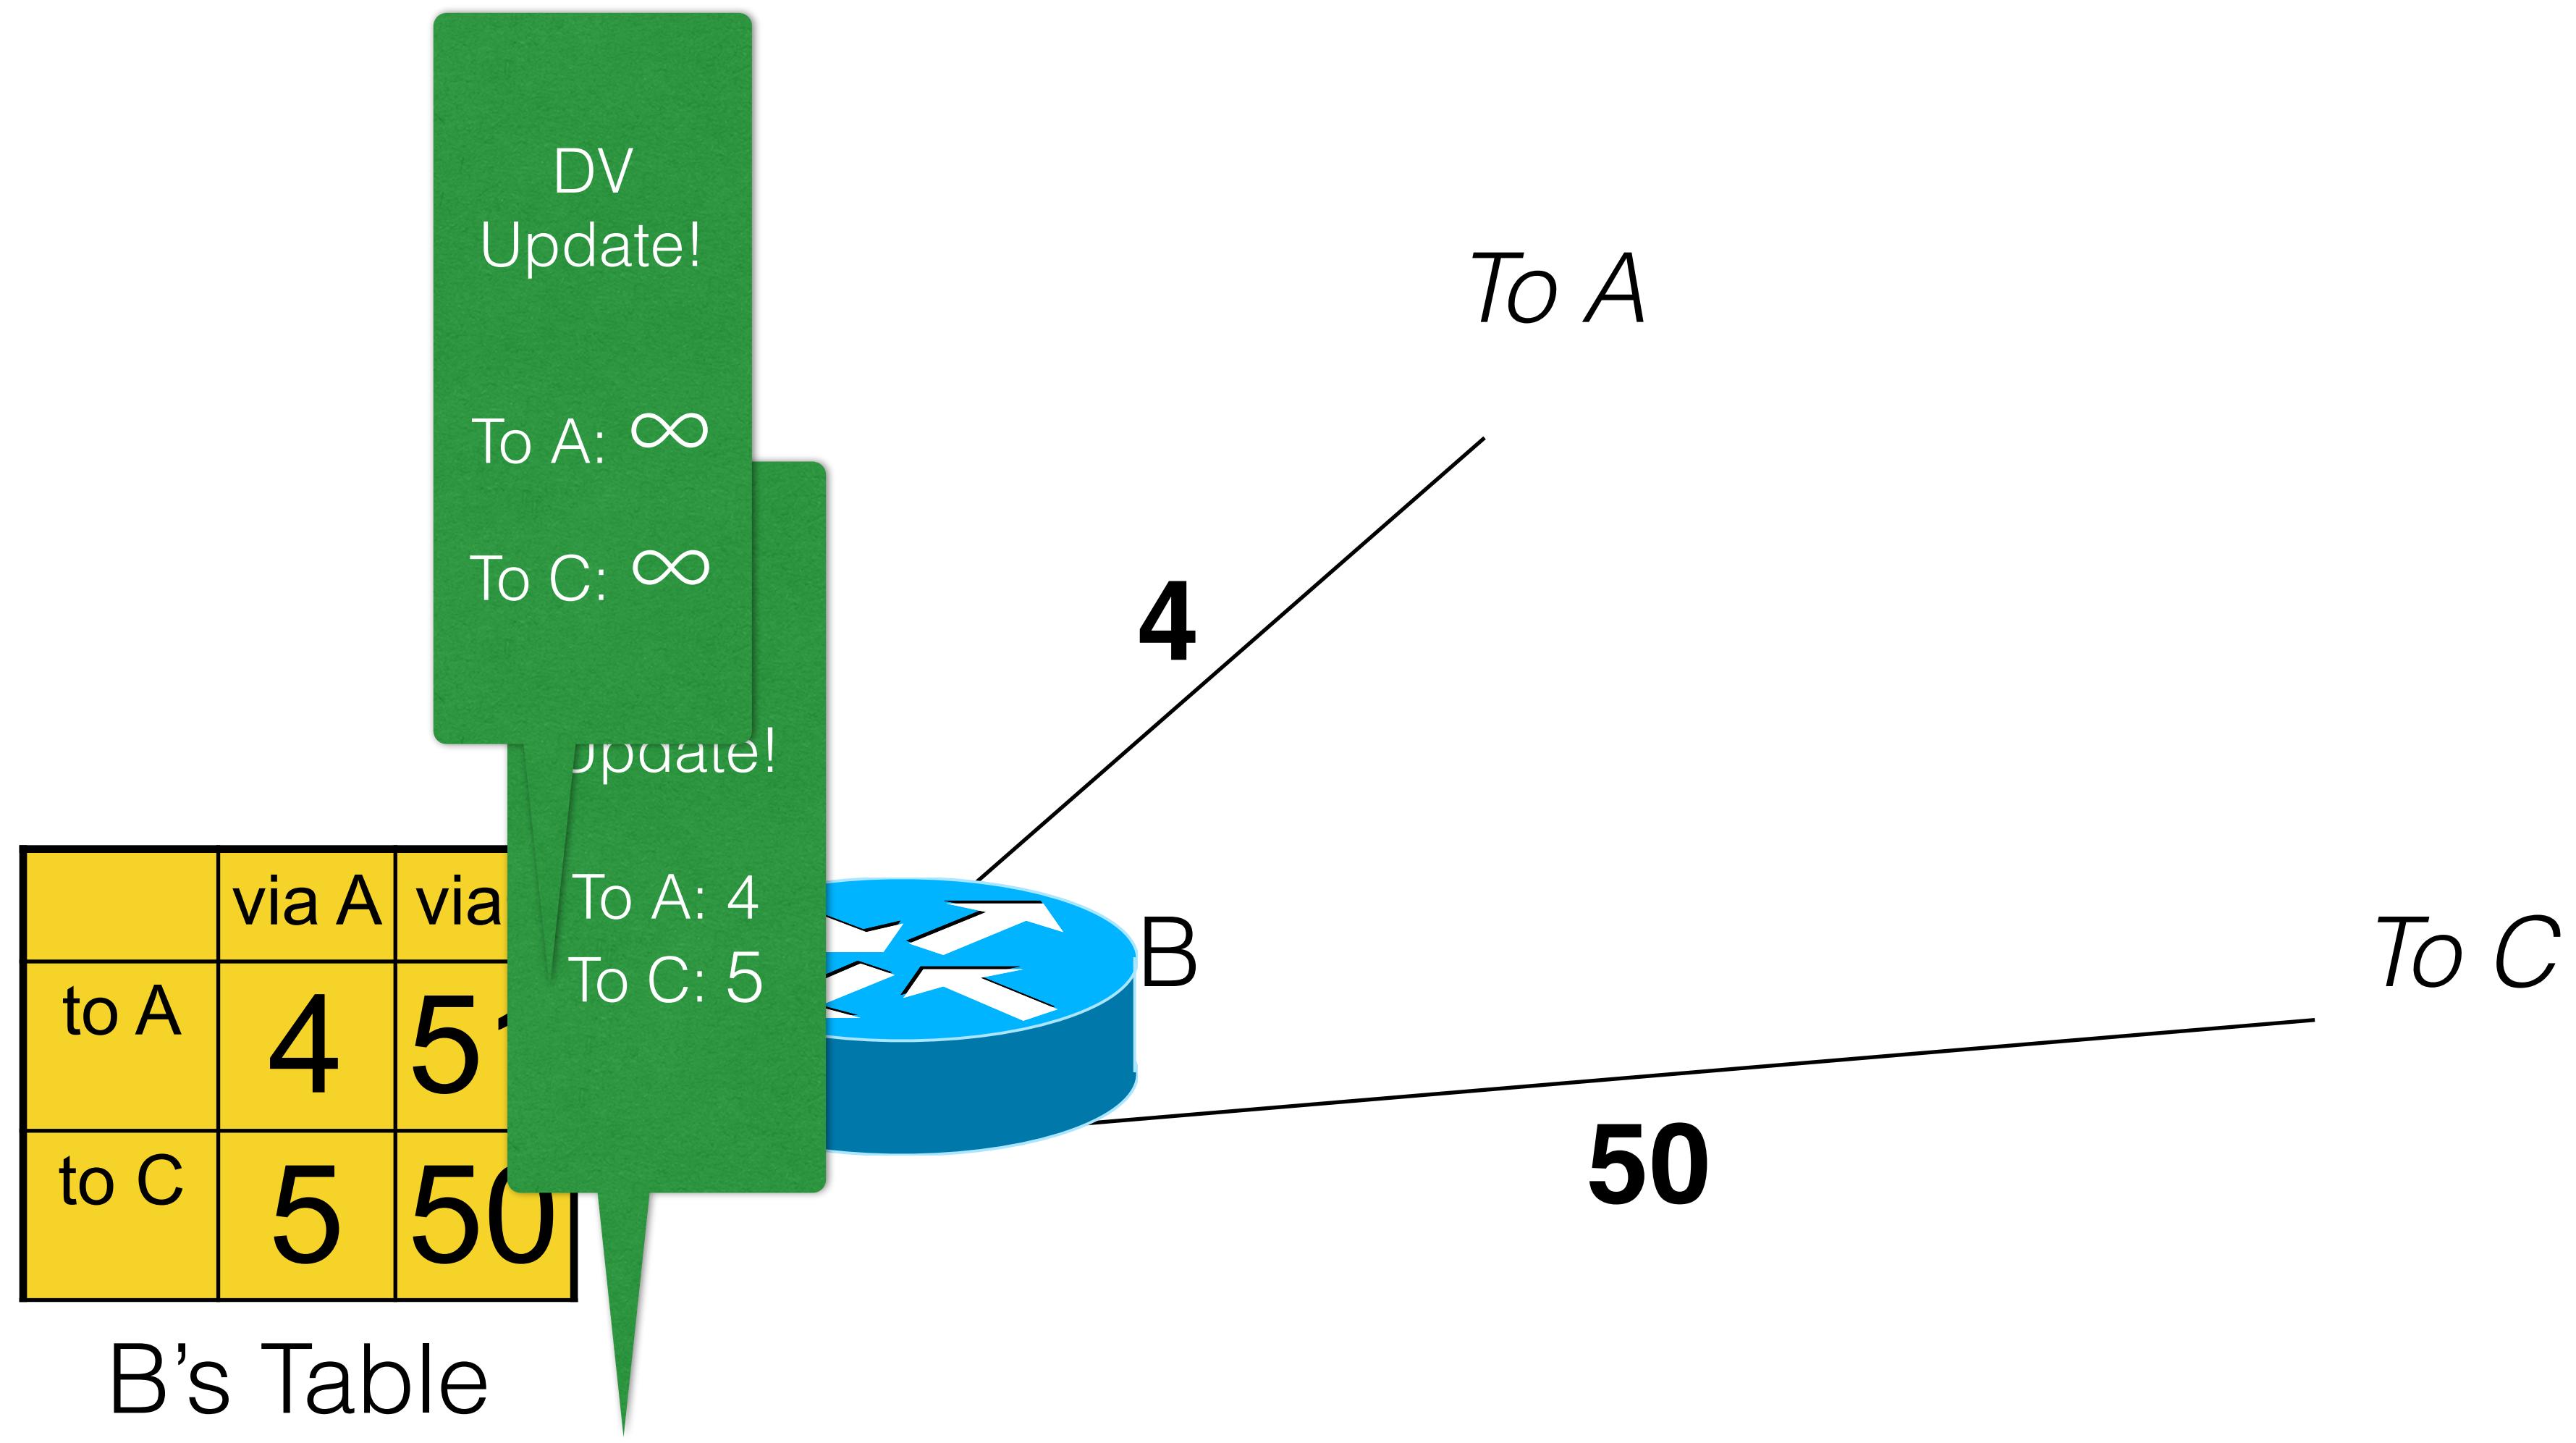
Running into trouble....

	via A	via C
to A	4	51
to C	5	50

B's Table



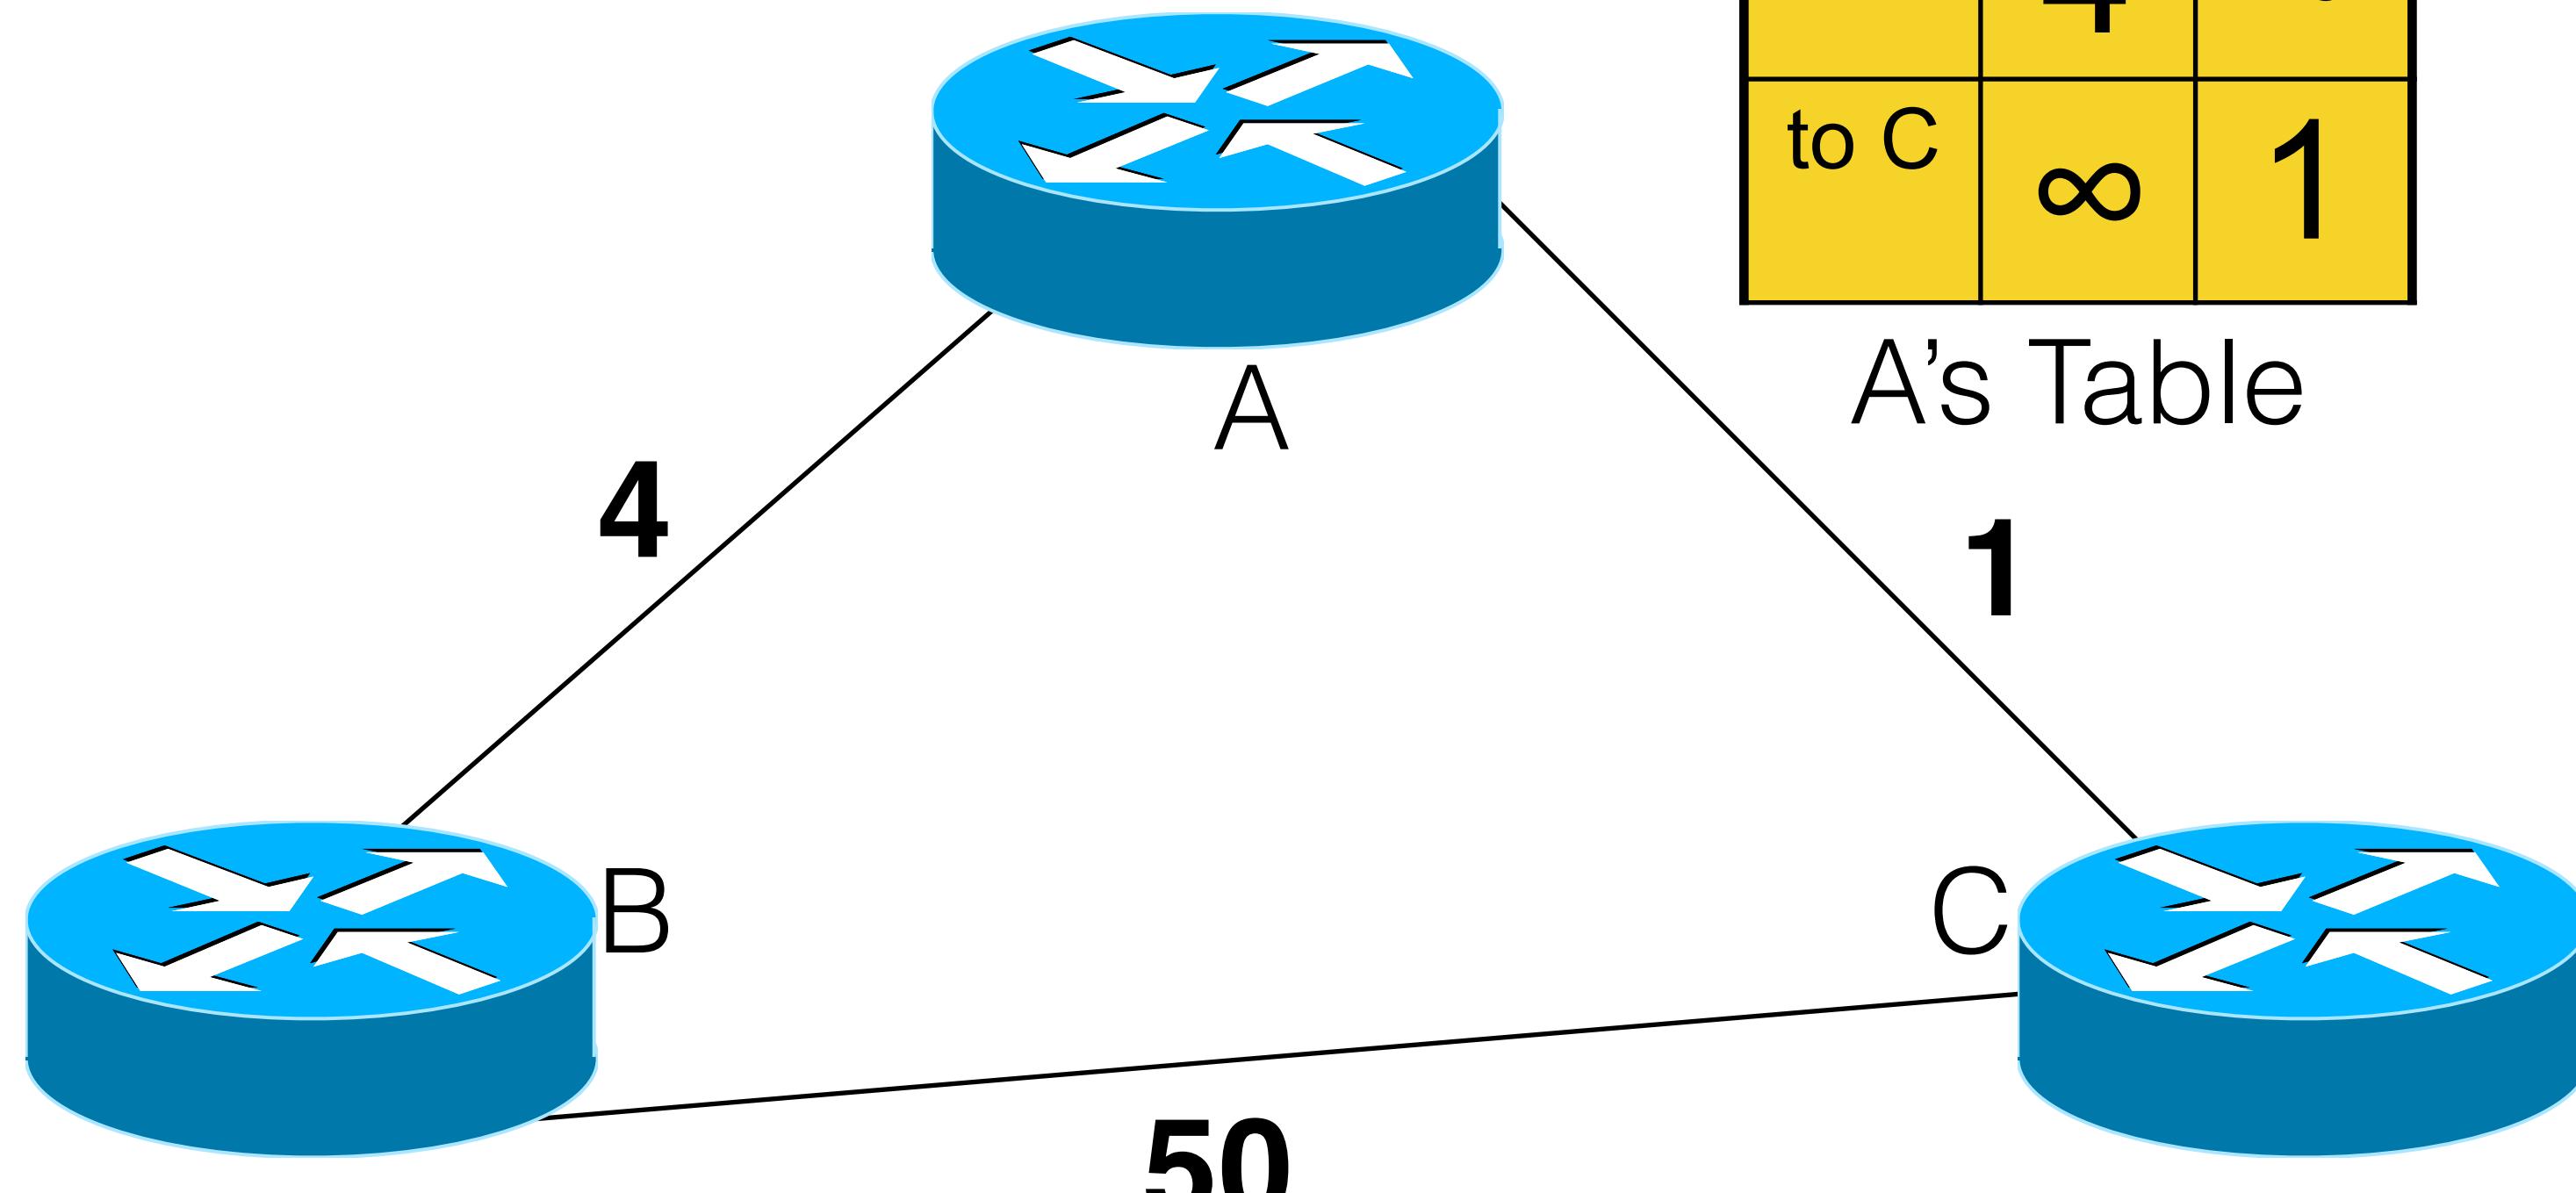
Zoom In



No Bad Loopy Routes!

	via A	via C
to A	4	51
to C	5	50

B's Table



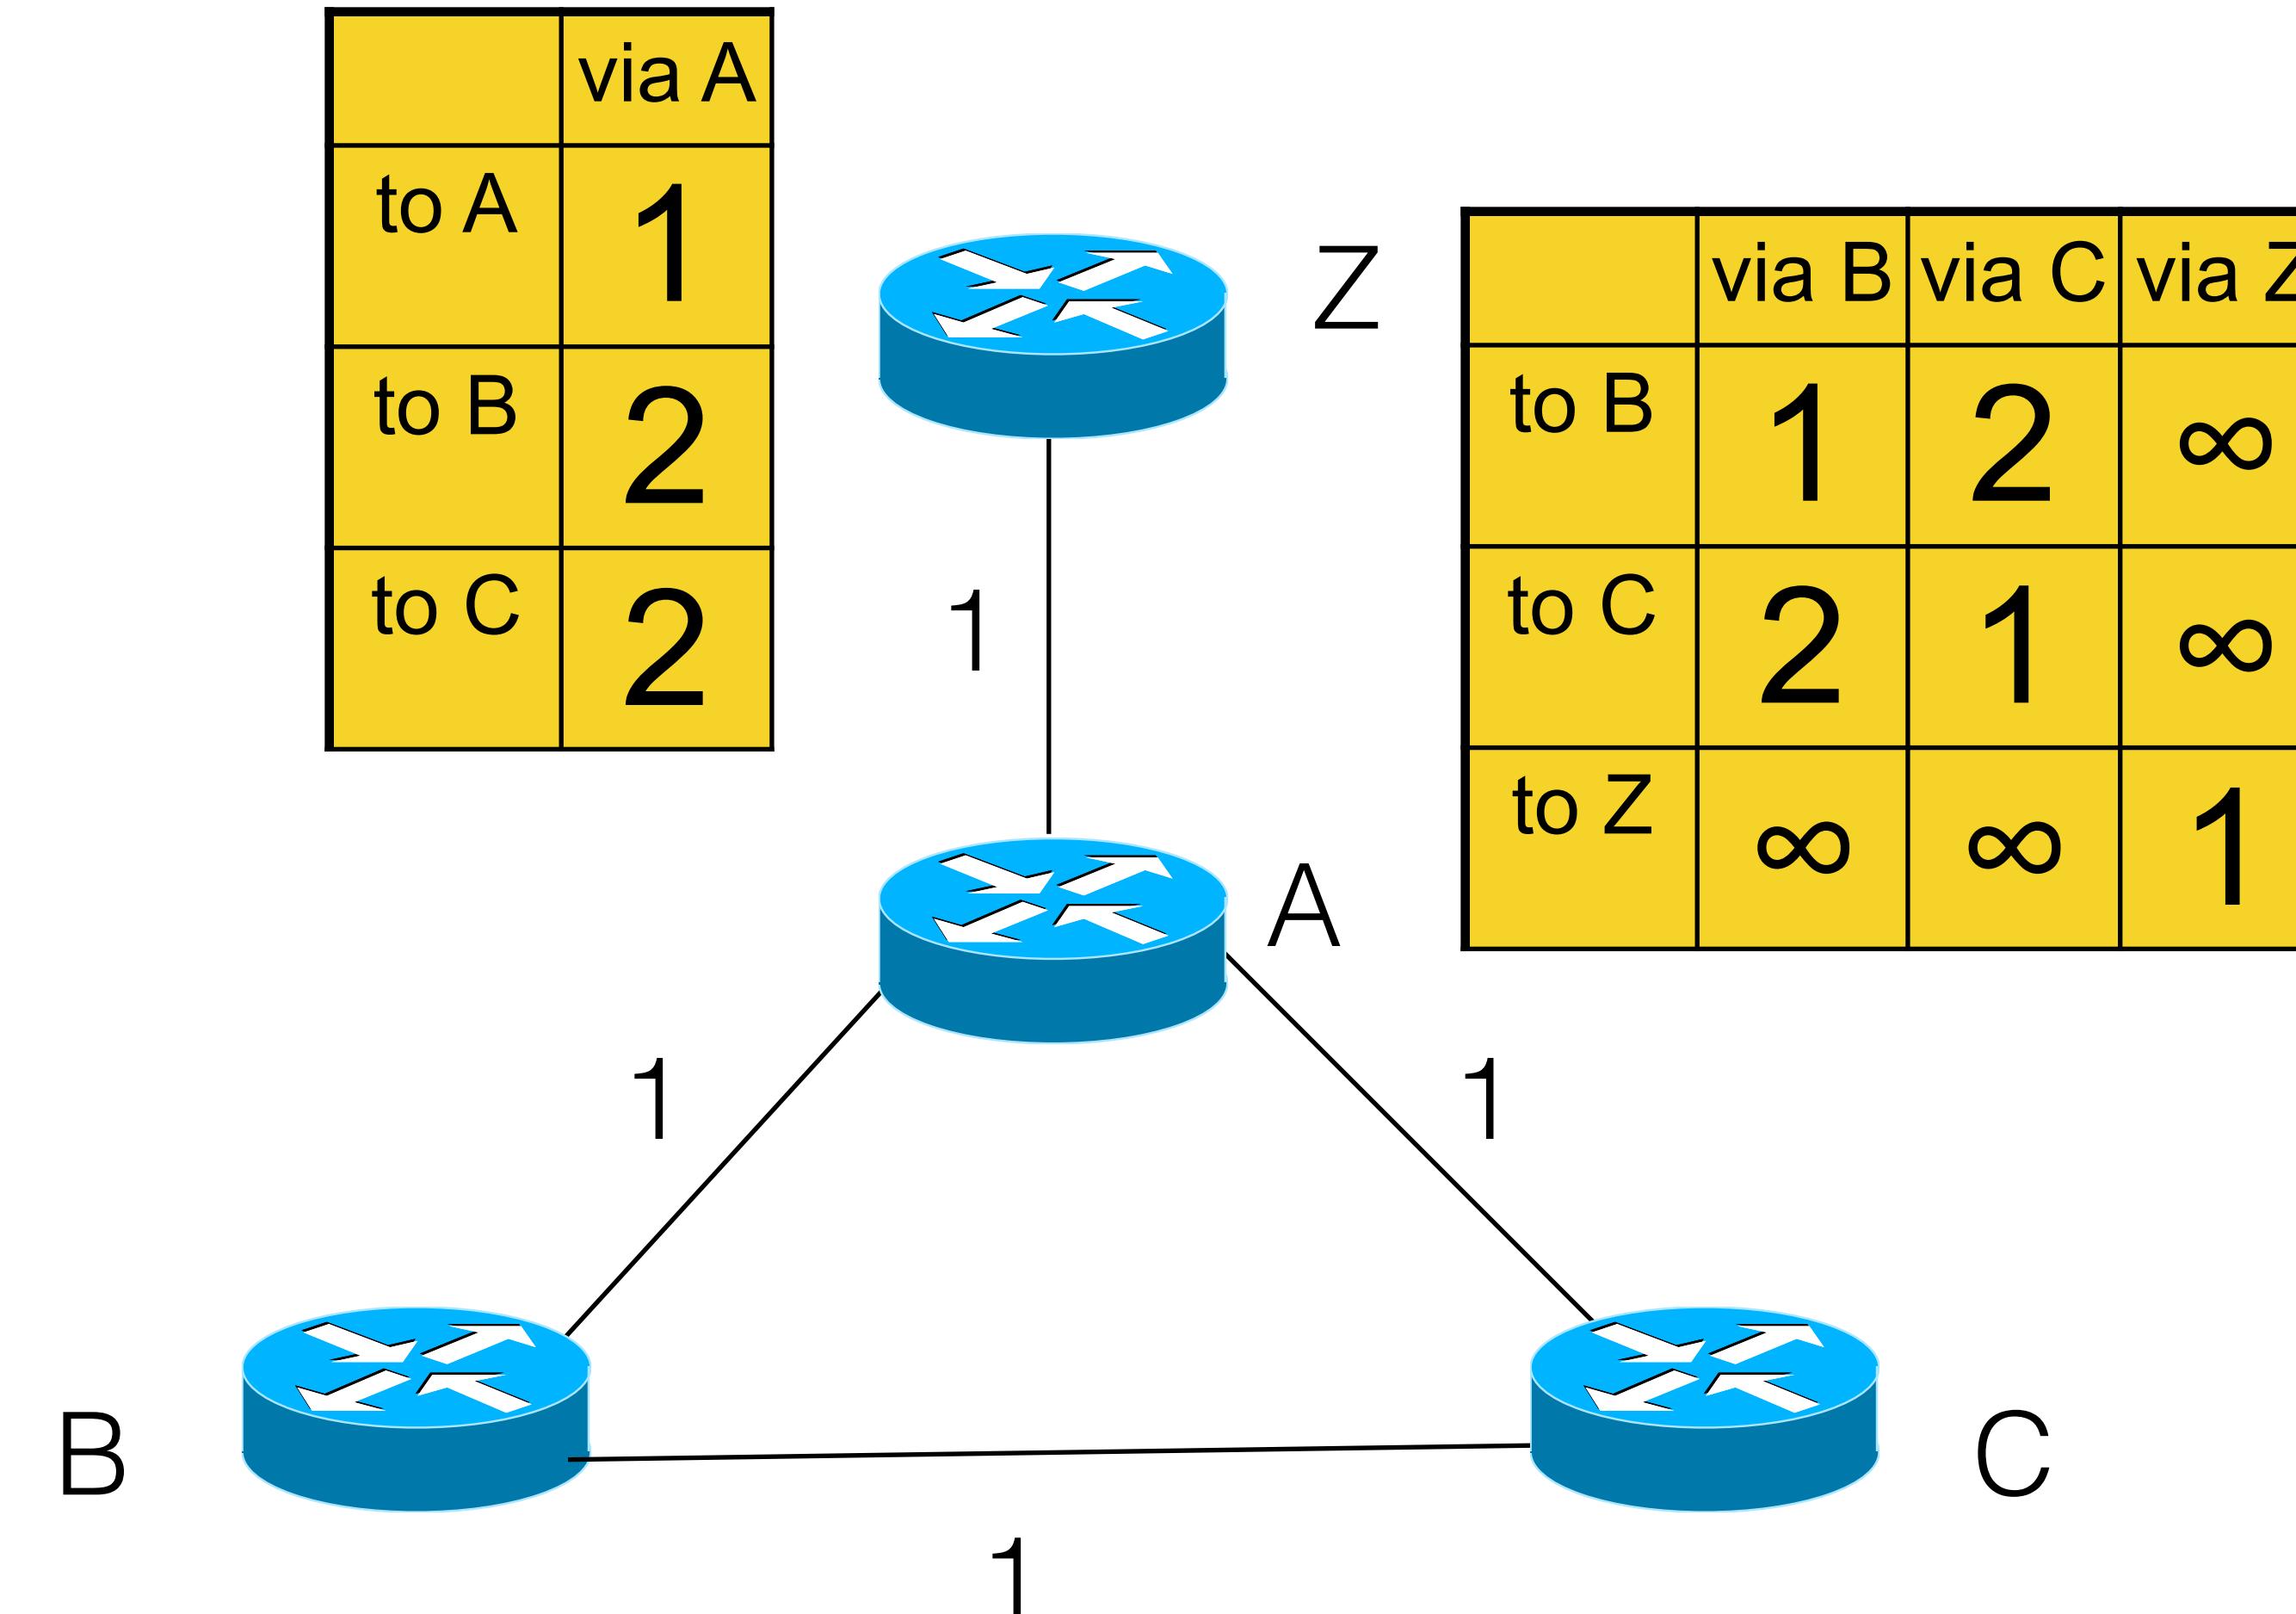
...for that graph.



A completely sad graph

	via A	via C
to A	1	2
to C	2	1
to Z	2	3

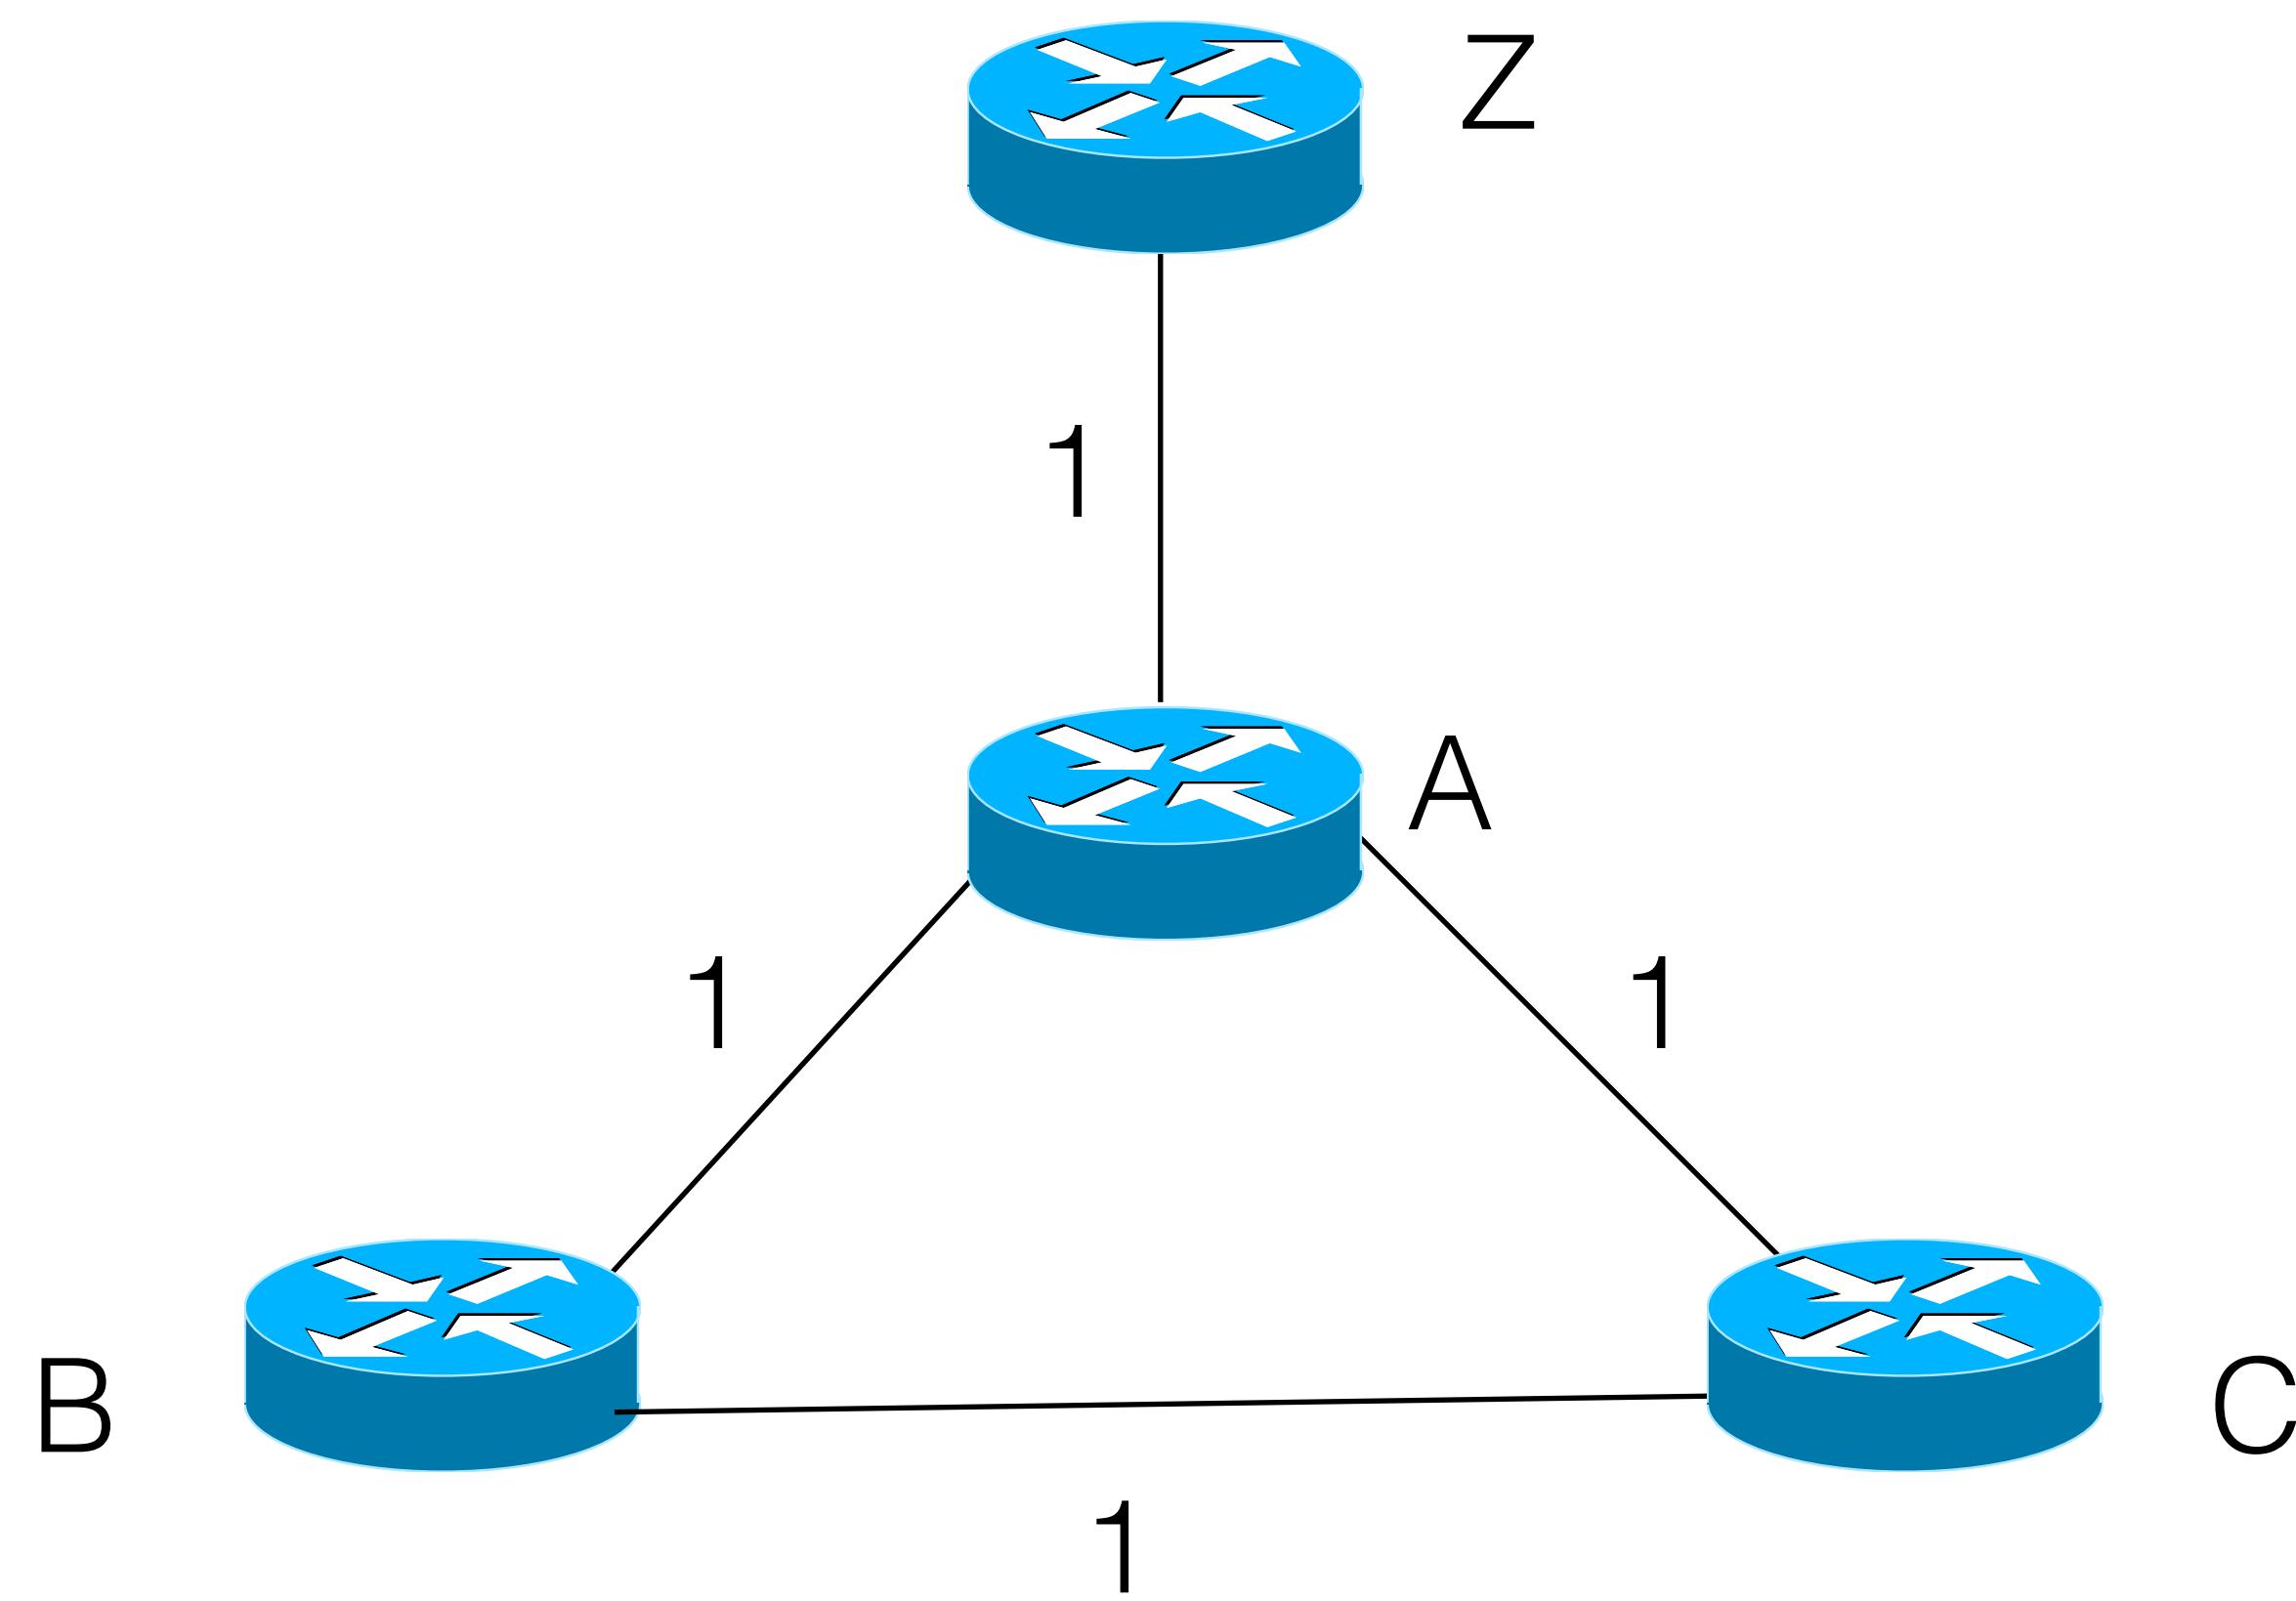
	via A
to A	1
to B	2
to C	2

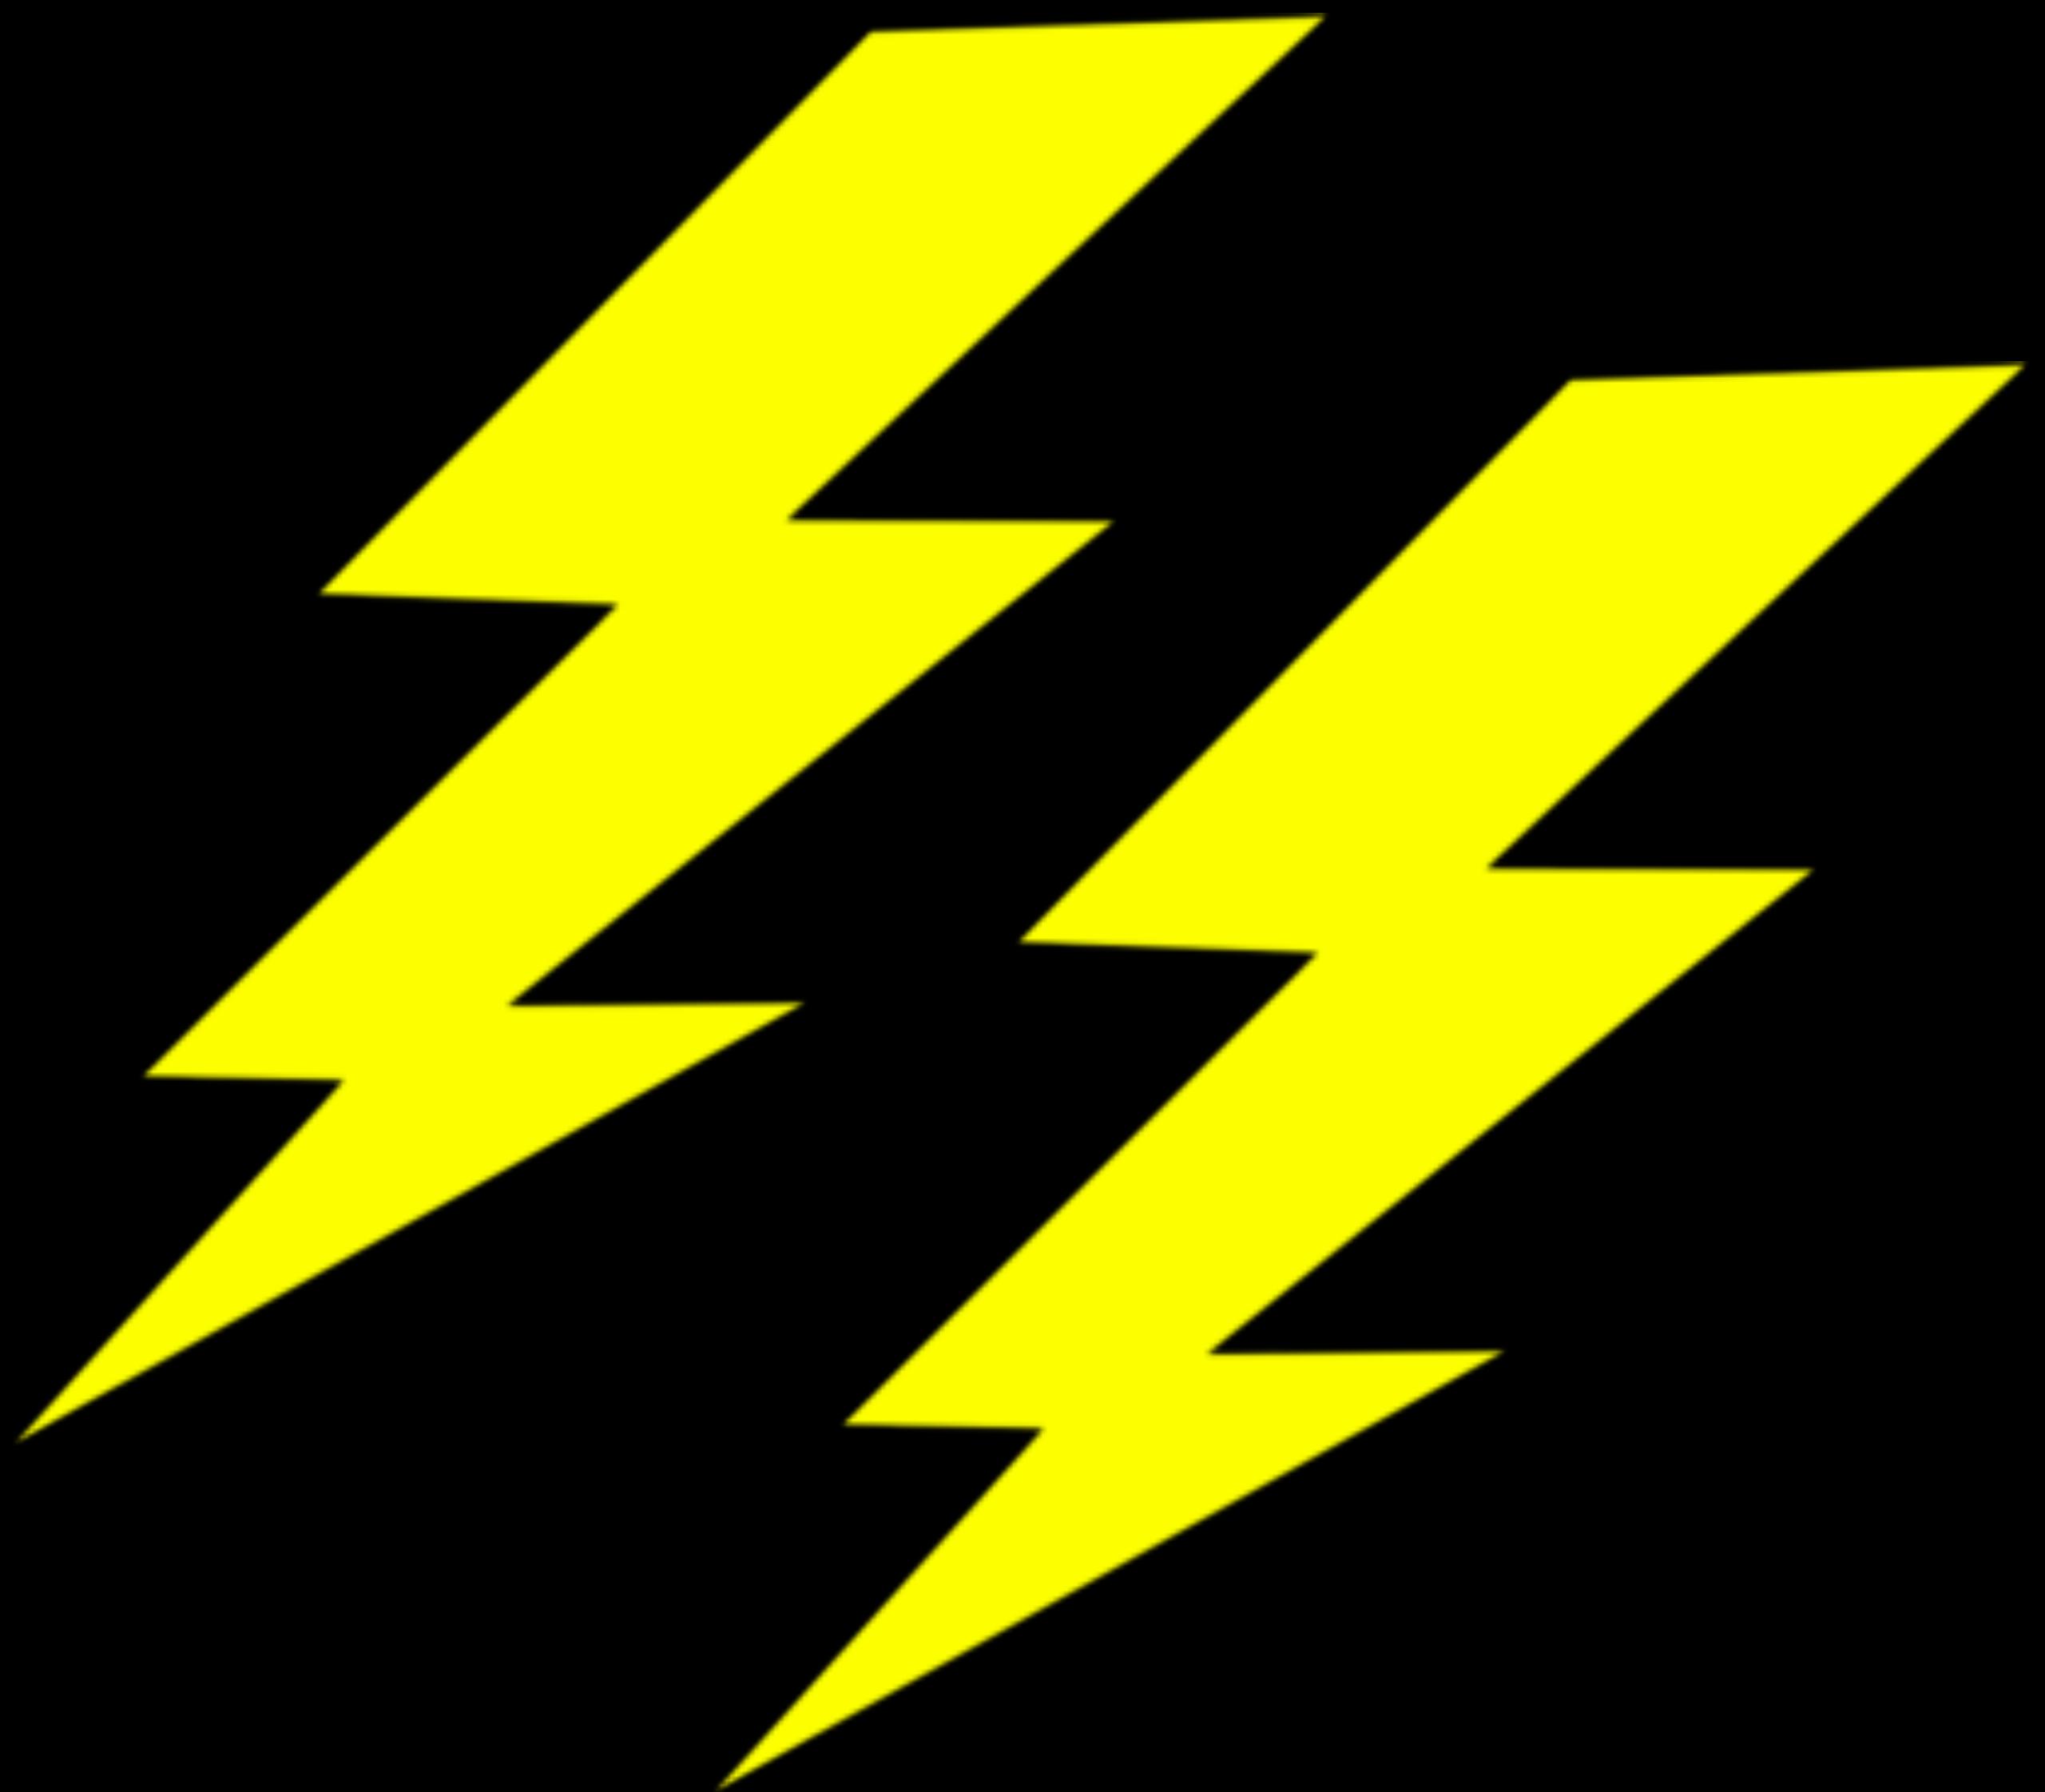


	via A	via B
to A	1	2
to B	2	1
to Z	2	3



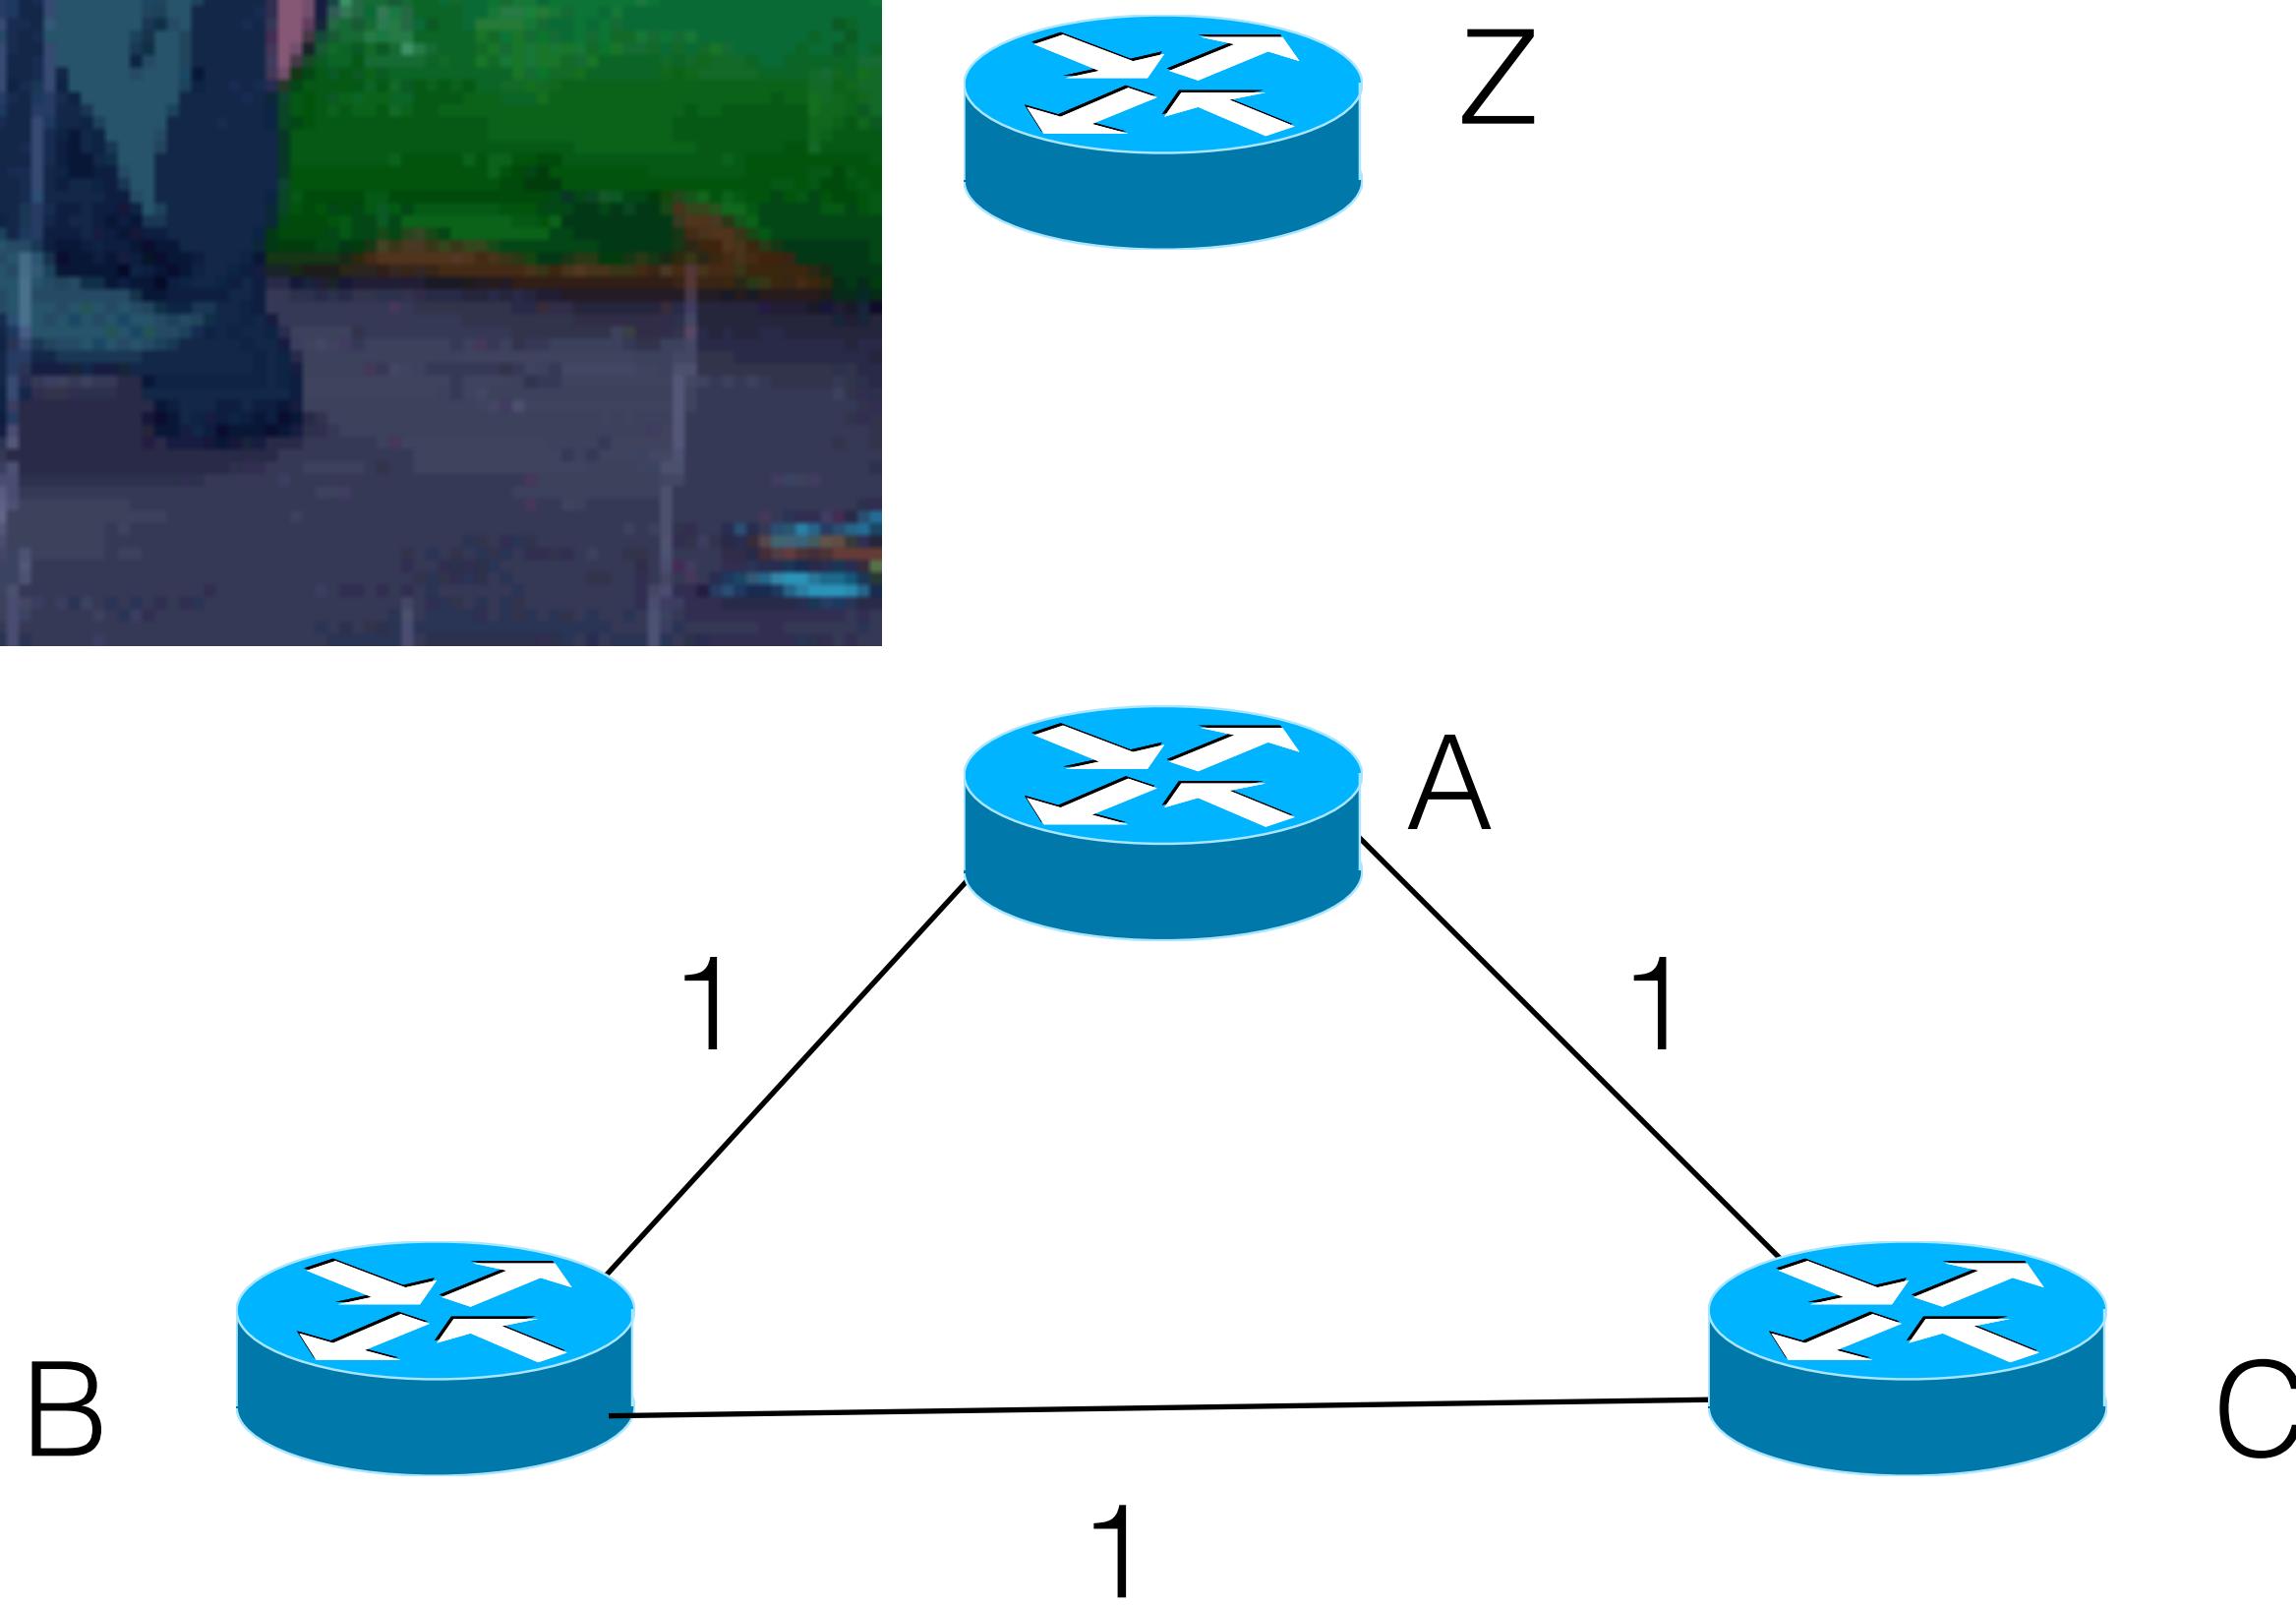
A completely sad graph





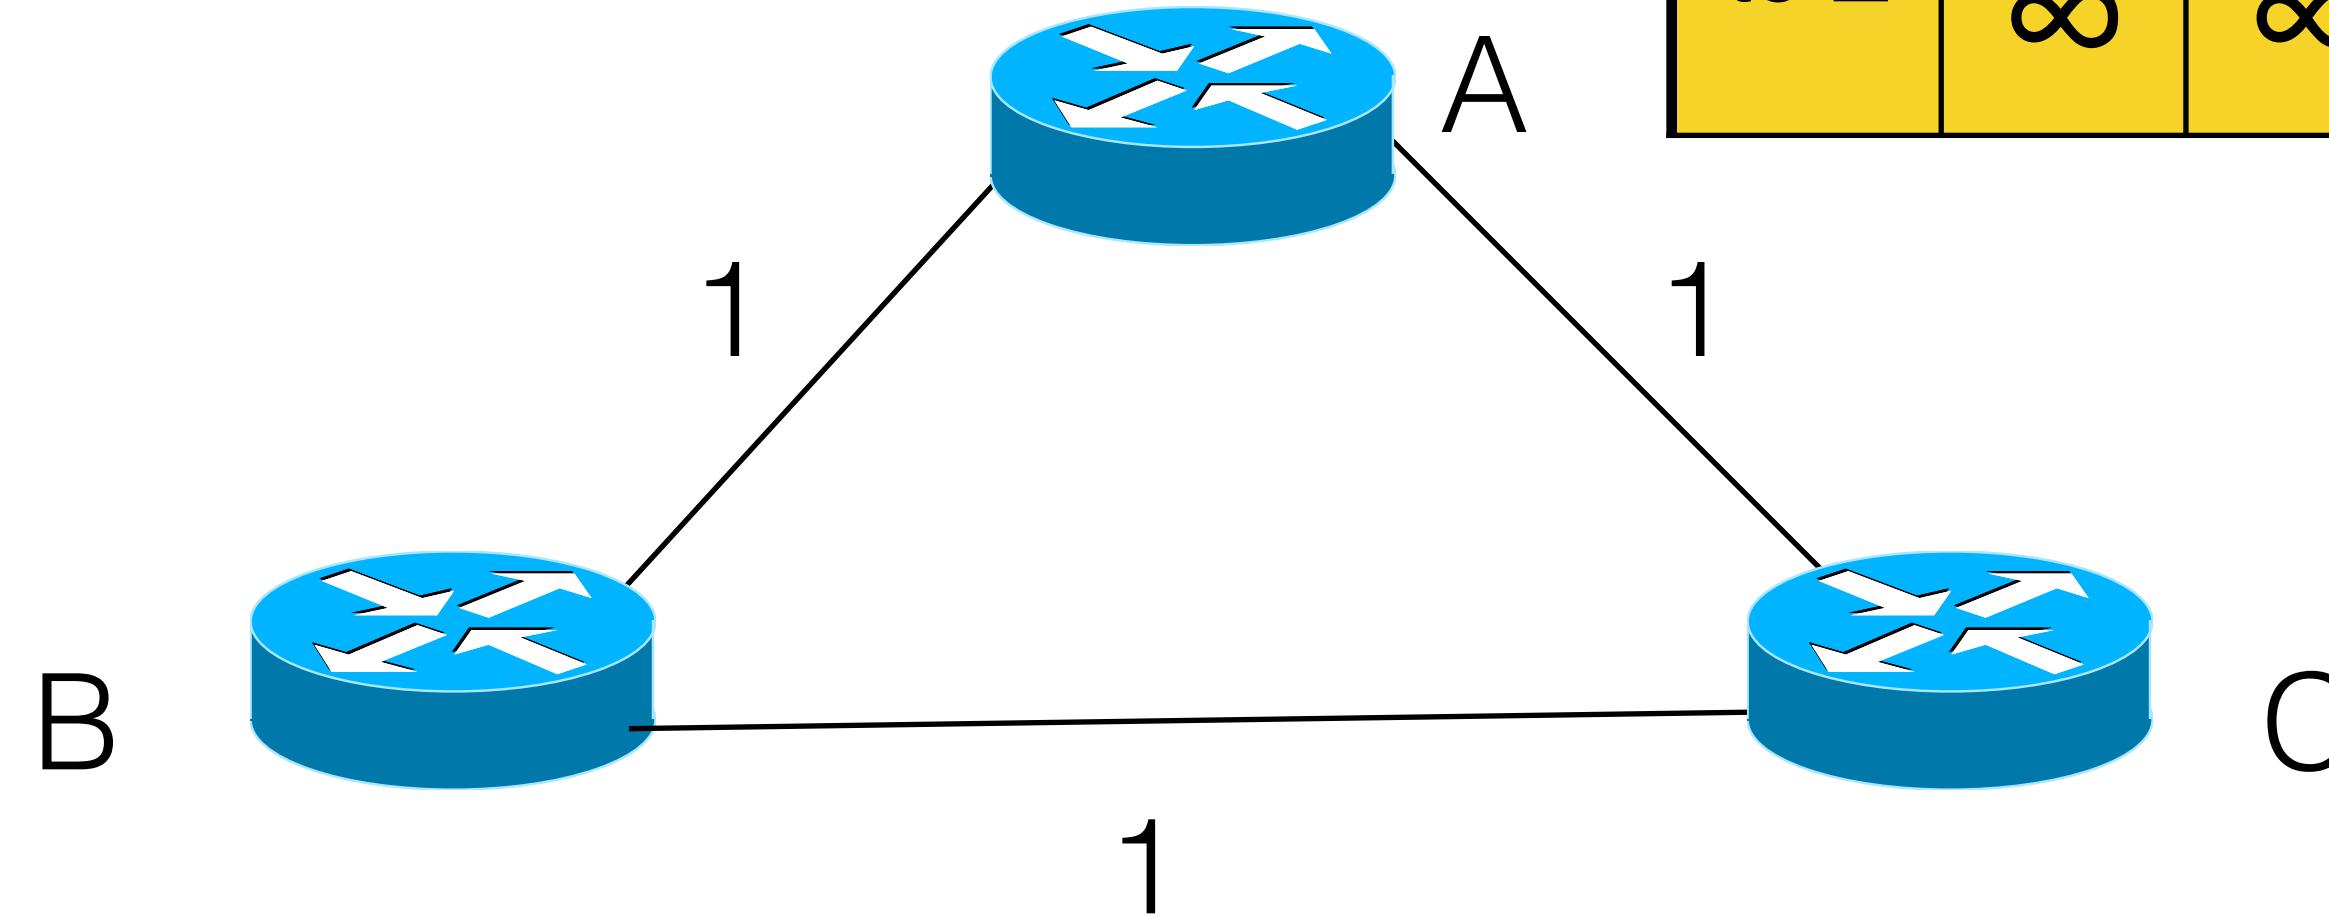


completely sad graph



A completely sad graph

	via A	via C
to A	1	2
to C	2	1
to Z	2	3



	via B	via C	via Z
to B	1	2	∞
to C	2	1	∞
to Z	∞	∞	∞

	via A	via B
to A	1	2
to B	2	1
to Z	2	3



A completely sad graph

	via A	via C
to A	1	2
to C	2	1
to Z	2	3

DV
Update!

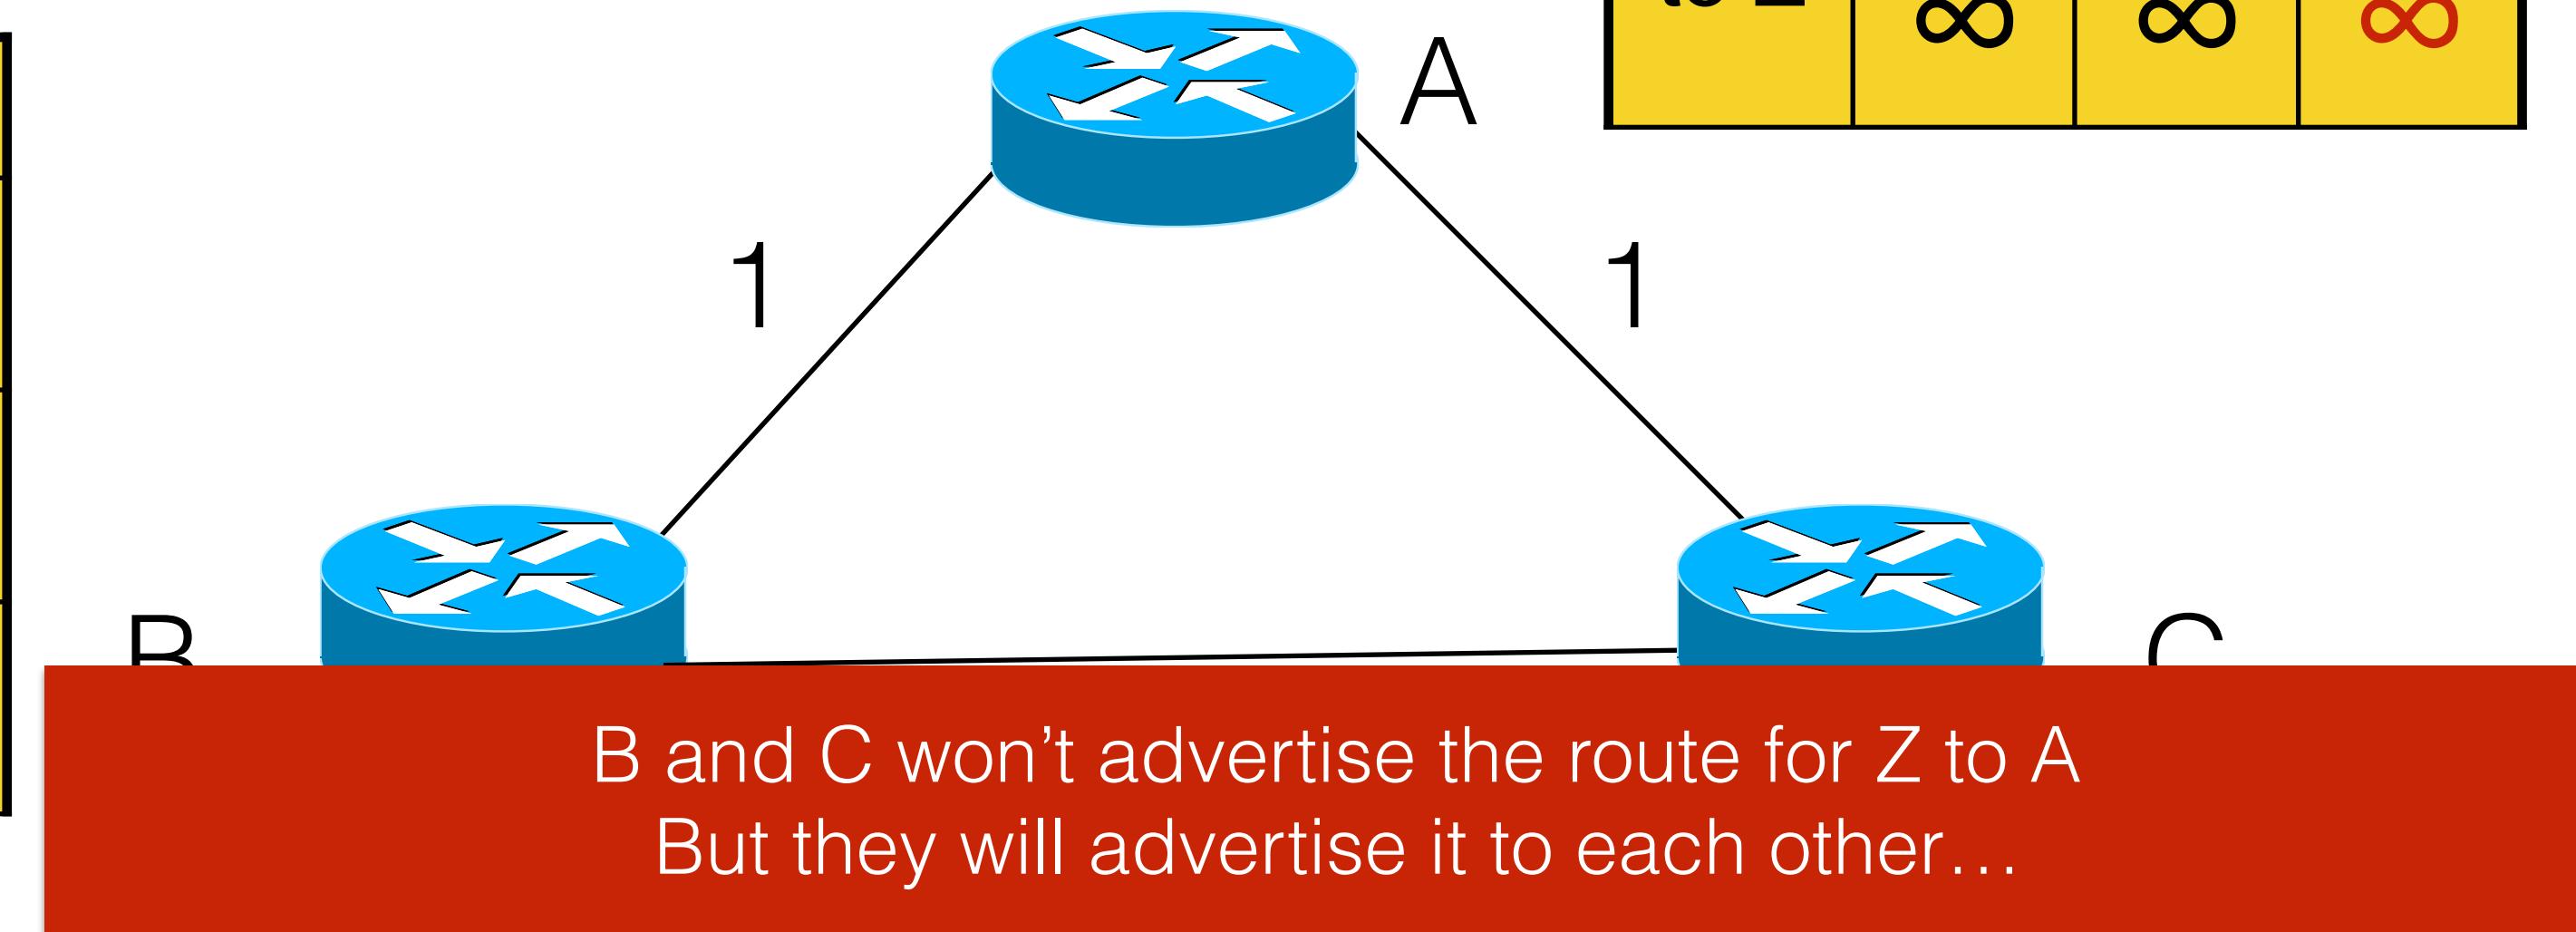
To B: 1
To C: 1
To Z: ∞

	via B	via C	via Z
to B	1	2	∞
to C	2	1	∞
to Z	∞	∞	∞

DV
Update!

To A: 1
To B: 1
To Z: 2

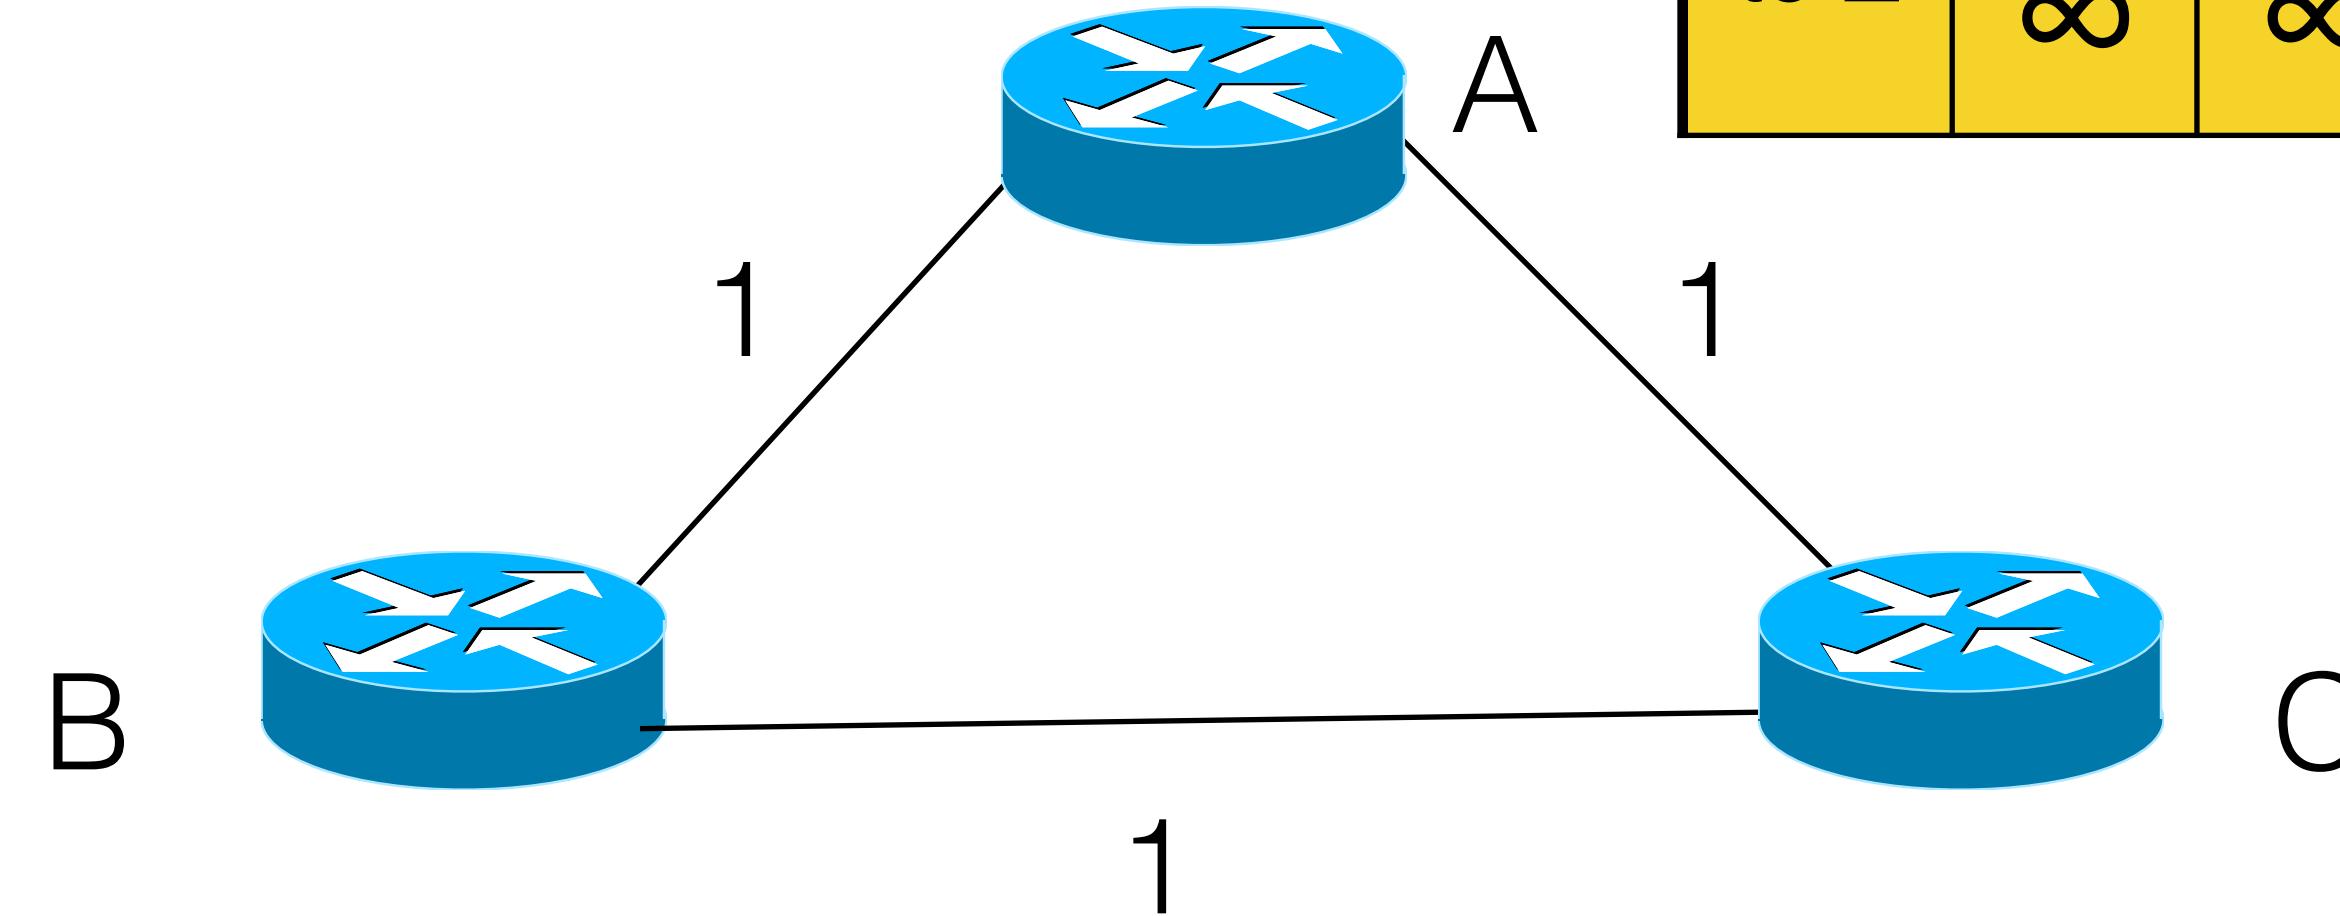
	via A	via B
to A	1	2
to B	2	1
to Z	2	3



B and C won't advertise the route for Z to A
But they will advertise it to each other...

A completely sad graph

	via A	via C
to A	1	2
to C	2	1
to Z	∞	3



	via B	via C	via Z
to B	1	2	∞
to C	2	1	∞
to Z	∞	∞	∞

	via A	via B
to A	1	2
to B	2	1
to Z	∞	3



A completely sad graph

	via A	via C
to A	1	2
to C	2	1
to Z	∞	3

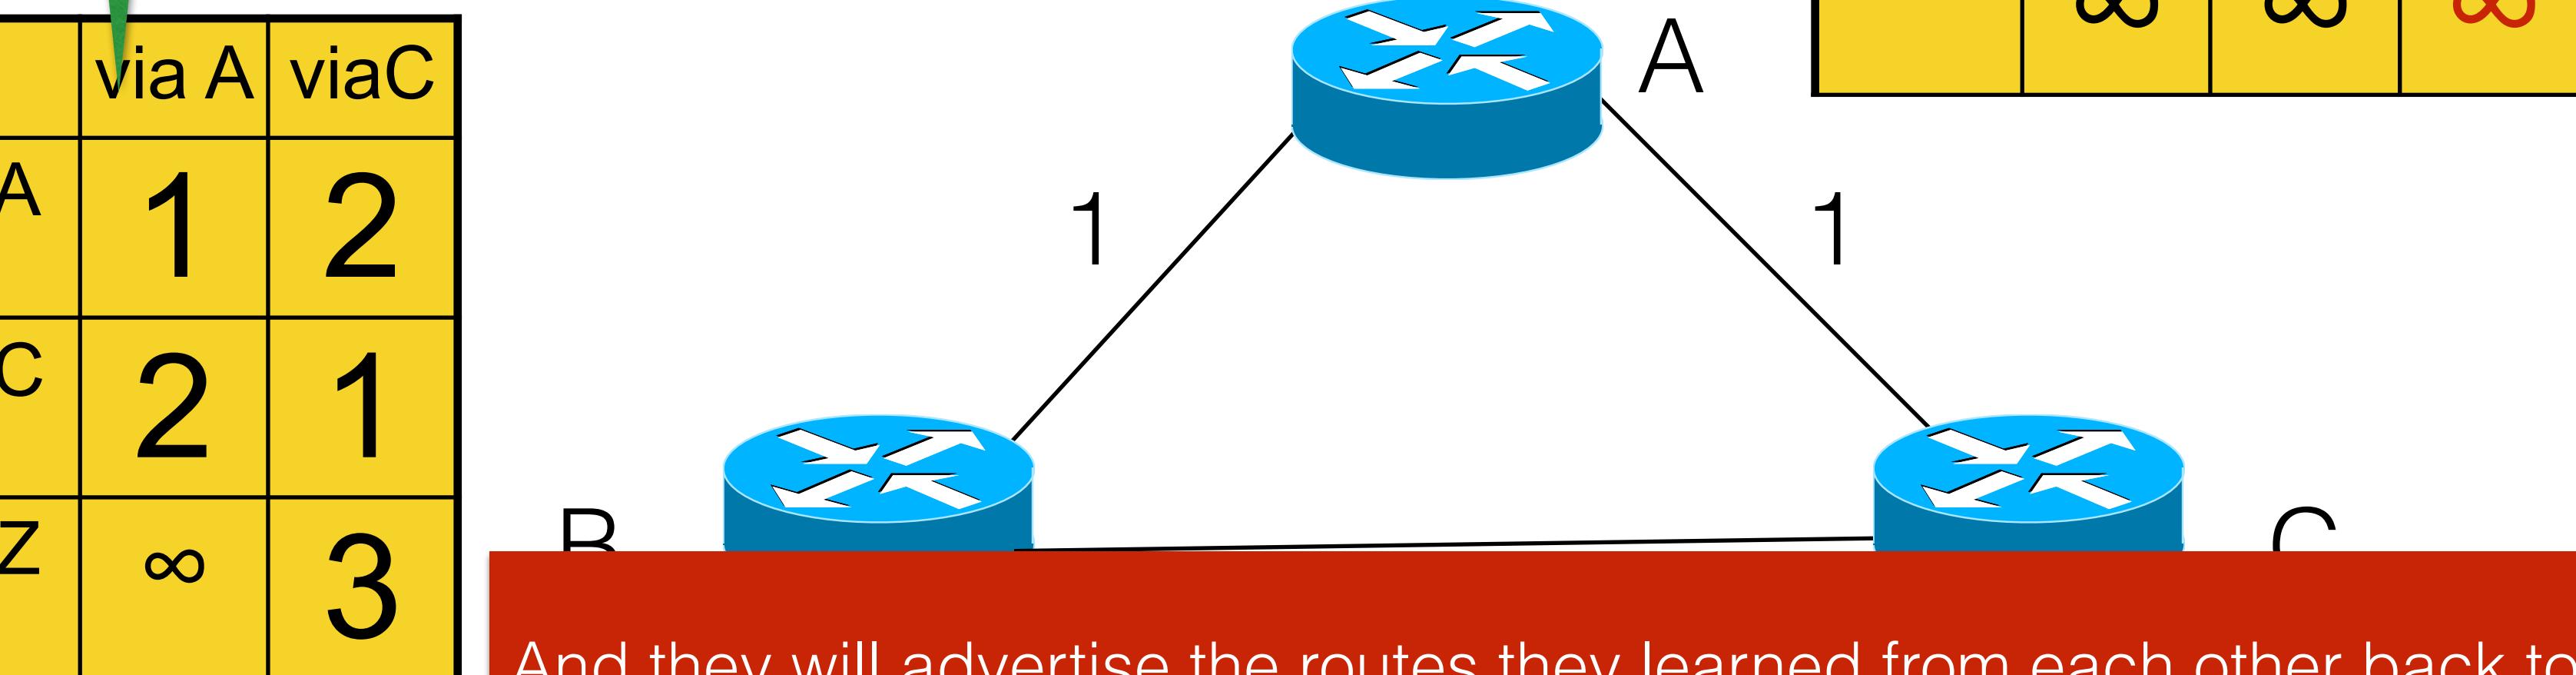
DV
Update!

To A: 1
To C: 1
To Z: 3

DV
Update!

To B: 1
To C: 1
To Z: ∞

	via B	via C	via Z
to B	1	2	∞
to C	2	1	∞
to Z	∞	∞	∞



DV
Update!

To A: 1
To B: 1
To Z: 3

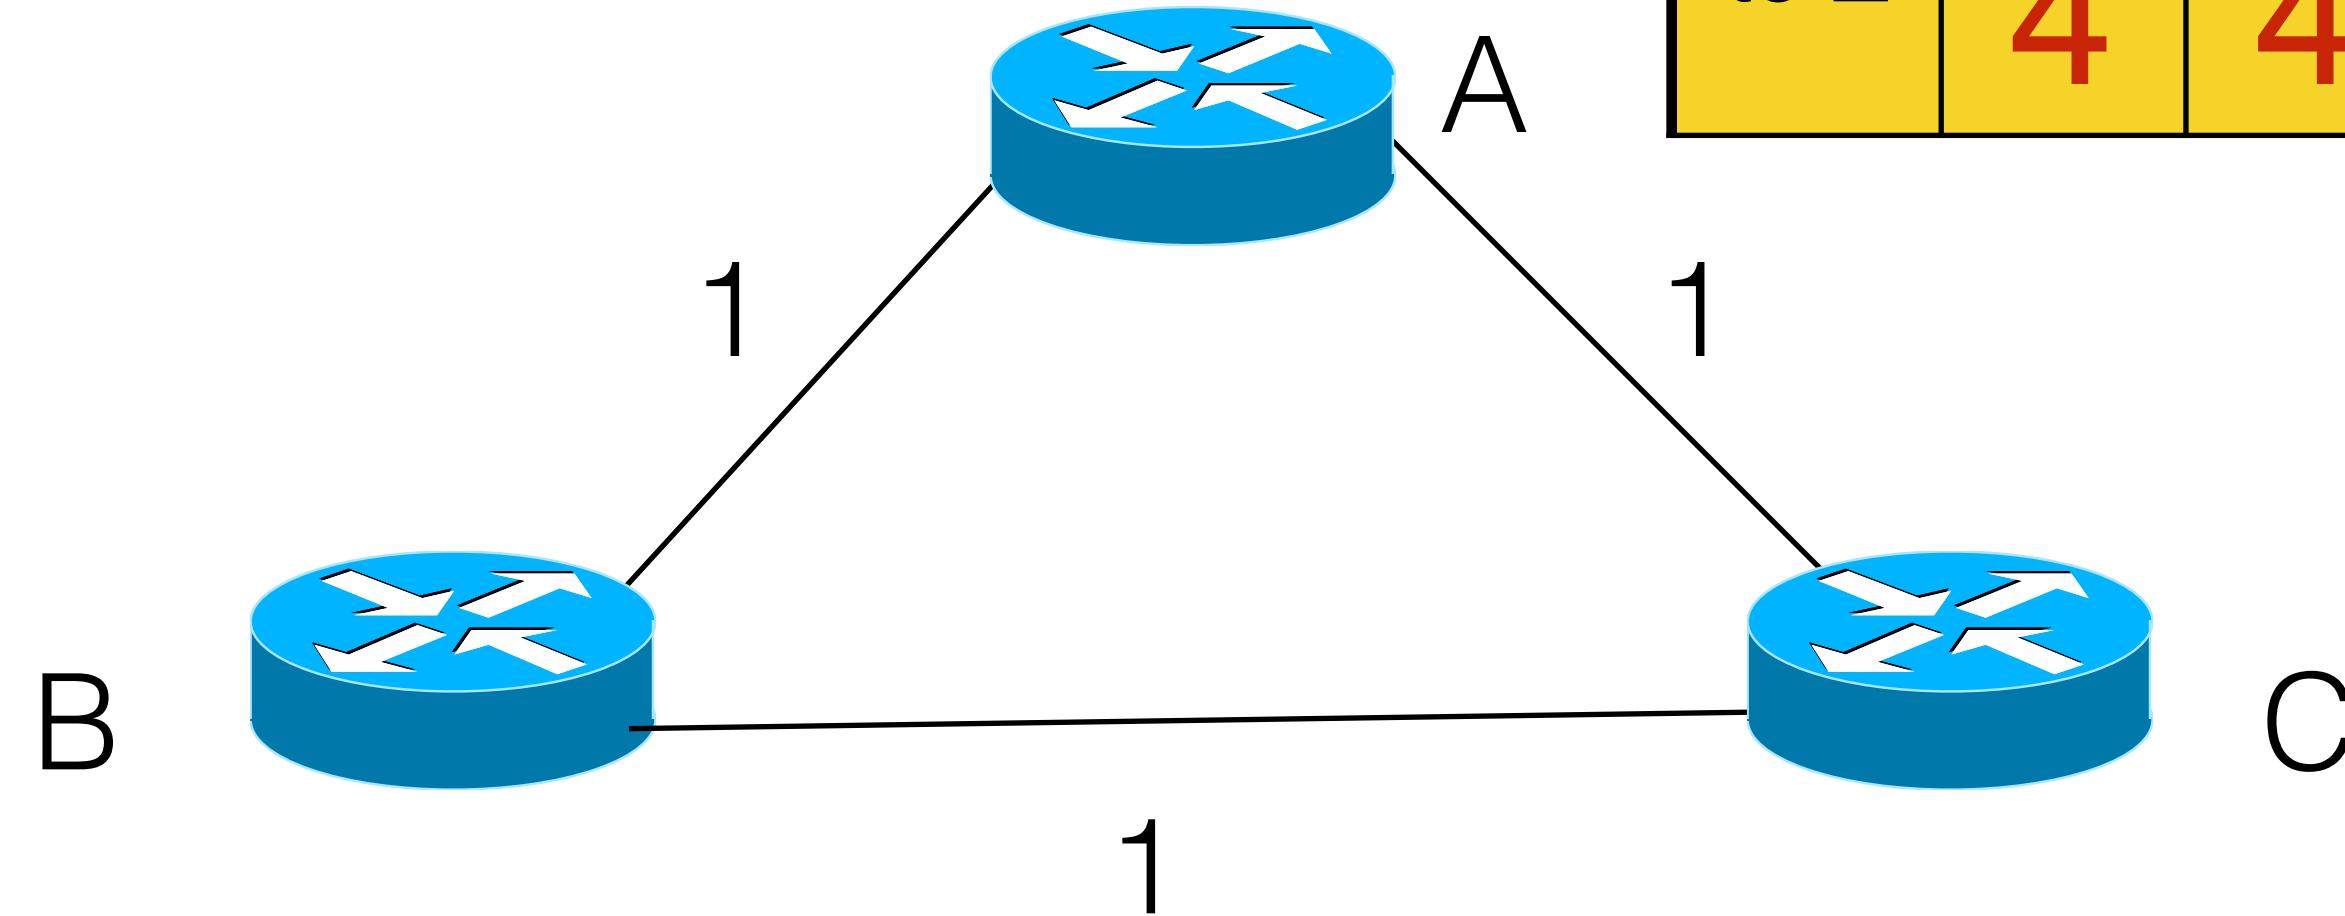
	via A	via B
to A	1	2
to B	2	1
to Z	∞	3



And they will advertise the routes they learned from each other back to A.

A completely sad graph

	via A	via C
to A	1	2
to C	2	1
to Z	∞	3



	via B	via C	via Z
to B	1	2	∞
to C	2	1	∞
to Z	4	4	∞

	via A	via B
to A	1	2
to B	2	1
to Z	∞	3



Here we go again...



In this case, we will count how
high?



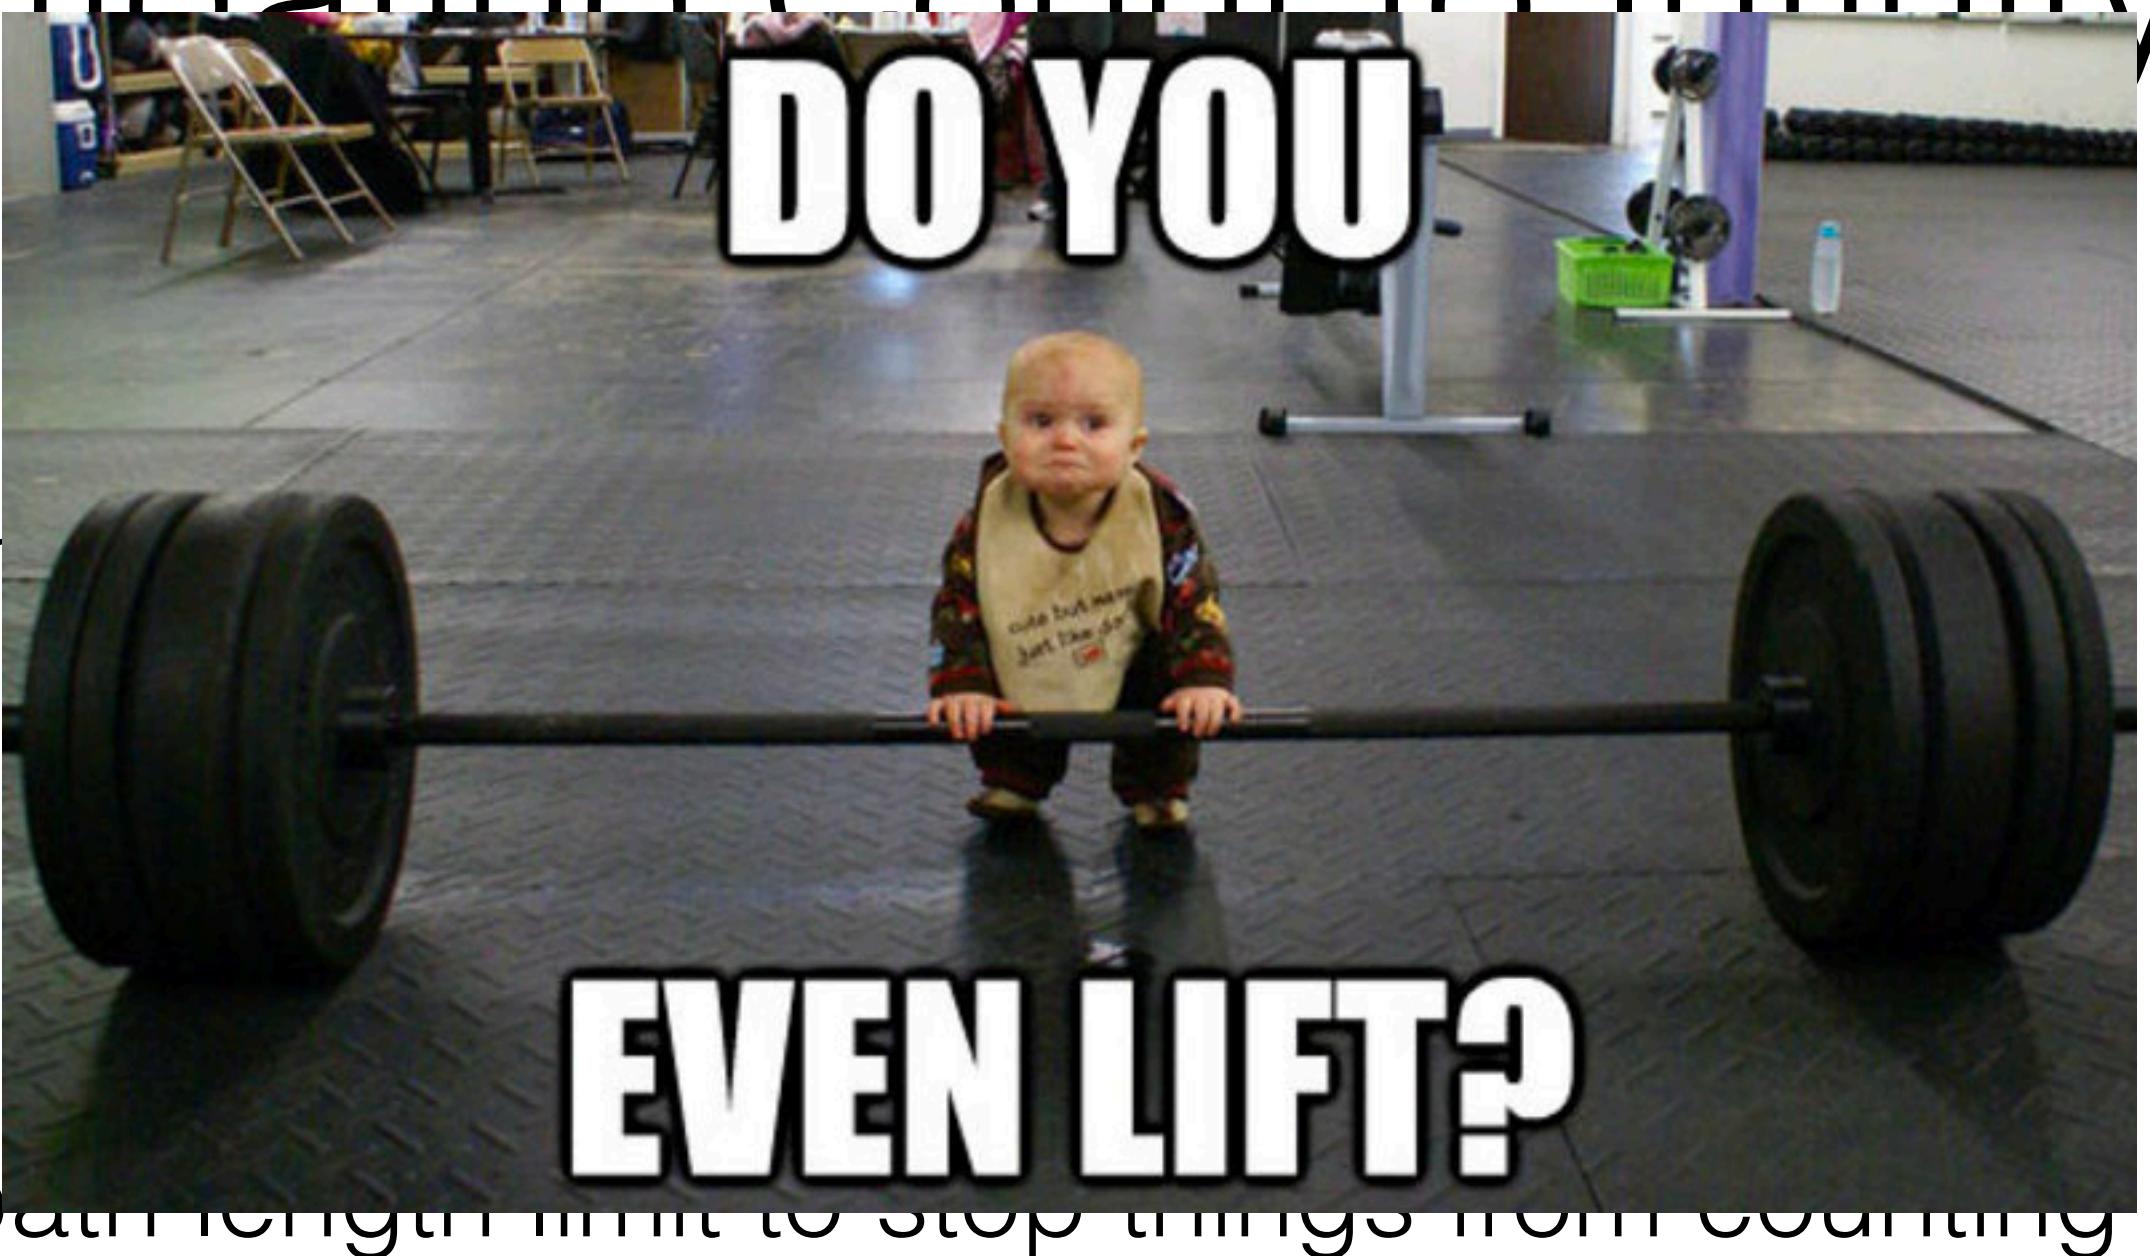
Three Techniques for Mitigating Count to Infinity

- Split Horizon/Poison Reverse
 - If I select a route I received from you in my distance vector, instead I will report a path length of INFINITY back to you.
- Maximum Path Lengths
 - Need to stop counting forever — set a path length limit to stop things from counting forever.
- Pushdown Timers



Three Techniques for Mitigating Count to Infinity

- Split Horizon/Poison Reverse
 - If I select a route I received from you in my routing table, I send back a path length of INFINITY back to you.
- Maximum Path Lengths
 - Need to stop counting forever... set a path length limit to stop things from counting
 - Sorry not sorry
 - I'm here to exercise your brain and make you mentally BUFF.
- Pushdown Timers
 - I'm not going to talk about these in lecture, instead I'm going to give you a tricky question about them, either on the homework or the midterm.



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Distance Vector
e.g RIP

Need to recompute
spanning tree if failure

Very slow recovery due
to count to infinity.

Yes

Yes

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

$O(\# \text{ switches} * \text{max node degree}) \ O(\#nodes)$

Need to run spanning
tree protocol before
routing

Need to run DV before
routing — takes length of
longest best path time.

Still sends new
connections everywhere.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Next

- I am going to teach you TWO MORE ROUTING ALGORITHMS
 - I know
 - We're not even halfway through routing algorithms we will learn in this course even...
- Then we will talk about some routing algorithms used in practice (briefly)
- Then I will (finally) get started talking about what you're all really here for...

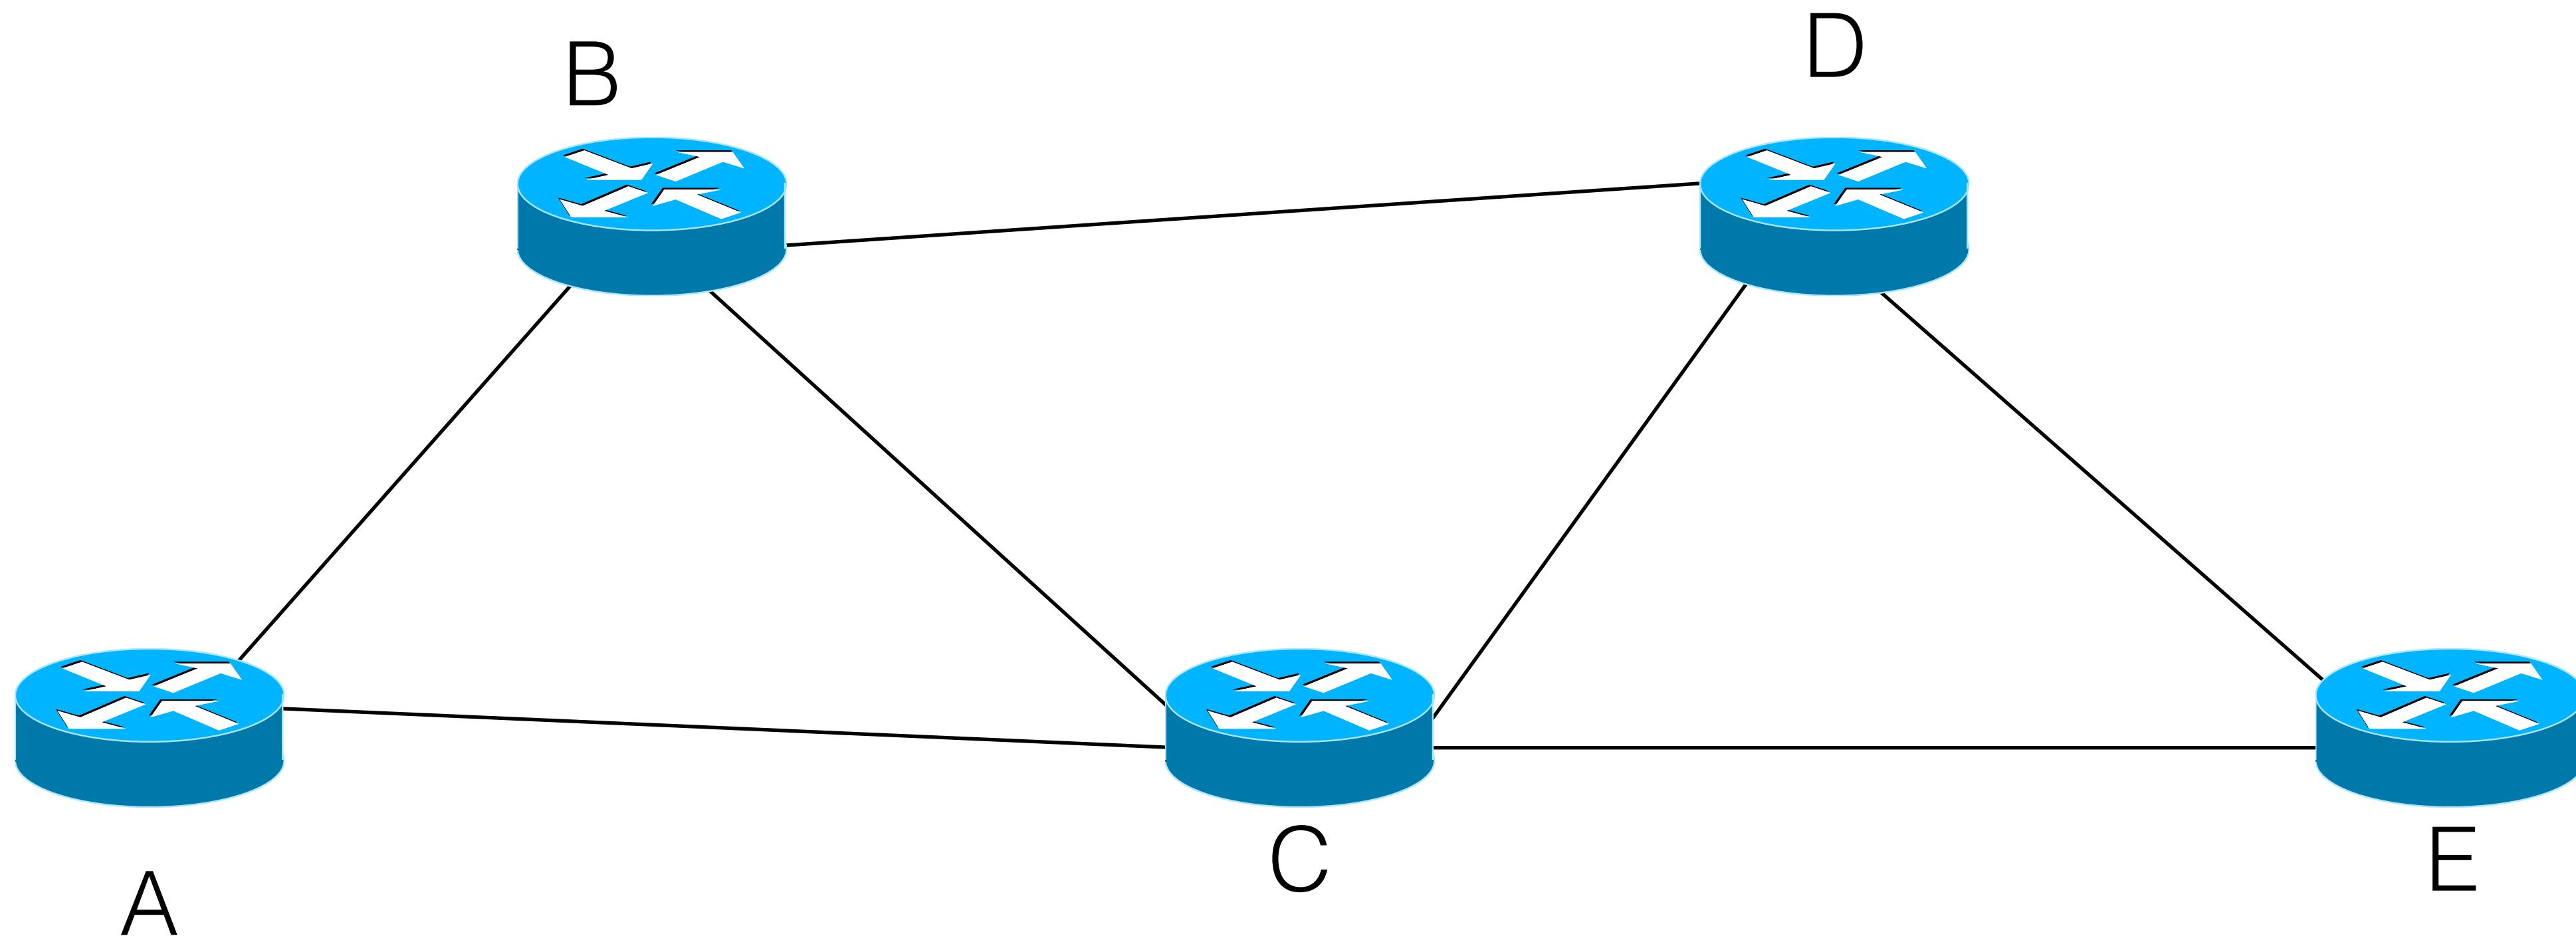


THE INTERNET



Our Newest Friend: Link State Algorithm (e.g. OSPF)

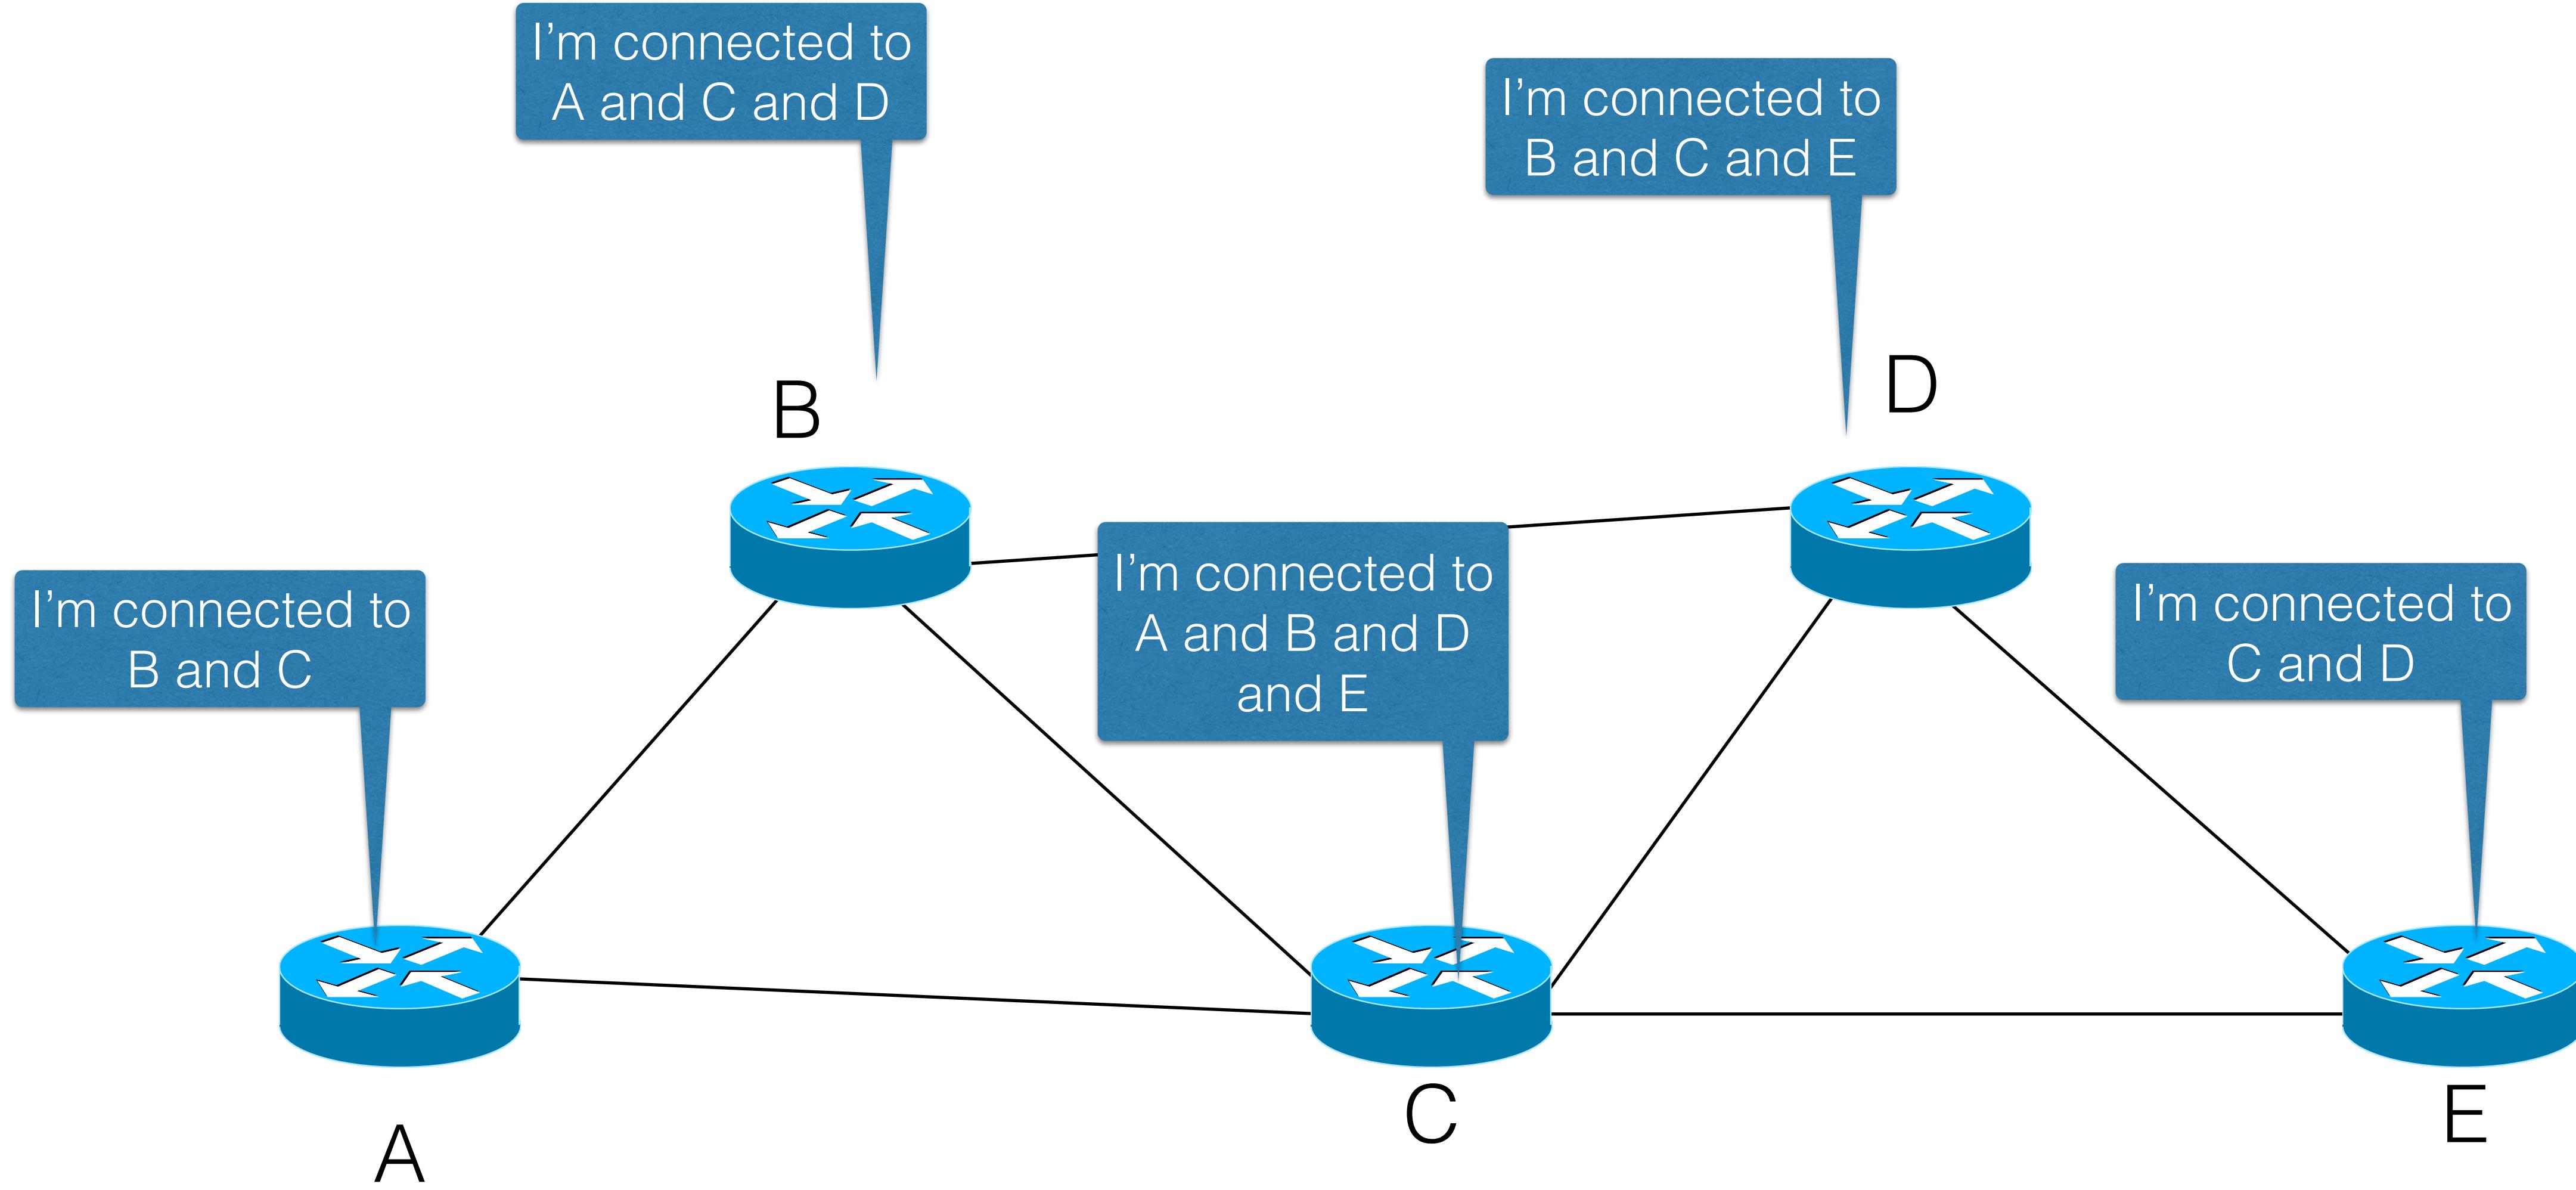




Link State, In a Nutshell

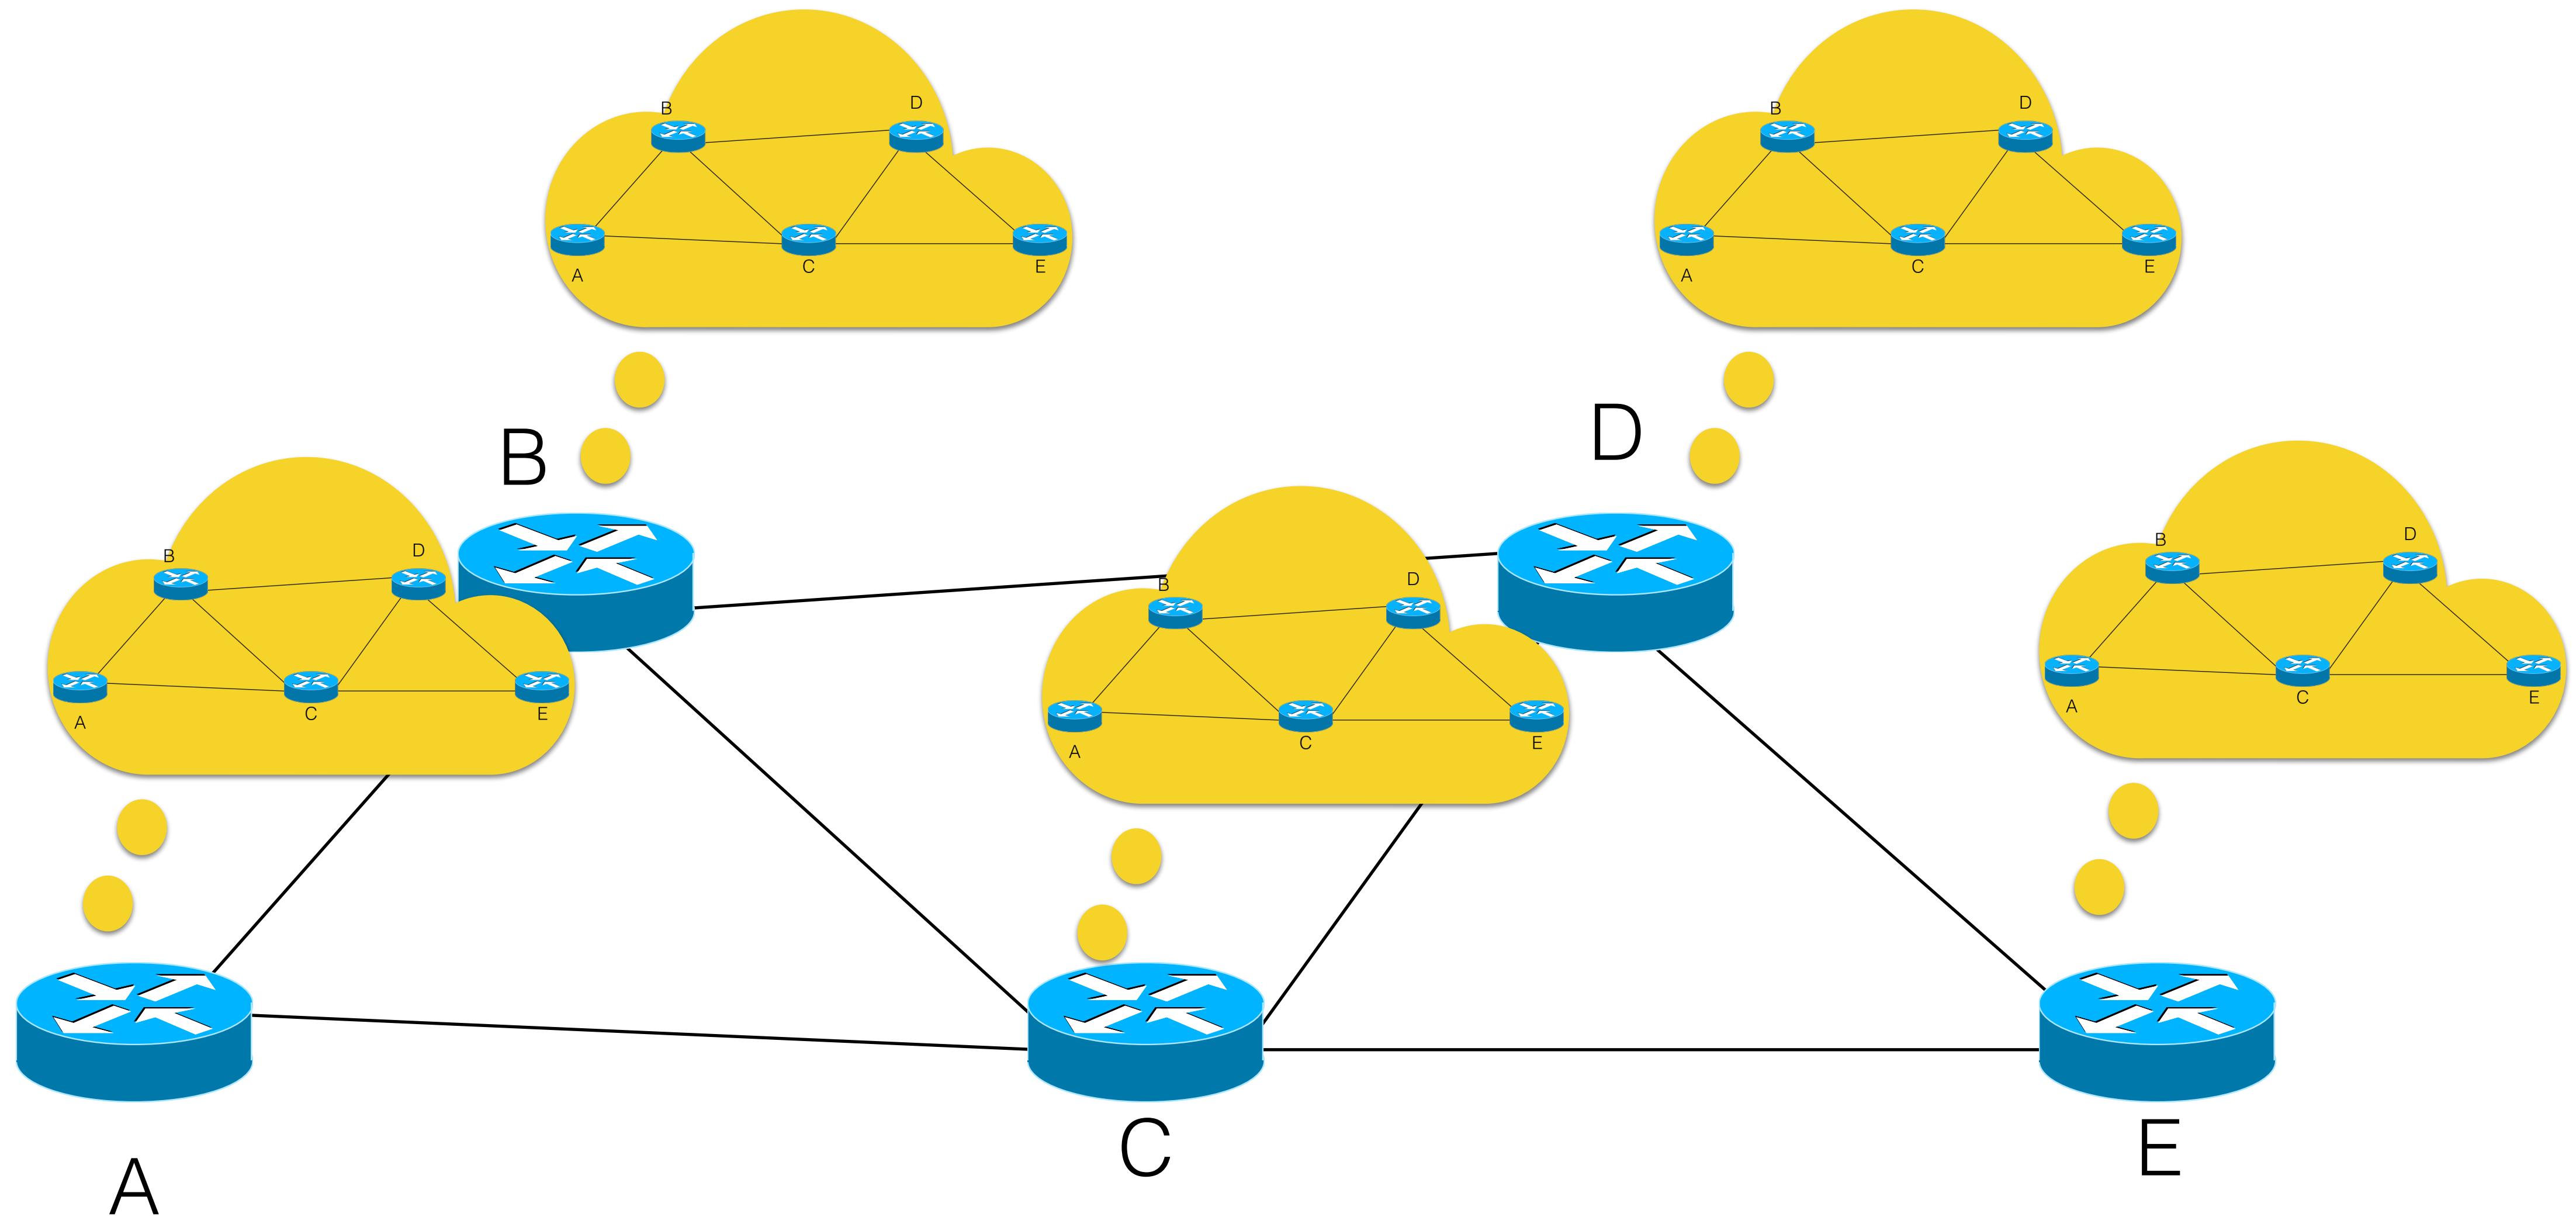
- Everyone knows who they are connected to directly.
- Every node *broadcasts* a list of who they are connected to (and with what link weight) to every other node in the network.
- Every node then can — locally — figure out what the entire network graph is.
- Each node then uses a shortest-path algorithm to find its shortest path to every possible destination.
- Each node then builds its own routing table.





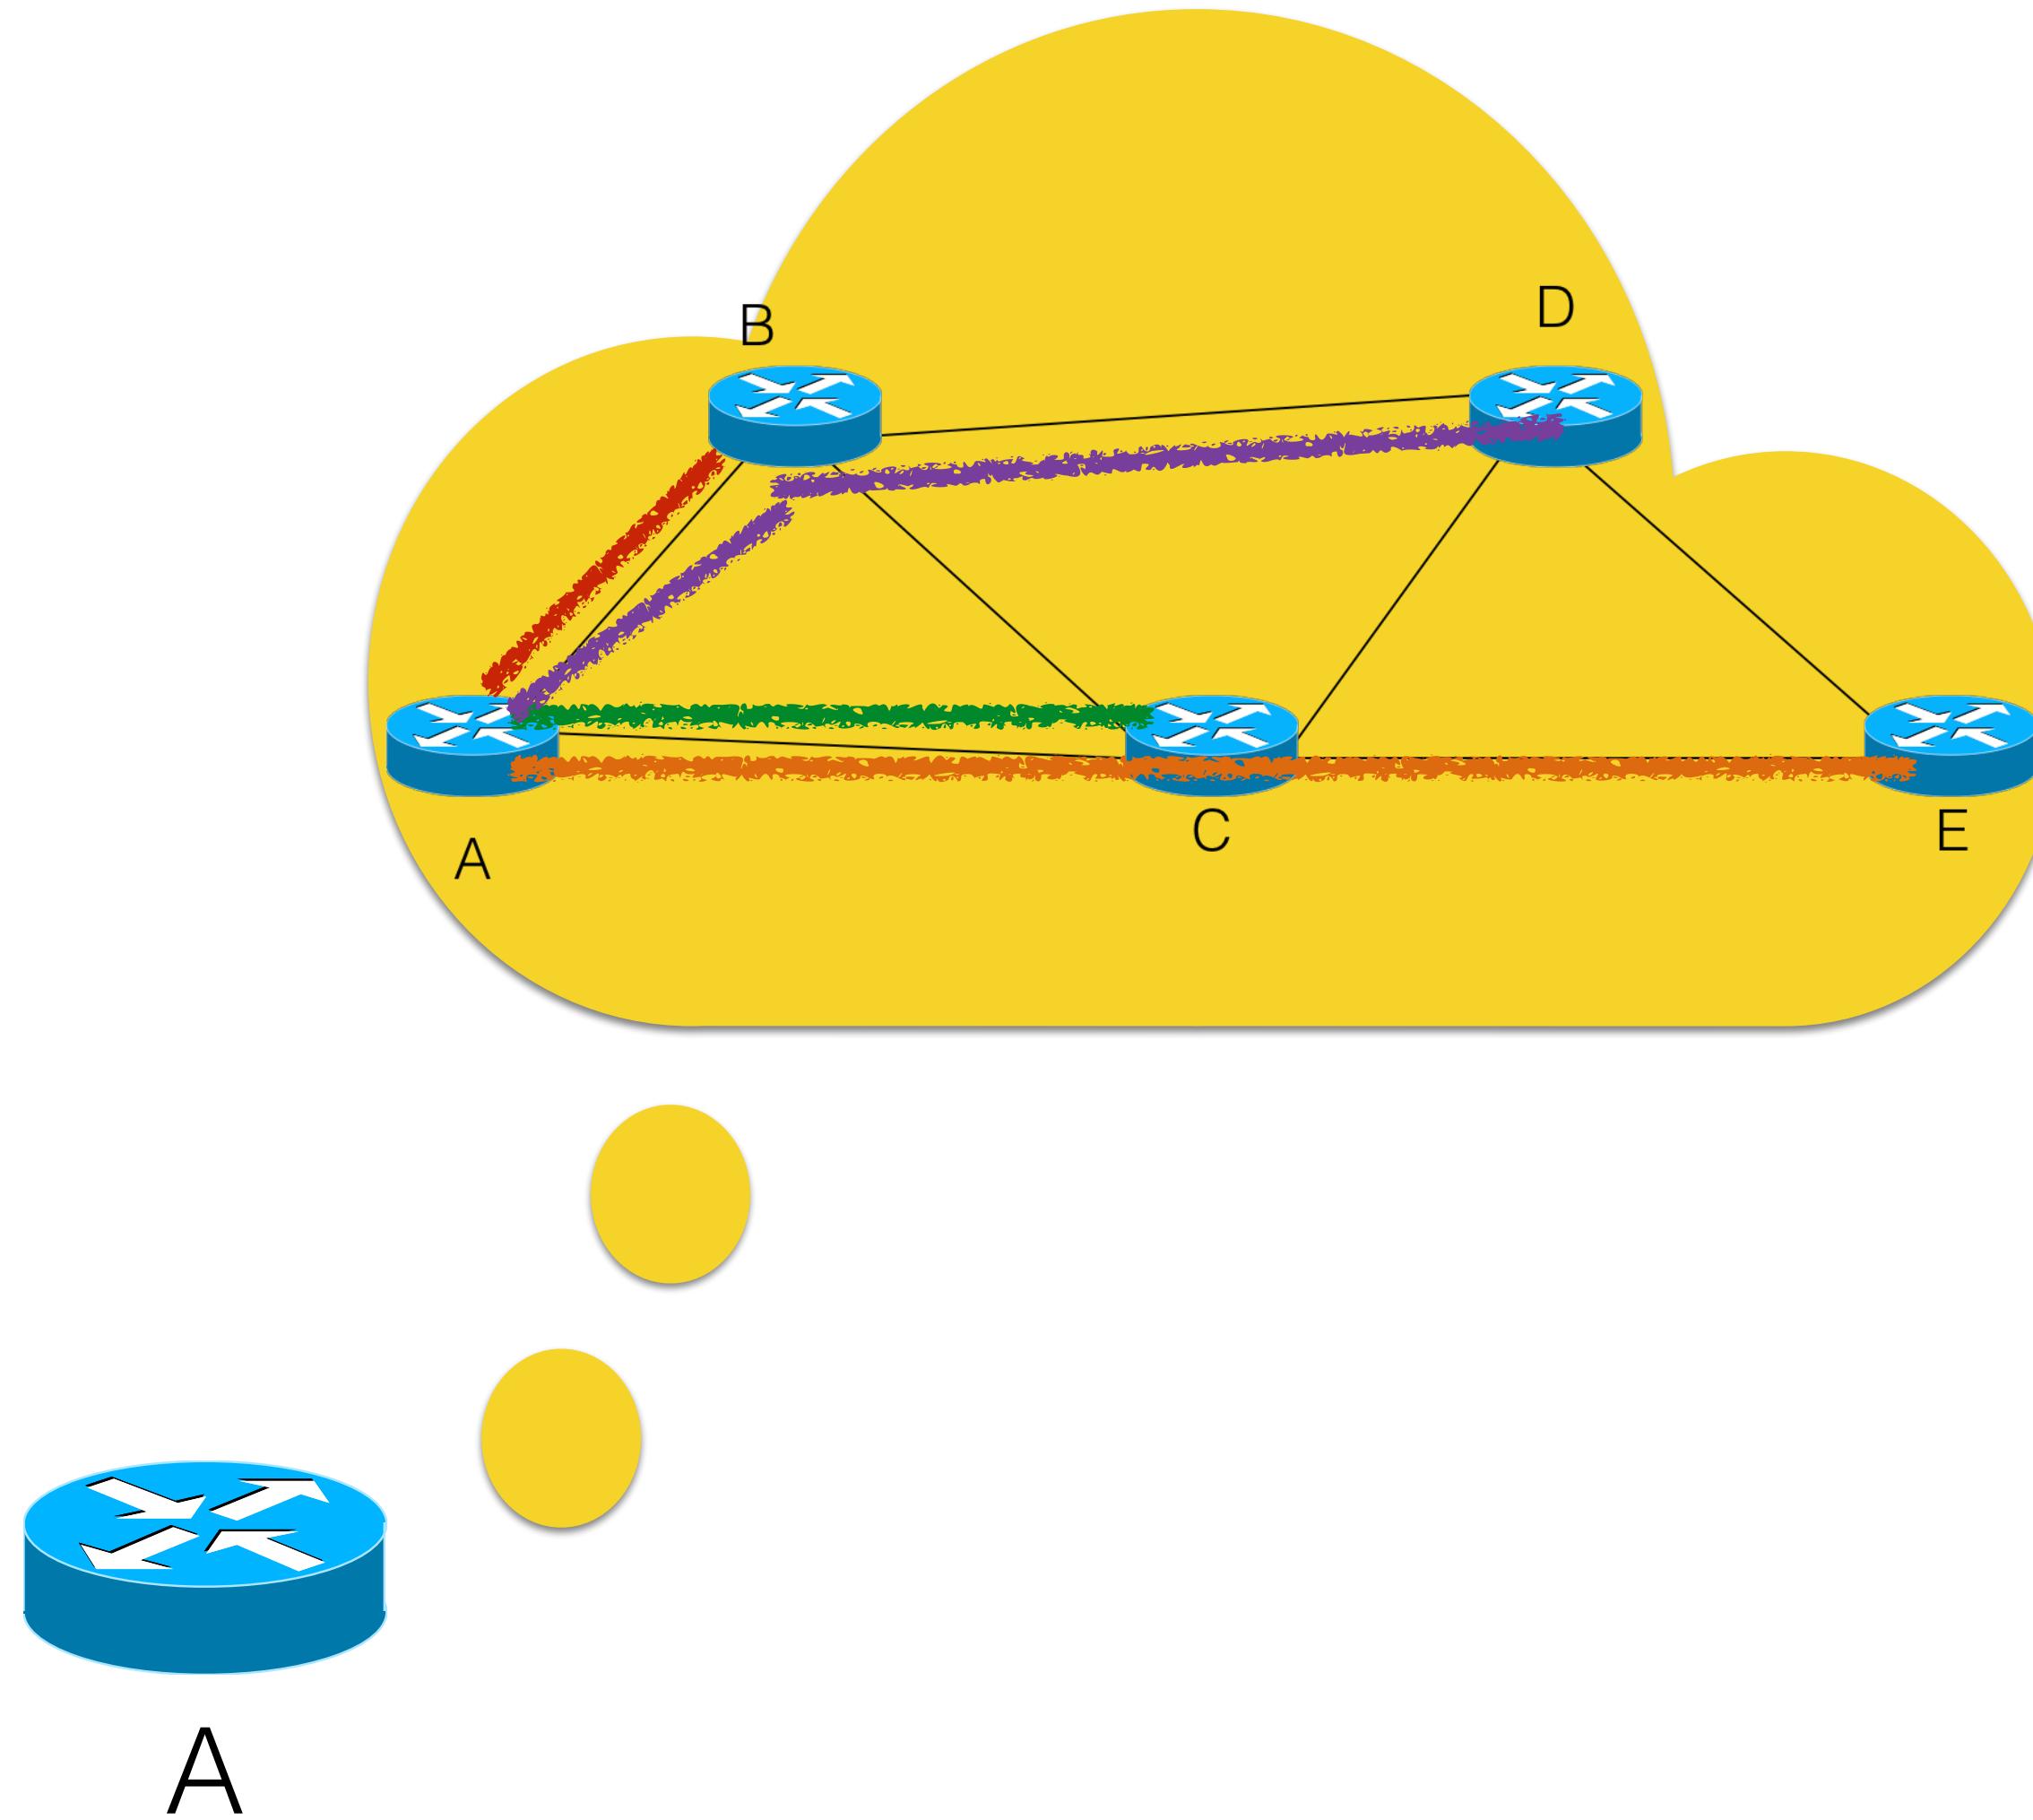
I think animating all of these messages being broadcast will give me carpal tunnel syndrome.
Imagine they are being broadcast.





Each node, through receiving these broadcast messages, can then figure out the entire network structure.





Once a router knows the entire graph structure, it can easily compute its shortest path to every other node in the network.



How do we compute A's
shortest path to all other nodes?



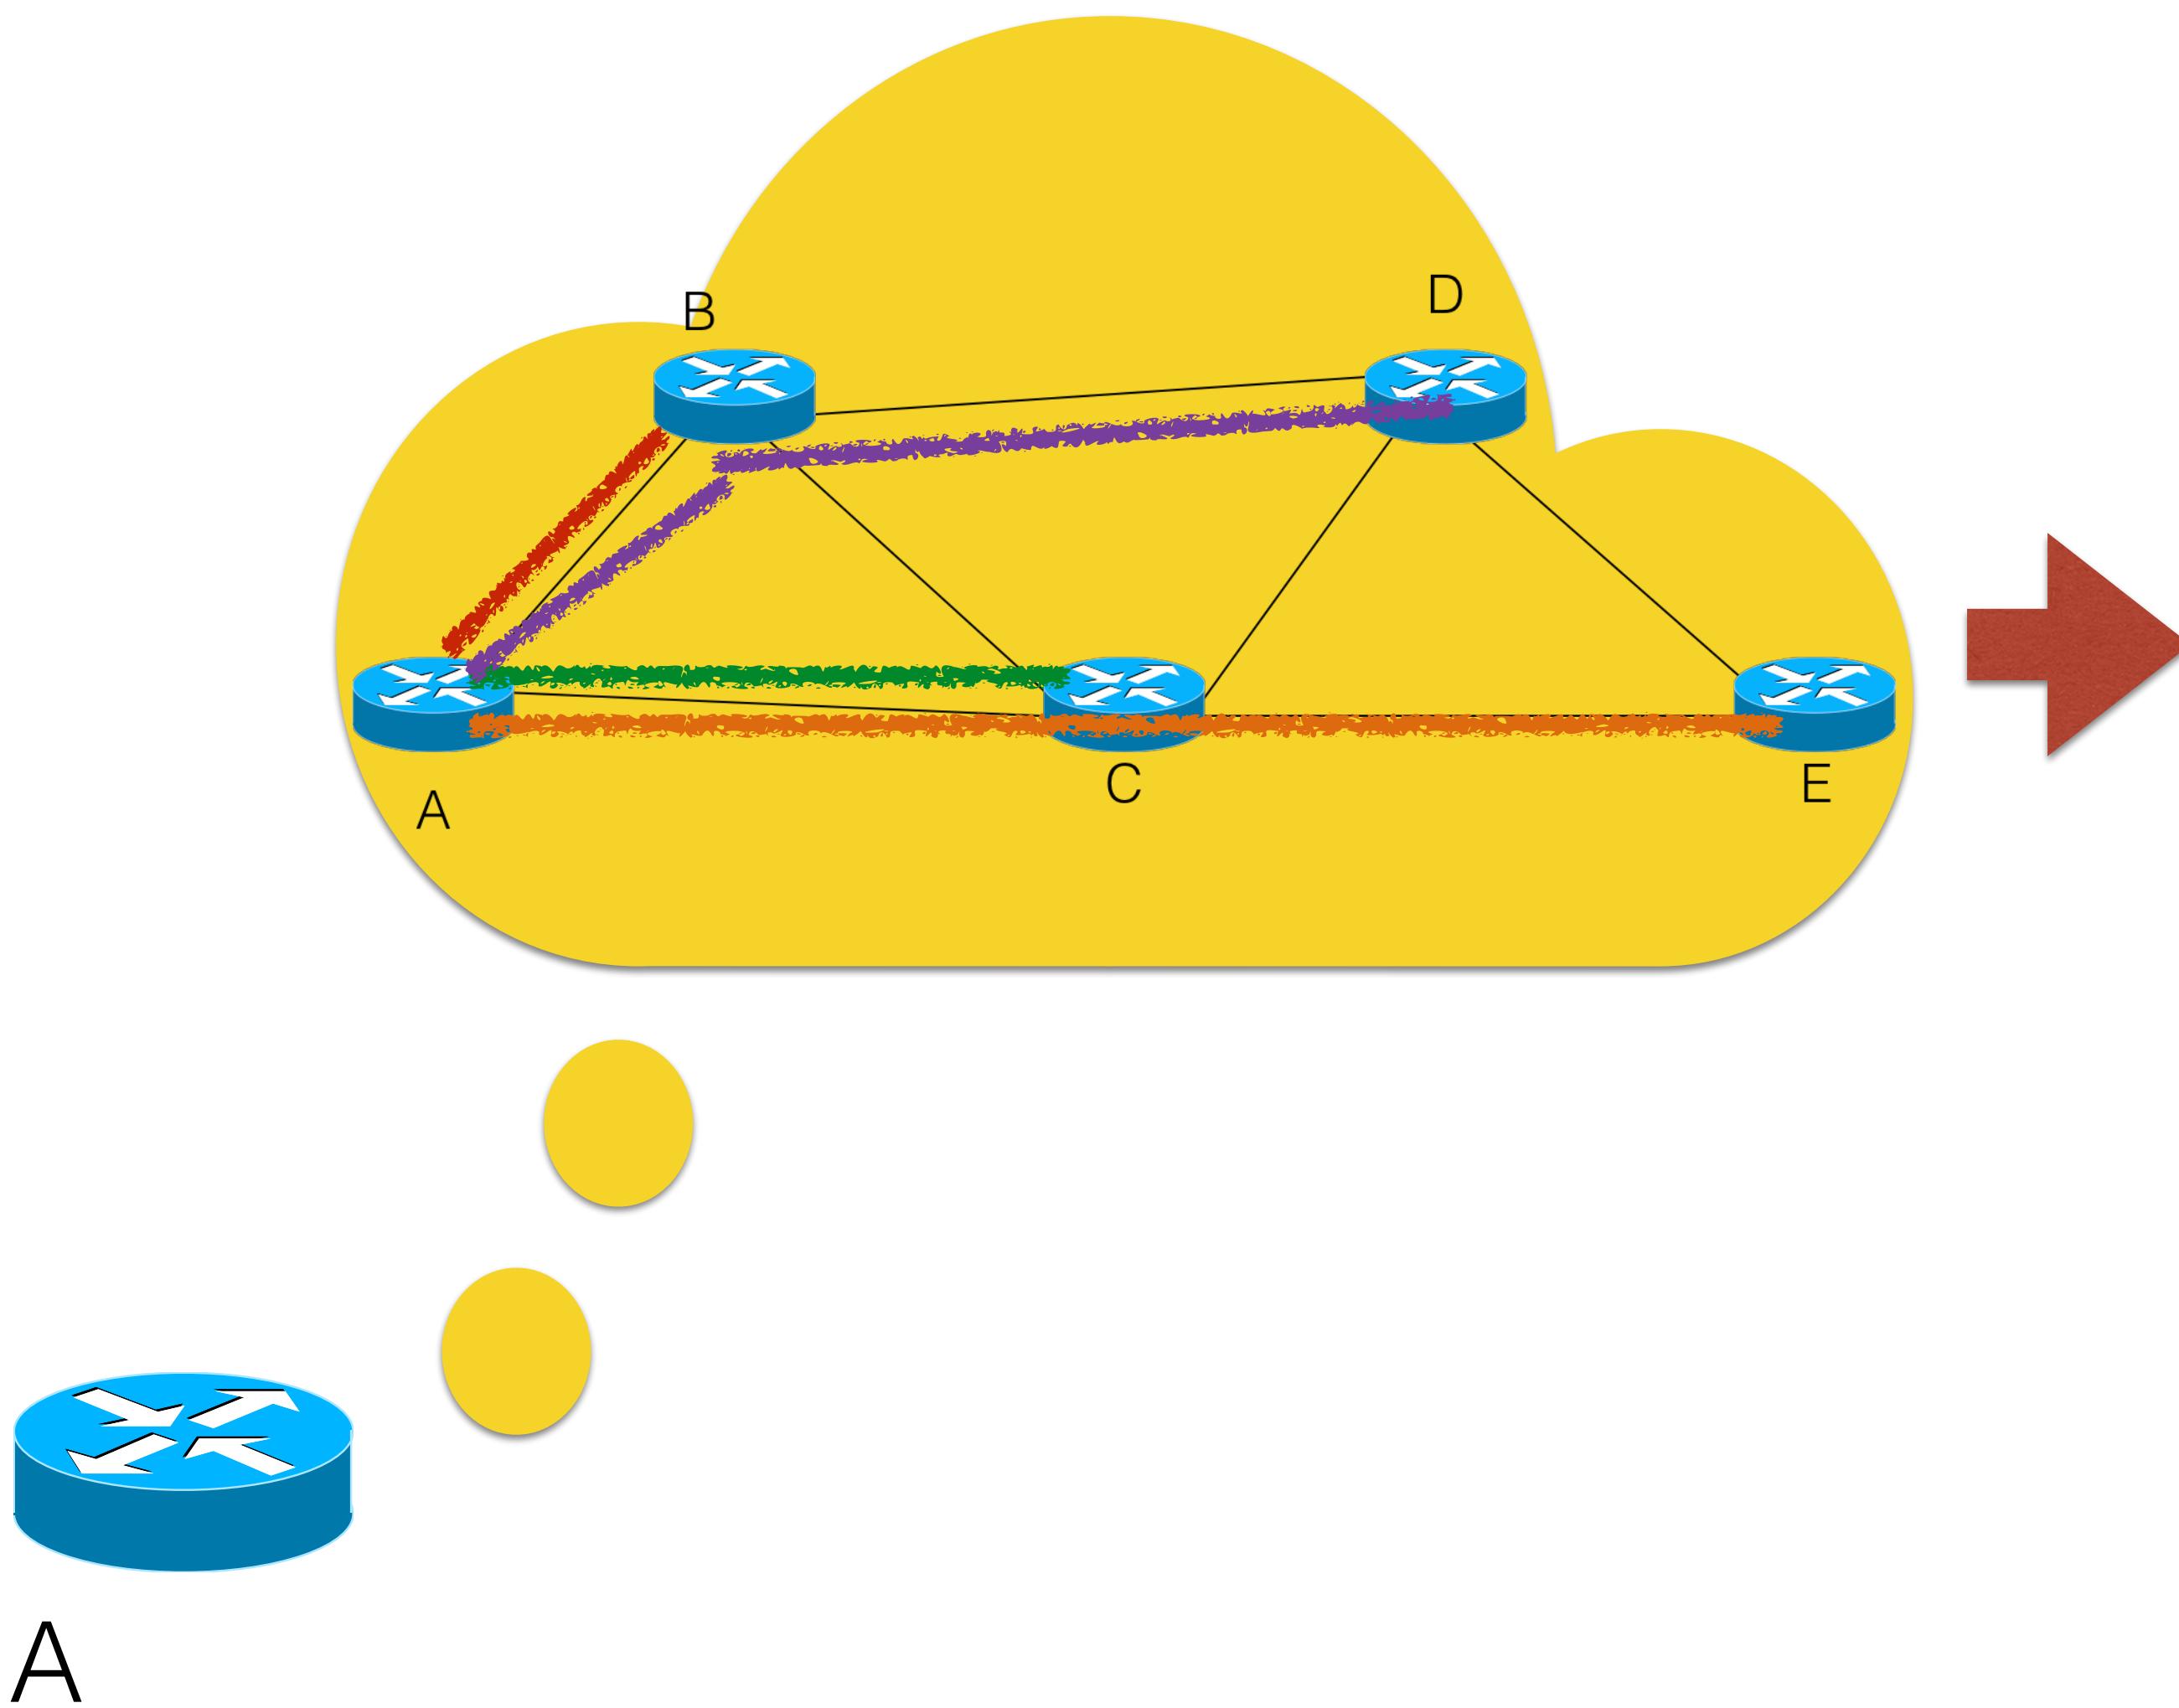
DIJKSTRA'S ALGORITHM



Edsger Dijkstra
Turing Award Winner
(11 May 1930 – 6 August 2002)

- Finds the shortest path from a source to all other destinations
- Runs in $O(|E| + |V| \log |V|)$ time
 - for E edges and V vertices
- You should have learned this 15-210, 15-251, or literally a billion other courses.
- *If you haven't learned it yet, come to OH and I'll teach it to you.*
 - *It's also on Wikipedia.*
 - *It's very cool — greedy algo!*





Once A knows its shortest path to every destination, it creates a routing table — what is the next hop for each destination?

To B? Forward to B
To C? Forward to C
To D? Forward to B
To E? Forward to C



That's like... it.



Broadcast Network w/ Learning Switches

Resilience

If there is a route, the
packet will reach dest!

Fully Distributed

Yes

State per Node

Learning Switch:
 $O(\#nodes)$

Convergence

No setup time at all!

Routing Efficiency

Broadcast Storms

Shortest Path?

Not Necessarily...

Broadcast Network w/ Learning Switches and Spanning Tree

Distance Vector
e.g RIP

Need to recompute
spanning tree if failure

Very slow recovery due
to count to infinity.

Yes

Yes

Learning Switch:
 $O(\#nodes) + \text{Path to Root: } O(\text{constant})$

$O(\# \text{ switches} * \text{max node degree}) + O(\#nodes)$

Need to run spanning
tree protocol before
routing

Need to run DV before
routing — takes length of
longest best path time.

Still sends new
connections everywhere.

Packets sent directly to
their destination.

Not Necessarily...

Yes



Distance Vector
e.g RIP

Resilience

Very slow recovery due
to count to infinity.

Fully Distributed

Yes

State per Node
(+ Routing Table)

$O(\# \text{ switches} * \text{max node degree})$

Convergence

Need to run DV before
routing — takes length of
longest best path time.

Routing Efficiency

Packets sent directly to
their destination.

Shortest Path?

Yes



	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	
Fully Distributed	Yes	
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	
Convergence	Need to run DV before routing — takes length of longest best path time.	
Routing Efficiency	Packets sent directly to their destination.	Packets sent directly to their destination.
Shortest Path?	Yes	Yes



Okay, I did the first two for you — fill in the rest of this table w/ your neighbor.



	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	
Fully Distributed	Yes	
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Routing Efficiency	Packets sent directly to their destination.	Packets sent directly to their destination.
Shortest Path?	Yes	Yes



	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	
Fully Distributed	Yes	
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Routing Efficiency	Packets sent directly to their destination.	Packets sent directly to their destination.
Shortest Path?	Yes	Yes



	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	
Fully Distributed	Yes	Yes
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Routing Efficiency	Packets sent directly to their destination.	Packets sent directly to their destination.
Shortest Path?	Yes	Yes



	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.
Fully Distributed	Yes	Yes
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Routing Efficiency	Packets sent directly to their destination.	Packets sent directly to their destination.
Shortest Path?	Yes	Yes



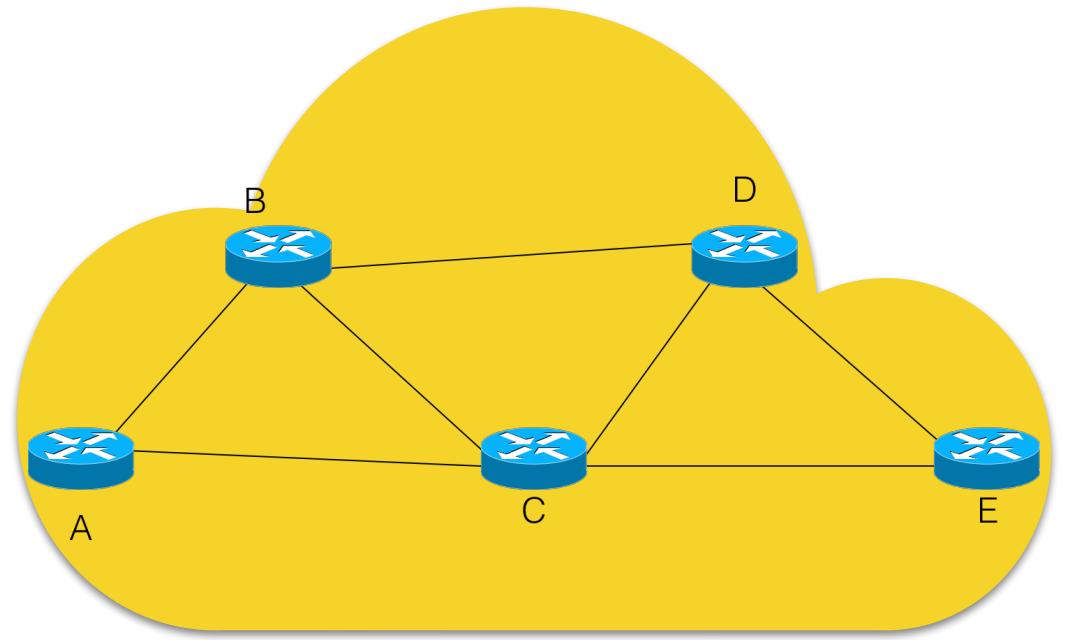
One more point of comparison:
Switch Control Plane Complexity



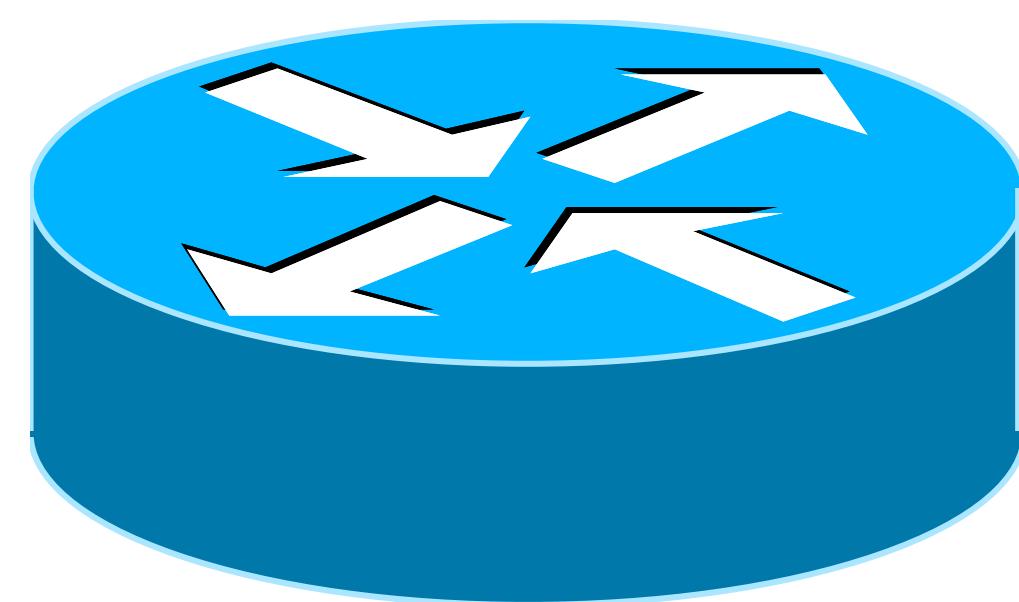
A switch has two components

- “Data Plane”
 - When a packet comes in, the data plane reads from a routing table and decides where to send the packet. Then it sends the packet out the correct port.
- “Control Plane”
 - This is the “brains” of the switch — the part that decides what to put into the routing table.

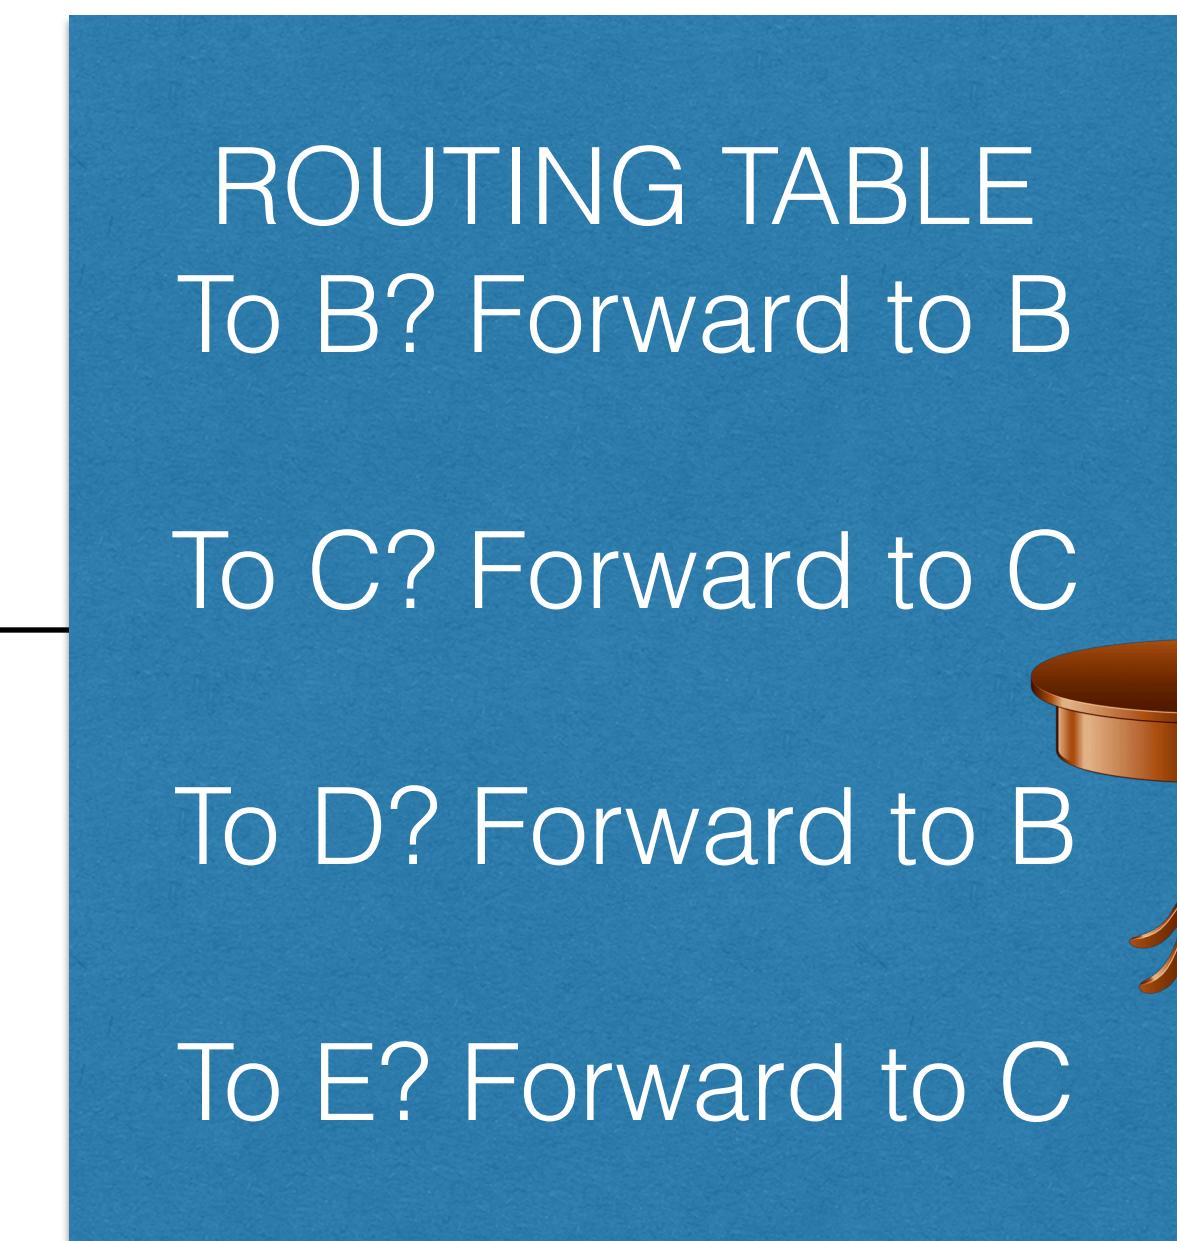




Send and receive updates
Decide what my new routes should be
WRITE to routing table



Control Plane
Data Plane



Receive packets.
READ from routing table to learn where I am supposed to send them
Send packet out of the correct port.



	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.
Fully Distributed	Yes	Yes
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Control Plane Complexity		

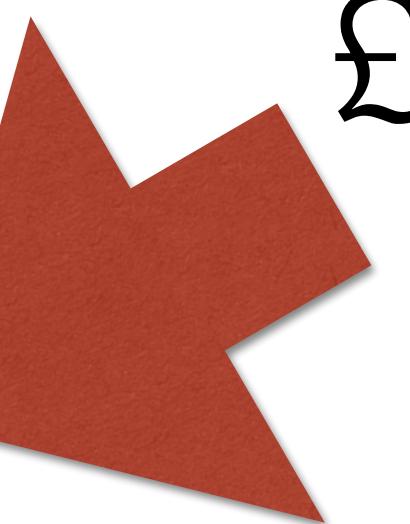


	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.
Fully Distributed	Yes	Yes
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Control Plane Complexity	Just select the “min” of all the updates I have heard from. (Dumb-ish Switch)	



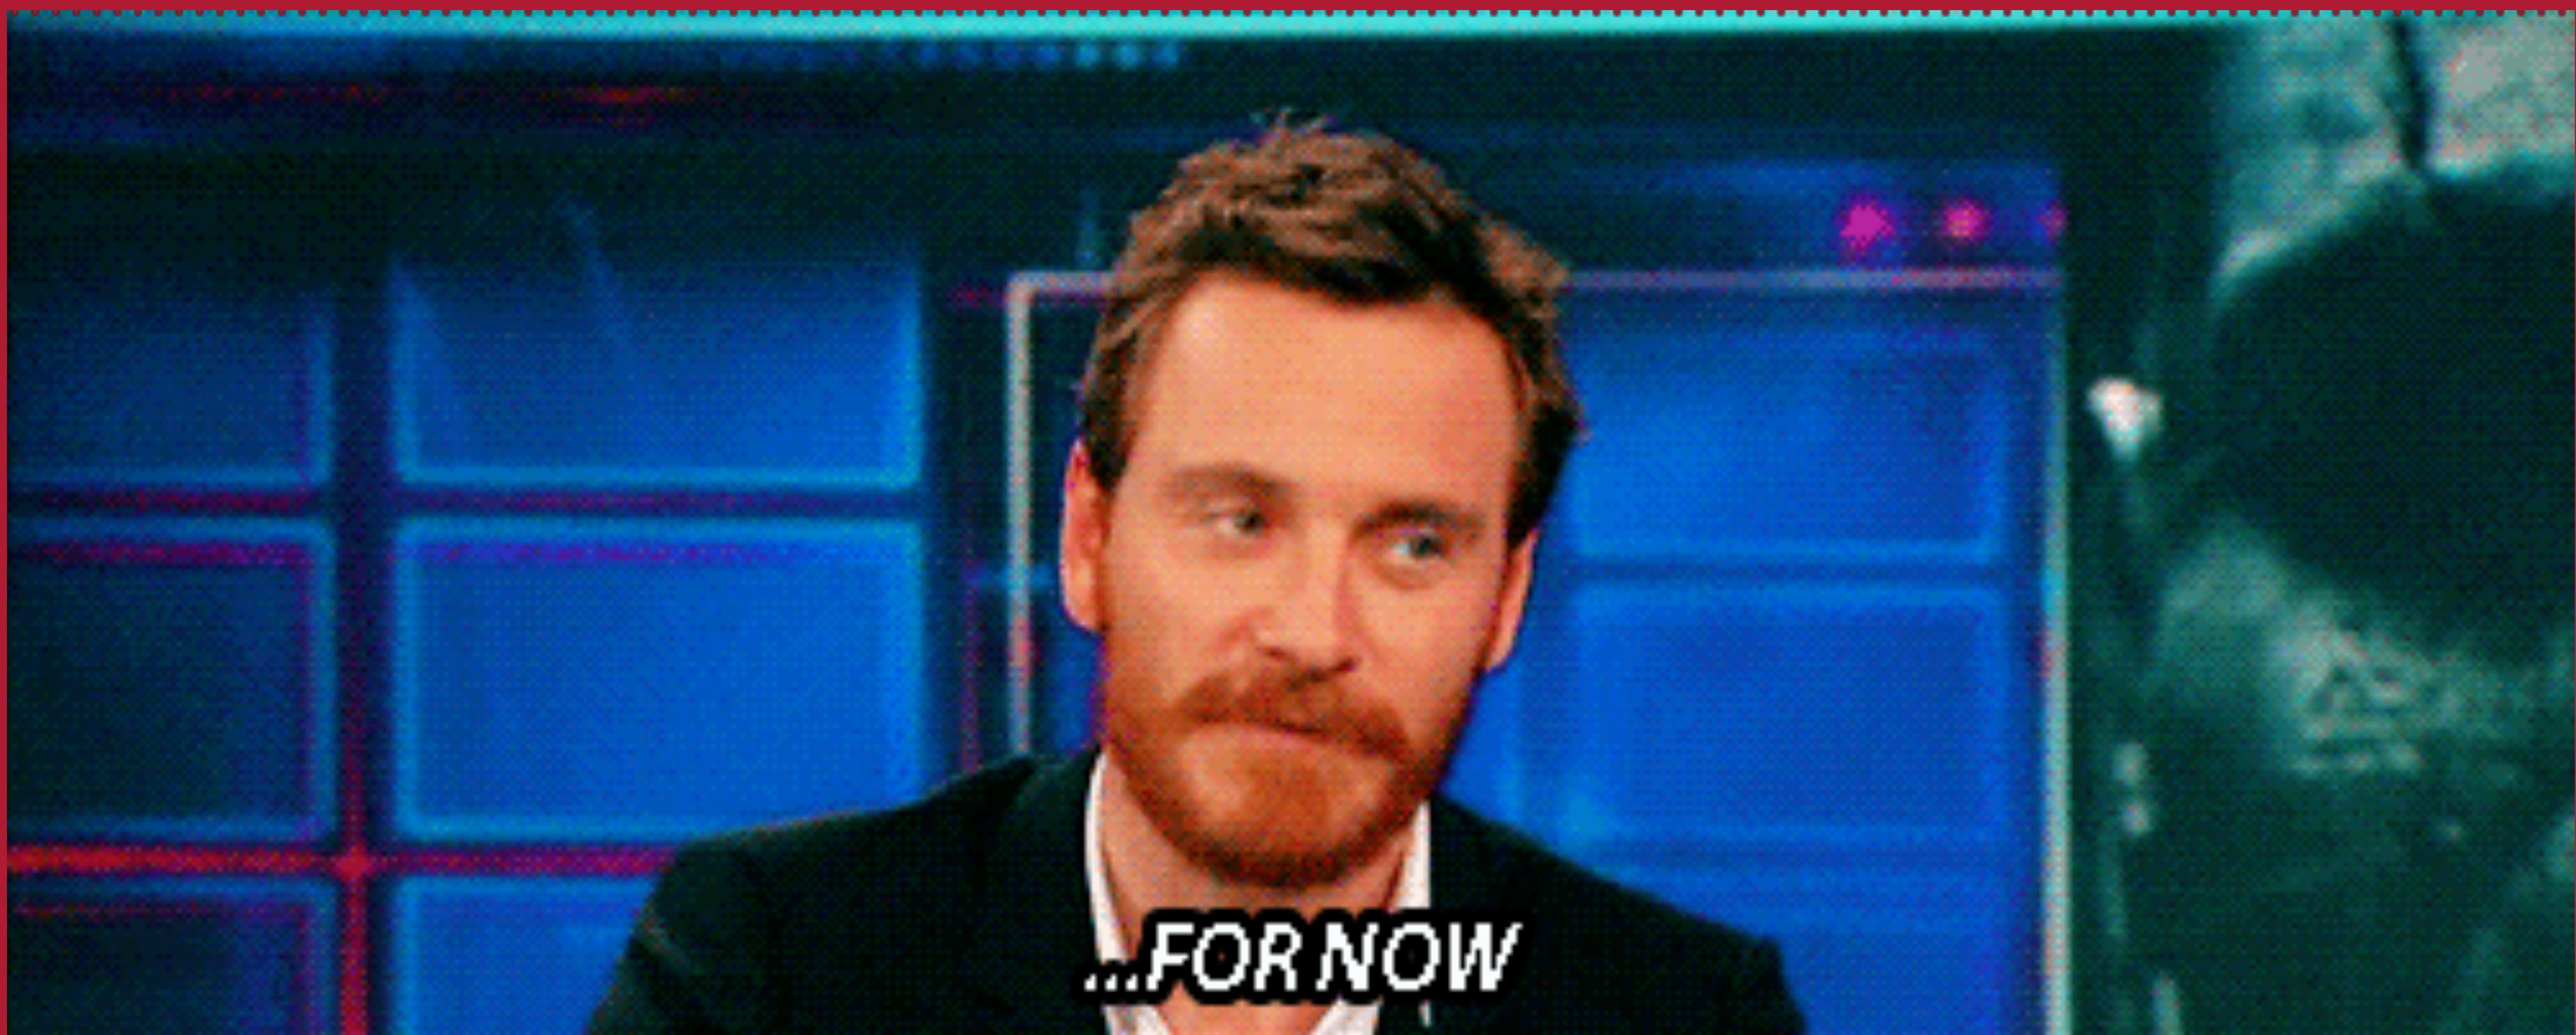
	Distance Vector e.g RIP	Link State e.g OSPF
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.
Fully Distributed	Yes	Yes
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$
Control Plane Complexity	Just select the “min” of all the updates I have heard from. (Dumb-ish Switch)	Rebuild network topology, run Dijkstra's algorithm over it.

£\$€



Last Routing Algorithm...





..FOR NOW



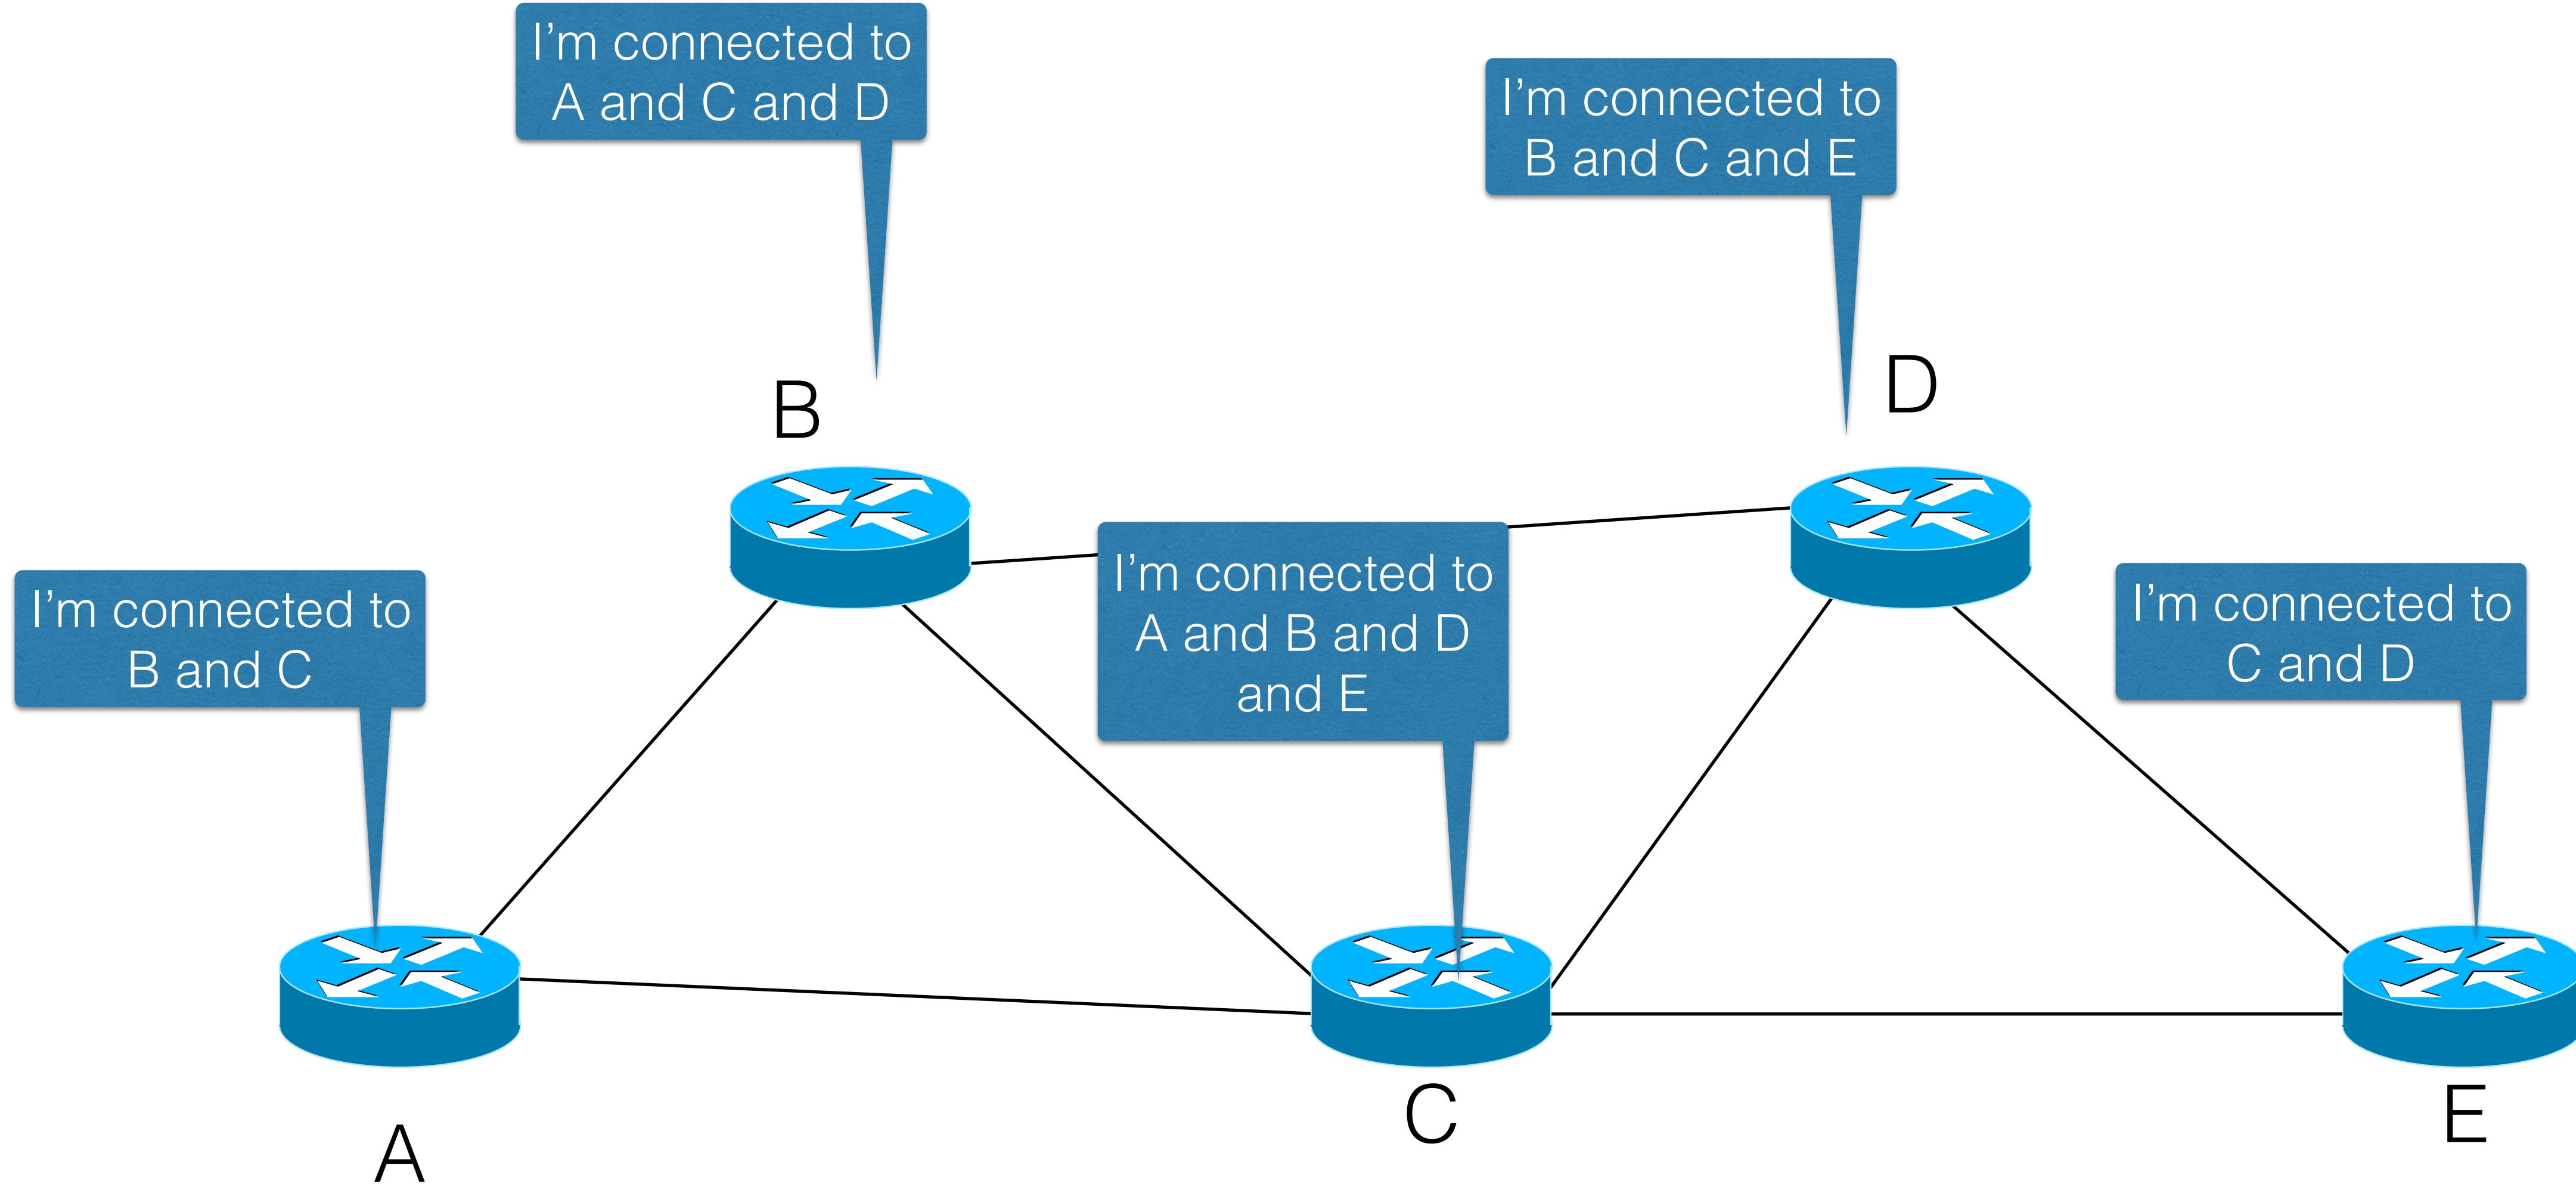
Centralized Routing (aka “Software-Defined Networking”)

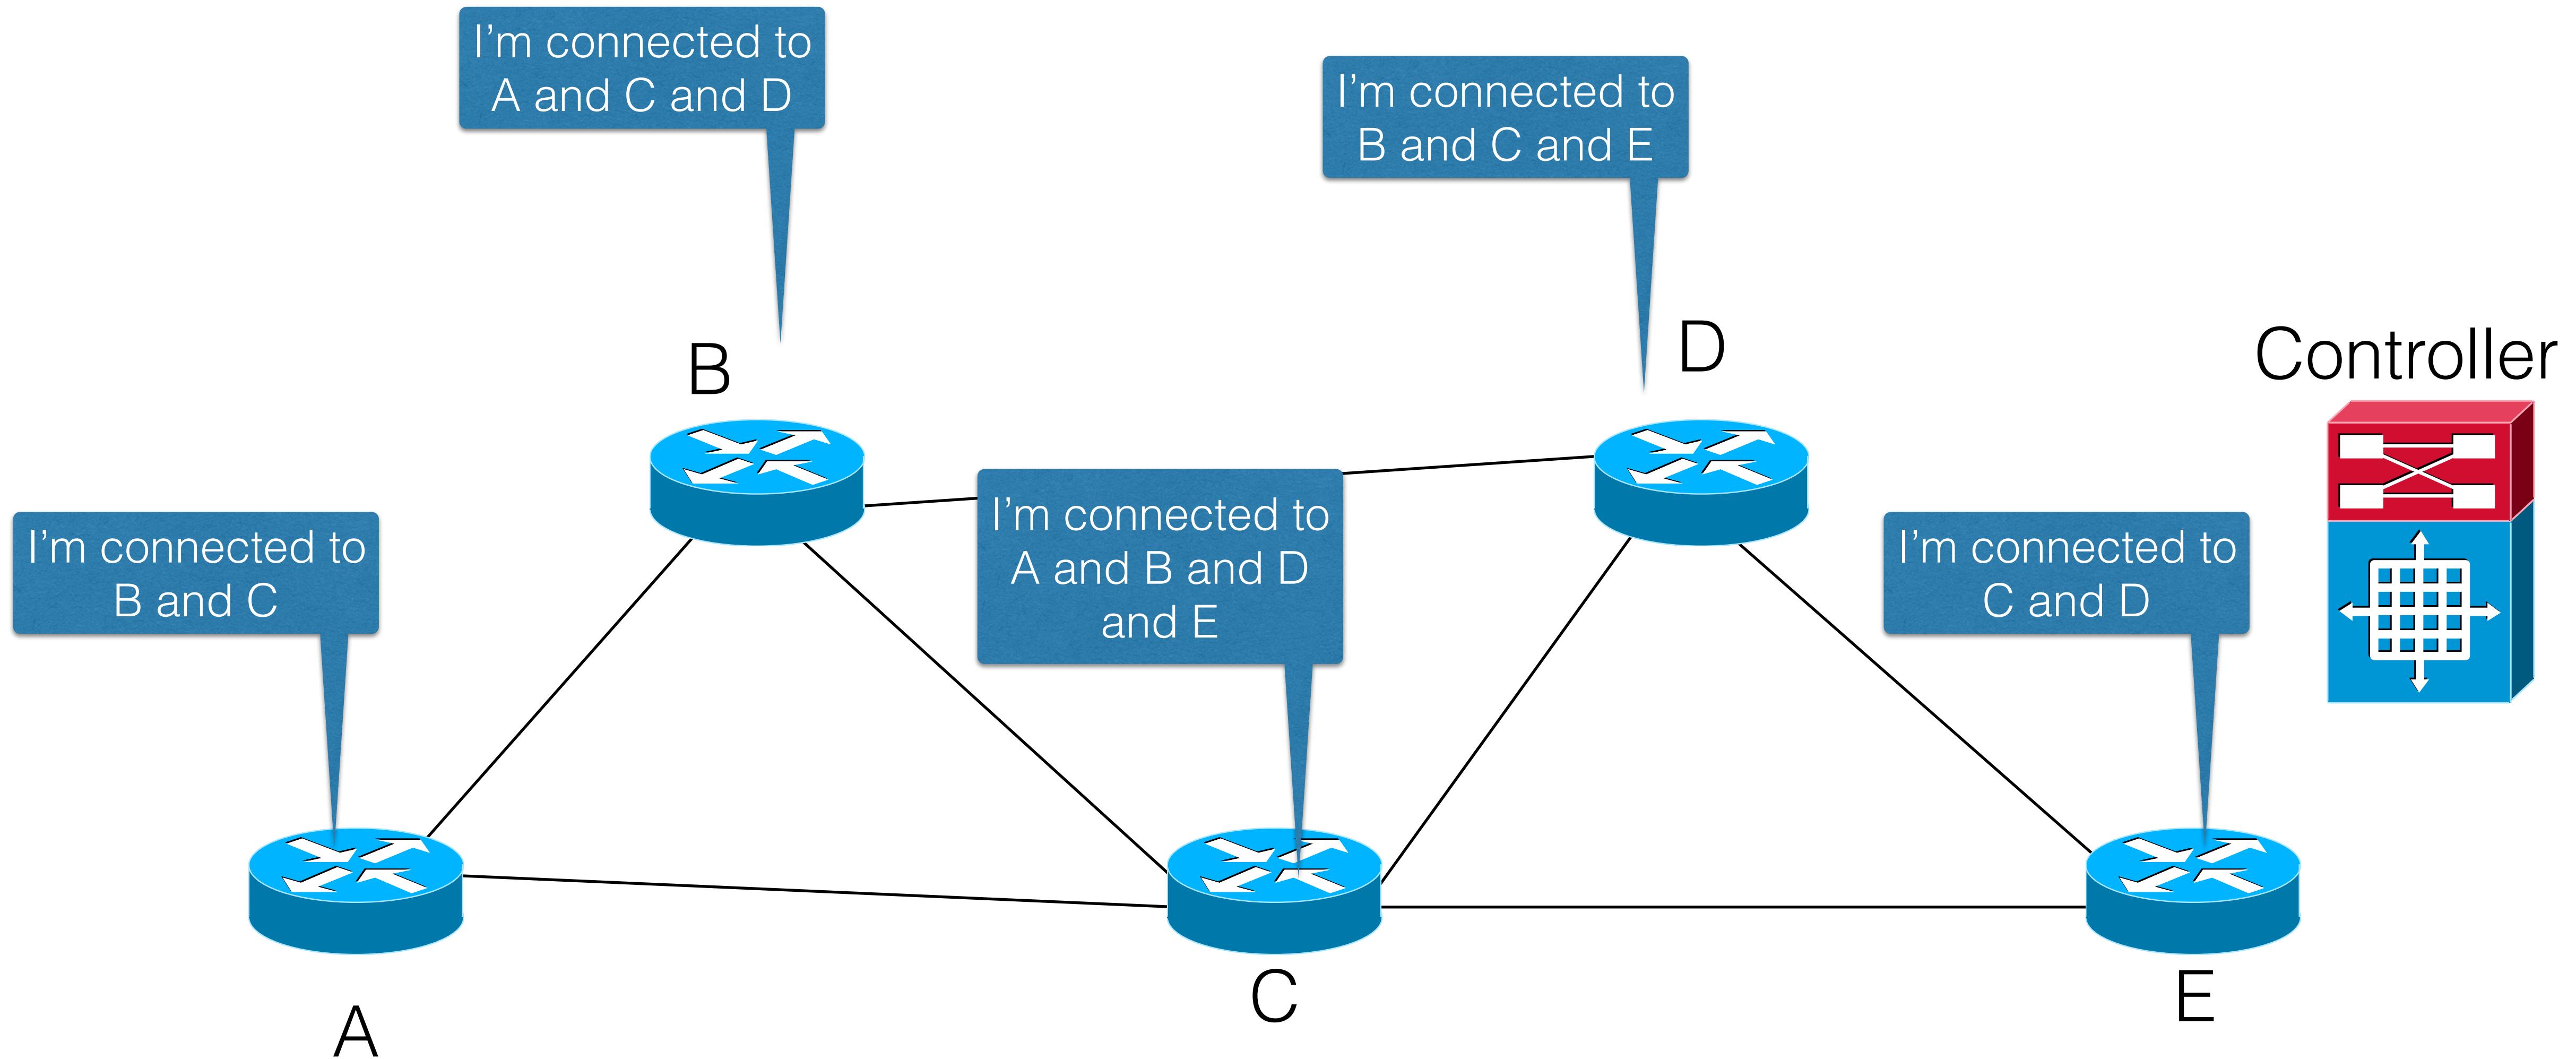


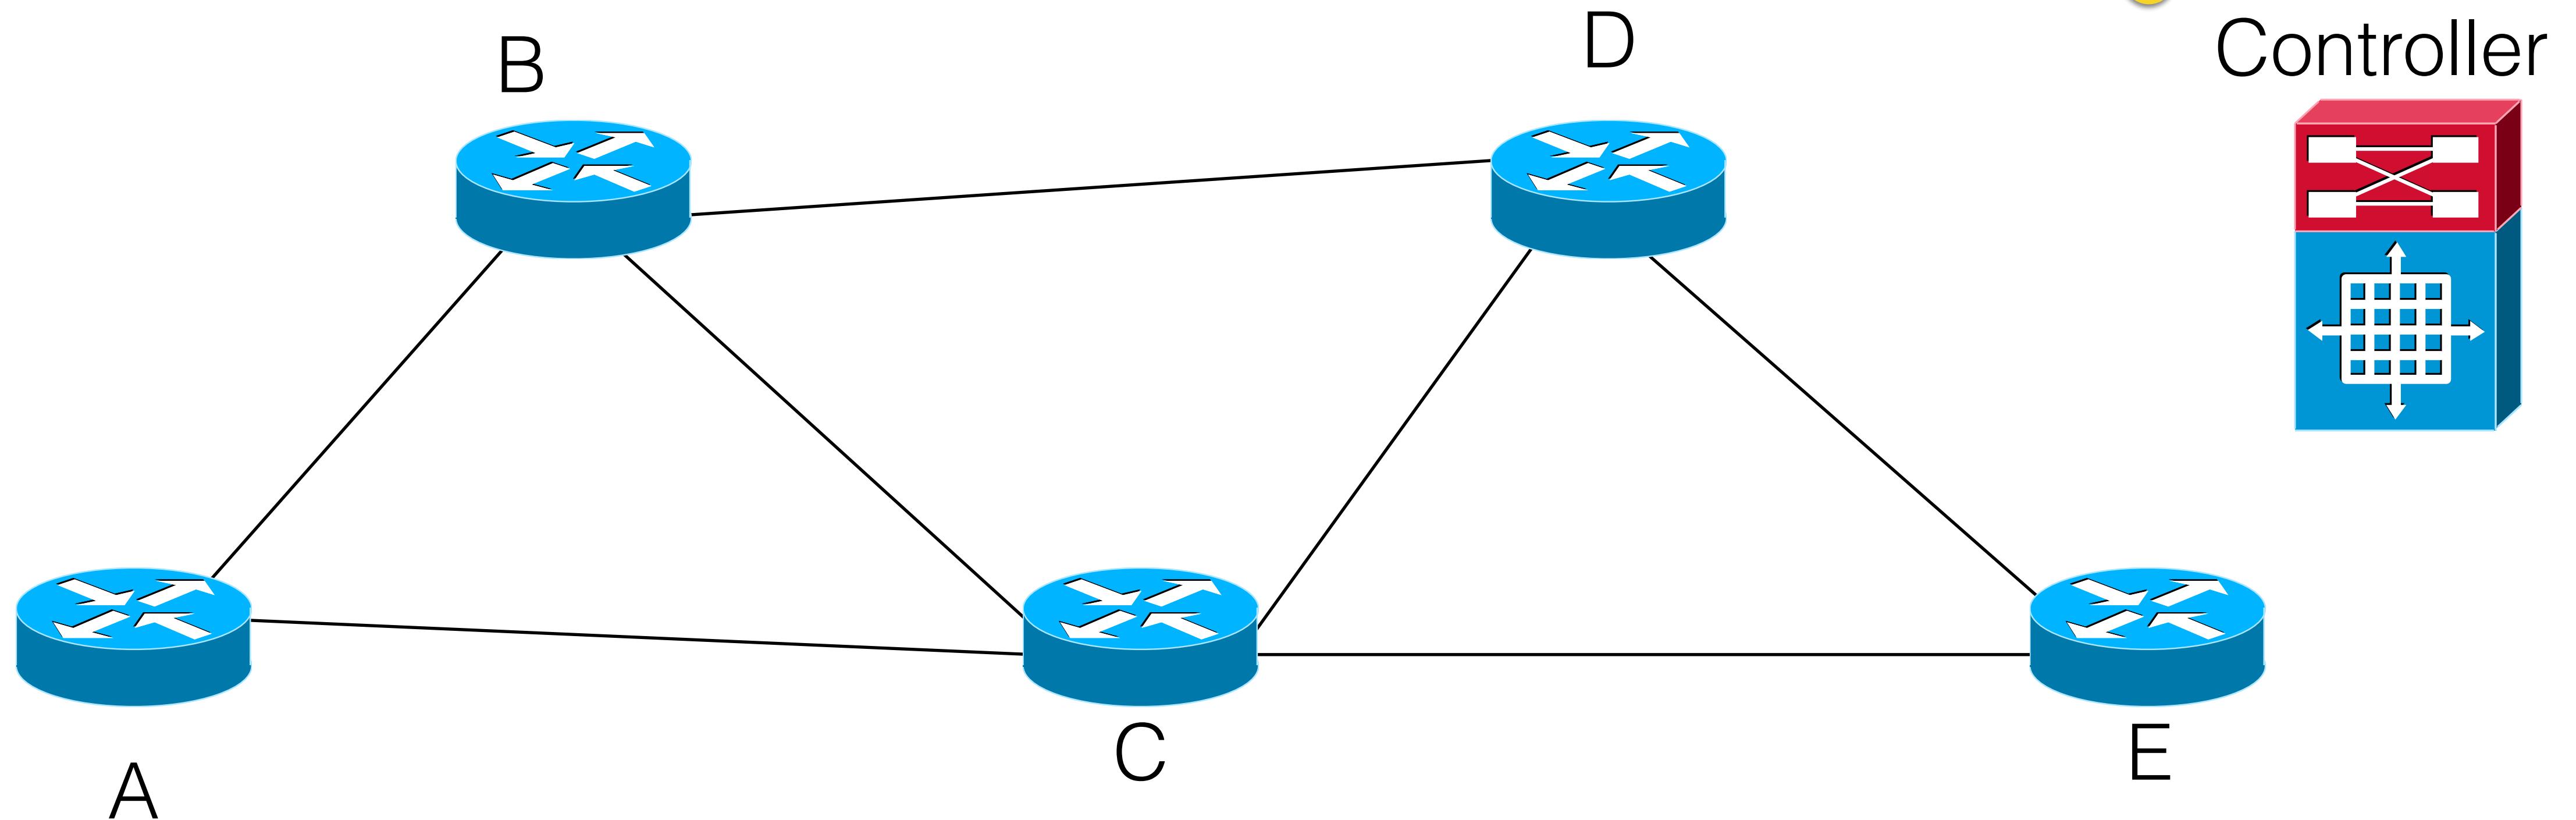
Centralized Routing, In a Nutshell

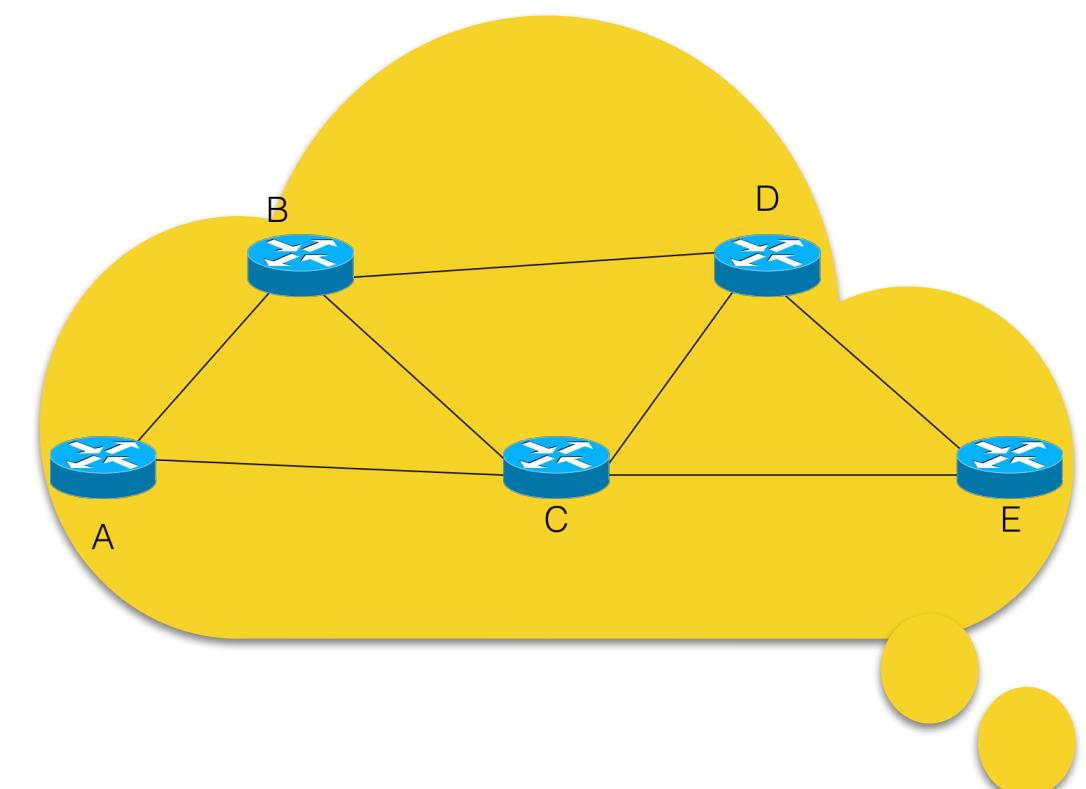
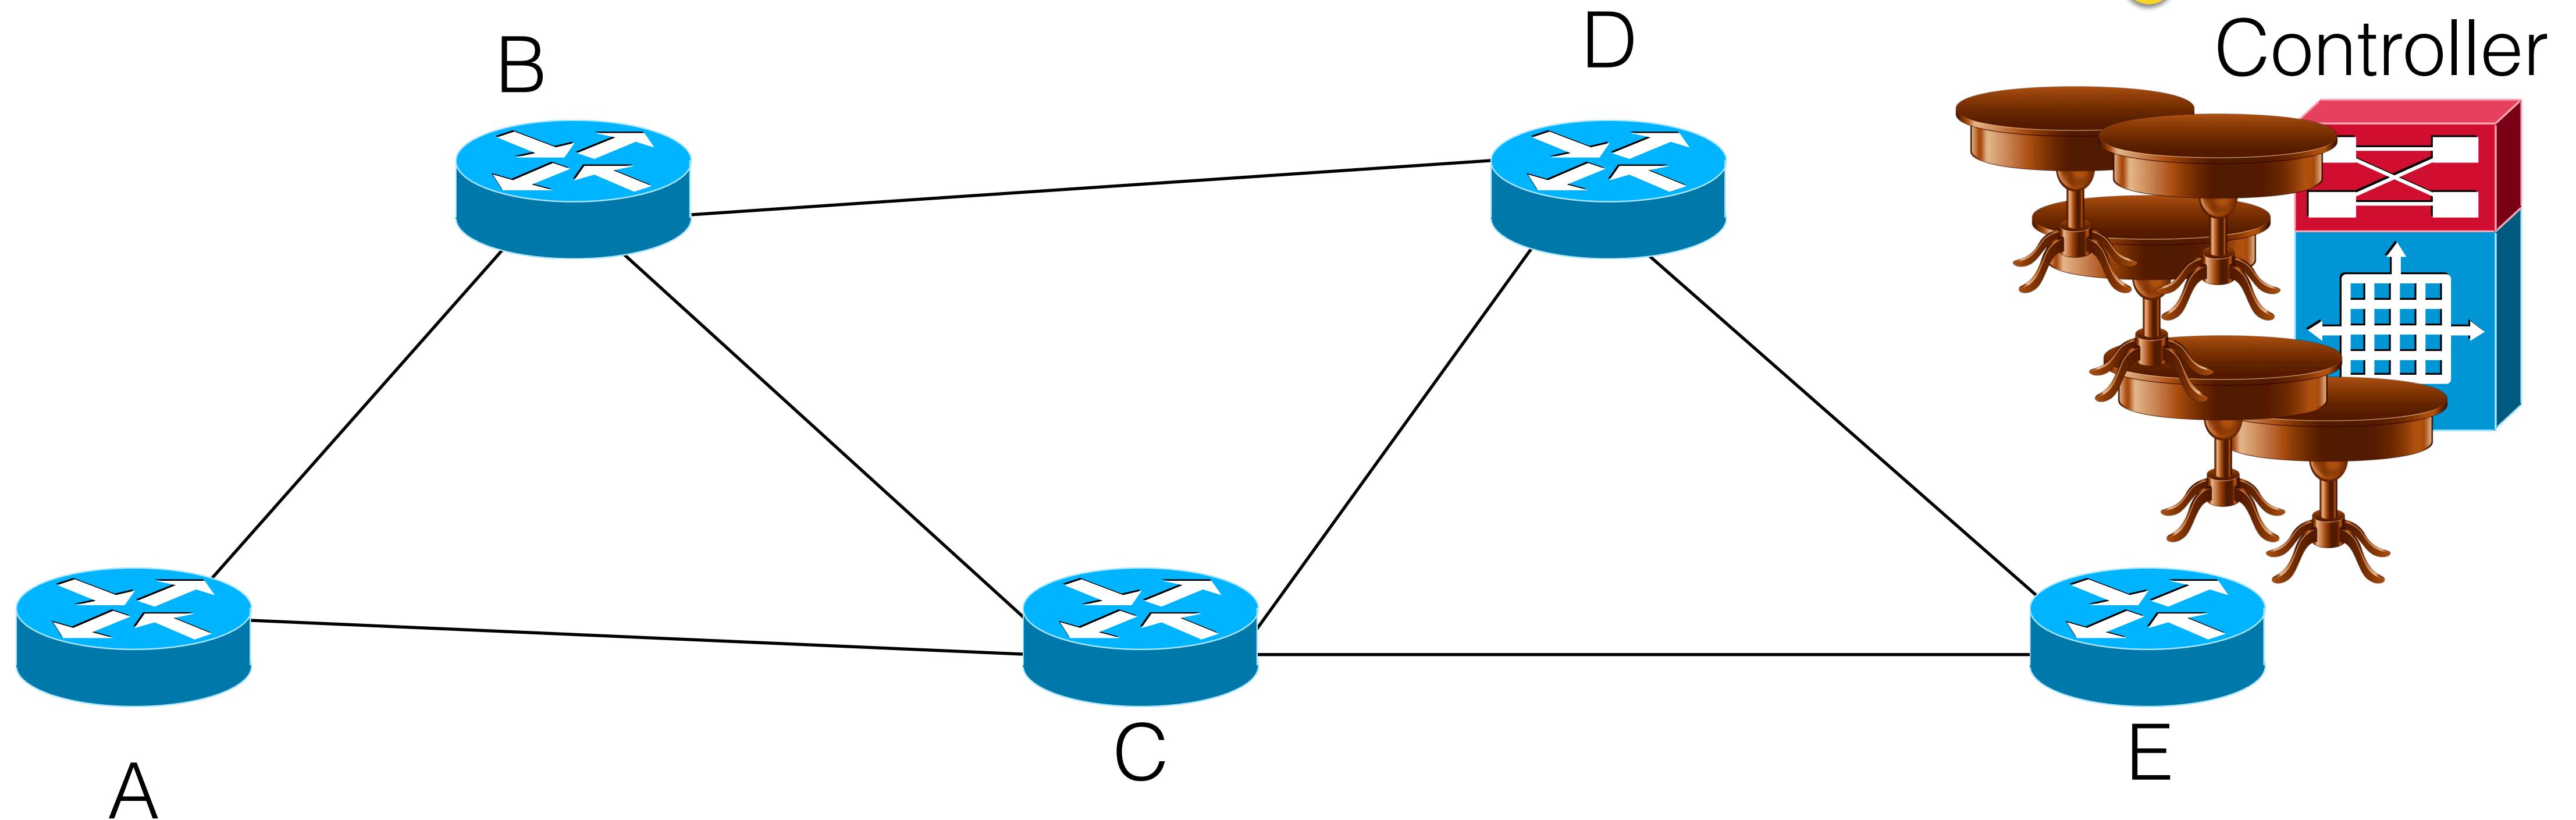
- Like Link State, every node knows who it is connected to.
- Instead of broadcasting to every other node, all nodes tell a special *controller* node who they are connected to.
- The controller computes the best routes for everyone.
- The controller then tells every node what entries to put in their routing tables.











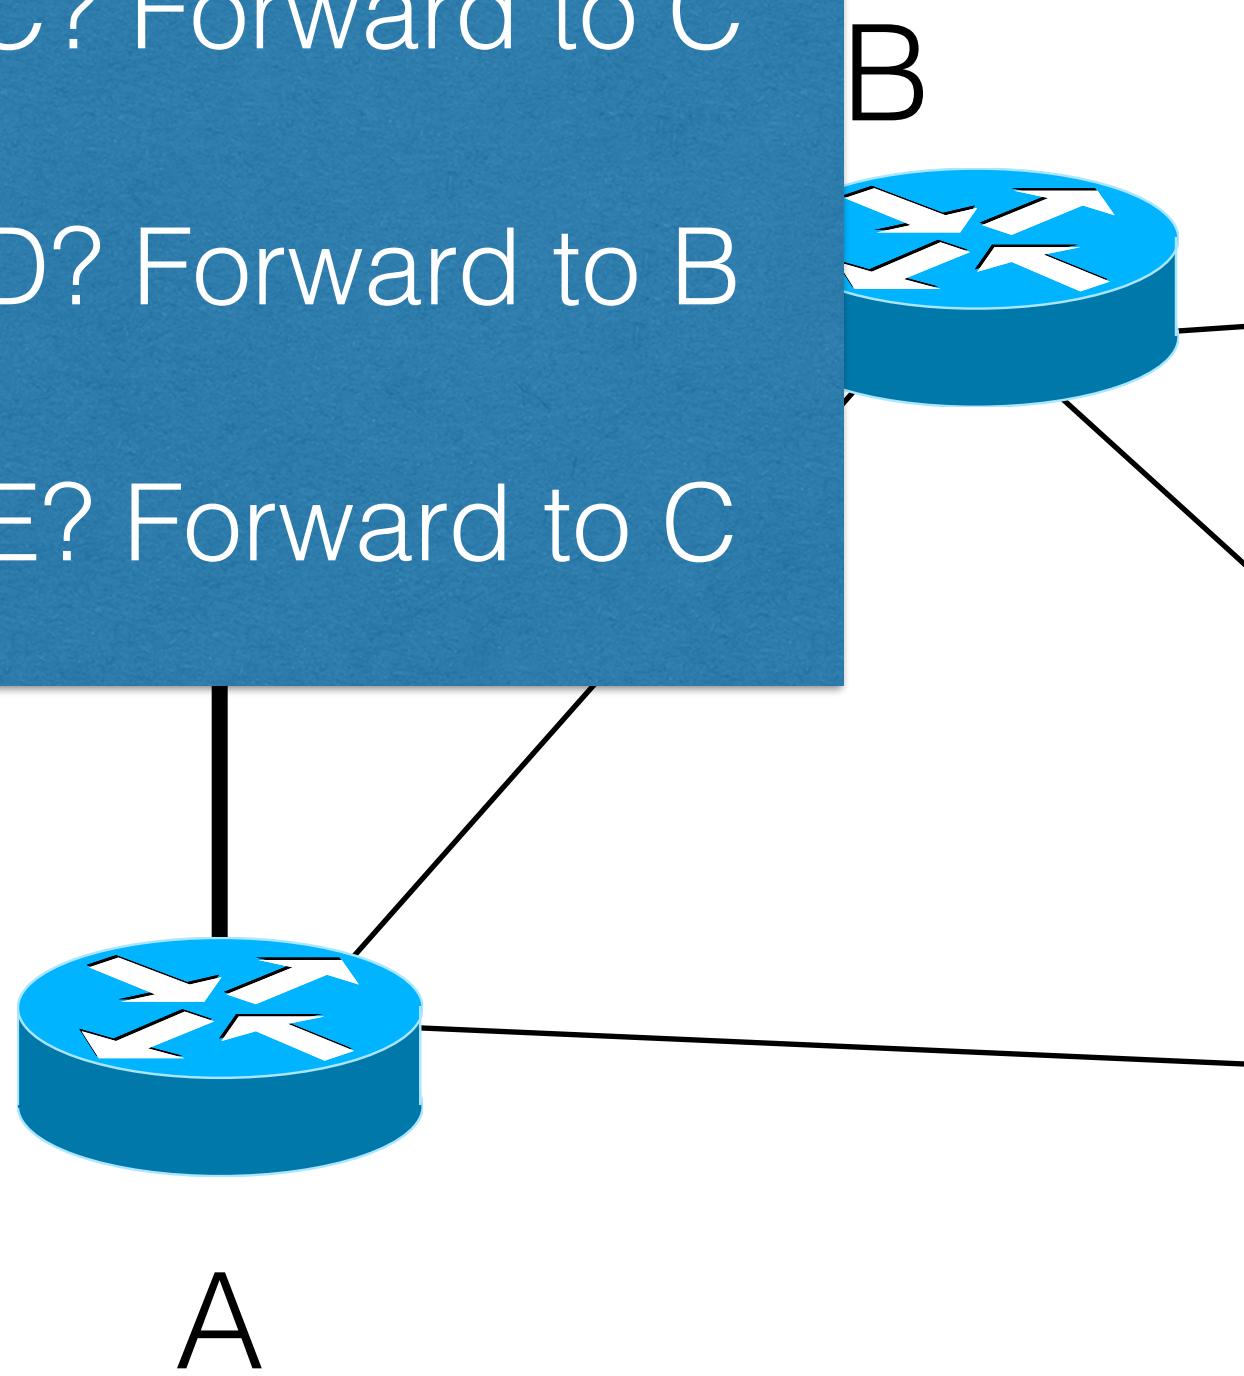
ROUTING TABLE

To B? Forward to B

To C? Forward to C

To D? Forward to B

To E? Forward to C



Now each switch remembers the new routing table.

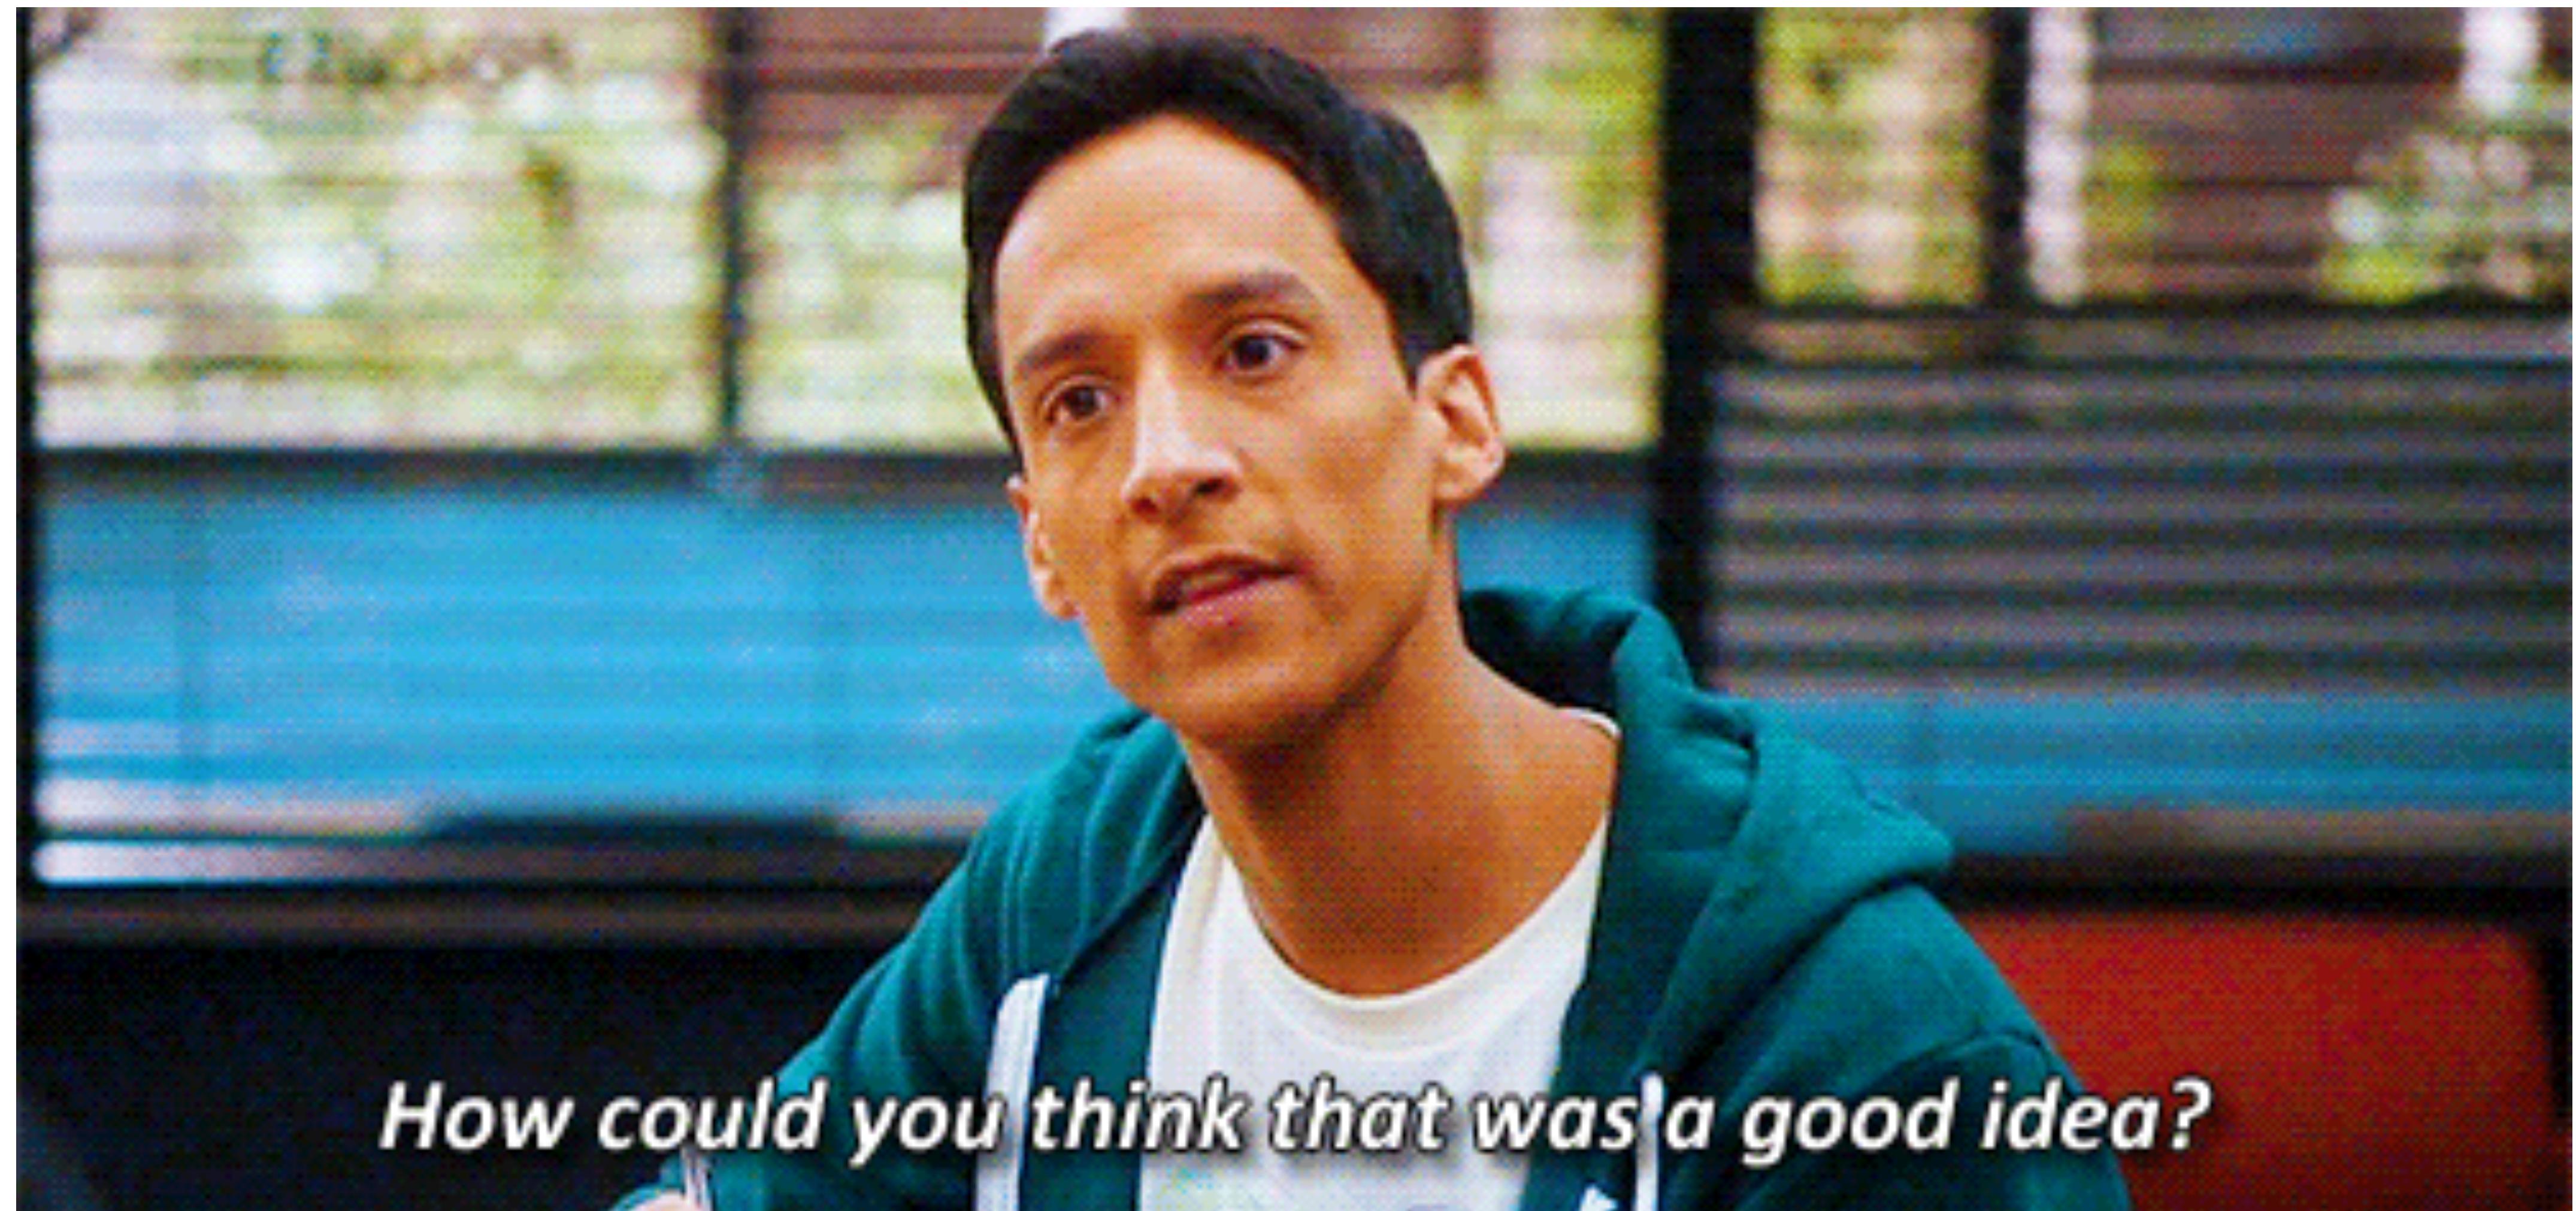
If a link or node fails, the switches notify the controller. The controller re-computes each node's route and sends the new routes out.

	Distance Vector e.g RIP	Link State e.g OSPF	Centralized e.g SDN
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.	Doesn't recover at all if controller(s) fail.
Fully Distributed	Yes	Yes	No
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$	$O(C)$ (Besides Routing Table)
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$	No distributed convergence
Control Plane Complexity	Just select the “min” of all the updates I have heard from. (Dumb-ish Switch)	Rebuild network topology, run Dijkstra’s algorithm over it.	EXTREMELY SIMPLE



	Distance Vector e.g RIP	Link State e.g OSPF	Centralized e.g SDN
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.	Doesn't recover at all if controller(s) fail.
Fully Distributed	Yes	Yes	No
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$	$O(C)$ (Besides Routing Table)
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$	No distributed convergence
Control Plane Complexity	Just select the “min” of all the updates I have heard from. (Dumb-ish Switch)	Rebuild network topology, run Dijkstra’s algorithm over it.	EXTREMELY SIMPLE



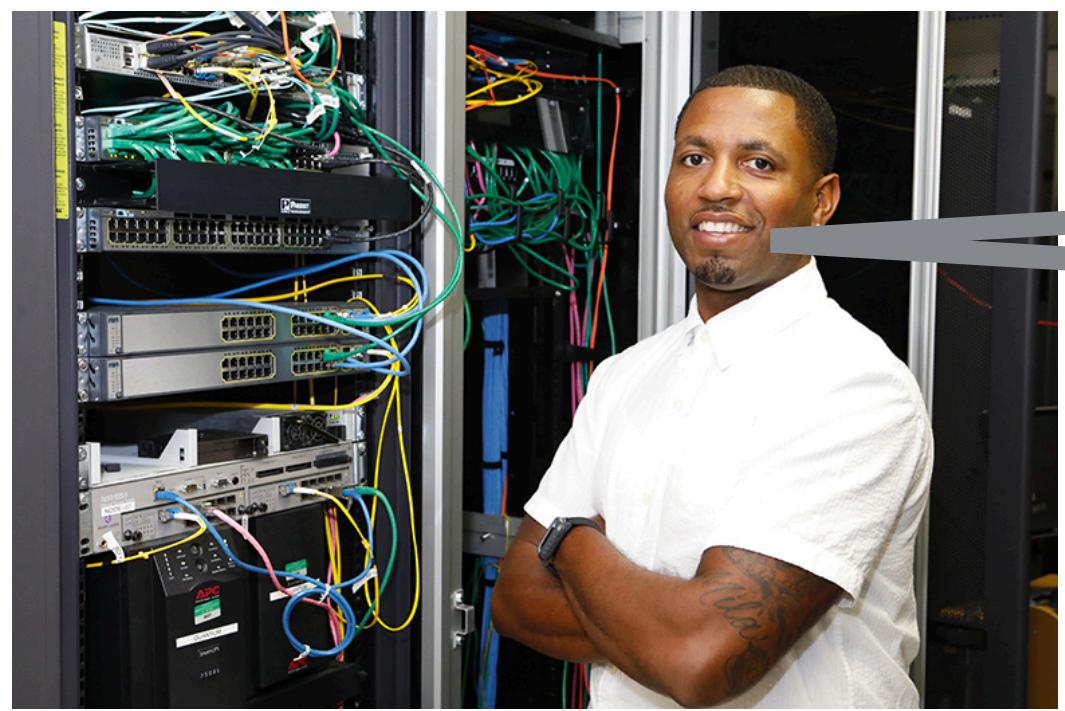


How could you think that was a good idea?

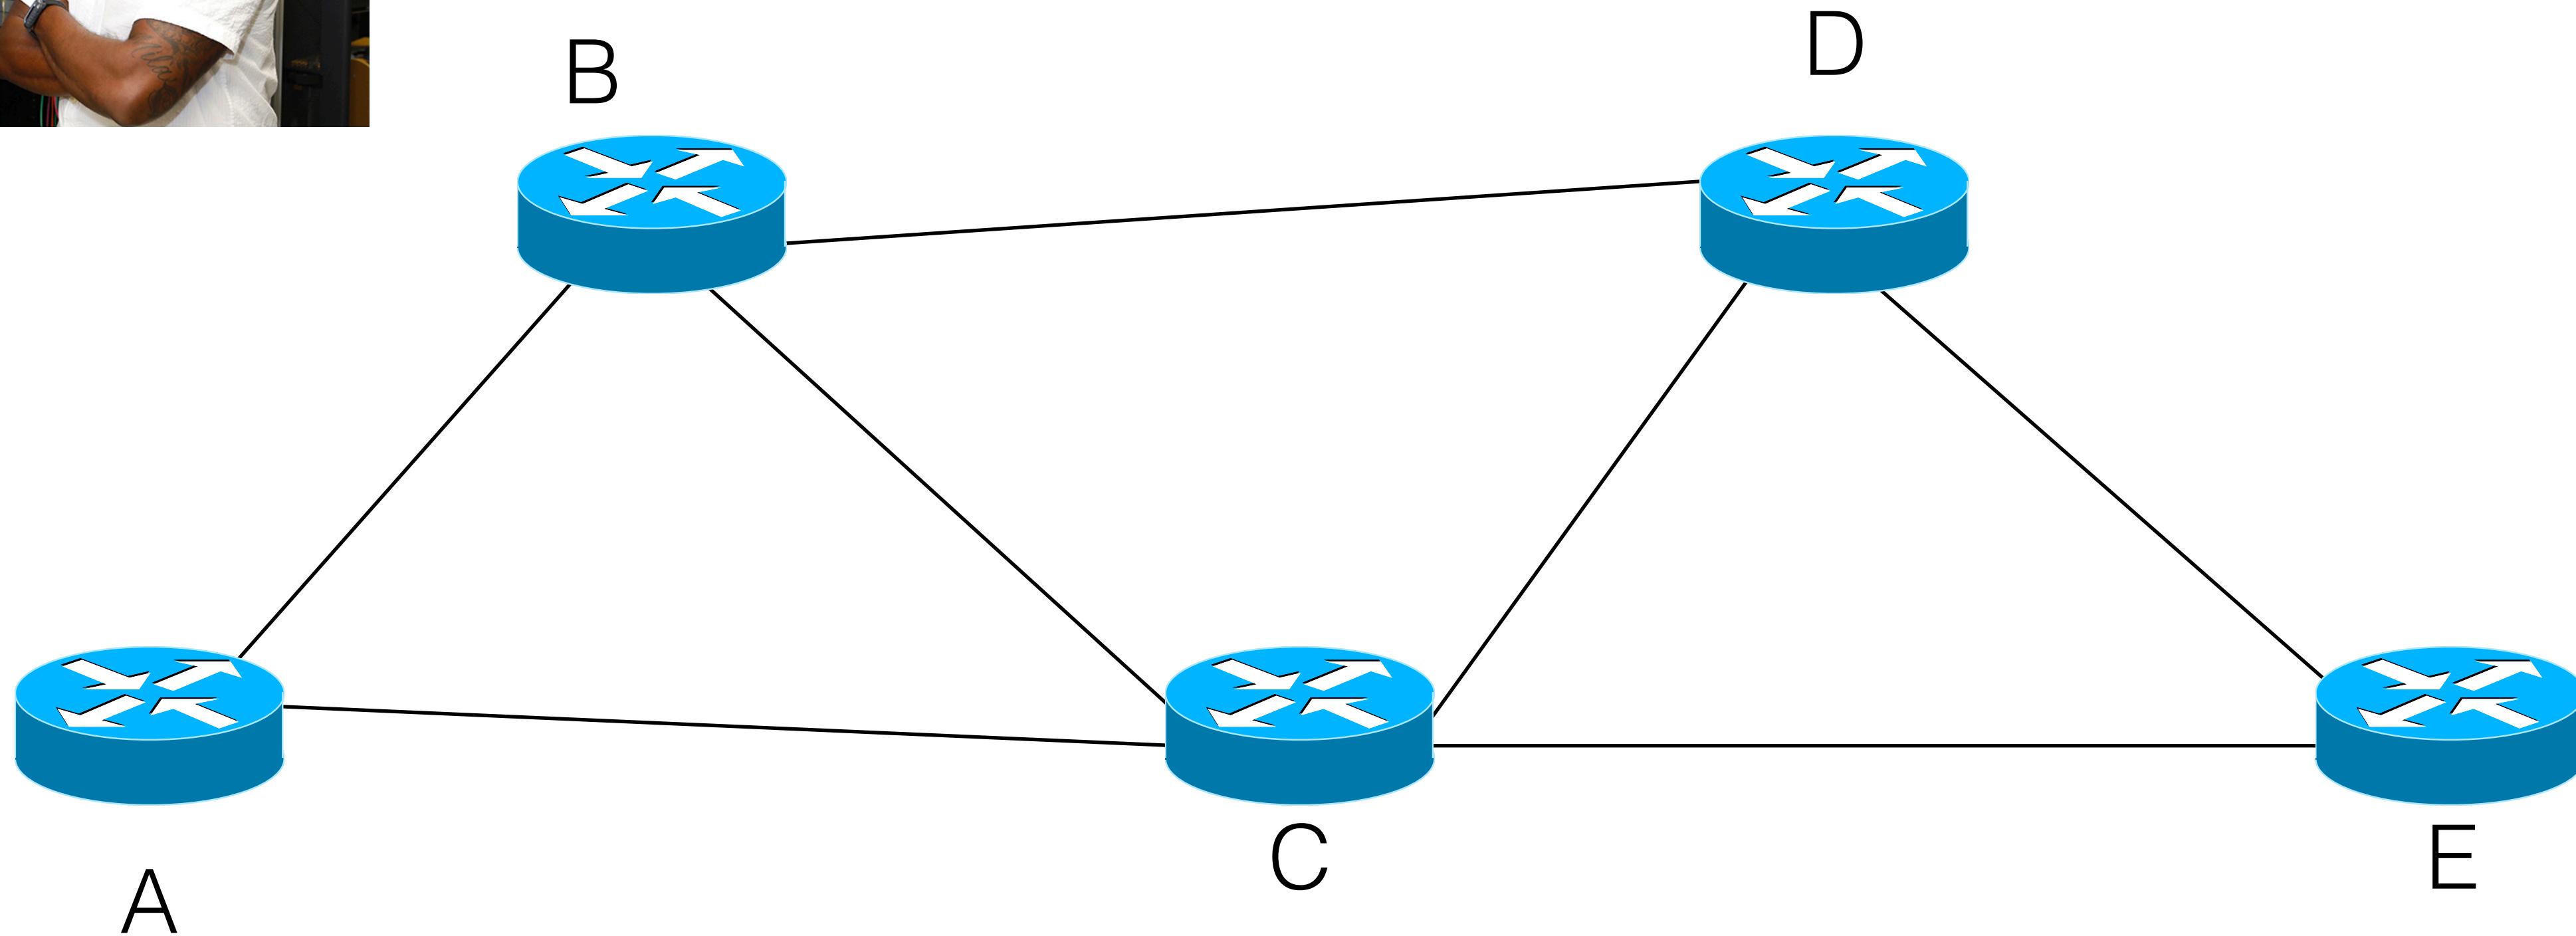


Fun Fact: Centralized Routing is considered “state of the art” — why in the world would people choose this over other designs that are fundamentally more resilient??





I want my network to work normally, EXCEPT B should NOT be allowed to communicate with E.

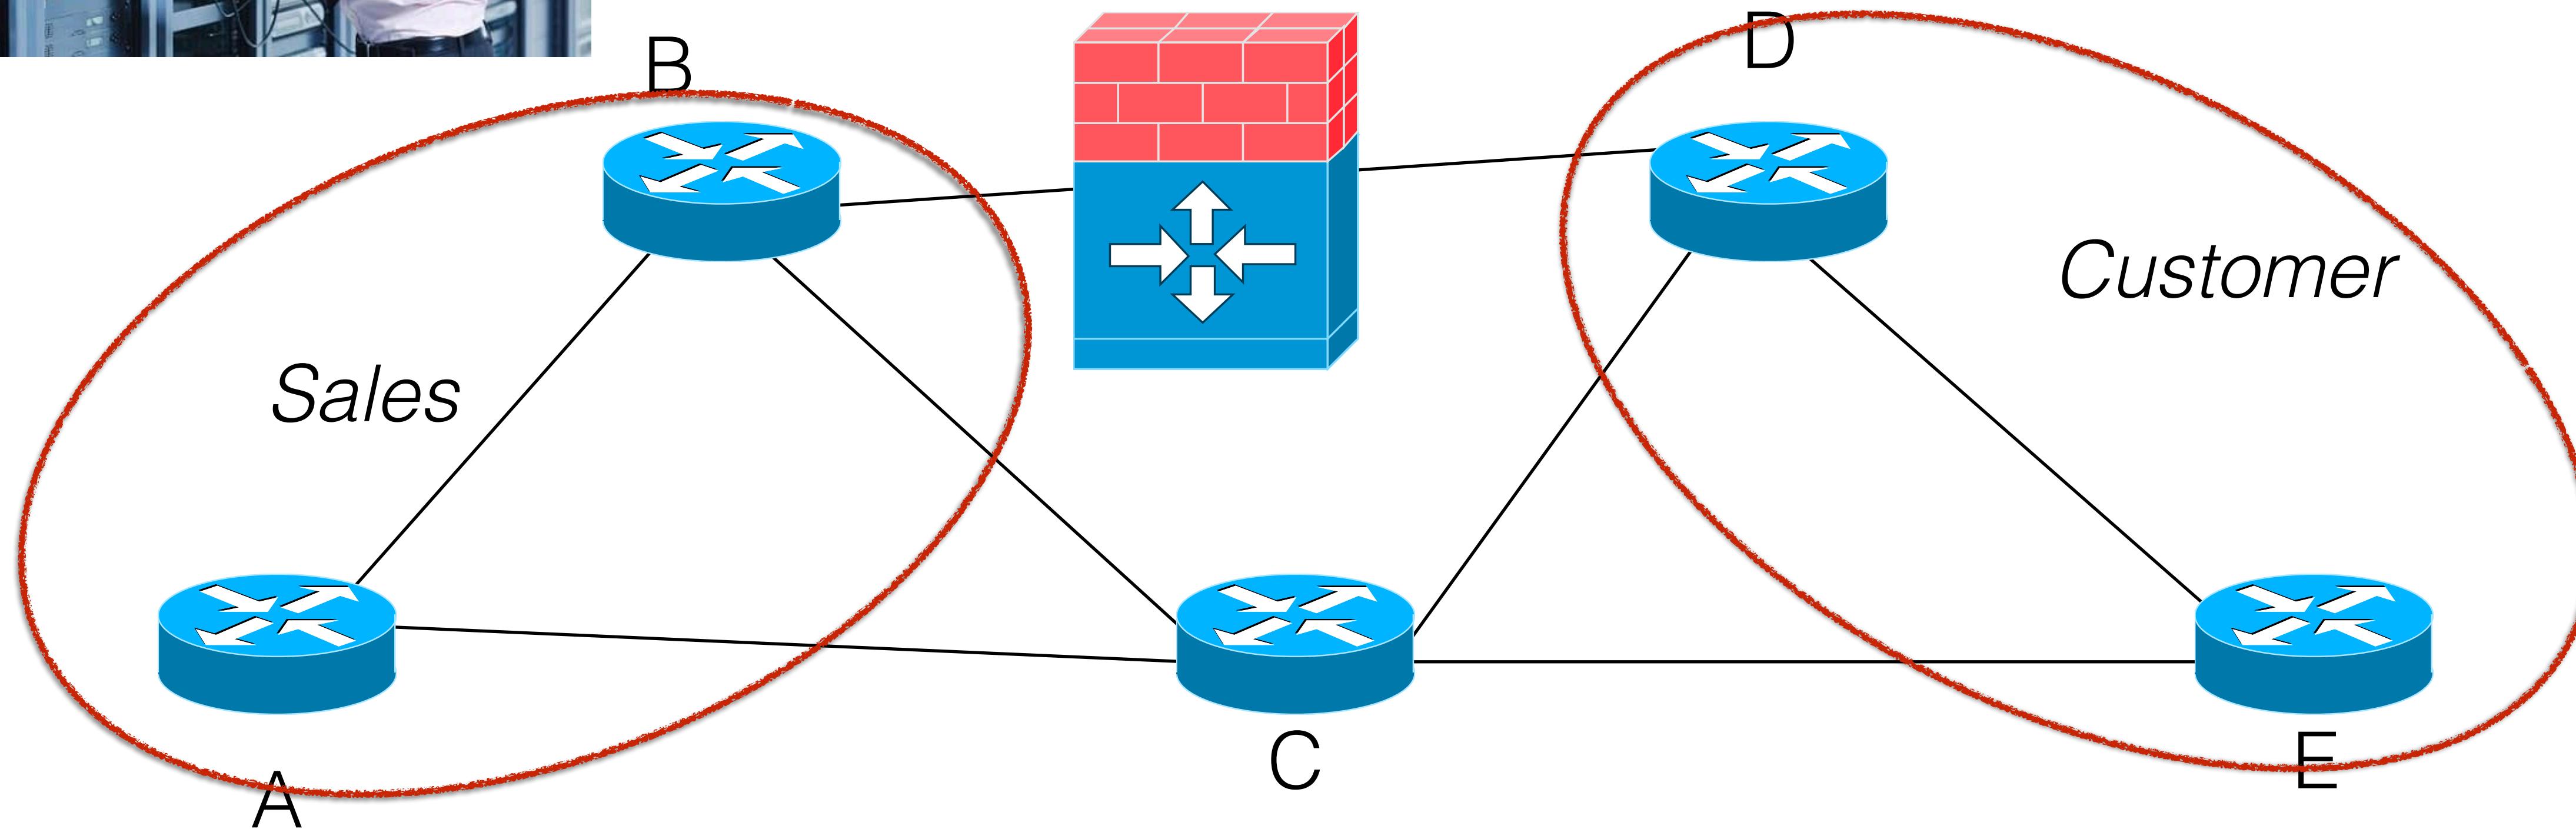


Traditional routing algorithms are designed to achieve global reachability — but can't enforce *policy requirements*.





I need to make sure all traffic going from my sales network to my customer network goes through a firewall.



Traditional routing algorithms are designed to achieve global reachability — but can't enforce *policy requirements*.



Network Policy

- You want to tell the network an “exception” or a “special case”
 - Something to do other than “Let everyone talk to each other!”
- With fully distributed algorithms, you have to distribute the policy
 - And different nodes have to behave differently! You might even need to configure each node specially, depending on the policies.
- With a centralized controller, you configure the controller with your policy. The controller makes the decisions, and the switches don’t have to be configured specially to apply the policy.
 - They just receive their routing tables from the controller.



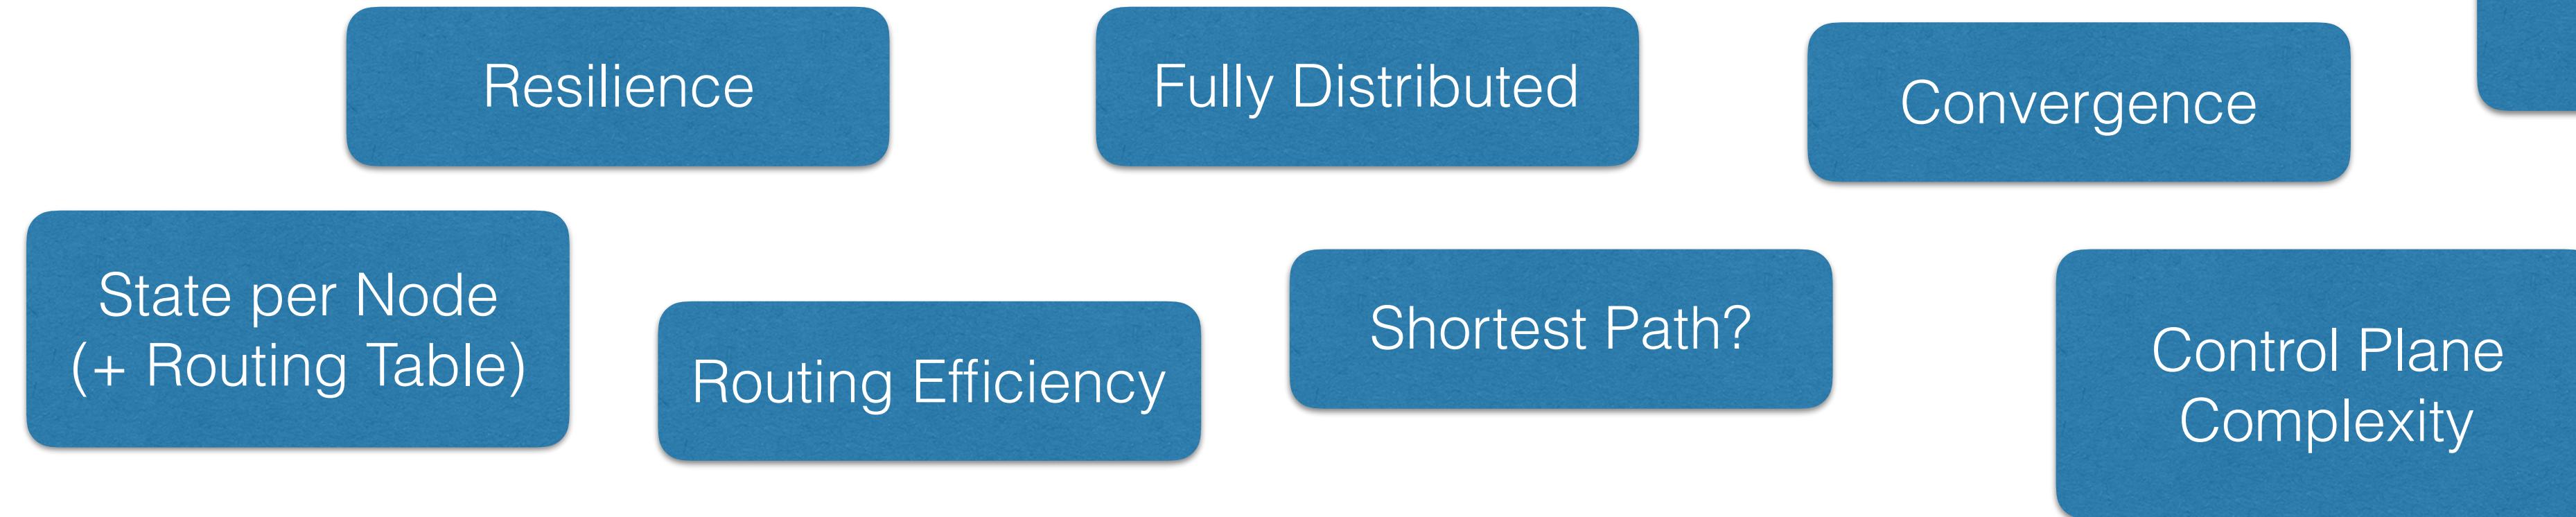
	Distance Vector e.g RIP	Link State e.g OSPF	Centralized e.g SDN
Resilience	Very slow recovery due to count to infinity.	Re-Run Dijkstra and you're good to go.	Doesn't recover at all if controller(s) fail.
Fully Distributed	Yes	Yes	No
State per Node (+ Routing Table)	$O(\# \text{ switches} * \text{max node degree})$	$O(\# \text{ edges})$	$O(C)$
Convergence	Need to run DV before routing — takes length of longest best path time.	Flood network w/ updates and then run Dijkstra: $O(E + V \log V)$	No distributed convergence
Control Plane Complexity	Just select the “min” of all the updates I have heard from. (Dumb-ish Switch)	Rebuild network topology, run Dijkstra’s algorithm over it.	EXTREMELY SIMPLE
Network Policy Support	Hard	Hard	Easy



So what algorithm is best?



Well, it depends...



Do I want cheap, dumb switches, or smart ones?

Is my network big or small?

Do I have network policies to enforce?

Do I need my network to survive a hurricane? earthquake? natural disaster?



In Practice

- My rack of servers in my research lab just uses broadcast routing!
- There are only six machines, connected by one switch.
- Most small setups like this just use broadcast routing.

Your TA Ranysha installing new network cards in the machine room.



In Practice

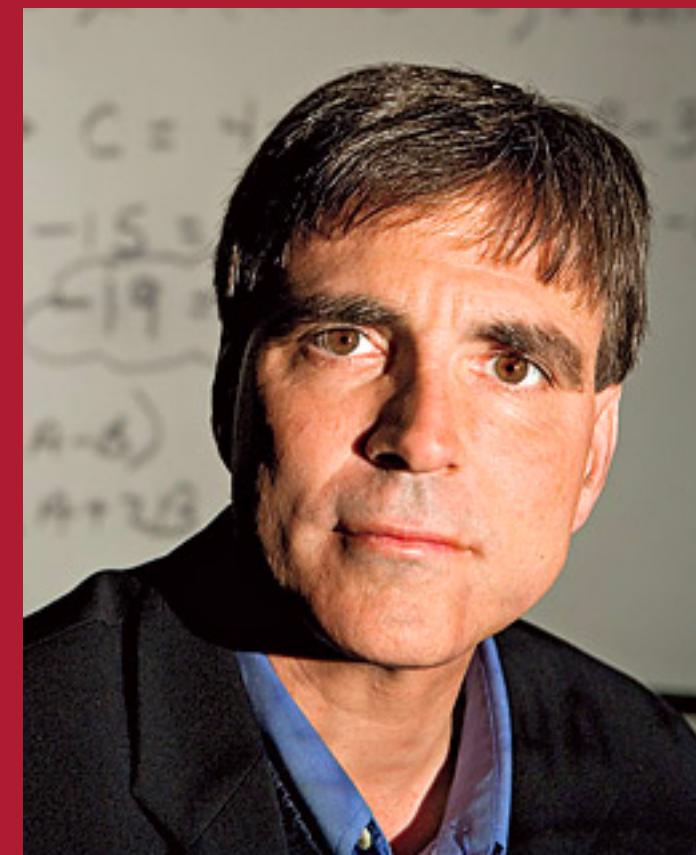


- Google's network spans across the whole world
 - This is called a “WAN” — a “wide area network”
 - Still administered by one organization — so it's one network (not the INTERnet). But it's very big.
- This network is called B4 and it uses a COMBINATION of Link State Routing (OSPF) and Centralized Routing (SDN)
- Just for fun: you can read about this network here: <https://dl.acm.org/citation.cfm?id=2486019>



Systems Engineering Wisdom

“Engineering isn’t about perfect solutions. It’s about doing the best you can with limited resources.”



— Randy Pausch
CMU Professor, ACM Fellow
(October 23, 1960 – July 25, 2008)



You have survived basic routing!

- We will learn a few more routing algorithms later in this class.
 - But now you are ready to move on to learn about...
 - THE INTERNET!!!!
 - Any questions before we move on?

